

UNIVERSITÉ LIBRE DE BRUXELLES

1 Faculté des Sciences
2 Département d’Informatique
3 Année Académique 2019/2020

4 Thèse présentée en vue de l’obtention du grade de Docteur en Sciences

5 ALGORITHMS AND DATA STRUCTURES 6 FOR 7 3SUM AND FRIENDS

8 *by*

9 Aurélien Ooms

10 Jury de thèse:

- 11 Prof. Stefan Langerman (Université libre de Bruxelles, Président)
- 12 Prof. Samuel Fiorini (Université libre de Bruxelles, Secrétaire)
- 13 Prof. Jean Cardinal (Université libre de Bruxelles, Promoteur)
- 14 Prof. John Iacono (Université libre de Bruxelles)
- 15 Prof. Moshe Lewenstein (Bar Ilan University)
- 16 Prof. Nabil Mustafa (ESIEE Paris)

17

To the Profane

18 Solving a problem consists in mapping a given data to an output in some
19 automated way. In Computer Science, we solve problems with computers:
20 we organize the data with data structures, and process those structures
21 with algorithms. Solving problems consumes resources, for example, time.
22 Problems are considered harder if they take more resources to solve.

23 We, computer scientists, are lazy. We do not want to solve the same
24 problem over and over with ad-hoc solutions depending on the data. By
25 automated we really mean that the solutions we design work for any data,
26 and in particular, they work for any data size “ n ”. Ideally, the resource
27 consumption of our solutions should scale well with this parameter. This is
28 why we study the behaviour of such solutions when applied to large¹ inputs.

29 Geometry is about space, points, curves, surfaces . . . We study geometric
30 problems: The given data is geometric, and the question about the data is
31 geometric. The proposed algorithms and data structures therefore exploit
32 geometry.

33 This thesis studies questions about geometric problems that are simple
34 to formulate. However simple they are, nobody has managed to answer them.
35 One of the problems asks to decide, given n points in the plane, whether
36 three of them lie on a common line. Of course, we can solve this problem.
37 The question of interest here is how fast it can be solved.

38 It is easy to solve this problem by testing all possible candidates but that
39 is quite inefficient. There is a more involved but standard construction that
40 takes significantly less time to execute but that is still considered inefficient.
41 As of today, we do not know of any reason why this problem would be much
42 harder than simply reading the data from begin to end.

43 Does this matter at all? It turns out simple problems appear to be
44 bottlenecks of more complex ones: unless we manage to improve our under-

1 *Gigantic multiplied by colossal multiplied by staggeringly huge is the sort of concept we're trying to get across here.* – Douglas Adams, The Restaurant at the End of the Universe.

45 standing of the simple ones, we are stuck with inefficient solutions for all of
46 them.

47 In this thesis, we expose novel algorithms and data structures related to
48 the problem stated above. Hopefully, this contribution will one day, directly
49 or indirectly, help us solve some of the interesting open questions about this
50 problem.

51

Token of Appreciation

52

TBD

Contents

53

54	To the Profane	i
55	Token of Appreciation	iii
56	Table of Contents	vii
57	On Notation	viii
58	0 In a Dozen Pages	1
59	I Without Proof	14
60	1 Models of Computation	15
61	1.1 Algorithms	15
62	1.1.1 Random Access Machines	16
63	1.1.2 Computation and Decision Trees	16
64	1.2 Data Structures	19
65	1.2.1 Encodings	19
66	2 History	21
67	2.1 3SUM & k -SUM	21
68	2.1.1 Variants	22
69	2.1.2 Point Location	23
70	2.1.3 Information Theoretic Lower Bound	23
71	2.1.4 Higher Lower Bounds	24
72	2.1.5 Uniform Algorithms	25
73	2.1.6 Nonuniform Algorithms	27
74	2.2 GPT & 3POL	28
75	2.2.1 Variants	29
76	2.2.2 Reductions from k -SUM	29

77	2.2.3	Lower Bounds and Order Types	31
78	2.2.4	Algorithms	33
79	2.2.5	More on Order Types	33
80	2.2.6	Encodings	34
81	2.2.7	The Intermediate Problem	36
82	2.2.8	Combinatorics	37
83	3	Contributions	39
84	3.1	Meiser Applied to k -SUM	39
85	3.2	Grønlund and Pettie Applied to 3POL	42
86	3.3	Slightly Subquadratic Encodings for GPT	44
87	3.4	Better Encodings for 3SUM	47
88	4	Developments	49
89	4.1	Better Nonuniform Algorithms for k -SUM	49
90	4.1.1	Using Vertical Decomposition	49
91	4.1.2	Using Inference Dimension	50
92	4.2	Timothy Chan Strikes Again	51
93	5	Open Questions	52
94	5.1	About Algorithms	52
95	5.2	About Encodings	54
96	II	The Computational Geometer’s Toolbox	55
97	6	Arrangements	56
98	6.1	Counting Cells	56
99	6.2	Pseudolines	57
100	6.3	Zone Theorem	58
101	6.4	Cell Decomposition	59
102	6.4.1	Bottom Vertex Triangulation	59
103	6.4.2	Vertical Decomposition	60
104	7	Chirotopes	62
105	7.1	Duality	62
106	7.2	Canonical Labelings	64

107	8 Divide and Conquer	66
108	8.1 Epsilon Nets and Cuttings	66
109	8.1.1 Range Spaces	66
110	8.1.2 VC-dimension	67
111	8.1.3 Epsilon Nets	68
112	8.1.4 Hyperplanes in Linear Dimension	68
113	8.1.5 Cuttings	69
114	8.1.6 Algebraic Range Spaces	70
115	8.1.7 Derandomization	70
116	8.2 Hierarchical Cuttings	70
117	9 Existential Theory of the Reals	73
118	9.1 Cylindrical Algebraic Decomposition	73
119	III Algorithms	75
120	A Solving k-SUM using Few Linear Queries	76
121	A.1 Meiser Solves k -SUM	77
122	A.1.1 Query Complexity	77
123	A.1.2 Time Complexity	81
124	A.1.3 Query Size	85
125	A.2 Missing Details	86
126	A.2.1 Keeping Queries Linear in Algorithm 1	86
127	A.2.2 Algebraic Computation Trees	88
128	A.2.3 Uniform Random Sampling	89
129	A.2.4 Proof of Lemma A.5	90
130	B Subquadratic Algorithms for Algebraic 3SUM	93
131	B.1 First Subquadratic Algorithms for 3POL	93
132	B.1.1 Nonuniform Algorithm for Explicit 3POL	94
133	B.1.2 Uniform Algorithm for Explicit 3POL	101
134	B.1.3 Nonuniform Algorithm for 3POL	104
135	B.1.4 Uniform Algorithm for 3POL	110
136	B.2 Subproblems	115
137	B.2.1 Polynomial Batch Range Searching	115
138	B.2.2 Polynomial Dominance Reporting	118

139	B.3 Applications	123
140	B.3.1 GPT for Points on Curves	124
141	B.3.2 Incidences on Unit Circles	126
142	B.3.3 Points Spanning Unit Triangles	128
143	IV Data Structures	129
144	C Subquadratic Encodings for Point Configurations	130
145	C.1 Encoding Order Types via Hierarchical Cuttings	131
146	C.2 Sublogarithmic Query Complexity	142
147	C.3 Higher-Dimensional Encodings	149
148	D Encoding 3SUM	155
149	D.1 Representation by Numbers	155
150	D.2 Space-Optimal Representation	158
151	D.3 Subquadratic Space and Constant Query Time	158
152	Bibliography	161
153	List of Contributions and Open Questions	176

154

On Notation

155

This short chapter hopes to lift any ambiguity in the notation used.

156

Big-Oh

157

To express the asymptotic behaviour of functions representing resource complexities (time and space) we use the standard *big-oh* notation (see for instance [55, Chapter 3]). For brevity, we add a few tweaks.

160

- The notation $\tilde{O}(\cdot)$ ignores polylogarithmic factors. For instance, we have $n^3 \log^2 n = \tilde{O}(n^3)$.

162

- The symbol ϵ (not to be mistaken for ε) denotes a positive real number that can be made as small as desired. Writing $T(n) = O(n^{12/7+\epsilon})$ means that for any fixed $\delta > 0$, $T(n) = O(n^{12/7+\delta})$, where the constant of proportionality may depend on δ . In particular, $n^2 \log^2 n = O(n^{2+\epsilon})$.

166

- The notation $O_{p_1, p_2, \dots}(f(n))$ means that the constant of proportionality depends on the parameters p_i .

168

- When we write $O(f(n_1, n_2, \dots))$ we assume that one of the variables n_i is the one going towards infinity in the big-oh definition, while the others are monotone functions of n_i (increasing or decreasing depending on the context).

172

Because this asymptotic notation takes little care of constant factors, logarithms are in base two unless otherwise indicated.

174

Sets

175

We denote by \mathbb{R}^d the d -dimensional Euclidean space and try to be consistent with the set notation used to represent the subsets of this space.

177

- A point is indicated by a lowercase letter, for instance a vertex p .

- 178 • A set of points is indicated by an uppercase letter, for instance an
179 hyperplane H , a cell C , or a simplex S .
 - 180 • A set of sets of points is indicated by a calligraphic uppercase letter,
181 for instance a set of hyperplanes \mathcal{H} , a net \mathcal{N} , or an arrangement $\mathcal{A}(\mathcal{H})$.
- 182 For finite sets, we sometimes use the short notation $[n] = \{1, 2, \dots, n\}$
183 and describe a set of cardinality n as a n -set.

0

In a Dozen Pages

186 This thesis is a compilation of the contributions from four papers:

- 187 A “*Solving k -SUM Using Few Linear Queries*” with Jean Cardinal and
188 John Iacono [41].
- 189 B “*Subquadratic Algorithms for Algebraic 3SUM*” with Luis Barba, Jean
190 Cardinal, John Iacono, Stefan Langerman, and Noam Solomon [20].
- 191 C “*Subquadratic Encodings for Point Configurations*” with Jean Cardinal,
192 Timothy Chan, John Iacono, and Stefan Langerman [38].
- 193 D “*Encoding 3SUM*” with Sergio Cabello, Jean Cardinal, John Iacono,
194 Stefan Langerman, and Pat Morin [37].

195 Paper A was presented at CG:YRF 16 and ESA 16. Paper B was presented
196 at EuroCG 17 and SoCG 17. It is published in DCG. Paper C was presented
197 at FWCG 17, EuroCG 18, and SoCG 18. It is published in JoCG. Paper D
198 was presented at EuroCG 19.

199 We begin this thesis with an overview of the studied topic. We explain
200 the context in which those papers were written and expose the contributions
201 contained in each of them.

202 Degeneracy Testing Problems

203 In this thesis, we study *Degeneracy Testing Problems*: an instance of size
204 n of such a problem is a single point in high-dimensional euclidean space
205 $q \in \mathbb{R}^{O(n)}$. Such an instance is called *general* if and only if it passes a series
206 of algebraic tests (usually $n^{O(1)}$ of them). If it fails one of the tests, it is

207 called *degenerate*. Our goal is to determine how fast an instance can be
208 classified as general or degenerate.

209 The terminology is justified because most instances are general: the set
210 of degenerate instances is a zero-measure subset of the input space. It also
211 makes sense to visualize the input space as the euclidean space: the algebraic
212 tests naturally induce a partition of the input space into semialgebraic sets.
213 Solving the problem therefore amounts to locate the input point q in this
214 partition of space. Our goal is thus to determine how fast this input point
215 can be located.

216 Degeneracy testing problems are easy decision problems because there
217 are only a finite number of candidate tests to try. The ones we study can all
218 be solved by brute-force in polynomial time because the number of tests is
219 polynomial. We show how, in some cases, this naive approach is definitely
220 subsumed by divide and conquer techniques exploiting the geometry of the
221 setting.

222 GPT

223 Let us illustrate by giving a first example of a degeneracy testing problem.
224 We begin with a definition:

225 **Definition 1.** A set of n points in \mathbb{R}^d is in general position if and only if
226 every $(d + 1)$ -subset spans the entire space. A point set that is not in general
227 position is called *degenerate*.

228 The General Position Testing problem (GPT) asks to decide if a given set
229 of n point is in general position. We can solve this problem by brute-force in
230 $O(n^{d+1})$ time. We can do it an order of magnitude faster by constructing the
231 dual arrangement of hyperplanes in $O(n^d)$ time. On the one hand, improving
232 on this slightly better solution appears to be non-trivial: there exists a class
233 of algorithms that cannot do better even though they exploit one of the
234 core structures of the problem, the chirotope axioms. On the other hand,
235 information theory only gives a decision tree lower bound of $\Omega(n \log n)$. A
236 popular conjecture is that no $o(n^d)$ time real-RAM algorithm exists for this
237 problem.

238 Nonuniform Algorithms

239 Another model of computation in which no $o(n^d)$ time algorithm for
240 GPT is known is the algebraic computation tree model. In this model,
241 an algorithm is a rooted tree whose internal nodes are either arithmetic
242 operations or sign tests on real variables, and whose leaves are the result of
243 the computation.

244 This model is more generous than the real-RAM model in the sense that
245 all computations that can be carried out by only knowing the input size incur
246 no cost. Because a computation tree has a fixed size, we need a different
247 tree for each input size. Therefore, we say that this model is *nonuniform*
248 since it allows to have a distinct algorithm for each input size.

249 This thesis considers both uniform algorithms in the real-RAM and word-
250 RAM models of computation and nonuniform algorithms in the algebraic
251 computation tree, algebraic decision tree, and linear decision tree models of
252 computation. For a given task, the complexity of the nonuniform algorithm
253 is less than the complexity of the uniform algorithm. While a nonuniform
254 algorithm is rarely practical, designing those at least means making progress
255 on the question of whether a sensible computation tree lower bound can be
256 derived.

257 3SUM

258 The 3SUM problem also falls in the category of degeneracy testing
259 problems. This problem asks to decide whether a given set of n numbers
260 contains a triple whose sum is zero. We can solve this problem by brute-force
261 in $O(n^3)$ time, and in $O(n^2)$ time with a slightly more clever algorithm.

262 However toyish 3SUM may look like, it is considered one of several key
263 problems in P: many geometric problems reduce from it in subquadratic time.
264 Hence, a conjectured quadratic lower bound on 3SUM implies a conditional
265 lower bound on all those more practical problems.

266 Like for GPT, there exist lower bounds for 3SUM in restricted models
267 of computation: 3SUM cannot be decided in $o(n^2)$ time if the only way we
268 inspect the input is by testing for the sign of weighted sums of three input
269 numbers.

270 Even before this lower bound was known, it was conjectured that a
271 quadratic lower bound would hold in other models of computation like the

272 real-RAM model.

273 A first stab at the conjecture was made when it was proven that for
 274 integer input numbers, it is possible to beat the conjectured lower bound
 275 by a few logarithmic factors [19]. However, it remained open whether such
 276 improvements were possible for real inputs.

277 Eventually, in a breakthrough paper, Grønlund and Pettie gave a sub-
 278 quadratic uniform algorithm that shaves a root of a logarithmic factor from
 279 quadratic time [91]. Since then more roots of logarithmic factors have been
 280 shaved [79, 81]. To this day, it is still conjectured that, for all $\delta > 0$, 3SUM
 281 requires $\Omega(n^{2-\delta})$ time to solve in the real-RAM model.

282 k -SUM

283 The paper of Grønlund and Pettie also discusses the following general-
 284 ization of the 3SUM problem: “For a fixed k , given a set of n real numbers,
 285 decide whether there exists a k -subset whose elements sum to zero.” This
 286 problem is called the k -SUM problem.

287 Obviously, the 3SUM problem is the k -SUM problem where $k = 3$.
 288 Moreover, there is a simple reduction from k -SUM to 2SUM when k is even
 289 and to 3SUM when k is odd. Those reductions yield a $O(n^{\frac{k}{2}} \log n)$ time
 290 real-RAM algorithm for k even and a $O(n^{\frac{k+1}{2}})$ time real-RAM algorithm for
 291 k odd.

292 In their paper, in addition to the slightly subquadratic uniform algorithm
 293 for 3SUM, Grønlund and Pettie give a strongly subquadratic nonuniform
 294 algorithm for 3SUM. The algorithms runs in time $\tilde{O}(n^{3/2})$, and, because
 295 of the aforementioned reduction, immediately yields an improved $\tilde{O}(n^{\frac{k}{2}})$
 296 nonuniform time complexity for k -SUM when k is odd.

297 As for uniform time complexity we do not know whether this nonuniform
 298 improvement can be transferred to the real-RAM model: we do not know of
 299 any real-RAM $o(n^{\frac{k+1}{2}})$ time algorithm for k -SUM when k is odd.

300 The k -SUM problem reduces to the following point location problem:
 301 “Given a input point $q \in \mathbb{R}^n$, locate q in the arrangement of $\binom{n}{k}$ hyperplanes of
 302 equation $x_{i_1} + x_{i_2} + \dots + x_{i_k} = 0$.” Applying the best nonuniform algorithms for
 303 point location in arrangements of hyperplanes by Meyer auf der Heide [114]
 304 and Meiser [113] yields linear decision trees of depth $n^{O(1)}$ for k -SUM, where
 305 the constant of proportionality in the big-oh does not depend on k .

306 **Paper A** Our first contribution is a finer analysis of Meiser’s algorithm
 307 that shows that the depth of his decision tree when applied to k -SUM is
 308 actually $O(n^3 \log^2 n)$. On top of that, we show how to implement a variant
 309 of this decision tree in the real-RAM model so that its uniform running time
 310 is $n^{\frac{k}{2} + O(1)}$ while keeping the nonuniform running time unchanged. Note that
 311 a naive implementation of Meiser’s algorithm has a uniform running time of
 312 $n^{k+O(1)}$.

313 The approach taken by Meiser’s algorithm follows the prune and search
 314 method: Take a sample of the hyperplanes for which we do not yet know
 315 the relative location of the input point. Locate the input point with respect
 316 to this sample by brute-force. This amounts to identifying the cell of the
 317 sample’s arrangement which contains the input point. Refine the location
 318 of the input point in the sample by partitioning the containing cell into
 319 low complexity subcells. Whenever this low complexity subcell is located
 320 completely on one side of an hyperplane, the input point is located on the
 321 same side, and so we can discard this hyperplane without a single query to
 322 the input point. Since some hyperplanes may intersect the low complexity
 323 subcell, rince and repeat.

324 The location refinement is necessary because this is the only way we can
 325 guarantee a bound on the number of hyperplanes we have to recurse on.
 326 Using the theory of ε -nets, we can show that, for a sample of reasonable size,
 327 any simplex that is not intersected by a sample hyperplane is not intersected
 328 by more than a constant fraction of the set of hyperplanes from which
 329 the sample is drawn (with high probability). We define the refined subcell
 330 of the algorithm presented above to be the simplex of the bottom-vertex
 331 triangulation of the sample’s arrangement that contains the input point.

332 Sorting

333 Before continuing, let us go back to the origins of those different problems.
 334 The link between them will be made even clearer.

335 Sorting is one of the oldest and most relevant data management problems.
 336 It is the archetypal computation tree problem. Usually presented, sorting is
 337 about permutations in *arrays*, but we do not like that. We use a different
 338 abstraction:

339 **Problem 1** (Sorting). Given n numbers $q_1, q_2, \dots, q_n \in \mathbb{R}$, for each pair $1 \leq$

340 $i < j \leq n$, decide whether $q_i < q_j$.

341 In other words, we see the sorting problem as an information retrieval
342 problem: how many comparisons do we need to make in order to *know* the
343 answer to all comparisons. The usual sorting algorithms rearrange an array
344 of input numbers into sorted order. Another way to think about it is that
345 they compute the permutation that sends the input array to a sorted version
346 of itself. This permutation is a data structure such that given any index in
347 the input array, we can query for the corresponding index in the sorted array
348 (called the *rank* of the element). To retrieve the result of a comparison of two
349 elements of the input array we can compare their ranks in this permutation.
350 The usual way of defining the sorting problem restricts the data structure
351 that should be used to encode the $\binom{n}{2}$ comparisons. The way we model the
352 sorting problem lifts this restriction because it does not ask to structure this
353 information in a nice way.

354 The sorting problem also has its decision problem variant: Element
355 Uniqueness.

356 **Problem 2** (Element Uniqueness). Given n numbers $q_1, q_2, \dots, q_n \in \mathbb{R}$, decide
357 whether there exists $1 \leq i < j \leq n$, such that $q_i = q_j$.

358 It is easy to see that Sorting amounts to locating the point q in the
359 arrangement of hyperplanes of equations $x_i - x_j = 0$, and that Element
360 Uniqueness reduces both to and from the 2SUM problem in linear time.
361 Actually, under reasonable assumptions on the computation model, sorting
362 and element uniqueness are the same problems: if all questions we ask about
363 the input are linear in the input numbers, then proving that the input does
364 not contain any duplicate entries requires to sort the input. The same
365 statement carries over to the k -SUM with respect to its “sorting version”:
366 computing the sign of all $\binom{n}{k}$ sums of k input numbers. Therefore, in those
367 models of computation, the sorting problem is equivalent to 2SUM. Because
368 of this relationship between Sorting and the k -SUM problem, we see why
369 the better understanding of k -SUM is a natural next move.

370 Among all the problems this thesis touches on, Sorting and Element
371 Uniqueness are the simplest and best understood. We know $\Omega(n \log n)$ lower
372 bounds in many nonuniform models of computation and we also know a long
373 list of simple real-RAM algorithms for those problems whose running times

374 match those lower bounds. For all that matters here, those problems are
375 solved.

376 A Zoo of Generalizations

377 Obviously, the k -SUM problem is not the only possible way to generalize
378 Sorting.

379 Hopcroft's problem asks whether given n points and n hyperplanes in \mathbb{R}^d ,
380 one of the points lies on one of the hyperplanes. When $d = 1$, this problem is
381 Element Uniqueness. Finding the location ("above"/"below") of each point
382 with respect to each hyperplane generalizes Sorting.

383 Orthogonal Vectors (OV) is a problem that has gained popularity in the
384 emerging field of fine grained complexity theory [154]. The problem is to
385 decide whether a set of n d -dimensional vectors contains an orthogonal pair.
386 It is easy to see that this problem is equivalent to Hopcroft's problem in
387 \mathbb{R}^{d-1} .

388 The dominance reporting problem asks, given n points in \mathbb{R}^d , to report
389 all pairs of points such that the first dominates the other in all dimensions.
390 Once again, for $d = 1$, this problem is Sorting because it asks for the answer
391 to all comparisons of the type $p_i \leq q_i$.

392 Sorting $X + Y$ is also a canonical problem in P : given two sets X and
393 Y of n numbers each, sort the set $\{x + y : x \in X, y \in Y\}$. Sorting $X + Y$
394 reduces linearly to the sorting version of 4SUM because it asks for the sign
395 of all comparisons of the type $x + y \leq x' + y'$.

396 We already saw that GPT and Sorting belong to the same family of
397 high-dimensional point location problems. There is a good reason for that:
398 when $d = 1$, GPT is Element Uniqueness. In one-dimensional space, GPT
399 asks whether any two points are the same. Picturing this space as the
400 (horizontal) real line, we see that the "sorting version" of GPT asks to
401 compute for each pair of points which one is on the "left" of the other which
402 simply amounts to Sorting the one-dimensional input points.

403 Intermediate Problems

404 A perspicacious reader may have frowned at the previous paragraphs
405 wondering whether this embryo of classification brings more insight than
406 confusion. "Sure!", one may say, "Many problems involve sorting, and

407 therefore, according to this dubious definition, are generalizations of it.”

408 However, the author begs to differ.

409 As a first example, take one of the best known algorithms for 3SUM. This
 410 algorithm reduces an instance of 3SUM to a sorting phase and a searching
 411 phase [91]. The sorting phase consists in answering all questions of the type
 412 $x_i + x_{i'} \leq x_j + x_{j'}$ with some restriction on the indices. Well, it turns out
 413 that this phase is exactly an instance of the sorting version of the 2SUM
 414 problem where the input numbers are the differences $x_i - x_j$.¹ Moreover, to
 415 make the uniform algorithm practical, they implement the resolution of the
 416 2SUM instance via dominance reporting. Note that this sorting problem is
 417 similar to the Sorting $X + Y$ problem where the answers to most questions
 418 are not cared for.

419 As a second example, we have the GPT problem. For this problem,
 420 the fact that the one-dimensional version is Sorting does not seem to give
 421 any insight on how to apprehend the more general problem, even in two
 422 dimensions. In order to make some progress in that direction, we need to
 423 capture more precisely to which family of problems GPT belongs. For that,
 424 we look at the algebra behind the problem.

425 GPT asks whether for any choice of $\binom{n}{d+1}$ input points p_i with coordinates
 426 $(p_{i,1}, p_{i,2}, \dots, p_{i,d})$, the determinant

$$\det \begin{pmatrix} 1 & p_{1,1} & p_{1,2} & \cdots & p_{1,d} \\ 1 & p_{2,1} & p_{2,2} & \cdots & p_{2,d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & p_{d+1,1} & p_{d+1,2} & \cdots & p_{d+1,d} \end{pmatrix}$$

428 is zero. This determinant is a degree- d $(d^2 + d)$ -variate polynomial. In
 429 particular, when $d = 2$, the determinant is a degree-2 6-variate polynomial.
 430 The GPT testing problem then amounts to testing whether the coordinates
 431 of any combination of the input points yields a root of that polynomial.
 432 We therefore consider the more general d -dimensional- k -POL ($dD-k$ -POL)
 433 problem: Given a dk -variate constant degree polynomial F and a set S of
 434 n points in \mathbb{R}^d , decide whether $F(S^k)$ contains any zeroes. For instance,
 435 2D-3POL generalizes GPT with $d = 2$ where the constant degree polynomial

¹This observation generalizes to the k -SUM problem: any k -SUM instance can be reduced to a larger $(k - 1)$ -SUM instance followed by a searching phase.

436 is the 3×3 determinant mentioned above. Moreover, 1D-3POL generalizes
 437 3SUM where the polynomial is simply the sum function. Equally interesting
 438 is the fact that 2D-2POL generalizes Hopcroft’s problem with $d = 2$ where
 439 the polynomial is the dot product.

440 **Paper B** In paper B, we generalize Grønlund and Pettie’s approach to
 441 solve 1D-3POL (or more simply, 3POL) in subquadratic time. Our approach
 442 is essentially the same in that it reduces 3POL to a sorting phase and a
 443 searching phase, the sorting phase being an instance of 2D-2POL.² Again,
 444 the implementation of the uniform algorithm solves this instance of 2D-2POL
 445 using dominance reporting (a generalization of it).

446 This result illustrates why a better understanding of the landscape of
 447 problems surrounding GPT helps to identify intermediate problems whose
 448 resolution marks progress towards the question of whether GPT admits
 449 subquadratic algorithms. Figure 1 depicts this landscape.

450 Encodings

451 Naive algorithms for Element Uniqueness, 3SUM, and k -SUM would
 452 search all possible combinations of 2, 3, or k input numbers for a match and
 453 reach horrible running times: respectively $O(n^2)$, $O(n^3)$, and $O(n^k)$.

454 The way better uniform algorithms for those problems work is by a
 455 combination of sorting and searching. The sorting part constructs a data
 456 structure and the searching part uses this data structure to answer the
 457 question at hand. The achieved running time is a balance between the cost
 458 of sorting and the cost of searching. This results in better running times
 459 than the naive “search only” solutions.

460 For instance, an efficient algorithm for the Element Uniqueness problem
 461 would first sort the input, then scan it for duplicates among adjacent numbers
 462 in this sorted order. The data structure that is constructed is the permutation
 463 we talked about earlier. This structure achieves a good compressing ratio by
 464 encoding the answer to all $O(n^2)$ pairwise comparisons of two input numbers
 465 in $O(n \log n)$ bits.

²Note that according to the implied definition, d D- k -SUM is equivalent to k -SUM. Therefore, the 2SUM instance in the sorting phase of their 3SUM algorithm is an instance of 2D-2SUM in disguise.

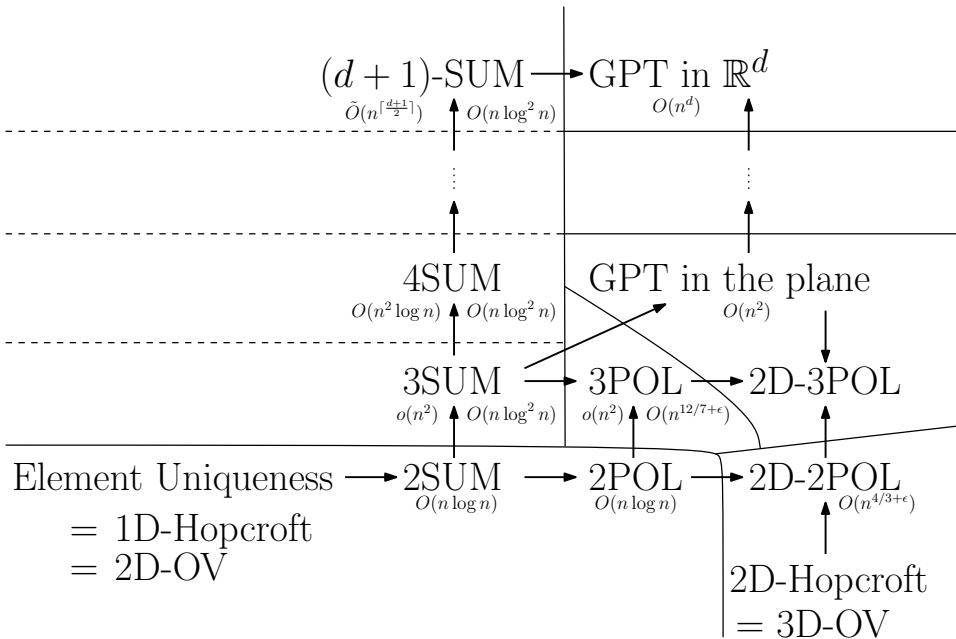


Figure 1. The decision problems surrounding GPT. Arrows indicate linear time reductions. $A = B$ indicates that A and B are the same problems. Known upper bounds are indicated next to the problem's name. When the uniform and nonuniform upper bounds differ, we place the best uniform upper bound on the right. Solid lines indicate apparent time-complexity differences between the best known nonuniform algorithms. Dotted lines indicate apparent time-complexity differences between the best known uniform algorithms.

466 For this reason Element Uniqueness is comparatively simpler than the
467 3SUM and k -SUM problems. For a nonuniform algorithm, constructing this
468 structure is sufficient to solve the problem: Once the construction is over, we
469 can discard the input because all the information we need is encoded in the
470 data structure. Since the nonuniform models of computation we consider
471 do not care about computations not involving the input, the searching part
472 does not cost anything.

473 The case of 3SUM and k -SUM is more complex: The sorting phase only
474 encodes a fraction of all possible queries. This leaves a significant amount
475 of work for the searching phase. Therefore, both the sorting phase and the
476 searching phase contribute to the nonuniform cost of those algorithms.

477 The case of GPT is also interesting. The naive algorithm queries $O(n^{d+1})$
478 input tuples while the better algorithm constructs the dual arrangement in
479 $O(n^d)$ time. As in the case of Element Uniqueness, this dual arrangement
480 is the data structure that encodes the answer to all $O(n^{d+1})$ queries while
481 achieving a significant space gain.

482 While the goal sought in the design of algorithms is to find the best
483 possible balance between those sorting and searching phases, a different
484 question becomes evident: “What is the most resource efficient data structure
485 encoding all the combinatorial information carried by the input?”. For
486 this question, resource efficiency can be measured in three ways: space
487 requirements, construction time, and query time. If one only cares about
488 space, a trivial answer points its head out: if there are only X combinatorial
489 types then each type can be encoded with $\lceil \log X \rceil$ bits. However, this solution
490 is unlikely to yield good construction time or good query time.

491 In the case of Element Uniqueness, this question has been answered long
492 ago. Sorting the input in $O(n \log n)$ time will construct a permutation using
493 $O(n \log n)$ bits that can be queried for any pairwise comparison in $O(1)$ time.
494 However, the question is still widely open for 3SUM, k -SUM, and GPT.

495 **Papers C & D** In Paper C we design the first subquadratic space data
496 structure for encoding the combinatorial type of a two-dimensional GPT
497 instance. This data structure can be constructed in quadratic time and
498 queries are answered in sublogarithmic time. Those results can be adapted
499 to work for higher-dimensional GPT to yield sub- $O(n^d)$ space data structures
500 with good construction and query times.

501 Since 3SUM reduces to GPT, the results of Paper C can be applied to
 502 encode the combinatorial type of 3SUM instances. However, since 3SUM
 503 is much better understood than GPT we should aim for better encodings.
 504 This is exactly what we do in Paper D. By filling the gaps in the partial
 505 data structure used in Grønlund and Pettie’s algorithm [91], we design an
 506 encoding that uses $O(n^{3/2} \log n)$ bits, can be constructed in $O(n^2)$ time and
 507 answers queries in constant time.

508 The Secret Ingredient

509 What makes the success of our methods are now-standard geometric
 510 divide-and-conquer tools: ε -nets and cuttings.

511 The idea of an ε -net is simple, yet powerful. Given a sample of m
 512 hyperplanes in \mathbb{R}^d , construct a sample of those hyperplanes of size $f(d, \varepsilon)$
 513 uniformly at random. Then, with high probability, any simplex that is not
 514 intersected by an hyperplane of the sample is intersected by at most εm
 515 hyperplanes of the original set. The sample is called a *net* because it does
 516 not let the fat simplices through: a simplex intersected by more than εm
 517 hyperplanes must be intersected by one of the sample. This yields an efficient
 518 point location tool: construct a sample, find the cell of its arrangement that
 519 contains the query point, find the simplex of its triangulation that contains
 520 the query point, and recurse on the fraction of hyperplanes that intersect it.

521 Triangulate the whole arrangement of the sample and you get a ε -cutting:
 522 a partition of space into simplices such that each simplex is intersected by
 523 at most εm hyperplanes. We use cuttings in applications where multiple
 524 queries are made on the same set of hyperplanes.

525 Wrapping it up

526 Since publication of those papers, a few developments have surfaced.

527 Ezra and Sharir [72] show how trading simplices of the bottom-vertex
 528 triangulation for prisms of the vertical decomposition in Meiser’s algorithm
 529 yields a shallower decision tree of depth $O(n^2 \log n)$. Essentially, the im-
 530 provement over our result in Paper A lies in the fact that, for vertical
 531 decomposition, the sample size can be taken to be an order of magnitude
 532 smaller.

533 In a breakthrough paper, Kane, Lovett, and Moran [103] give a linear

534 decision tree of depth $O(n \log^2 n)$ for k -SUM, almost matching the $\Omega(n \log n)$
535 lower bound. This improves both on Paper A and Ezra and Sharir [72].

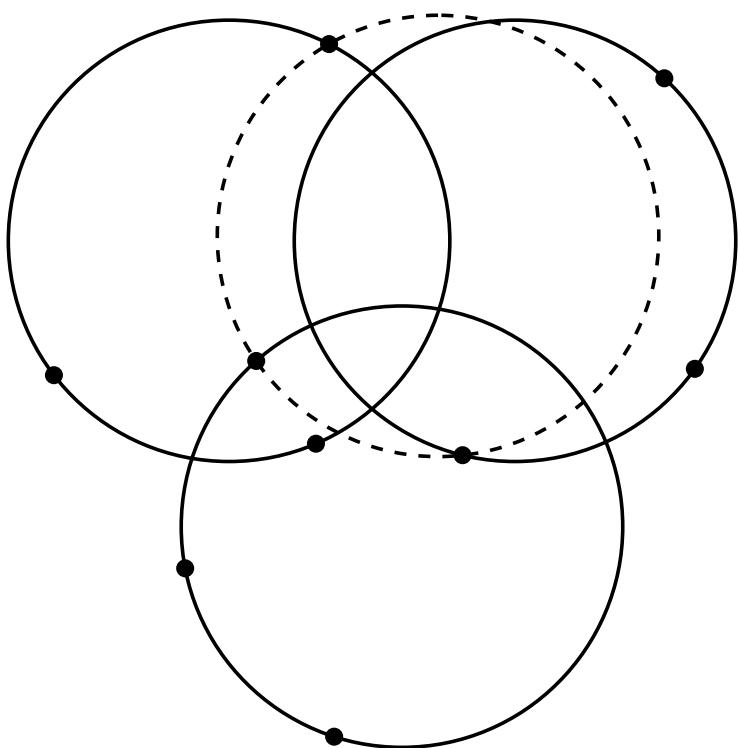
536 In [46], Chan shaves more logarithmic factors from the time complexity of
537 uniform algorithms for 3SUM and 3POL. While we focused on applications
538 that solve 3SUM-hard geometric problems with one-dimensional input in
539 Paper B, he shows how the ideas that work for 3POL also work for some
540 3SUM-hard geometric problems with two-dimensional data.

I

541

Without Proof

542



1

543

544

Models of Computation

545 In Theoretical Computer Science, we simplify reality in order to make
546 our job easier: we create models of what we think computers are and use
547 mathematics to reason about them. This allows us, not only to show, but to
548 *prove* that our algorithms are correct and efficient, according to those models.
549 Hopefully we manage to capture the essence of what practical computation
550 is about, making our lemmas and theorems relevant in practice.

551 We have to start this thesis somewhere, and, for our results to make
552 sense, it is only natural that we begin by defining the models we use.

553 1.1 Algorithms

554 We study two classes of algorithms: uniform and nonuniform.

555 *Uniform* algorithms are considered practical. They can be implemented
556 on a real computer because they take care of data management issues in an
557 automated way.

558 *Nonuniform* algorithms ignore the data management aspect of practical
559 computation. They do not give details on how to structure the intermediate
560 results of their computation. Those results are assumed to be instantly
561 accessible without any organization. Since we consider problem instances
562 of arbitrarily large size, this is not possible in practice: processor units can
563 only hold a fixed amount of data at any given time. The data has to be
564 stored somewhere in some structured way.

565 The naming uniform simply means that the algorithm (its description) is
566 the same for all input sizes. This is what is generally expected from algorithm
567 design: to output a finite size description of an automated problem solving
568 method that works for any instance size. Nonuniform algorithms on the

569 other hand are allowed to have distinct descriptions for each input size.
570 Each nonuniform algorithm can be seen as an infinite sequence of uniform
571 algorithms $A = A_1, A_2, \dots$, each A_n hardcoding the data management part
572 in its description. This description therefore is allowed to grow with n .
573 Designing a nonuniform algorithm A amounts to describing a method that
574 given n outputs A_n .

575 In this thesis, we study the uniform *random access machine* (RAM)
576 models real-RAM and word-RAM, and the nonuniform decision tree models
577 algebraic computation tree, algebraic decision tree, and linear decision tree.
578 Those are described in the following paragraphs.

579 1.1.1 Random Access Machines

580 The real-RAM and word-RAM models have the RAM in common. They
581 both assume a memory whose access cost is constant and a constant number
582 of standard operations on the numbers they manipulate.

583 The real-RAM model is a classic of Computational Geometry. The
584 input is assumed to consist of n real numbers and those numbers can
585 be manipulated exactly at constant cost using arithmetic and comparison
586 operators. This makes sense from the geometer's point of view since geometric
587 data is best abstracted as real numbers. For data management purposes,
588 the model also allows to manipulate $\log n$ -bits integers with arithmetic,
589 comparison, and bitwise operators but does not allow the conversion from a
590 real number to an integer.

591 The word-RAM model models computers as we know them more closely.
592 It considers an input consisting of n integers called words and allows a
593 constant number of standard unary or binary operators to be executed in
594 constant time. Those words have bitsize $w \geq \log n$.

595 1.1.2 Computation and Decision Trees

596 A decision tree is a constant-degree rooted tree whose internal nodes
597 encode binary queries to an omniscient oracle and whose leaves encode the
598 result of the computation. The complexity of the tree is defined to be the
599 length of its longest root to leaf path called the *depth* of the tree.

600 A decision tree cannot inspect the input directly. To make progress, it
601 inquires about the input through “yes/no” questions asked to the oracle.

602 The oracle answers those questions by “yes” or “no” honestly and accurately.

603 An identification problem consists in, given an input of size n , distinguishing among $f(n)$ input classes, that is, to put the instance in the right box. A decision problem is an identification where $f(n) = 2$.

606 Of course, for every identification problem with $f(n)$ input classes, there
 607 is a decision tree that solves the problem by asking $O(\log f(n))$ “yes/no”
 608 questions of the form: “Is the input in the following range of input classes?”
 609 This is best possible: every query carries at most one additional bit of
 610 information, and there are at least $\lceil \log f(n) \rceil$ bits necessary to label each
 611 input class. This lower bound is called the Information Theoretic Lower
 612 Bound, or ITLB, and holds for any decision tree model.

613 The queries the above decision tree makes are too powerful to make the
 614 study of this kind of model interesting in general. Moreover, practically,
 615 those queries have no structure a priori and it is not at all obvious how they
 616 can be encoded. For a decision tree model to make sense, we have to restrict
 617 the kind of queries that can be made.

618 Linear Decision Trees

619 A first natural way to restrict the complexity of the queries of our decision
 620 tree model is to only allow linear queries on the input. This will be sufficient
 621 to solve problems such as Sorting, Element Uniqueness, 3SUM, k -SUM, and
 622 subset-sum.

623 In the *s-linear decision tree model*, queries consists in testing the sign
 624 of a linear function on at most s numbers q_{i_1}, \dots, q_{i_s} of the input q_1, \dots, q_n .
 625 Such a query is called a *s-linear query* and can be written as

$$626 \quad \alpha_1 q_{i_1} + \cdots + \alpha_s q_{i_s} \stackrel{?}{\leq} \alpha_0$$

627 Each question is defined to cost a single unit. All other operations can be
 628 carried out for free but may not examine the input vector q . We refer to
 629 n -linear decision trees simply as linear decision trees.

630 Note that this model generalizes the Comparison Tree Model commonly
 631 used for Sorting and Element Uniqueness, where queries have the form
 632 $q_i \stackrel{?}{\leq} q_j$.

633 **Algebraic Decision Trees**

634 The linear decision tree model is not powerful enough for more complicated
 635 problems. General Position Testing for instance involves discovering the sign
 636 of quadratic polynomials of the input, which is impossible in general with
 637 linear queries only.

638 The (bounded-degree) *algebraic decision tree (ADT)* [128, 146, 155] is
 639 an algebraic generalization of linear decision trees. An algebraic decision
 640 tree performs constant-cost branching operations that amount to test ~~the~~
 641 sign of a constant-degree polynomial of the input numbers.

642 In this model, for an input q_1, \dots, q_n , a query can be written as

$$643 \quad p(q_1, \dots, q_n) \stackrel{?}{\leq} 0,$$

644 where $p \in \mathbb{R}[x_1, \dots, x_n]$ is a polynomial of constant degree. Again, operations
 645 not involving the input are free.

646 **Algebraic Computation Trees**

647 Computation trees differ from decision trees in that instead of exclusively
 648 asking “yes/no” questions, they perform arithmetic operations on variables
 649 whose results can be remembered and reused as operands, but only inspected
 650 using comparison queries. This allows, for instance, to make decision based
 651 on arbitrary polynomial expressions of the input with a sane way of counting
 652 the complexity of constructing such an expression. This would not be allowed
 653 in the algebraic decision tree model.

654 The internal nodes of an *algebraic computation trees* are labeled with
 655 arithmetic ($r \leftarrow o_1 \text{ op } o_2, \text{op} \in \{+, -, \times, \div\}$) and branching ($z : 0$) operations.
 656 Any internal node labeled $r \leftarrow o_1 \text{ op } o_2$ has outdegree one and is such that
 657 either $o_k = q_i$ for some i or there exists an ancestor $o_k \leftarrow x \text{ op } y$ of this node,
 658 and any internal node labeled $z : 0$ has outdegree three and is such that
 659 either $z = q_i$ for some i or there exists an ancestor $z \leftarrow x \text{ op } y$ of this node.
 660 Similarly to the previously described decision tree models, leafs of the tree
 661 correspond to input classes.

662 Using this model can be interpreted as only counting the operations
 663 that involve the input, that is, members of the input or results of previous
 664 operations involving the input. All other arithmetic operations are for free.

665 1.2 Data Structures

666 Data Structures are ubiquitous in algorithmics. Given some data, we
667 want to store it in computer memory in a way that 1) allows us to efficiently
668 access the information we want to extract from this data and 2) does not
669 require too much storage space. The term *structure* is by opposition to a
670 randomly ordered stream of data. ??

671 Usually, a third important feature of those structures is that they can
672 be also updated efficiently if the underlying data changes. We do not
673 study or use that aspect in this thesis, instead we only require that those
674 data structures can be constructed from scratch in an efficient way. Data
675 structures allowing efficient updates are usually named *dynamic* while the
676 ones we study are *static* data structures.

677 Those intuitive concepts are summarized as follows:

678 **Preprocessing Time** Given some data, the time it takes to construct the
679 corresponding data structure in the given computation model.

680 **Space** The amount of space the data structure consumes in the given
681 computation model. Can be expressed in bits, words, memory cells, ...

682 **Query Time** Given a data structure, the time consumed by a single query
683 in the given computation model.

684
685 In §1.2.1 we give a precise definition of a particular type of data structure
686 we study: encodings. ?

687 1.2.1 Encodings

688 A data structure is called *succinct* if its space usage is close to the ITLB.
689 This concept was first introduced by Jacobson in his PhD thesis [101]. Since
690 then, it has been extensively studied. Raman et al. studied the dynamic
691 implementation of such data structures and gave the first rank-select succinct
692 data structures [130, 131]. Those data structures are even efficient in practice
693 as shown by Vigna with their implementations [153].

694 In this thesis, we also try to get good bounds on the space used by the
695 data structures we design. However, in most cases we are still very far from

696 the ITLB. We make all our data structure design problems fit in a single
697 framework dubbed *instance encoding* so that their space and query time
698 requirements can be easily compared.

699 **Definition 2.** For fixed k and given a function $f : [n]^k \rightarrow [O(1)]$, we define
700 a $(S(n), Q(n))$ -encoding of f to be a string of $S(n)$ bits such that, given this
701 string and any $i \in [n]^k$, we can compute $f(i)$ in $Q(n)$ time in the word-RAM
702 model with word size $w \geq \log n$.

2

703

History

704

705 In this chapter we define and recount the history of the different problems
706 on which we made progress: 3SUM and k -SUM (§2.1), and GPT and 3POL
707 (§2.2). 

708 Note that we expose the state of knowledge as it was when we started
709 this thesis back in 2015. For developments during the years 2016 to 2019
710 see §4.

711 2.1 3SUM & k -SUM

712 The 3SUM problem is defined as follows: given n distinct real numbers,
713 decide whether any three of them sum to zero.

714 **Problem 3** (3SUM). Given n numbers $q_1 < q_2 < \dots < q_n \in \mathbb{R}$, decide
715 whether there exist $i < j < k \in [n]$ such that $q_i + q_j + q_k = 0$.

716 The seminal paper by Gajentaan and Overmars showed the crucial role of
717 3SUM in understanding the complexity of several problems in computational
718 geometry [80].

719 A popular conjecture is that no $n^{2-\Omega(1)}$ -time real-RAM algorithm solves
720 3SUM.

721 **Conjecture 2.1.** *There is no $n^{2-\Omega(1)}$ -time real-RAM algorithm for 3SUM.*

722 This conjecture has been used to show conditional lower bounds for
723 problems in P, notably in computational geometry with problems such
724 as GeomBase, general position testing (GPT) [80] and Polygonal Contain-
725 ment [21], and more recently for string problems such as Local Alignment [3]

and Jumbled Indexing [18], as well as dynamic versions of graph problems [2, 105, 125], triangle enumeration and Set Disjointness [105].

The k -SUM problem is a straightforward generalization of the 3SUM problem: given n distinct real numbers, decide whether k of them sum to zero.

Problem 4 (k -SUM). Given n numbers $q_1 < q_2 < \dots < q_n \in \mathbb{R}$, decide whether there exist $i_1 < i_2 < \dots < i_k \in [n]$ such that $\sum_{j=1}^k q_{i_j} = 0$.

It has been known for long that k -SUM is W[1]-hard and proved recently to be W[1]-complete [1]. The k -SUM problem is also a fixed-parameter version of the subset-sum problem, a standard NP -complete problem.

Similarly to 3SUM, the k -SUM problem has proved to be a computational bottleneck in high dimensional convex hull construction or general position testing [70].

For all those reasons, 3SUM and k -SUM are considered as key subjects of an emerging theory of complexity-within-P (or “fine-grained” complexity), on par with other problems such as all-pairs shortest paths, orthogonal vectors, boolean matrix multiplication, and conjectures such as the Strong Exponential Time Hypothesis [4, 42, 98, 117, 126].

2.1.1 Variants

For the sake of simplicity, we will sometimes consider the following definition of 3SUM:

Problem 5 (3SUM variant). Given three sets A , B , and C of n real numbers, decide whether any triple of numbers in $A \times B \times C$ sums to zero.

A similar variant can be defined for k -SUM:

Problem 6 (k -SUM variant). Given k sets S_i of n real numbers, decide whether there exists $(s_1, s_2, \dots, s_k) \in S_1 \times S_2 \times \dots \times S_k$ such that $\sum_{i=1}^k s_i = 0$.

The k -SUM problem can be further generalized to the linear degeneracy testing problem (k -LDT) where we allow coefficients other than zero and one.

Problem 7 (k -LDT). Given n input numbers $q_1 < \dots < q_n \in \mathbb{R}$ and constants $p_0, \dots, p_n \in \mathbb{R}$ decide whether there exists $i_1, i_2, \dots, i_k \in [n]$ such that $\sum_{j=1}^k p_j q_{i_j} = p_0$.

758 Our algorithms apply to this more general problem with only minor
 759 changes.

760 2.1.2 Point Location

761 The Point Location problem is a classic problem in Computational
 762 Geometry.

763 **Problem 8** (Point Location). Given a set \mathcal{H} of m hyperplanes and a query
 764 point $q \in \mathbb{R}^d$, find the cell $C \in \mathcal{A}(\mathcal{H})$ such that $q \in C$.

765 Because the bounds in Theorem 6.1 are attained when \mathcal{H} is in general
 766 position, there is a lower bound of $\Omega(d \log m)$ on the depth of decision trees
 767 solving this problem.

768 For our purpose, it is useful to see the k -SUM problem as a point location
 769 problem in \mathbb{R}^n where the coordinates of q are the input numbers and where
 770 \mathcal{H} is the set of all $\binom{n}{k}$ hyperplanes of equation $\sum_{j=1}^k x_{i_j} = 0$. Locating q
 771 in $\mathcal{A}(\mathcal{H})$ amounts to deciding in n -dimensional space, for each hyperplane
 772 $H \in \mathcal{H}$, whether q lies on, above, or below H . Since q lies on some H if and
 773 only if the k -SUM instance is degenerate, this constitutes a valid reduction.
 774 We emphasize that in this reduction the set of hyperplanes depends only on
 775 k and n and not on the actual input vector q .

776 Sorting can also be seen as a point location problem in \mathbb{R}^n . In this case \mathcal{H}
 777 is the set of all $\binom{n}{2}$ hyperplanes of equation $x_i = x_j$. The arrangement $\mathcal{A}(\mathcal{H})$
 778 has exactly $n!$ n -dimensional cell that correspond to the $n!$ permutations the
 779 input might have. Identifying the cell containing the input point reveals its
 780 permutation, hence solving the sorting problem.

781 See [144, Section 34.6] for more on point location in high dimension. 

782 2.1.3 Information Theoretic Lower Bound

783 The ITLB for sorting is the logarithm of the number of possible permutations
 784 of the input, that is $\lceil \log n! \rceil = \Omega(n \log n)$, for *any* decision tree. This
 785 does not hold for element uniqueness: in an arbitrarily powerful decision
 786 tree model, using a single query, we can find out whether the input contains
 787 duplicates. This matches the useless lower bound of $\log 2 = 1$.

788 This observation holds for any decision problem: if the problem amounts
 789 to retrieve a single bit of information, the “yes/no” answer, we cannot we
 790 derive any useful lower bound for arbitrary decision tree models.

791 Restricting our model to only allow linear queries, we can derive a more
 792 useful lower bound for element uniqueness.

793 **Lemma 2.1.** *In the linear decision tree model, any algorithm solving the
 794 Element Uniqueness problem must sort its input (if it has no duplicates).*

795 *Proof.* Assume that the input q has no duplicates. Then, looking at q as a point
 796 in \mathbb{R}^n , q is contained in a d -cell of the arrangement of hyperplanes of equation
 797 $x_i = x_j$. Sorting the input amounts to identifying this cell. Assume we have
 798 solved the element uniqueness problem for this instance and further assume, by
 799 contradiction, that we have not identified the cell containing q . In the model
 800 we consider, each query/answer pair corresponds to a halfspace containing q .
 801 Because we have not identified the d -cell containing q , it must be that points
 802 from at least two d -cells of the arrangement are contained in the intersection of
 803 the halfspaces. Because halfspaces are convex sets, their intersection is convex
 804 too, and thus connected. In this connected set, a path from q to any point from
 805 the other d -cell must at some point intersect one of the hyperplanes $x_i = x_j$
 806 that separates the two cells. An adversary can therefore update q to be this
 807 intersection point without changing any of the previous answers, even though
 808 this new point contains duplicates, a contradiction. \square

809 Because of the decision tree lower bound on Sorting, we have *The following* 

810 **Corollary 2.2.** *Any linear decision tree solving the Element Uniqueness
 811 problem has depth $\Omega(n \log n)$.*

812 The same conclusion can be drawn for the 3SUM and k -SUM problems
 813 since Element Uniqueness reduces to them.

814 **Corollary 2.3.** *Any linear decision tree solving the k -SUM problem has
 815 depth $\Omega(n \log n)$.*

816 2.1.4 Higher Lower Bounds

817 In 1999, Erickson showed that we cannot solve k -SUM or k -LDT in
 818 subquadratic time in the k -linear decision tree model [69]:

819 **Theorem 2.4** (Erickson [69]). *The depth of a k -linear decision tree solving
 820 the k -LDT problem is $\Omega(n^{\lceil \frac{k}{2} \rceil})$.*

The proof uses an adversary argument which can be explained geometrically. As we already observed, we can solve k -LDT problems by modeling them as point location problems in an arrangement of hyperplanes. Solving one such problem amounts to determining which cell of the arrangement contains the input point. The adversary argument of Erickson [69] is that there exists a cell having $\Omega(n^{\lceil \frac{k}{2} \rceil})$ boundary facets and in this model point location in such a cell requires testing each facet.

In 2005, Ailon and Chazelle slightly extended the range of query sizes for which a nontrivial lower bound could be established, elaborating on Erickson's technique [15]. They study s -linear decision trees to solve the k -SUM problem when $s > k$. In particular, they give an additional proof for the $\Omega(n^{\lceil \frac{k}{2} \rceil})$ lower bound of Erickson and generalize the lower bound for the s -linear decision tree model when $s > k$. Note that the exact lower bound given by Erickson for $s = k$ is $\Omega((nk^{-k})^{\lceil \frac{k}{2} \rceil})$ while the one given by Ailon and Chazelle is $\Omega((nk^{-3})^{\lceil \frac{k}{2} \rceil})$. Their result improves therefore the lower bound for $s = k$ when k is large. The lower bound they prove for $s > k$ is the following

Theorem 2.5 (Ailon and Chazelle [15]). *The depth of a s -linear decision tree solving the k -LDT problem is*

$$\Omega(nk^{-3})^{(1-\epsilon_k)(\frac{2k-s}{2})/\lceil \frac{s-k+1}{2} \rceil},$$

where $\epsilon_k > 0$ tends to 0 as $k \rightarrow \infty$.

This lower bound breaks down when $k = \Omega(n^{\frac{1}{3}})$ or $s \geq 2k$ and the cases where $k < 6$ give trivial lower bounds. For example, in the case of 3SUM with $s = 4$ we only get an $\Omega(n)$ lower bound.

2.1.5 Uniform Algorithms

On the real-RAM and word-RAM, it is easy to solve any k -SUM instance in time $O(n^{k/2} \log n)$ for k even and $O(n^{\frac{k+1}{2}})$ for k odd, and hence 3SUM in time $O(n^2)$. Those algorithms inspect the input using k -linear queries exclusively and the lower bounds of Erickson and Ailon and Chazelle in this model of computation essentially match their complexity.

Because fixing two of the numbers a and b in a triple only allows for one solution to the equation $a + b + x = 0$, an instance of 3SUM has at most

853 n^2 degenerate triples. An instance giving a matching lower bound is for
 854 example the set of n integers $\{\frac{1-n}{2}, \frac{3-n}{2}, \dots, \frac{n-1}{2}\}$ (for odd n) with $\frac{3}{4}n^2 + \frac{1}{4}$
 855 degenerate triples. One might be tempted to think that the number of
 856 “solutions” to the problem would lower bound the complexity of algorithms
 857 for the decision version of the problem, as it is the case for this problem, and
 858 other problems, in restricted models of computation [68, 69]. This intuition
 859 is incorrect.

860 In 2008, Baran, Demaine, and Pătrașcu gave the first subquadratic
 861 algorithms for 3SUM [19]. They design subquadratic Las Vegas algorithms
 862 for 3SUM on integer and rational numbers in the circuit RAM, word RAM,
 863 external memory, and cache-oblivious models of computation. The idea of
 864 their approach is to exploit the parallelism of the models, using linear and
 865 universal hashing. However, since their algorithms do not handle real inputs,
 866 this did not settle the question of subquadratic algorithms in full generality.

867 In 2014, Grønlund and Pettie gave the first subquadratic algorithms for
 868 real-input 3SUM [91]. Using an old trick due to Fredman [77], they prove the
 869 existence of a linear decision tree solving the 3SUM problem using a strongly
 870 subquadratic number of linear queries. The classical quadratic algorithm for
 871 3SUM uses 3-linear queries while the decision tree of Grønlund and Pettie
 872 uses 4-linear queries and requires $O(n^{3/2}\sqrt{\log n})$ of them.

873 **Theorem 2.6** (Grønlund and Pettie [91]). *There is a 4-linear decision tree
 874 of depth $O(n^{3/2}\sqrt{\log n})$ for 3SUM.*

875 They show how to adapt this decision tree to run in subquadratic time
 876 in the real-RAM model. They design two subquadratic 3SUM real-RAM
 877 algorithms. A deterministic one running in $O(n^2/(\log n/\log\log n)^{\frac{2}{3}})$ time
 878 and a randomized one running in $O(n^2(\log\log n)^2/\log n)$ time with high
 879 probability.

880 In 2015, Chan and Lewenstein designed strongly subquadratic word-RAM
 881 algorithms for a high-dimensional variant of integer 3SUM with applications
 882 to jumbled indexing [47]. This result generalizes the folk wisdom that 3SUM
 883 on small integers can be solved in nearlinear time using FFT’s.

884 The same year, Freund [79] and Gold and Sharir [81] improved on the
 885 results of Grønlund and Pettie [91]. Freund [79] gave a deterministic algo-
 886 rithm for 3SUM running in $O(n^2 \log\log n/\log n)$ time. Gold and Sharir [81]
 887 gave another deterministic algorithm for 3SUM with the same running time

and shaved off the $\sqrt{\log n}$ factor in the decision tree complexities of 3SUM and k -SUM given by Grønlund and Pettie.

Theorem 2.7 (Freund [79], Gold and Sharir [81]). *There is a $O(\frac{n^2 \log \log n}{\log n})$ -time real-RAM algorithm for 3SUM.*

2.1.6 Nonuniform Algorithms

The nonuniform algorithm of Grønlund and Pettie for 3SUM in the 4-linear decision tree model generalizes to k -SUM [91]. In the $(2k - 2)$ -linear decision tree model, only $O(n^{\frac{k}{2}} \sqrt{\log n})$ queries are required for odd values of k . Putting this in perspective with the lower bounds of Erickson and Ailon and Chazelle, this indicates that increasing the size of the queries, thus making the model more powerful, does yield improvements on the depth of the minimal-height decision tree.

It has been well established that there exist nonuniform polynomial-time algorithms for the subset-sum and knapsack problems, even though those problems are NP-complete. In 1984, Meyer auf der Heide was the first to use a point location algorithm to solve the knapsack problem in the linear decision tree model in polynomial time [114]. He thereby answered a question raised by Dobkin and Lipton [58, 59], Yao [156] and others. However, if one uses this algorithm to locate a point in an arbitrary arrangement of hyperplanes the running time is increased by a factor linear in the greatest coefficient in the equations of all hyperplanes. In 1993, a second algorithm was described by Meiser [113], and is derived from a point location data structure for arrangements of hyperplanes using *bottom vertex triangulation* and ε -nets (see §6.4.1 and §8.1.3). The complexity of Meiser's point location algorithm is polynomial in the dimension, logarithmic in the number of hyperplanes and does not depend on the coefficients in the equations of the hyperplanes. A useful complete description of this algorithm is given by Bürgisser et al. [36, Section 3.4].

When applied to k -SUM, those algorithms can be cast as the construction of a n -linear decision tree, even though k is constant. They both yield $n^{O(1)}$ -time nonuniform algorithms for k -SUM where the constant of proportionality does not depend on k , thus exhibiting the potential superiority of n -linear queries.

921 2.2 GPT & 3POL

922 In the plane, the General Position Testing problem (GPT) asks whether,
 923 given n points, three of them are collinear [?]. Because of some elementary
 924 algebra, the following form is equivalent:

925 **Problem 9** (GPT in the plane). Given n points $(x_1, y_1), \dots, (x_n, y_n) \in \mathbb{R}^2$,
 926 decide whether there exist $i < j < k \in [n]$ such that

$$927 \det \begin{pmatrix} 1 & x_i & y_i \\ 1 & x_j & y_j \\ 1 & x_k & y_k \end{pmatrix} = 0.$$

928 Note that this determinant is a degree-two polynomial in six variables

$$929 x_i y_j - y_i x_j - x_i y_k + y_i x_k + x_j y_k - y_j x_k.$$

930 In Computational Geometry this determinant is better known as the
 931 (counterclockwise) orientation $\nabla(i, j, k) \in \{-, 0, +\}$ of the points i , j , and
 932 k with coordinates (x_i, y_i) , (x_j, y_j) , and (x_k, y_k) and is defined as the sign
 933 of this determinant. It is called *orientation* or *sidedness* because for three
 934 given points this sign gives the orientation of the triangle those points define.
 935 It is often written as

$$936 \nabla(i, j, k) = \text{sign}((x_j - x_i)(y_k - y_i) - (y_j - y_i)(x_k - x_i)),$$

937 with only five subtractions and two multiplications.

938 The problem generalizes to higher dimension: given n points in \mathbb{R}^d , decide
 939 whether any $d+1$ of them lie on a common hyperplane. Again, using algebra:

940

941 **Problem 10.** Given n points $p_i = (p_{i,1}, p_{i,2}, \dots, p_{i,d}) \in \mathbb{R}^d$, decide whether
 942 there exist $i_1 < i_2 < \dots < i_{d+1} \in [n]$ such that

$$943 \det \begin{pmatrix} 1 & p_{i_1,1} & p_{i_1,2} & \dots & p_{i_1,d} \\ 1 & p_{i_2,1} & p_{i_2,2} & \dots & p_{i_2,d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & p_{i_{d+1},1} & p_{i_{d+1},2} & \dots & p_{i_{d+1},d} \end{pmatrix} = 0.$$

944 Again, this determinant is a degree d polynomial in $d^2 + d$ variables.

945 Later in this document we mention GPT without specifying the dimension.
 946 The problem is then assumed to be in the plane. When we consider a different
 947 parameterization we explicitly mention it.

948 **2.2.1 Variants**

949 Similarly to the 3SUM problem, we can study a variant of GPT where
 950 the tested triples come from different sets.

951 **Problem 11** (GPT variant in the plane). Given three sets A , B , and C of n
 952 points in \mathbb{R}^2 , decide whether any triple of points in $A \times B \times C$ is collinear.

953 This variant is useful for a particular reduction from 3SUM. A similar
 954 variant can be defined for the d -dimensional problem.

955 **2.2.2 Reductions from k -SUM**

956 This section is divided into two parts. The first part exposes a folk
 957 reduction from k -SUM to GPT showing GPT in \mathbb{R}^d is as hard as $(d+1)$ -
 958 SUM. The second part exposes other interesting connections between k -SUM
 959 and GPT that can be drawn through Vandermonde determinants, yielding
 960 equivalent hardness proofs for GPT [70].

961 **Folklore**

962 Given an instance of the three-set variant of 3SUM, we can reduce it to an
 963 instance of GPT where the input sets are on three horizontal lines.

964 **Observation 2.8.** *Given three sets A , B , and C of n real numbers, let*

$$\begin{aligned} 965 \quad A' &= \{a' = (a, 0) : a \in A\}, \\ 966 \quad B' &= \{b' = (b, 2) : b \in B\}, \\ 967 \quad C' &= \{c' = (-c/2, 1) : c \in C\}. \end{aligned}$$

969 Then for any $a \in A$, $b \in B$, and $c \in C$, we have that $a + b + c = 0$ if and only
 970 if a' , b' , and c' are collinear.

971 *Proof.* The equation of a line defined by points $a' \in A'$ and $b' \in B'$ is

$$\begin{aligned} 972 \quad y &= \frac{2}{b-a}x - \frac{2}{b-a}a \\ 973 \quad &\iff x = a + \frac{y(b-a)}{2}, \end{aligned}$$

975 which simplifies to $x = \frac{a+b}{2}$ when $y = 1$. □

976 Therefore, solving GPT is at least as hard as solving 3SUM.

?? ↗

977 In the dual (see §7.1), this reduction is even easier to apprehend and
978 trivially generalizes to d -dimensional GPT.

979 **Observation 2.9.** Given k sets S_i of n real numbers, in \mathbb{R}^{k-1} let S'_i , $1 \leq i \leq$
980 $k-1$, be the set of n hyperplanes $x_i = s$ for $s \in S_i$ and let S'_k be the set of n
981 hyperplanes $\sum_{i=1}^{k-1} x_i + s = 0$ for $s \in S_k$. Then for any $s_i \in S_i$, $1 \leq i \leq k$, we
982 have that $\sum_{i=1}^k s_i = 0$ if and only if the corresponding hyperplanes intersect.

983 *Proof.* Trivially. □

984 This proves that GPT in \mathbb{R}^d is as hard as $(d+1)$ -SUM.

985 **Points on the Weird Moment Curve**

986 When $d = 1$, GPT is exactly the element uniqueness problem.

987 **Problem 12** (GPT on the real line). Given n points $q_1, \dots, q_n \in \mathbb{R}$, decide
988 whether there exist $i < j \in [n]$ such that

$$989 \det \begin{pmatrix} 1 & q_i \\ 1 & q_j \end{pmatrix} = q_j - q_i = 0.$$

990 One seemingly stupid way to solve this problem is to consider the Van-
991 dermonde matrix

$$992 V(x) = \begin{pmatrix} 1 & x_1 & x_1^2 & \dots & x_1^{n-1} \\ 1 & x_2 & x_2^2 & \dots & x_2^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 & \dots & x_n^{n-1} \end{pmatrix},$$

993 and to compute its determinant

$$994 |V(q)| = \det \begin{pmatrix} 1 & q_1 & q_1^2 & \dots & q_1^{n-1} \\ 1 & q_2 & q_2^2 & \dots & q_2^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & q_n & q_n^2 & \dots & q_n^{n-1} \end{pmatrix} = \prod_{1 \leq i < j \leq n} (q_j - q_i).$$

995 This can be done in $O(n \log n)$ ring multiplications by defining $p(x) =$
 996 $\prod_{i=0}^{n-1} (x - q_i)$ and computing

$$\begin{aligned} 997 \det(V(q))^2 &= \text{disc}_x(p(x)) \\ 998 &= (-1)^{\binom{n}{2}} \text{res}_x(p(x), \text{diff}_x(p(x))) \\ 999 &= (-1)^{\binom{n}{2}} \prod_{i=0}^{n-1} \text{diff}_x(p(x))(q_i), \\ 1000 \end{aligned}$$

1001 using fast polynomial interpolation to first find the coefficients of $p(x)$
 1002 and then using fast polynomial evaluation to evaluate $\text{diff}_x(p(x))$ at all
 1003 q_i [6, 106, 107, 148].

1004 Consider points on the moment curve $m_d(t) = (t, t^2, \dots, t^d)$. Given a
 1005 GPT instance consisting of points $m_d(q_i)$ on this curve the determinant in
 1006 Problem 10 becomes

$$1007 \det \begin{pmatrix} 1 & q_1^1 & q_1^2 & \cdots & q_1^d \\ 1 & q_2^1 & q_2^2 & \cdots & q_2^d \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & q_{d+1}^1 & q_{d+1}^2 & \cdots & q_{d+1}^d \end{pmatrix} = |V(q)|,$$

1008 and solving the problem amounts to deciding Element Uniqueness.

1009 Consider points on the *weird* moment curve $\omega_d(t) = (t, t^2, \dots, t^{d-1}, t^{d+1})$.
 1010 If our input only consists of points $\omega_d(q_i)$ on this curve then the determinant
 1011 becomes

$$1012 \det \begin{pmatrix} 1 & q_0^1 & q_0^2 & \cdots & q_0^{d-1} & q_0^{d+1} \\ 1 & q_1^1 & q_1^2 & \cdots & q_1^{d-1} & q_1^{d+1} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 1 & q_d^1 & q_d^2 & \cdots & q_d^{d-1} & q_d^{d+1} \end{pmatrix} = |V(q)| \cdot \sum_{i=1}^{d+1} q_i,$$

1013 which is the predicate we want to test for $(d + 1)$ -SUM [70] (modulo the
 1014 $|V(q)|$ factor which is nonzero for distinct input numbers).

1015 In the plane, this reduction from instances of 3SUM to points on the
 1016 cubic curve $\omega_2(t) = (t, t^3)$ is also part of folklore.

1017 2.2.3 Lower Bounds and Order Types

1018 Even though 3SUM and k -SUM reduce to GPT, it does not necessarily
 1019 mean that the lower bounds for those problems carry over. The lower bounds

1020 in §2.1.3 are in the linear decision tree model and we cannot hope to solve
 1021 GPT in those models because they do not allow us to answer quadratic
 1022 polynomial queries on the input. That being said, solving an instance of GPT
 1023 as in Observation 2.8 with orientation predicates is the same as trying to
 1024 solve the original 3SUM instance with 3-linear queries. Hence, the quadratic
 1025 lower bound of Erickson [69] carries over and GPT requires $\Omega(n^2)$ time in a
 1026 model where one is only allowed orientation queries. More generally, GPT
 1027 in \mathbb{R}^d requires $\Omega(n^{\lceil \frac{d+1}{2} \rceil})$ time in a model where one is only allowed to query
 1028 the orientation of d -simplexes of the input points.

1029 However, neither those lower bounds, nor the weaker ones in the n -linear
 1030 decision tree model, carry over in the more generous algebraic decision tree
 1031 model and algebraic computation tree model. To get a superlinear lower
 1032 bound for GPT in those models we need work a little bit harder.

1033 Given an instance of GPT, its *order type* is defined to be the map that
 1034 sends each triple of input points to its orientation.

1035 **Definition 3** (Order Type of a Point Set). Given a set of n labeled points $P =$
 1036 $\{p_1, p_2, \dots, p_n\}$, we define the *order type* of P to be the function $\chi: [n]^3 \rightarrow$
 1037 $\{-, 0, +\}: (a, b, c) \mapsto \nabla(p_a, p_b, p_c)$ that maps each triple of point labels to the
 1038 orientation of the corresponding points, up to isomorphism.

1039 **Definition 4** (Realizable Order Type). When considering an arbitrary func-
 1040 tion $f: [n]^3 \rightarrow \{-, 0, +\}$ we say that f is a realizable order type if there is a
 1041 n -set $P \subset \mathbb{R}^2$ such that its order type is f .

1042 If we consider the problem of identifying the order type of a set of points,
 1043 we can achieve a lower bound that holds in any arbitrary decision tree model.
 1044 It suffices ~~say~~ to bound the number of possible order types.

1045 **Theorem 2.10** (Goodman and Pollack [87], Alon [16]). *There are $2^{\Theta(n \log n)}$*
 1046 *realizable order types.*

1047 With the following corollary

1048 **Corollary 2.11.** *A decision tree identifying the order type of a point sets*
 1049 *has depth $\Omega(n \log n)$.*

1050 The question now is whether we can do the same as in the Element
 1051 Uniqueness versus Sorting case: show that the decision problem is as hard
 1052 as the identification problem. Ben-Or [25] shows exactly that.

1053 **Theorem 2.12** (Ben-Or [25]). *Given a set $W \subset \mathbb{R}^d$ having N connected
 1054 components, the depth of any algebraic computation tree or bounded-degree
 1055 algebraic decision tree deciding whether an input point $q \in \mathbb{R}^d$ is in W is
 1056 $\Omega(\log N - d)$.*

1057 Because this Theorem applies to both W and its complement, we can
 1058 replace N by $\max\{\#W, \#(\mathbb{R}^d - W)\}$ in the statement, where $\#X$ denote
 1059 the number of connected components of the set X .

1060 **Corollary 2.13.** *Any algebraic computation tree or bounded-degree algebraic
 1061 decision tree that solves GPT has depth $\Omega(n \log n)$.*

1062 *Proof.* Let $W \subset \mathbb{R}^{2n}$ be the set defined by the union of the $\binom{n}{3}$ algebraic
 1063 varieties of equation

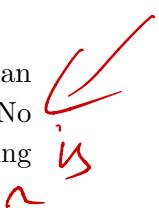
$$1064 \quad (x_j - x_i)(y_k - y_i) - (y_j - y_i)(x_k - x_i) = 0.$$

1065 Solving GPT on input $q \in \mathbb{R}^{2n}$ amounts to deciding whether $q \in W$. The set
 1066 $\mathbb{R}^{2n} - W$ has $2^{\Omega(n \log n)}$ connected components by Theorem 2.10. \square

1067 Note that both Theorem 2.10 and Theorem 2.12 rely on the Petrovskii-Olešnik-Thom-Milnor (POTM) Theorem that bounds the number of
 1068 connected components of algebraic varieties [23, 115, 151].

1070 2.2.4 Algorithms

1071 By brute-force, GPT can be solved in $O(n^3)$ time. Quadratic time can
 1072 be achieved by constructing the dual arrangement [93, Theorem 24.4.1]. No
 1073 subquadratic algorithm known, uniform or nonuniform. Essentially, nothing
 1074 known.



1075 2.2.5 More on Order Types

1076 A great deal of the literature in computational geometry deals with
 1077 the notion of order type [7–14, 16, 17, 28, 30, 31, 67, 71, 73–75, 82, 85–
 1078 90, 94, 100, 104, 108, 112, 121, 138–140, 149]. The order type of a point set
 1079 has been further abstracted into combinatorial objects known as (rank-three)
 1080 *oriented matroids* [75]. The *chirotope axioms* define consistent systems of
 1081 signs of triples [28]. From the topological representation theorem [30], all
 1082 such *abstract order types* correspond to pseudoline arrangements, while,

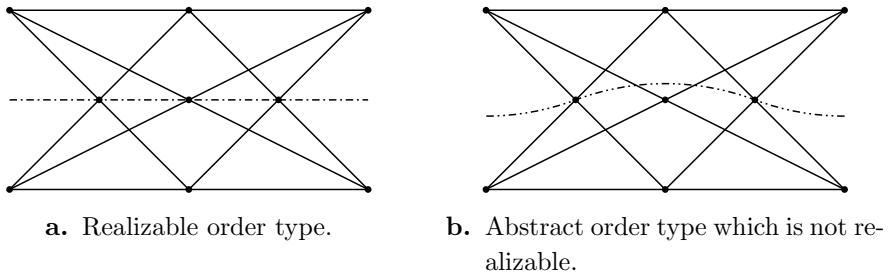


Figure 2.1. Pappus’s configuration.

from the standard projective duality, order types of point sets correspond to straight line arrangements. See Chapter 6 of The Handbook for more details [139].

When the order type of a pseudoline arrangement can be realized by an arrangement of straight lines, we call the pseudoline arrangement *stretchable*. As an example of a nonstretchable arrangement, Levi gives Pappus’s configuration where eight triples of concurrent straight lines force a ninth, whereas the ninth triple cannot be enforced by pseudolines [108] (see Figure 2.1). Ringel shows how to convert the so-called “non-Pappus” arrangement of Figure 2.1 (b) to a simple arrangement while preserving nonstretchability [140]. All arrangements of eight or fewer pseudolines are stretchable [84], and the only nonstretchable simple arrangement of nine pseudolines is the one given by Ringel [138]. More information on pseudoline arrangements is available in Chapter 5 of The Handbook [83].

Figure 2.1 shows that not all pseudoline arrangements are stretchable. Indeed, most are not: there are $2^{\Theta(n^2)}$ abstract order types [73] and only $2^{\Theta(n \log n)}$ realizable order types (Theorem 2.10).

Theorem 2.14 (Felsner [73]). *There are $2^{\Theta(n^2)}$ abstract order types.*

This discrepancy stems from the algebraic nature of realizable order types, as illustrated by the main tool used in the upper bound proofs (the POTM Theorem [23, 115, 151]).

2.2.6 Encodings

At SoCG’86, Bernard Chazelle asked [89]:

“How many bits does it take to know an order type?”

1107 This question is of importance in Computational Geometry for the
1108 following two reasons: First, in many algorithms dealing with sets of points
1109 in the plane, the only relevant information carried by the input is the
1110 combinatorial configuration of the points given by the orientation of each
1111 triple of points in the set (clockwise, counterclockwise, or collinear) [61, 67,
1112 104]. Second, computers as we know them can only handle numbers with
1113 finite description and we cannot assume that they can handle arbitrary real
1114 numbers without some sort of encoding. The study of *robust* algorithms
1115 is focused on ensuring the correct solution of problems on finite precision
1116 machines. Chapter 41 of The Handbook of Discrete and Computational
1117 Geometry is dedicated to this issue [158].

1118 Regarding encoding the function χ of Definition 3 (as in Definition 2),
1119 Theorem 2.14 together with information theory imply that $\Theta(n^2)$ bits are
1120 necessary and sufficient for abstract order types whereas $\Theta(n \log n)$ bits
1121 are necessary and sufficient for realizable order types. Optimal encodings
1122 matching those bounds can be produced by a simple enumeration algo-
1123 rithm. However, it is unclear how the original information can be efficiently
1124 reconstructed from those encodings. On the other hand, storing all $\binom{n}{3}$
1125 orientations in a lookup table to render this information accessible seems
1126 wasteful.

1127 Another obvious idea for encoding the order type of a point set is to store
1128 the coordinates of the points, and answer orientation queries by computing
1129 the corresponding determinant. While this should work in many practical
1130 settings, it cannot work for all point sets. Perles’s configuration shows that
1131 some configuration of points, containing collinear triples, forces at least one
1132 coordinate to be irrational [92] (see Figure 2.2). It is easy to see that order
1133 types of points in general position can always be represented by rational (or
1134 integer) coordinates. However, it is well known that some configurations
1135 require doubly exponential coordinates, hence coordinates with exponential
1136 bitsizes if represented in the binary numeral system [90].

1137 Goodman and Pollack defined λ -matrices which can encode abstract
1138 order types using $O(n^2 \log n)$ bits [85] and can be constructed in $O(n^2)$
1139 time [63]. They asked if the space requirements could be moved closer to the
1140 information-theoretic lower bounds. Everett, Hurtado, and Noy complained
1141 that this encoding does not allow a fast decoding for individual triples [71].
1142 Knuth and Streinu independently gave new encodings of size $O(n^2 \log n)$

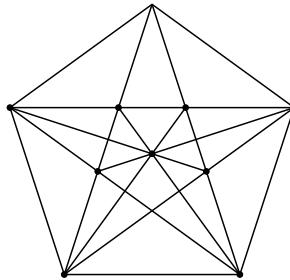


Figure 2.2. Perles’s configuration.

that allow orientation queries in constant time [104, 149].¹ Felsner and Valtr showed how to encode abstract order types optimally in $O(n^2)$ bits via the wiring diagram of their corresponding allowable sequence [73, 74] (as defined in [82]). Aloupis et al. gave an encoding of size $O(n^2)$ that can be computed in $O(n^2)$ time and that can be used to test for the isomorphism of two distinct point sets in the same amount of time [17].

2.2.7 The Intermediate Problem

On par with GPT we consider an algebraic generalization of the 3SUM problem: we replace the sum function by a constant-degree polynomial in three variables $F \in \mathbb{R}[x, y, z]$ and ask to determine whether there exists a *degenerate* triple (a, b, c) of input numbers such that $F(a, b, c) = 0$. We call this problem the *3POL problem*.

Problem 13 (3POL). Let $F \in \mathbb{R}[x, y, z]$ be a trivariate polynomial of constant degree, given three sets A , B , and C , each containing n real numbers, decide whether there exist $a \in A$, $b \in B$, and $c \in C$ such that $F(a, b, c) = 0$.

In addition to generalizing 3SUM, this problem can model particular instances of GPT: consider points in the plane that are constrained to lie on three parameterized polynomial curves of constant degree, then an algorithm for 3POL can determine whether any three of those points are collinear. Note that this generalizes the reductions from 3SUM found in §2.2.2, where the input points either lie on three horizontal lines or on a single cubic curve.

¹We attract the attention of the reader on the fact that we express size in bits. Other authors, [71] and [149] in particular, express size in number of words, which is off by at least a logarithmic factor.

1164 In fact, this problem is not new: it has already been studied before from a
1165 combinatorics point of view as exposed in the next section.

1166 2.2.8 Combinatorics

1167 In a series of results spanning fifteen years, Elekes and Rónyai [65], Elekes
1168 and Szabó [66], Raz, Sharir and Solymosi [136], and Raz, Sharir and de
1169 Zeeuw [133] give upper bounds on the number of degenerate triples for
1170 the 3POL problem. For the particular case $F(x, y, z) = f(x, y) - z$ where
1171 $f \in \mathbb{R}[x, y]$ is a constant-degree bivariate polynomial, Elekes and Rónyai [65]
1172 show that the number of degenerate triples is $o(n^2)$ unless f is *special*. Special
1173 for f means that f has one of the two special forms

$$\text{1174 } f(u, v) = h(\varphi(u) + \psi(v)) \quad \text{or} \quad f(u, v) = h(\varphi(u) \cdot \psi(v)),$$

1175 where h, φ, ψ are univariate polynomials of constant degree. It must be noted
1176 that the 3SUM problem falls in the special category since, in that case, f is
1177 the sum function. Elekes and Szabó [66] later generalized this result to a
1178 broader range of functions F using a wider definition of specialness. Raz,
1179 Sharir and Solymosi [136] and Raz, Sharir and de Zeeuw [133] improved
1180 both bounds to $O(n^{11/6})$. They translated the problem into an incidence
1181 problem between points and constant-degree algebraic curves. Then, they
1182 showed that unless f (or F) is special, these curves have low multiplicities.
1183 Finally, they applied a theorem due to Pach and Sharir [122] bounding the
1184 number of incidences between the points and the curves.

1185 **Theorem 2.15** (Raz, Sharir and de Zeeuw [133]). *Let A, B, C be n -sets of
1186 real numbers and $F \in \mathbb{R}[x, y, z]$ be a polynomial of constant degree, then*

$$\text{1187 } |Z(F) \cap (A \times B \times C)| = O(n^{11/6}),$$

1188 unless F has some group related form.²

1189 Raz, Sharir and de Zeeuw [133] also look at the number of degenerate
1190 triples for the General Position Testing problem when the input is restricted
1191 to points lying on a constant number of constant-degree algebraic curves.

²Because our results do not depend on the meaning of *group related form*, we do not bother defining it here. We refer the reader to Raz, Sharir and de Zeeuw [133] for the exact definition.

1192 **Theorem 2.16** (Raz, Sharir and de Zeeuw [133]). *Let C_1, C_2, C_3 be three
1193 (not necessarily distinct) irreducible algebraic curves of degree at most d in
1194 \mathbb{C}^2 , and let $S_1 \subset C_1, S_2 \subset C_2, S_3 \subset C_3$ be finite subsets. Then the number of
1195 proper collinear triples in $S_1 \times S_2 \times S_3$ is*

1196 $O_d(|S_1|^{1/2}|S_2|^{2/3}|S_3|^{2/3} + |S_1|^{1/2}(|S_1|^{1/2} + |S_2| + |S_3|)),$

1197 unless $C_1 \cup C_2 \cup C_3$ is a line or a cubic curve.

1198 Nassajian Mojarrad, Pham, Valculescu and de Zeeuw [120] and Raz,
1199 Sharir and de Zeeuw [134] proved bounds for versions of the problem where
1200 F is a 4-variate polynomial.

3

1201

Contributions

1202

1203 In this chapter we expose our contributions. All sections have ta-
1204 bles that list past results and our contributions together with a
1205 reference. When the result is a contribution from this thesis, this
1206 reference has the format *Section(Contribution)*. The section ref-
1207 erence points to the details and proofs to attain the result and
1208 the contribution reference points to the theorem statement.

1209 3.1 Meiser Applied to k -SUM

1210 In Paper A we focus on the computational complexity of k -SUM. We
1211 recall its definition.

1212 **Problem 4** (k -SUM). Given n numbers $q_1 < q_2 < \dots < q_n \in \mathbb{R}$, decide
1213 whether there exist $i_1 < i_2 < \dots < i_k \in [n]$ such that $\sum_{j=1}^k q_{i_j} = 0$.

1214 In §A.1.1, we show the existence of an n -linear decision tree of depth
1215 $\tilde{O}(n^3)$ for k -SUM using a careful implementation of Meiser's algorithm [113].
1216 Although the high-level algorithm itself is not new, we refine the implemen-
1217 tation and analysis for the k -SUM problem.¹ Meiser presented his algorithm
1218 as a general method of d -dimensional point location in the arrangement of
1219 m hyperplanes that yielded a $\tilde{O}(d^4 \log m)$ -depth algebraic computation tree;
1220 when viewing the k -SUM problem as a point location problem, with $d = n$
1221 and $m = O(n^k)$, Meiser's algorithm can be applied to this problem, yielding
1222 a $\tilde{O}(n^4)$ -depth algebraic computation tree. We show that while the original

¹After submitting our results, we learned from a personal communication with Hervé Fournier that a similar analysis for arbitrary hyperplanes appears in his PhD thesis [76] (in French).

1223 algorithm was cast as a nonuniform polynomial-time algorithm, it can be
 1224 implemented in the linear decision tree model with an $\tilde{O}(n^3)$ upper bound.
 1225 Moreover, this result implies the same improved upper bound on the depth
 1226 of algebraic computation trees for the k -SUM problem, as shown in §A.2.2.

1227 There are two subtleties to this result. The first is inherent to the
 1228 chosen complexity model: even if the number of queries to the input is
 1229 small (in particular, the degree of the polynomial complexity is invariant on
 1230 k), the time required to *determine which queries should be performed* may
 1231 be arbitrary. In a naïve analysis, we show it can be trivially bounded by
 1232 $\tilde{O}(n^{k+2})$. In §A.1.2 we improve on this and present an algorithm to choose
 1233 which decisions to perform whereby the running time can be reduced to
 1234 $\tilde{O}(n^{\frac{k}{2}+8})$. Hence, we obtain an $\tilde{O}(n^{\frac{k}{2}+8})$ time randomized algorithm in the
 1235 RAM model expected to perform $\tilde{O}(n^3)$ linear queries on the input.

1236 **Contribution 1.** *There exist linear decision trees of depth $O(n^3 \log^2 n)$*
 1237 *solving the k -SUM and the k -LDT problems. Furthermore, for the k -SUM*
 1238 *problem there exists an $\tilde{O}(n^{\lceil \frac{k}{2} \rceil + 8})$ Las Vegas algorithm in the word-RAM*
 1239 *model expected to perform $O(n^3 \log^2 n)$ linear queries on the input. This*
 1240 *algorithm also solves the k -LDT problem within the same time bounds pro-*
 1241 *vided the coefficients of the underlying linear equation are constant rational*
 1242 *numbers.*

1243 For those analyses, we consider algorithms in the standard word-RAM
 1244 model with $\Theta(\log n)$ -size words, but in which the input $q \in \mathbb{R}^n$ is accessible
 1245 *only* via a linear query oracle. Hence we are not allowed to manipulate the
 1246 input numbers directly. The complexity is measured in two ways: by counting
 1247 the total number of queries, just as in the linear decision tree model, and by
 1248 measuring the overall running time, taking into account the time required
 1249 to determine the sequence of linear queries. This two-track computation
 1250 model, in which the running time is distinguished from the query complexity,
 1251 is commonly used in results on comparison-based sorting problems where
 1252 analyses of both runtime and comparisons are of interest [39, 40, 147].

1253 The second issue we address is that the linear queries in the above
 1254 algorithm may have size n , that is, they may use all the components of the
 1255 input. The lower bound of Erickson shows that if the queries are of minimal
 1256 size, the number of queries cannot be a polynomial independent of k such as
 1257 what we obtain, so non-minimal query size is clearly essential to a drastic

Table 3.1. Complexities of past and new algorithms for the k -SUM problem.

For our new algorithms, the number of blocks is a parameter that allows us to change the query size (see §A.1.3). The origin of the constant in the exponent of the time complexity is due to Lemma A.5. We conjecture it can be reduced, though substantial changes in the analysis will likely be needed to do so.

Reference	Blocks	Query Size	Queries	Time
Folklore	—	k	$\tilde{O}(n^{\lceil \frac{k}{2} \rceil})$	$\tilde{O}(n^{\lceil \frac{k}{2} \rceil})$
Meiser [113]	—	n	$\tilde{O}(n^3)$ §A.1.1	$\tilde{O}(n^{k+2})$ §A.1.1
[81, 91]	—	$2k - 2$	$\tilde{O}(n^{\frac{k}{2}})$	$\tilde{O}(n^{\frac{k}{2}})$
· (1)	1	n	$\tilde{O}(n^3)$ §A.1.1	$\tilde{O}(n^{\lceil \frac{k}{2} + 8 \rceil})$ §A.1.2
§A.1.3 (2)	b	$k \lceil \frac{n}{b} \rceil$	$\tilde{O}(b^{k-4} n^3)$	$\tilde{O}(b^{\lfloor \frac{k}{2} - 9 \rfloor} n^{\lceil \frac{k}{2} + 8 \rceil})$
§A.1.3 (A.7)	$b = n^{\Theta(1)}$	$o(n)$	$\tilde{O}(n^3)$	$\tilde{O}(n^{\lceil \frac{k}{2} + 8 \rceil})$
§A.1.3 (A.8)	$b = \Theta(n^\alpha)$	$O(n^{1-\alpha})$	$\tilde{O}(n^{3+(k-4)\alpha})$	$\tilde{O}(n^{(1+\alpha)\frac{k}{2} + 8.5})$

reduction in the number of queries needed. This gives rise to the natural question as to what is the relation between query size and number of queries. In particular, one natural question is whether queries of size less than n would still allow the problem to be solved using a number of queries that is a polynomial independent of k . We show that this is possible; in §A.1.3, using a blocking scheme, we introduce a family of algorithms exhibiting an explicit tradeoff between the number of queries and their size.

Contribution 2. *For any integer $b > 0$, there exists a $k \lceil \frac{n}{b} \rceil$ -linear decision tree of depth $\tilde{O}(b^{k-4} n^3)$ solving the k -SUM problem. Moreover, this decision tree can be implemented as an $\tilde{O}(b^{\lfloor \frac{k}{2} \rfloor - 9} n^{\lceil \frac{k}{2} \rceil + 8})$ Las Vegas word-RAM algorithm.*

We expose two corollaries of this contribution: We show that we can restrict our algorithms to use $o(n)$ -linear queries while keeping the same complexity bounds, up to polylogarithmic factors. We also give a range of tradeoffs for $O(n^{1-\alpha})$ -linear decision trees. Although the proposed algorithms still involve nonconstant-size queries, this is the first time such tradeoffs are explicitly tackled. Table 3.1 summarizes our results.

1275 3.2 Grønlund and Pettie Applied to 3POL

1276 In Paper B we focus on the computational complexity of 3POL. Since
 1277 3POL contains 3SUM, an interesting question is whether a generalization
 1278 of Grønlund and Pettie’s 3SUM algorithm exists for 3POL. If this is true,
 1279 then we might wonder whether we can “beat” the $O(n^{11/6}) = O(n^{1.833\dots})$
 1280 combinatorial bound of Raz, Sharir and de Zeeuw [133] with nonuniform algo-
 1281 rithms (see §2.2.8). We give a positive answer to both questions: we design
 1282 a uniform $O(n^2(\log \log n)^{3/2}/(\log n)^{1/2})$ -time algorithm and a nonuniform
 1283 $O(n^{12/7+\varepsilon}) = O(n^{1.7143})$ -time algorithm for 3POL. To prove our uniform
 1284 result, we present a fast algorithm for the Polynomial Dominance Reporting
 1285 (PDR) problem, a far reaching generalization of the Dominance Reporting
 1286 problem. As the algorithm for Dominance Reporting and its analysis by
 1287 Chan [44] is used in fast algorithms for all-pairs shortest paths, $(\min,+)$ -
 1288 convolutions, and 3SUM, we expect this new algorithm will have more
 1289 applications.

1290 To make the exposition of our results simpler, we study two different
 1291 variants of the 3POL problem. The first variant is the 3POL problem as
 1292 defined in the previous chapter (§2.2.7). We recall its definition here.

1293 **Problem 13** (3POL). Let $F \in \mathbb{R}[x, y, z]$ be a trivariate polynomial of constant
 1294 degree, given three sets A , B , and C , each containing n real numbers, decide
 1295 whether there exist $a \in A$, $b \in B$, and $c \in C$ such that $F(a, b, c) = 0$.

1296 The second variant is a special case of the 3POL problem where we
 1297 restrict the trivariate polynomial F to have the form $F(a, b, c) = f(a, b) - c$.
 1298 We call it the explicit 3POL problem because the dependency on the third
 1299 variable is explicitly given:

1300 **Problem 14** (explicit 3POL). Let $f \in \mathbb{R}[x, y]$ be a bivariate polynomial of
 1301 constant degree, given three sets A , B , and C , each containing n real numbers,
 1302 decide whether there exist $a \in A$, $b \in B$, and $c \in C$ such that $c = f(a, b)$.

1303 Similarly to Grønlund and Pettie [91], we consider both nonuniform
 1304 and uniform models of computation. For the nonuniform model, Grønlund
 1305 and Pettie consider linear decision trees. Because linear decision trees
 1306 are not powerful enough to even solve the problems we are looking at,
 1307 in this contribution, we consider *bounded-degree algebraic decision trees*

Table 3.2. Algorithmic results for 3POL and its explicit variant.

Reference	Model	Complexity
§B.1.1 (3), §B.1.3 (5)	ADT	$O(n^{12/7+\epsilon})$
§B.1.2(4), §B.1.4(6)	real-RAM	$O(n^2(\log \log n)^{3/2}/(\log n)^{1/2})$

1308 (ADT), an algebraic generalization of linear decision trees (see §1.1.2). For
 1309 the uniform model we consider the real-RAM model with only the four
 1310 arithmetic operators (see §1.1.1), the same model as in [91].

1311 In §B.1.1 and §B.1.2, we first design uniform and nonuniform algorithms
 1312 for explicit 3POL.

1313 **Contribution 3.** *Explicit 3POL can be solved in $O(n^{12/7+\epsilon})$ time in the
 1314 bounded-degree algebraic decision tree model.*

1315 **Contribution 4.** *Explicit 3POL can be solved in $O(\frac{n^2(\log \log n)^{3/2}}{(\log n)^{1/2}})$ time in
 1316 the real-RAM model.*

1317 In §B.1.3 and §B.1.4, we show that those algorithms can be adapted to
 1318 solve the more general 3POL problem.

1319 **Contribution 5.** *3POL can be solved in $O(n^{12/7+\epsilon})$ time in the bounded-
 1320 degree algebraic decision tree model.*

1321 **Contribution 6.** *3POL can be solved in $O(\frac{n^2(\log \log n)^{3/2}}{(\log n)^{1/2}})$ time in the real-
 1322 RAM model.*

1323 Table 3.2 gives a summary of our results.

1324 **Applications** Our results can be applied to many algebraic degeneracy
 1325 testing problems, such as GPT, a well known 3SUM-hard problem for
 1326 which no subquadratic algorithm is known (see §2.2). Raz, Sharir and
 1327 de Zeeuw results on the 3POL problem [133] can be applied to obtain a
 1328 combinatorial bound of $O(n^{11/6})$ on the number of collinear triples when
 1329 the input points are known to be lying on a constant number of polynomial
 1330 curves, provided those curves are neither lines nor cubic curves. A corollary
 1331 of our first result is that GPT where the input points are constrained to lie
 1332 on $o((\log n)^{1/6}/(\log \log n)^{1/2})$ constant-degree polynomial curves (including

lines and cubic curves) admits a subquadratic real-RAM algorithm and a strongly subquadratic bounded-degree algebraic decision tree. Interestingly, both reductions from 3SUM to GPT on 3 lines (map a to $(a, 0)$, b to $(b, 2)$, and c to $(-\frac{c}{2}, 1)$) and from 3SUM to GPT on a cubic curve (map a to (a^3, a) , b to (b^3, b) , and c to (c^3, c)) construct such special instances of GPT (see §2.2.2 for details on those reductions). This constitutes the first step towards closing the major open question of whether GPT can be solved in subquadratic time. To further convince the reader of the expressive power of the 3POL problem, we also give reductions from the problem of counting triples of points spanning unit circles (§B.3.2), from the problem of counting triples of points spanning unit area triangles (§B.3.3), and from the problem of counting collinear triples in any dimension (§B.3.1).

A Remark The algorithms we present manipulate polynomial expressions. In computational geometry, it is customary to assume the real-RAM model can be extended to allow the computation of roots of constant degree polynomials. We distance ourselves from this practice and take particular care of using the real-RAM model and the bounded-degree algebraic decision tree model with only the four arithmetic operators (see Chapter 9 on the Existential Theory of the Reals).

3.3 Slightly Subquadratic Encodings for GPT

In Paper C, we are interested in *compact* encodings for order types: we wish to design data structures using as few bits as possible that can be used to quickly answer orientation queries of a given abstract or realizable order type. There exist various ways to encode abstract order types optimally (see §2.2.6). However, it is not known how to decode the orientation of one triple from any of those optimal encodings in, say, sublinear time. Moreover, since the information-theoretic lower bound for realizable order types is only $\Omega(n \log n)$, we must ask if this space bound is approachable for those order types while keeping orientation queries reasonably efficient. Our contributions rely heavily on the hierarchical cuttings of Chazelle [48] described in detail in §8.2.

For ease of presentation, we state our results in terms of the encoding definition of §1.2.1 that we restate here.

1366 **Definition 2.** For fixed k and given a function $f : [n]^k \rightarrow [O(1)]$, we define
 1367 a $(S(n), Q(n))$ -encoding of f to be a string of $S(n)$ bits such that, given this
 1368 string and any $i \in [n]^k$, we can compute $f(i)$ in $Q(n)$ time in the word-RAM
 1369 model with word size $w \geq \log n$.

1370 In §C.1, we give the first optimal encoding for abstract order types that
 1371 allows efficient query of the orientation of any triple: the encoding is a data
 1372 structure that uses $O(n^2)$ bits of space with queries taking $O(\log n)$ time in
 1373 the word-RAM model with word size $w \geq \log n$.

1374 **Contribution 7.** *All abstract order types have an encoding using $O(n^2)$ bits
 1375 of space and allowing for queries in $O(\log n)$ time.*

1376 Our encoding is far from being space-optimal for realizable order types.
 1377 We show that its construction can be easily tuned to only require slightly
 1378 subquadratic space in this case.

1379 **Contribution 8.** *All realizable order types have a $O(\frac{n^2(\log \log n)^2}{\log n})$ -bits en-
 1380 coding allowing for queries in $O(\log n)$ time.*

1381 In §C.2, we further refine our encoding to reduce the query time to
 1382 $O(\log n / \log \log n)$.

1383 **Contribution 9.** *All abstract order types have an encoding using $O(n^2)$ bits
 1384 of space and allowing for queries in $O(\frac{\log n}{\log \log n})$ time.*

1385 **Contribution 10.** *All realizable order types have a $O(\frac{n^2 \log^\epsilon n}{\log n})$ -bits encoding
 1386 allowing for queries in $O(\frac{\log n}{\log \log n})$ time.*

1387 In the realizable case, we give quadratic upper bounds on the preprocess-
 1388 ing time required to compute an encoding in the real-RAM model.

1389 **Contribution 11** (Contributions 17 and 18). *In the real-RAM model and
 1390 the constant-degree algebraic decision tree model, given n real-coordinate
 1391 input points in \mathbb{R}^2 we can compute the encoding of their order type as in
 1392 Contributions 7, 8, 9, and 10 in $O(n^2)$ time.*

1393 In §C.3 we generalize our encodings to chirotopes of point sets in higher
 1394 dimension.

Table 3.3. Past results and new contributions on order type encoding.

Reference	Query Time	Space (bits)	Preprocessing
Trivial	$O(1)$	$O(n^3)$	$O(n^3)$
Enumeration	$2^{\Omega(n \log n)}$	$O(n \log n)$	$2^{\Omega(n \log n)}$
λ -matrices [85]	?	$O(n^2 \log n)$	$O(n^2)$
Numerical [90]	$O(n)$	$2^{\Theta(n)}$?
Permutations [104, 149]	$O(1)$	$O(n^2 \log n)$	$O(n^2)$
Wiring diagrams [73, 74]	$O(n^2)$	$O(n^2)$	$O(n^2)$
Canonical Labeling [17]	?	$O(n^2)$	$O(n^2)$
Abstract §C.1,§C.2 (9)	$O(\frac{\log n}{\log \log n})$	$O(n^2)$	$O(n^2)$
Realizable §C.1,§C.2 (8)	$O(\log n)$	$O(\frac{n^2 \log^2 \log n}{\log n})$	$O(n^2)$
Realizable §C.1,§C.2 (10)	$O(\frac{\log n}{\log \log n})$	$O(\frac{n^2}{\log^{1-\epsilon} n})$	$O(n^2)$
Realizable in \mathbb{R}^d §C.3 (12)	$O(\frac{\log n}{\log \log n})$	$O(\frac{n^d \log^2 \log n}{\log n})$	$O(n^d)$

1395 **Contribution 12.** All realizable chirotopes of rank $k \geq 4$ have an encoding
 1396 using $O(\frac{n^{k-1} (\log \log n)^2}{\log n})$ bits of space and allowing for queries in $O(\frac{\log n}{\log \log n})$
 1397 time.

1398 **Contribution 13.** In the real-RAM model and the constant-degree algebraic
 1399 decision tree model, given n real-coordinate input points in \mathbb{R}^d we can compute
 1400 the encoding of their chirotope as in Theorem 12 in $O(n^d)$ time.

1401 Table 3.3 gives a summary of our results.

1402 **A Remark** Our data structure is the first subquadratic encoding for
 1403 realizable order types that allows efficient query of the orientation of any
 1404 triple. It is not known whether a subquadratic algebraic computation
 1405 tree exists for identifying the order type of a given point set. Any such
 1406 computation tree would yield another subquadratic encoding for realizable
 1407 order types, although the obtained encoding might not be query-efficient. We
 1408 see the design of compact encodings for realizable order types as a subgoal
 1409 towards subquadratic nonuniform algorithms for this related problem, a
 1410 major open problem in Computational Geometry. Note that pushing the
 1411 preprocessing time below quadratic would yield such an algorithm.

1412 3.4 Better Encodings for 3SUM

1413 In Paper D, we consider the following problem: given three sets of real
1414 numbers, output a word-RAM data structure from which we can efficiently
1415 recover the sign of the sum of any triple of numbers, one in each set.

1416 This problem is similar to the problem of encoding the order type of a
1417 finite set of points studied in Paper C. While this contribution showed that
1418 it was possible to achieve slightly subquadratic space and logarithmic query
1419 time, we show here that we can encode the 3SUM *type* of n numbers using
1420 $\tilde{O}(n^{3/2})$ bits to allow for constant time queries in the word-RAM. We also
1421 study lower and upper bounds of other natural 3SUM type encodings.

1422 As there are only $O(n^3)$ queries, a table of size $n^3 \log_2 3 + O(1)$ bits
1423 suffices to give constant query time [60]. This can be improved to $O(n^2 \log n)$
1424 bits of space by storing for each pair (i, j) the values $k_{<}(i, j) = \max\{0\} \cup$
1425 $\{k : a_i + b_j + c_k < 0\}$ and $k_{>}(i, j) = \min\{n+1\} \cup \{k : a_i + b_j + c_k > 0\}$. For a
1426 query (i, j, k) , we compare k against the values $k_{<}(i, j)$ and $k_{>}(i, j)$ to recover
1427 $\chi(i, j, k)$ in $O(1)$ time. All $k_{<}(i, j)$ and $k_{>}(i, j)$ can be computed in $O(n^2)$
1428 time via the classic quadratic time algorithm for 3SUM.

1429 One seemingly simple representation is to store the numbers in A , B and
1430 C ; however these are reals and thus we need to make them representable
1431 using a finite number of bits. In §D.1 we show that a minimal integer
1432 representation of a 3SUM instance may require $\Theta(n)$ bits per value, which
1433 would give rise to a $O(n)$ query time and $O(n^2)$ space, which is far from
1434 impressive.

1435 **Contribution 14.** *Every 3SUM instance has an equivalent integer instance
1436 where all values have absolute value at most $2^{O(n)}$. Furthermore, there exists
1437 an instance of 3SUM where all equivalent integer instances require numbers
1438 at least as large as the n th Fibonacci number and where the standard binary
1439 representation of the instance requires $\Omega(n^2)$ bits.*

1440 In [38] the problem of given a set of n lines, to create an encoding of
1441 them so that the orientation of any triple (the *order type*) can be determined
1442 was studied; our problem is a special case of this where the lines only have
1443 three slopes. Can we do better for the case of 3SUM? We answer this in the
1444 affirmative.

1445 Again, for ease of presentation, we state our results in terms of the
1446 encoding definition of §1.2.1 that we restate here.

Table 3.4. Past results and new contributions on 3SUM type encoding.

Reference	Query time	Space (bits)	Preprocessing
Trivial	$O(1)$	$O(n^3)$	$O(n^3)$
Almost trivial	$O(1)$	$O(n^2 \log n)$	$O(n^2)$
Order type (Paper C)	$O(\log n)$	$O(\frac{n^2 \log^2 \log n}{\log n})$	$O(n^2)$
Order type (Paper C)	$O(\frac{\log n}{\log \log n})$	$O(\frac{n^2}{\log^{1-\epsilon} n})$	$O(n^2)$
Numerical §D.1(14)	$O(n)$	$O(n^2)$	$n^{O(1)}$
Space-optimal §D.2(15)	$n^{O(1)}$	$O(n \log n)$	$n^{O(1)}$
Query-optimal §D.3(16)	$O(1)$	$\tilde{O}(n^{3/2})$	$O(n^2)$

1447 **Definition 2.** For fixed k and given a function $f : [n]^k \rightarrow [O(1)]$, we define
1448 a $(S(n), Q(n))$ -encoding of f to be a string of $S(n)$ bits such that, given this
1449 string and any $i \in [n]^k$, we can compute $f(i)$ in $Q(n)$ time in the word-RAM
1450 model with word size $w \geq \log n$.

1451 In §D.2 we show how to use an optimal $O(n \log n)$ bits of space with a
1452 polynomial query time.

1453 **Contribution 15.** All 3SUM types have a $O(n \log n)$ -bits $n^{O(1)}$ -querytime
1454 encoding.

1455 Finally, in §D.3 we show how to use $\tilde{O}(n^{3/2})$ space to achieve $O(1)$ -time
1456 queries.

1457 **Contribution 16.** All 3SUM types have a $\tilde{O}(n^{3/2})$ -bits $O(1)$ -querytime
1458 encoding.

1459 Table 3.4 gives a summary of past and new results.

4

1460

Developments

1461

1462 Since the publication of our results, new developments have surfaced.

4.1 Better Nonuniform Algorithms for k -SUM

1464 Meiser's algorithm solves a broader class of problems than the k -SUM
1465 problems. In its general formulation, this algorithm solves the point location
1466 problem (Problem 8): to locate an input point in a fixed arrangement of m
1467 hyperplanes in R^d .

1468 Our analysis of Meiser's algorithm yields a nonuniform upper bound
1469 of $O(d^3 \log d \log m)$ for this problem. However, this is far from optimal.
1470 The information theoretic lower bound obtained through Buck's theorem
1471 (Theorem 6.1) is only $\Omega(d \log m)$.

1472 Two new developments have surfaced. The first reduces the nonuni-
1473 form complexity of Meiser's algorithm for the point location problem from
1474 $O(d^3 \log d \log m)$ down to $O(d^2 \log m)$. The second invents a completely novel
1475 technique to reduce the nonuniform complexity of the same problem down to
1476 $O(d \log d \log m)$, provided the coefficients of the hyperplanes are constants.

1477 We briefly sketch those two developments in the next subsections.

4.1.1 Using Vertical Decomposition

1479 In 2016, Ezra and Sharir [72] improve the decision tree depth of Meiser's
1480 algorithm by using vertical decomposition (see §6.4.2) as the cell refinement
1481 subroutine.

1482 A big player in the complexity of our algorithm is the size of the sample
1483 needed for the ε -net. It turns out that with vertical decomposition the sample
1484 size can be significantly reduced: $O(d)$ instead of $O(d^2 \log d)$ for simplices. A

sample size of $O(d \log d)$ can easily be attained using existing results on the VC-dimension of certain range spaces. One of the key results in [72] is to get the sample size down to $O(d)$.

For instance, when applying this improved algorithm to the k -SUM problem or the subset-sum problem, the change in the sample size causes Meiser's decision tree to be much shallower than our implementation: the depth of the tree is reduced by a $n \log n$ factor in both cases leading to overall complexities of $O(n^2 \log n)$ and $O(n^3)$ respectively.

4.1.2 Using Inference Dimension

In 2017, Kane, Lovett, and Moran [103] proved that the point location problem can be solved in $O(d \log d \log m)$ linear queries.

Their decision tree is also a prune and search algorithm following the generic approach of Meiser: sample, locate, prune, recurse. The key ingredient is a new tool dubbed *inference dimension* that takes the role of the VC-dimension in Meiser's algorithm.

The major catch with this new decision tree is that it will only work for sparse hyperplanes: hyperplanes with a small hamming weight. This is precisely the case for the application of point location to the k -SUM and subset-sum problems. For those applications, the decision trees in [103] are shallower than the decision trees in [72] by a factor of $n / \log n$: $O(n \log^2 n)$ for k -SUM and $O(n^2 \log n)$ for subset-sum.

Another feature of this algorithm is that it only uses *comparison queries*. When applied to k -SUM, our algorithm and the one in [72] are n -linear decision trees, while the one in [103] is a $2k$ -linear decision tree.

In 2018, the authors of [103] published another paper generalizing their decision tree to work for arbitrary hyperplanes [102]. However, the depth they obtain is much worse in that case: $O(d^4 \log d \log m)$ for arbitrary hyperplanes. Moreover, it requires the use of *generalized comparison queries* which are more complex than the comparison queries in [103].

In 2019, Hopkins, Kane, and Lovett [99] provide another decision tree for point location of depth $O(d \log^2 m)$ when the set of hyperplanes is drawn from a distribution with weak restrictions. This new algorithm only uses comparison queries.

1518 4.2 Timothy Chan Strikes Again

1519 In 2017, Timothy Chan presented a $O((n^2/\log^2 n)(\log \log n)^{O(1)})$ -time
1520 real-RAM algorithm for 3SUM [46], shaving a logarithmic factor from previ-
1521 ous solutions [79, 81].

1522 His technique relies on cuttings in near-logarithmic dimension on an
1523 augmented real-RAM model with constant time nonstandard operations on
1524 $\Theta(\log n)$ bits words. Cuttings essentially replace dominance reporting in
1525 Grønlund and Pettie’s algorithm [91]. This is inspired by another one of his
1526 papers on APSP [45].

1527 This new technique for shaving off logarithmic factors can be applied to
1528 other problems, such as (median, +)-convolution, (median, +)-matrix multi-
1529 plication, and 3SUM-hard problems in computational geometry including
1530 explicit 3POL.

5

Open Questions

1533 These are a few interesting questions that are left unanswered.

1534 5.1 About Algorithms

1535 The best known linear decision tree for k -SUM has depth $O_k(n \log^2 n)$
 1536 while the lower bound is $\Omega_k(n \log n)$ [103] (see §4.1.2). Can we match the
 1537 information theoretic lower bound for k -SUM with linear decision trees?

1538 **Open Question 1.** *Is there a $O_k(n \log n)$ -depth LDT for k -SUM?*

1539 The same question can be asked about the point location problem. The
 1540 best linear decision tree for arbitrary hyperplanes has depth $O(d^2 \log m)$ and
 1541 the lower bound is $\Omega(d \log m)$ [72] (see §4.1.1).

1542 **Open Question 2.** *Is there a $O(d \log m)$ -depth LDT for point location?*

1543 In Paper A we show how to efficiently implement Meiser's algorithm in
 1544 the word-RAM model when applied to k -SUM. Can the same be done for
 1545 the algorithms of Ezra and Sharir [72] and Kane, Lovett, and Moran [103]?

1546 **Open Question 3.** *Is there a word-RAM algorithm for k -SUM running in
 1547 time $n^{\frac{k}{2} + O(1)}$ accessing the input through $n^{3 - \Omega(1)}$ n -linear queries only?*

1548 We also have not managed to match the running time of the best known
 1549 uniform algorithm (see §2.1.5).

1550 **Open Question 4.** *Is there a word-RAM algorithm for k -SUM running in
 1551 time $\tilde{O}(n^{\lceil \frac{k}{2} \rceil})$ accessing the input through $n^{o(k)}$ n -linear queries only?*

1552 Grønlund and Pettie [91] showed how to solve 3SUM in subquadratic time
 1553 in the real-RAM model. However, so far, their technique and its subsequent
 1554 improvements [46, 79, 81], have failed to generalize to uniform algorithms
 1555 for k -SUM.

1556 **Open Question 5.** *Is there a $o(n^{\lceil \frac{k}{2} \rceil})$ -time real-RAM algorithm for k -SUM
 1557 with $k \geq 4$?*

1558 The best known algorithms for 3SUM are only *slightly* subquadratic,
 1559 shaving a polylogarithmic factor from quadratic runtime. In Paper B, we
 1560 show how to adapt one of those algorithms to solve the 3POL problem within
 1561 the same time bound. The best lower bound we have for those problems is
 1562 $\Omega(n \log n)$. The current conjecture is that no significant improvement can be
 1563 made with respect to known algorithms (see §2.1). We have to ask if this is
 1564 indeed true.

1565 **Open Question 6.** *Is there a $n^{2-\Omega(1)}$ -time real-RAM algorithm for 3SUM?*

1566 The question can also be asked more generally for 3POL.

1567 **Open Question 7.** *Is there a $n^{2-\Omega(1)}$ -time real-RAM algorithm for 3POL?*

1568 The best known nonuniform algorithm for 3SUM is a 6-linear decision tree
 1569 of depth $O(n \log^2 n)$ (see §4.1.2), the nonuniform algorithms in [79, 81, 91]
 1570 are 4-linear decision trees of depth $\tilde{O}(n^{3/2})$, and we know that there is no
 1571 3-linear decision tree of depth $o(n^2)$ for 3SUM [69] (see §2.1). We want to
 1572 know whether better 4-linear decision trees exists for 3SUM.

1573 **Open Question 8.** *Is there a $o(n^{\frac{3}{2}})$ -depth 4-linear decision tree for 3SUM?*

1574 If this is the case, then we would also like to know if this model of
 1575 computation suffers from more severe limitations than the 6-linear decision
 1576 tree model.

1577 **Open Question 9.** *Is there a $n \log^{O(1)} n$ -depth 4-linear decision tree for
 1578 3SUM?*

1579 The inference dimension method introduced in [103] is essentially a pigeon
 1580 hole argument applied to linear combinations of linear equations with small
 1581 coefficients. At first sight, this kind of argument does not seem to apply to

polynomial equations, even when the polynomials are of degree two. One can of course apply a Veronese embedding [95, 96] to linearize the equations but this has the effect of blowing up the dimension which directly impacts the efficiency of the method. Could some generalization of inference dimension, or any other method, improve on the $O(n^{12/7+\epsilon})$ nonuniform time bound we obtain in Paper B?

Open Question 10. *Is there a $o(n^{12/7})$ -depth bounded-degree algebraic decision tree or algebraic computation tree for 3POL?*

In §B.3.1 we show how 3POL algorithms can be used to solve particular instances of GPT in subquadratic time. Despite all our efforts, we still lack a subquadratic algorithm for GPT in the general case, even a nonuniform one. Like 3SUM and 3POL, the best known lower bound is $\Omega(n \log n)$ (see §2.2.3).

Open Question 11. *Is there a $o(n^2)$ -depth bounded-degree algebraic decision tree or algebraic computation tree for GPT?*

5.2 About Encodings

In Paper C we show how to get compact encodings for abstract order types with slightly sublogarithmic query time. Is it possible to bring this query time down to constant, or even just strongly sublogarithmic?

Open Question 12. *Is there a $O(n^2)$ -bits $O(\log^{1-\epsilon} n)$ -querytime encoding for abstract order types?*

In the same paper, we manage to get slightly subquadratic space encodings for realizable order types. Is it possible to get this down to strongly subquadratic while keeping the query time reasonable, say polynomial?

Open Question 13. *Is there a $n^{2-\Omega(1)}$ -bits $n^{O(1)}$ -querytime encoding for realizable order types?*

In Paper D we modify Grønlund and Pettie’s nonuniform algorithm to obtain a 3SUM type encoding with a strongly subquadratic space requirement and efficient queries. Is there any way to exploit the nearlinear decision tree in [103] to obtain a better space bound while keeping query time efficient?

Open Question 14. *Is there a $o(n^{3/2})$ -bits $\log^{O(1)} n$ -querytime encoding for 3SUM types?*

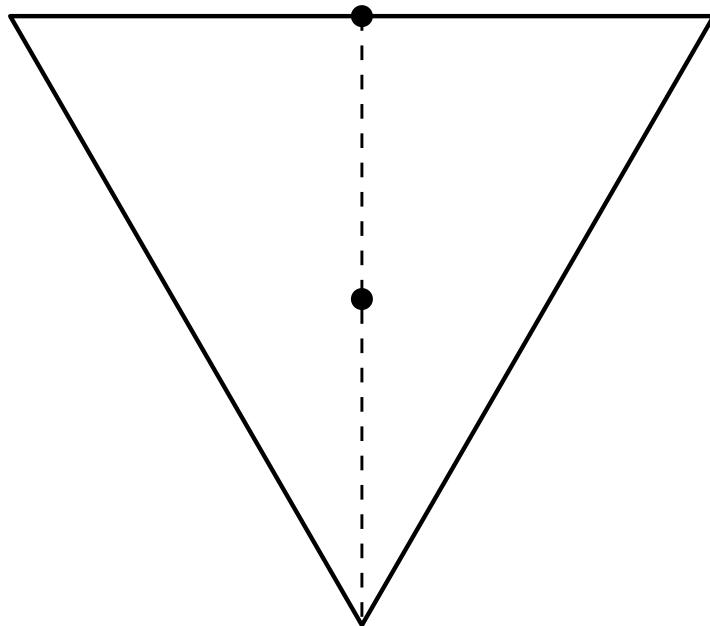
1614

II

1615

The Computational Geometer's Toolbox

1616



6

1617

Arrangements

1619 An *hyperplane* in \mathbb{R}^d is the set of points $q \in \mathbb{R}^d$ satisfying a linear
1620 equation of the form $\sum_{i=1}^d p_i q_i = 0$ for some tuple of coefficients $p \in \mathbb{R}^d$. An
1621 *affine* hyperplane has an additional independent term p_0 to get the equation
1622 $\sum_{i=1}^d p_i q_i = p_0$. In general we will drop the adjective “affine” and talk about
1623 hyperplanes without specifying whether $p_0 = 0$. In the plane an hyperplane
1624 is a line. In \mathbb{R}^3 an hyperplane is a plane. On the real line an hyperplane is a
1625 point.

1626 An hyperplane H partitions the space \mathbb{R}^d into two halfspaces, often
1627 denoted by H^- and H^+ . Naturally, H^- is defined to be the set of points
1628 such that $\sum_{i=1}^d p_i q_i < p_0$ and H^+ is defined to be the set of points such
1629 that $\sum_{i=1}^d p_i q_i > p_0$. When we want to be more precise as to whether the
1630 inequality is strict we can write $H^<$, H^{\leq} , $H^>$, or H^{\geq} . We also define H^0
1631 and $H^=$ as synonyms of H .

1632 A set of hyperplanes $\mathcal{H} = \{H_1, H_2, \dots, H_m\}$ partitions \mathbb{R}^d into convex
1633 regions called *cells*. This partition is denoted by $\mathcal{A}(\mathcal{H})$ and is called the
1634 *arrangement* of \mathcal{H} . Each sign vector $\sigma \in \{-, 0, +\}^m$ corresponds to a
1635 (potentially empty) cell of this arrangement defined as the set of points q
1636 such that $q \in H_i^{\sigma_i}$ for all $H_i \in \mathcal{H}$.

6.1 Counting Cells

1637 It is useful to understand the complexity of such an arrangement when
1638 the number m of hyperplanes grows and the dimension d is fixed. The goal
1639 is to bound the number of nonempty cells an arrangement can have. When
1640 $m \leq d$ each of the 3^m cells is nonempty, assuming general position. However,
1641 when we fix d and make m grow, we get a more reasonable behavior.

1643 **Theorem 6.1** (Buck [35], see also [93, Theorem 24.1.1 and Corollary 24.1.2]).

1644 Consider the partition of space defined by an arrangement of m hyperplanes
1645 in \mathbb{R}^d . The number of regions of dimension $k \leq d$ is at most

$$\binom{m}{d-k} \sum_{i=0}^k \binom{m-d+k}{i}.$$

1647 This bound is attained when the hyperplanes are in general position. The
1648 number of regions of all dimensions is $\Theta(m^d)$ in that case.

1649 This bound allows us to derive precise lower and upper bounds for
1650 algorithms manipulating those arrangements. See [93] for more details and
1651 generalizations.

1652 6.2 Pseudolines

1653 In the plane, line arrangements are generalized to pseudoline arrangements by dropping the “straight” nature of lines. A pseudoline in \mathbb{R}^2 is a
1654 simple curve connecting points at infinity. An arrangement of pseudolines is a
1655 collection of pseudolines that pairwise intersect exactly once. This definition
1656 can be made more formal by considering pairwise-intersecting simple closed
1657 curves in the projective plane instead. Some important differences between
1658 line and pseudoline arrangements have already been discussed in §2.2.5.

1660 Putting those differences aside, pseudoline arrangements share sufficient
1661 similarity with line arrangements as to allow a generic algorithmic treatment
1662 of both kinds. For instance, their arrangement complexities (§6.1) are
1663 asymptotically the same, the Zone Theorem (§6.3) holds in general, canonical
1664 labelings (§7.2) can be constructed by looking at the order type only (see
1665 §2.2.3 for a definition of order type), and generic cell decompositions (§6.4)
1666 allow the use of the divide-and-conquer methods of Chapter 8 in both cases.

1667 In Paper C we consider the problem of encoding order types of line
1668 or pseudoline arrangements. Strictly speaking, an order type consists of
1669 pure combinatorial information and does not require an embedding to be
1670 known. In particular, given an order type, it is $\exists\mathbb{R}$ -complete to decide
1671 whether a straigh embedding exists [118, 119]. For simplicity, the algorithms
1672 constructing the encodings in Paper C assume an embedding is given.

1673 For more on pseudoline arrangements see Chapter 5 of the Handbook
1674 on the subject [83]. For more on equivalent representations of pseudoline

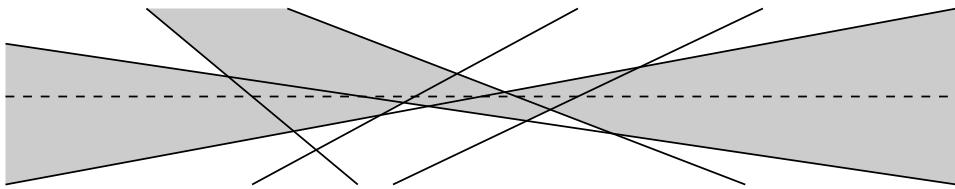


Figure 6.1. The zone defined by the dashed line in the two-dimensional arrangement of the plain lines is emphasized in light grey.

arrangements see Chapter 6 of the Handbook on oriented matroids [139].

6.3 Zone Theorem

The zone of a given pseudoline of an arrangement is the set of cells of the arrangement supported by that pseudoline. Figure 6.1 illustrates a zone in a two-dimensional arrangement of lines.

We define the complexity of each cell to be the number of its sides. We define the complexity of a zone to be the sum of the complexities of its cells. The Zone Theorem states that the complexity of any zone is linear.

Theorem 6.2 (Zone Theorem in the plane [27]). *Given an arrangement of $n + 1$ pseudolines, the sum of the numbers of sides in all the cells supported by one of the pseudolines is at most $\lfloor 9.5n \rfloor - 1$.¹*

This result is important because it allows optimal incremental construction of arrangements of line arrangements, a frequently used tool.

In [63] and in the first edition of [61], there were claims of generalization of this result to arrangements of hyperplanes in higher dimension. However, the proofs turned out to be flawed [64]. Edelsbrunner, Sturmfels, and Sharir were the first to provide a valid proof of the generalized result [64].

Theorem 6.3 (Zone Theorem [64]). *Given an arrangement of n hyperplanes in \mathbb{R}^d , the sum of the numbers of 0-faces, 1-faces, \dots , and $(d - 1)$ -faces in all the cells supported by one of the hyperplanes is $O(n^{d-1})$.*

They also cite valid proofs by Houle and Matoušek for a weaker version of the theorem: The Zone Theorem bounds the sum of complexities of the

¹Note that an earlier weaker (worse constant factor) linear bound is implied by a theorem in [50].

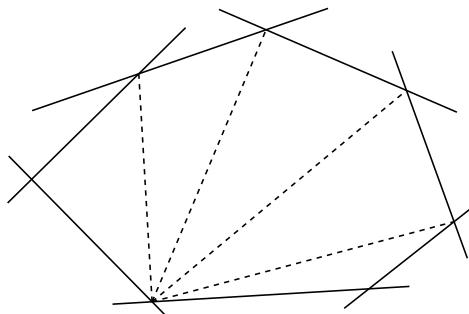


Figure 6.2. The bottom vertex triangulation of a cell in an arrangement of lines.

1697 cells in a zone. The full generalization of the theorem defines the complexity
 1698 of each cell to be the number of all 0-faces, 1-faces, ..., and $(d - 1)$ -faces
 1699 supporting the cell. The weak generalization only counts $(d - 1)$ -faces.
 1700 However, no publication of those proofs is known to the author.

1701 Note that the Zone Theorem implies the asymptotic upper bound of
 1702 $O(n^d)$ on the complexity of the arrangement (§6.1).

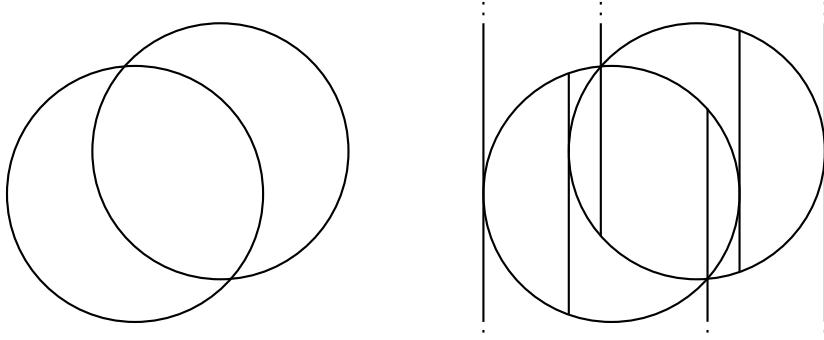
1703 6.4 Cell Decomposition

1704 Cells of arrangements may not behave nicely. In particular, they may have
 1705 arbitrary description complexity. In this section we give two decomposition
 1706 schemes that allow to partition cells into low-complexity subcells. Access to
 1707 such decompositions is an essential ingredient of divide-and-conquer methods
 1708 described in Chapter 8.

1709 Lookup the Handbook [93, §24.3.2], for a more general overview of cell
 1710 decomposition techniques.

1711 6.4.1 Bottom Vertex Triangulation

1712 Given an arrangement of hyperplanes in \mathbb{R}^d , its bottom vertex triangula-
 1713 tion is the partition of space obtained by triangulating each cell of the
 1714 arrangement recursively as follows: taking $d = 2$ as the base case (because
 1715 edges are already simplices), let p be the bottom-most vertex of the cell,
 1716 for each facet of the cell not containing p , for each $(d - 1)$ -dimensional
 1717 simplex S' of the bottom vertex triangulation of the facet in \mathbb{R}^{d-1} , construct
 1718 a d -dimensional simplex S that is the convex hull of S' and p . The bottom



- a. Some curves in \mathbb{R}^2 .
- b. Vertical decomposition of the curves in Figure 6.3a.

Figure 6.3. Vertical decomposition.

1719 vertex triangulation of an arrangement of lines in the plane is illustrated in
 1720 Figure 6.2.

1721 In Paper A we construct a simplex of the bottom vertex triangulation of a
 1722 ε -net as part of a point location procedure. In Paper B (§B.2.2) and Paper C
 1723 (all sections) we use hierarchical cuttings based on this decomposition in
 1724 order to construct an efficient point location data structure.

1725 See also [53] for a more thorough description and an application to
 1726 nearest neighbour data structures.

1727 **6.4.2 Vertical Decomposition**

1728 Given an arrangement of curves in \mathbb{R}^2 , its vertical decomposition is the
 1729 partition of space obtained by shooting a vertical segment from each vertex
 1730 of the arrangement until it hits a curve of the arrangement. The vertical
 1731 decomposition of an arrangement of two circles is illustrated in Figure 6.3.

1732 We use this decomposition in Paper B (§B.2.1) in order to apply a divide-
 1733 and-conquer algorithm of Matoušek [109] to an arrangement of curves. This
 1734 algorithm was originally designed to work for an arrangement of hyperplanes
 1735 using bottom-vertex triangulation decomposition. In Paper C (§C.1, §C.2),
 1736 we use this decomposition in order to encode the order type of an arrangement

1737 of pseudolines.

1738 Note that in the first application we only use vertical decomposition
1739 on a constant number of bounded-degree polynomial curves. In the second
1740 application we only care about the existence of such a decomposition. For
1741 those reasons, we do not discuss efficient construction of this decomposition.

1742 For the construction of the vertical decomposition of an arrangement
1743 of polynomial curves in \mathbb{R}^2 , we refer the reader to Pach and Sharir [123],
1744 Chazelle et al. [49], and Edelsbrunner et al. [62].

1745 **A Remark** This decomposition can be generalized to work for hypersur-
1746 faces in \mathbb{R}^d . Unfortunatly, the behaviour of such decompositions quickly
1747 degenerates with d . While bottom vertex triangulation gives a bound of
1748 $O(n^d)$ cells in \mathbb{R}^d and vertical decomposition gives a bound of $O(n^2)$ cells
1749 in \mathbb{R}^2 , we only know of a few upper bounds in fixed dimensions and some
1750 of them are worse than $O(n^d)$. This is why in our application of Meiser's
1751 algorithm we use the bottom vertex triangulation. However, as observed by
1752 Ezra and Sharir [72], the bad behaviour of vertical decompositions is not a
1753 bottleneck of Meiser's algorithm since we only need to look at a single cell
1754 of the decomposition. Moreover, the complexity of those cells is better than
1755 that of simplicial ones: vertical decomposition yields prisms supported by at
1756 most $2d$ hyperplanes of the arrangement whereas simplices of the bottom
1757 vertex triangulation are supported by $d^2 + d$ hyperplanes in the worst case.
1758 This has a direct impact on the VC-dimension of the range spaces defined
1759 by those objects and makes vertical decomposition the winning strategy for
1760 this particular algorithm.

7

1761

Chirotopes

1763 Chirotopes are the generalization of order types introduced in §2.2.3. We
1764 generalize Definitions 3 and 4 to point sets in \mathbb{R}^d .

1765 **Definition 5** (Chirotope of a Point Set in \mathbb{R}^d). Given a set of n points $p_i =$
1766 $(p_{i,1}, p_{i,2}, \dots, p_{i,d}) \in \mathbb{R}^d$, its (rank- $(d+1)$) chirotope is the function $\chi: [n]^{d+1} \rightarrow$
1767 $\{-, 0, +\}$ such that

$$\chi(i_1, i_2, \dots, i_{d+1}) = \det \begin{pmatrix} 1 & p_{i_1,1} & p_{i_1,2} & \cdots & p_{i_1,d} \\ 1 & p_{i_2,1} & p_{i_2,2} & \cdots & p_{i_2,d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & p_{i_{d+1},1} & p_{i_{d+1},2} & \cdots & p_{i_{d+1},d} \end{pmatrix},$$

1769 up to isomorphism.

1770 **Definition 6** (Realizable Chirotope). When considering an arbitrary function
1771 $f: [n]^{d+1} \rightarrow \{-, 0, +\}$ we say that f is a realizable chirotope if there is a n -set
1772 $P \subset \mathbb{R}^2$ such that its chirotope is f .

1773 In this chapter, we give details on two essential ingredients of our data
1774 structures in Paper C: the classic point-hyperplane duality and the canonical
1775 labeling of order types and chirotopes. The first one is necessary to be able to
1776 apply the encodings we obtain for line arrangements and hyperplane arrange-
1777 ments to point configurations. The second is used to get the preprocessing
1778 time of those encodings down to $O(n^d)$.

7.1 Duality

1780 Technically speaking, the encoding we describe for realizable chirotopes in
1781 Paper C encodes the chirotope of a given arrangement of lines or hyperplanes.

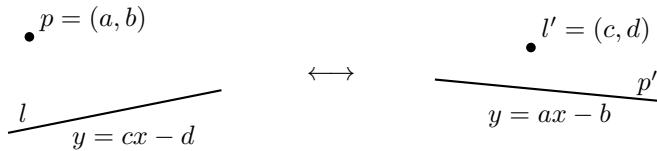


Figure 7.1. Order preserving duality: “ p is above l ” if and only if “ p' is above p' ”.

Moreover, for ease of presentation, we make the assumption that the vertices of this arrangement have finite coordinates. In the two-dimensional case, this is equivalent to having no two lines parallel. In these paragraphs, we give the details necessary to rigorously handle all realizable chirotopes, including degenerate ones. This is especially important in higher dimension, where the situation is a bit more complicated than in two dimensions.

In two dimensions, we wish to encode order types of point configurations. Since our encoding construction algorithm works with an arrangement of lines as input, we need a mapping from those primal points to their dual lines. This mapping should preserve the order type of the point configuration, hence it needs to be *order-preserving*. One such order-preserving duality is the mapping $(a, b) \leftrightarrow y = ax - b$ (see Figure 7.1).

To avoid parallel lines in the dual, it suffices to avoid intersection points at infinity. In the primal, this translates to avoiding two points of the configuration defining a vertical line, that is, with the same x coordinate. This is easily done by performing a tiny rotation in the primal. This (proper) rotation does not change the order type of the point set.

In higher dimension, the order-preserving point-line duality generalizes to the following order-preserving point-hyperplane duality: We map each d -dimensional point $(x_1, x_2, \dots, x_d) \in \mathbb{R}^d$ to the hyperplane $y_d = \sum_{i=1}^{d-1} x_i y_i - x_d$ and the hyperplane $x_d = \sum_{i=1}^{d-1} y_i x_i - y_d$ to the d -dimensional point $(y_1, y_2, \dots, y_d) \in \mathbb{R}^d$.

As before, we want hyperplanes to be non-parallel. In fact, we need an even stronger assumption: We want all linearly independent subsets of d hyperplanes to intersect in a point with finite coordinates. Having no intersection points at infinity in the dual means having no d points spanning a hyperplane parallel to the x_d axis in the primal. This is easy to avoid by applying tiny rotations in the primal. Again, those (proper) rotations do

1810 not change the chirotope of the point set.

1811 In dimension three and higher, one would think degenerate arrangements
 1812 lead to annoying nongeneral situations. However, those situations are easy
 1813 to handle with our technique: Degenerate subsets of hyperplanes are linearly
 1814 dependent. The determinant corresponding to a query asking about a
 1815 degenerate $d + 1$ subset is therefore zero. Our technique will identify those
 1816 degenerate queries and map them to the correct answer in a space-efficient
 1817 way.

1818 7.2 Canonical Labelings

1819 Given a point set, the composition of its order type χ with a permutation
 1820 ρ produces a new order type $\chi' = \chi \circ \rho$. This composition corresponds to a
 1821 relabeling of the point set. Aloupis et al. [17] defined the canonical labeling
 1822 $\rho^*(\chi)$ of an order type χ to be a permutation such that for all permutations
 1823 π we have $\rho^*(\chi \circ \pi) = \pi^{-1} \circ \rho^*(\chi)$. In other words, given two isomorphic
 1824 order types χ and χ' , we have $\chi \circ \rho^*(\chi) = \chi' \circ \rho^*(\chi')$, and $\rho^*(\chi')^{-1} \circ \rho^*(\chi)$ is
 1825 the isomorphism that sends χ to χ' .¹ They proved that the function ρ^* is
 1826 computable in $O(n^2)$ time. This first tool is useful to identify isomorphic
 1827 order types.

1828 They also showed that given any order type χ , a string $E(\chi)$ of $O(n^2)$
 1829 bits, called the representation of χ , can be computed in $O(n^2)$ time, such
 1830 that, if χ and χ' are two isomorphic order types, then $E(\chi) = E(\chi')$. This
 1831 second tool is useful to quickly compare two order types (a naive solution
 1832 would take $\Theta(n^3)$ time by first computing a canonical labeling, and then
 1833 comparing all triples).

1834 **Lemma 7.1** (Aloupis et al. [17]). *Given an order type presented as an oracle,
 1835 its canonical labeling of $O(n \log n)$ bits and its canonical representation of
 1836 $O(n^2)$ bits can be computed in $O(n^2)$ time in the word-RAM model.*

1837 Both tools generalize to chirotopes of point configurations in any dimension
 1838 d and, more generally, to chirotopes of rank $d + 1$.

1839 **Lemma 7.2** (Aloupis et al. [17]). *For all $d \geq 2$, given a rank- $(d + 1)$ chiro-
 1840 tope presented as an oracle, its canonical labeling of $O(n \log n)$ bits and its*

¹Sometimes, two order types χ and $-\chi$ are also considered to be isomorphic. See [17] for more details.

₁₈₄₁ canonical representation of $O(n^d)$ bits can be computed in $O(n^d)$ time in the
₁₈₄₂ word-RAM model.

8

1843

1844

Divide and Conquer

1845 A general divide-and-conquer paradigm in algorithm design is, given an
1846 instance of size n , to 1) divide it into at most a smaller instances of size at
1847 most $\frac{n}{b}$ in $p(n)$ preprocessing time, 2) solve those instances recursively, solving
1848 constant size instances by brute force, and 3) postprocess the solutions of
1849 those smaller instances to obtain the solution to the original one in $p(n)$
1850 time. An upper bound $T(n)$ on the time complexity of such an algorithm is
1851 the recurrence

$$1852 \quad T(n) = a T\left(\frac{n}{b}\right) + p(n).$$

1853 The asymptotic behavior of such a recurrence can then be derived from the
1854 Master Theorem [26, 55].

1855 In this chapter we expose a couple standard divide-and-conquer techniques
1856 in Computational Geometry that allow the implementation of this paradigm
1857 (or variants of it). Papers A, B, and C explicitly rely on those techniques.
1858 The subquadratic-space constant-querytime encoding in Paper D can also
1859 be interpreted as an ad-hoc implementation of those concepts (see §D.3).

1860 8.1 Epsilon Nets and Cuttings

1861 The efficient construction of ε -nets and cuttings is a necessary condition
1862 for many divide-and-conquer schemes in Computational Geometry. Those
1863 constructions are based on the concept of Vapnik-Chervonenkis dimension
1864 (VC-dimension) of a range space.

1865 8.1.1 Range Spaces

1866 A *range space* (or set system) consists of a *ground set* (or universe) and
1867 a family of subsets of this ground set called *ranges*. A simple example of a

range space is to take a finite set of points on the real line as the ground set and to take the subsets of points induced by intervals as the ranges. In general, we have the following definition.

Definition 7 (Range Space). A pair (X, \mathcal{R}) such that X is a set and $\mathcal{R} \subseteq 2^X$ is a family of subsets of X .

Note that, for simplicity, the examples, definitions, and lemmas we state here consider the case of finite ground sets. Those can be generalized to infinite universes.

In the worst case, \mathcal{R} could be equal to 2^X . If X is taken to be points in \mathbb{R}^d and the ranges are taken to be subsets induced by simple geometric objects, as in most geometric applications, this is not possible.

For instance, consider points and intervals on the real line. For any set of points, only $O(|X|^2)$ subsets can be induced by intervals. Another example is that of points in the plane and subsets induced by halfplanes. Because any halfplane can be moved to have exactly two points on its boundary, there are $O(|X|^2)$ such subsets. Those observations can be generalized. For that we need a few more definitions.

8.1.2 VC-dimension

A set is *shattered* by a family of ranges if each of its subsets can be induced by a range of the family.

Definition 8 (Shattered Set). Let (X, \mathcal{R}) be a range space. A subset of $T \subseteq X$ is said to be shattered by \mathcal{R} if every subset $T' \subseteq T$ is such that $T' = T \cap R$ for some $R \in \mathcal{R}$.

In our points and intervals example it is easy to see that a set of two distinct points can be shattered but a set of three points cannot. For the points and halfplanes example sets of three noncollinear points are shattered but not sets of four points.

The VC-dimension of a range space is defined to be the size of the largest shattered subset.

Definition 9 (VC-dimension). Let (X, \mathcal{R}) be a range space. Let v be the size of the largest $T \subseteq X$ such that T is shattered by \mathcal{R} . The VC-dimension of the range space is defined to be v .

1900 The Perles-Sauer-Shelah-Vapnik-Chervonenkis lemma gives a direct con-
 1901 nection between this parameter and the number of ranges.

1902 **Lemma 8.1** (Vapnik and Chervonenkis [152], Sauer [141], Shelah [143]). *Let
 1903 (X, \mathcal{R}) be a range space. If $|\mathcal{R}| > \sum_{i=0}^{k-1} \binom{|X|}{i}$ then the VC-dimension of the
 1904 range space is at least k . Conversely, if the VC-dimension of the range space
 1905 is v then $|\mathcal{R}| \leq \sum_{i=0}^{v-1} \binom{|X|}{i} = O(|X|^v)$.*

1906 In our examples, notice the discrepancy between the “identical” $O(|X|^2)$
 1907 bounds on the number of ranges and the different VC-dimensions (two and
 1908 three respectively).

1909 8.1.3 Epsilon Nets

1910 A net is a subset of the ground set such that it hits every large range.

1911 **Definition 10** (ε -net). Let (X, \mathcal{R}) be a range space and let ε be a real number
 1912 in the $[0, 1]$ interval. A subset $N \subseteq X$ is an ε -net for the range space if for
 1913 every $R \in \mathcal{R}$ such that $|R| > \varepsilon |X|$ we have that $N \cap R \neq \emptyset$.

1914 In our algorithms and data structures, we consider nets of range spaces
 1915 in \mathbb{R}^d . They are used as follows: construct an ε -net, partition the space
 1916 into a small number of ranges so that no range of the partition is hit by the
 1917 net. Then we have the guarantee that none of those ranges is large. For
 1918 many problems, such partitions of space with a small number of small ranges
 1919 usually lead to practical divide-and-conquer schemes. These partitions are
 1920 called *cuttings* (see §8.1.5).

1921 A very neat result is that, for constant VC-dimension, a reasonably small
 1922 ε -net can be constructed efficiently by random sampling.

1923 **Lemma 8.2.** *Let (X, \mathcal{R}) be a range space with VC-dimension v . Let N be
 1924 a uniform random sample of X of size $\Omega_v(r \log r)$. Then N is a $\frac{1}{r}$ -net for
 1925 (X, \mathcal{R}) with probability $1 - r^{-\Omega(1)}$.*

1926 See [52, Section 40.4] for more on range spaces and ε -nets.

1927 8.1.4 Hyperplanes in Linear Dimension

1928 Consider the range space where the ground set is a finite set of hyperplanes
 1929 \mathcal{H} in \mathbb{R}^n and where the ranges are the subsets of hyperplanes that can be
 1930 obtained by intersecting the ground set with a simplex.

1931 By combining a theorem due to Blumer et al. [29] with the results of
 1932 Meiser [113]¹, it is possible to obtain good bounds on ε -nets constructed
 1933 by random sampling. In Paper A we are interested in the dependency of
 1934 those bounds on the ambient dimension n . This is not explicitly tackled in
 1935 Lemma 8.2 since here the VC-dimension of this range space depends on n .

1936 **Lemma 8.3.** *For all real numbers $r > 1$ and $c \geq 1$, if we choose at least
 1937 $cn^2 \log nr \log r$ hyperplanes of \mathcal{H} uniformly at random and denote this selec-
 1938 tion \mathcal{N} then for any simplex intersected by more than $\frac{|\mathcal{H}|}{r}$ hyperplanes of
 1939 \mathcal{H} , with probability $1 - 2^{-\Omega(c)}$, at least one of the intersecting hyperplanes is
 1940 contained in \mathcal{N} .*

1941 The contrapositive states that if no hyperplane in \mathcal{N} intersects a given
 1942 simplex, then with high probability the number of hyperplanes of \mathcal{H} inter-
 1943 secting the simplex is at most $\frac{|\mathcal{H}|}{r}$.

1944 8.1.5 Cuttings

1945 A cutting in \mathbb{R}^d is a set of (possibly unbounded and/or non-full di-
 1946 mensional) bounded-complexity cells that together partition \mathbb{R}^d . For our
 1947 purposes, a cell is of bounded complexity if its boundary is defined by a
 1948 number of lines, pseudolines, or hyperplanes of some arrangement that solely
 1949 depends on the dimension, and not on the size of the arrangement. A
 1950 $\frac{1}{c}$ -cutting of a set of n elements is a cutting with the constraint that each of
 1951 its cells is intersected by at most $\frac{n}{c}$ elements.

1952 In constant dimension, a straightforward way to construct a $\frac{1}{c}$ -cutting for
 1953 a set of hyperplanes is to first construct a $\frac{1}{c}$ -net for the range space consisting
 1954 of hyperplanes as elements and simplices as ranges and then triangulate its
 1955 arrangement as in §6.4.1. Because none of the simplices of the triangulation
 1956 intersect the net, each simplex is intersected by at most a $\frac{1}{c}$ -fraction of the
 1957 hyperplanes. Through this construction we obtain $O(c^d \log^d c)$ simplicial
 1958 cells each intersected by at most $\frac{n}{c}$ hyperplanes.

1959 For hyperplanes in constant dimension, there exist various ways of con-
 1960 structing $\frac{1}{c}$ -cuttings of size $O(c^d)$, thereby removing the polylogarithmic
 1961 factor overhead (see for instance §8.2).

¹Note that Meiser used an older result due to Haussler and Welzl [97] and got an extra $\log n$ factor in the size of the ε -net. We thank Hervé Fournier for pointing this out.

1962 8.1.6 Algebraic Range Spaces

1963 In Paper B (§B.2.1) we use a similar technique to obtain a partition of
 1964 \mathbb{R}^2 into $O(c^2 \log^2 c)$ pseudotrapezoidal cells so that each cell is intersected by
 1965 at most a $\frac{1}{c}$ -fraction of some input polynomial curves. Because those curves
 1966 have bounded degree, we can design a range space with finite VC-dimension
 1967 and exploit it.

1968 In the same paper (§B.2.2) we use the technique of linearization [5, 157]
 1969 in order to tackle a more general problem in any constant dimension: given
 1970 a system of polynomial inequalities, we can define a single variable for each
 1971 monomial so that each inequality becomes linear. This has the effect of
 1972 greatly increasing the ambient dimension but allows to use the techniques
 1973 that normally only apply to linear ranges.

1974 8.1.7 Derandomization

1975 The ε -net construction based on sampling described above can be made
 1976 deterministic. Derandomization is usually achieved through the method of
 1977 conditional probabilities of Raghavan [129] and Spencer [145] (as in [48]).
 1978 This method yields a construction time that is linear in the size of the ground
 1979 set for ε -nets and ε -cuttings with constant ε .

1980 **Lemma 8.4** (Matoušek [110, Corollary 4.5]). *Let (X, \mathcal{R}) be a range space of
 1981 finite VC-dimension and let $r > 1$ be a constant. Under some reasonable
 1982 assumptions, we can compute a $\frac{1}{r}$ -net for (X, \mathcal{R}) of size $O(r \log r)$ in $O(|X|)$
 1983 time.*

1984 For other concrete examples of derandomization of nets and cuttings,
 1985 we refer the reader to Matoušek [111], Chazelle and Matoušek [51] and
 1986 Brönnimann et al. [34]. The Handbook has a nice overview of the subject [52,
 1987 Section 40.7]. See also [52, Section 40.1] for more on randomized divide-and-
 1988 conquer and [52, Section 40.6] for more on derandomization.

1989 8.2 Hierarchical Cuttings

1990 Compared to standard cuttings, hierarchical cuttings have the additional
 1991 property that they can be composed without multiplying the hidden constant
 1992 factors in the big-O notation [48]. In particular, they allow for $O(n^d)$ -space

1993 $O(\log n)$ -query d -dimensional point location data structures (for constant d).
1994 This is exploited in Paper B (§B.2.2) and Paper C (all sections).

1995 **Definition 11** (Hierarchical Cutting). Given n hyperplanes in \mathbb{R}^d , a ℓ -levels
1996 hierarchical cutting of parameter $r > 1$ for those hyperplanes is a sequence of
1997 ℓ levels labeled $0, 1, \dots, \ell - 1$ such that²

- 1998 • Level i has $O(r^{2i})$ cells,
- 1999 • Each of those cells is further partitioned into $O(r^2)$ subcells,
- 2000 • The collection of subcells is a $\frac{1}{r^{i+1}}$ -cutting for the set of hyperplanes,
- 2001 • The $O(r^{2(i+1)})$ subcells of level i are the cells of level $i + 1$.

2002 It is clear from reading through the various references that those hier-
2003 archical cuttings can be constructed for arrangements of pseudolines with
2004 the same properties: In [48], Chazelle first proves a vertex-count estimation
2005 lemma that only relies on incidence properties of line arrangements [48,
2006 Lemma 2.1]. Then he summons a lemma from [109] that relies on the finite
2007 VC-dimension of the range space at hand [48, Lemma 3.1]. Here the ground
2008 set is the set of pseudolines and the ranges are the subsets induced by
2009 intersections with pseudosegments. The VC-dimension of this range space is
2010 easily shown to be finite and is known to be at most 8: every arrangement of
2011 9 pseudolines contains a subset of 6 pseudolines in hexagonal formation [94],
2012 which cannot be shattered.³ Finally, he proves a lemma for the efficient
2013 construction of sparse ε -nets whose correctness again only relies on incidence
2014 properties of line arrangements [48, Lemma 3.2]. Using those three lemmas
2015 together with bottom vertex triangulation (§6.4.1) he is able to prove his
2016 main result:

2017 **Lemma 8.5** (Chazelle [48, Theorem 3.3]). *Given n hyperplanes in \mathbb{R}^d , for
2018 any real parameter $r > 1$, we can construct a ℓ -levels hierarchical cutting of
2019 parameter r for those hyperplanes in time $O(nr^{\ell(d-1)})$.*

2020 For pseudoline arrangements, bottom vertex triangulation can be traded
2021 for vertical decomposition (§6.4.2).

²In [48], Chazelle refers to this parameter as r_0 and uses r to mean r_0^ℓ . Here we drop the subscript for ease of presentation.

³This a quote from [32]. We could not access the original paper.

2022 **Lemma 8.6.** *Given n pseudolines, for any real parameter $r > 1$, we can
2023 construct a ℓ -levels hierarchical cutting of parameter r for those pseudolines
2024 in time $O(nr^\ell)$.*

2025 In particular, we will use those lemmas with $\ell = \lceil \log_r \frac{n}{t} \rceil$, for some
2026 parameter t , so that the last level of the hierarchy defines a $\frac{t}{n}$ -cutting whose
2027 cells are each intersected by at most t pseudolines (or hyperplanes).

2028 Note that in [48] the construction is described for constant parameter
2029 r . This restriction on the parameter is easily lifted: We can construct
2030 a hierarchical cutting with superconstant parameter r by constructing a
2031 hierarchical cutting with some appropriate constant parameter r' , and then
2032 skip levels that we do not need. This is used in §C.2 to reduce the query
2033 time from $O(\log n)$ to $O(\frac{\log n}{\log \log n})$.

9

2034

2035 Existential Theory of the Reals

2036 The problems we consider in Paper B require our algorithms to manip-
2037 ulate polynomial expressions and, potentially, their real roots. For that
2038 purpose, we rely on Collins’s cylindrical algebraic decomposition (CAD) [54].
2039 To understand the power of this method, and why it is useful for us, we give
2040 some background on the related concept of first-order theory of the reals.

2041 **Definition 12.** A Tarski formula $\phi \in \mathbb{T}$ is a grammatically correct formula
2042 consisting of real variables ($x \in \mathbb{R}$), universal and existential quantifiers on
2043 those real variables ($\forall, \exists: \mathbb{R} \times \mathbb{T} \rightarrow \mathbb{T}$), the boolean operators of conjunction
2044 and disjunction ($\wedge, \vee: \mathbb{T}^2 \rightarrow \mathbb{T}$), the six comparison operators ($<, \leq, =, \geq, >$
2045 , $\neq: \mathbb{R}^2 \rightarrow \mathbb{T}$), the four arithmetic operators ($+, -, *, /: \mathbb{R}^2 \rightarrow \mathbb{R}$), the usual
2046 parentheses that modify the priority of operators, and constant real numbers
2047 (\mathbb{R}). A Tarski sentence is a fully quantified Tarski formula. The first-order
2048 theory of the reals (FOTR) is the set of true Tarski sentences.

2049 Tarski [150] and Seidenberg [142] proved that FOTR is decidable. How-
2050 ever, the algorithm resulting from their proof has nonelementary complexity.
2051 This proof, as well as other known algorithms, are based on quantifier
2052 elimination, that is, the translation of the input formula to a much longer
2053 quantifier-free formula, whose validity can be checked. There exists a family
2054 of formulas for which any method of quantifier elimination produces a doubly
2055 exponential size quantifier-free formula [57].

2056 9.1 Cylindrical Algebraic Decomposition

2057 Collins’s CAD matches this doubly exponential complexity.

2058 **Theorem 9.1** (Collins [54]). FOTR can be solved in $2^{2^{O(n)}}$ time in the
2059 real-RAM model, where n is the size of the input Tarski sentence.

2060 See Basu, Pollack, and Roy [23] for additional details, Basu, Pollack,
2061 and Roy [22] for a singly exponential algorithm when all quantifiers are
2062 existential (existential theory of the reals, $\exists\mathbb{R}$), Caviness and Johnson [43]
2063 for an anthology of key papers on the subject, and Mishra [116] for a review
2064 of techniques to compute with roots of polynomials.

2065 Collins's CAD solves any *geometric* decision problem that does not
2066 involve quantification over the integers in time doubly exponential in the
2067 problem size. This does not harm our results as we exclusively use this
2068 algorithm to solve constant size subproblems. Geometric is to be understood
2069 in the sense of Descartes and Fermat, that is, the geometry of objects that
2070 can be expressed with polynomial equations. In particular, it allows us to
2071 make the following computations in the real-RAM and bounded-degree ADT
2072 models:

- 2073 1. Given a constant-degree univariate polynomial, count its real roots in
2074 $O(1)$ operations,
- 2075 2. Sort $O(1)$ real numbers given implicitly as roots of some constant-degree
2076 univariate polynomials in $O(1)$ operations,
- 2077 3. Given a point in the plane and an arrangement of a constant number
2078 of constant-degree polynomial planar curves, locate the point in the
2079 arrangement in $O(1)$ operations.

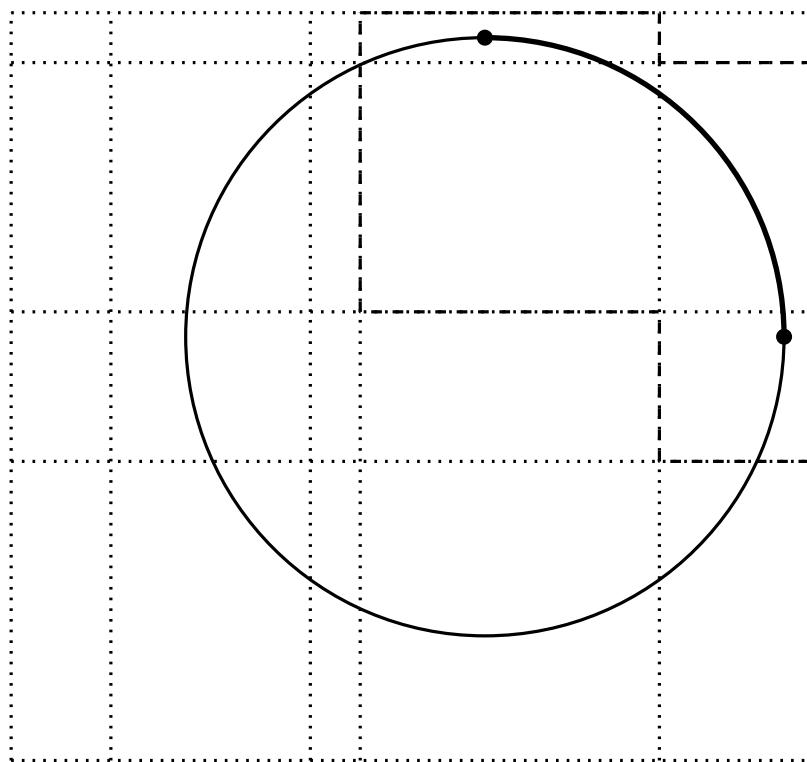
2080 Instead of bounded-degree algebraic decision trees as the nonuniform
2081 model we could consider decision trees in which each decision involves a
2082 constant-size instance of the decision problem in the first-order theory of
2083 the reals. The depth of a bounded-degree algebraic decision tree simulating
2084 such a tree would only be blown up by a constant factor.

III

2085

Algorithms

2086



A

2087

2088 Solving k -SUM using Few Linear 2089 Queries

2090

with Jean Cardinal and John Iacono

2091 The k -SUM problem is given n input real numbers to determine whether
2092 any k of them sum to zero. The problem is of tremendous importance in
2093 the emerging field of complexity theory within P , and it is in particular
2094 open whether it admits an algorithm of complexity $O(n^c)$ with $c < \lceil \frac{k}{2} \rceil$.
2095 Inspired by an algorithm due to Meiser (1993), we show that there exist linear
2096 decision trees and algebraic computation trees of depth $O(n^3 \log^2 n)$ solving
2097 k -SUM. Furthermore, we show that there exists a randomized algorithm
2098 that runs in $\tilde{O}(n^{\lceil \frac{k}{2} \rceil + 8})$ time, and performs $O(n^3 \log^2 n)$ linear queries on
2099 the input. Thus, we show that it is possible to have an algorithm with a
2100 runtime almost identical (up to the +8) to the best known algorithm but
2101 for the first time also with the number of queries on the input a polynomial
2102 that is independent of k . The $O(n^3 \log^2 n)$ bound on the number of linear
2103 queries is also a tighter bound than any known algorithm solving k -SUM,
2104 even allowing unlimited total time outside of the queries. By simultaneously
2105 achieving few queries to the input without significantly sacrificing runtime
2106 vis-à-vis known algorithms, we deepen the understanding of this canonical
2107 problem which is a cornerstone of complexity-within- P .

2108 We also consider a range of tradeoffs between the number of terms
2109 involved in the queries and the depth of the decision tree. In particular,
2110 we prove that there exist $o(n)$ -linear decision trees of depth $\tilde{O}(n^3)$ for the
2111 k -SUM problem.

2112 A.1 Meiser Solves k -SUM

2113 Our main contribution is an efficient implementation of an existing
 2114 algorithm by Meiser [113] when applied to the k -SUM problem.

2115 This section is divided into three subsections: §A.1.1 gives the outline of
 2116 the algorithm and analyses its complexity in the linear decision tree model.
 2117 §A.1.2 gives a second slightly more complex implementation and analysis of
 2118 this algorithm in the word-RAM model. §A.1.3 gives a simple tweak that
 2119 can be applied to those algorithms in order to reduce the complexity of the
 2120 queries involved.

2121 Missing details are found in §A.2.

2122 A.1.1 Query Complexity

2123 In this section and the next, we prove the following first result.

2124 **Contribution 1.** *There exist linear decision trees of depth $O(n^3 \log^2 n)$
 2125 solving the k -SUM and the k -LDT problems. Furthermore, for the k -SUM
 2126 problem there exists an $\tilde{O}(n^{\lceil \frac{k}{2} \rceil + 8})$ Las Vegas algorithm in the word-RAM
 2127 model expected to perform $O(n^3 \log^2 n)$ linear queries on the input. This
 2128 algorithm also solves the k -LDT problem within the same time bounds pro-
 2129 vided the coefficients of the underlying linear equation are constant rational
 2130 numbers.*

2131 Algorithm outline

2132 For a fixed set of hyperplanes \mathcal{H} and given input vertex q in \mathbb{R}^n , Meiser's
 2133 algorithm allows us to determine the cell of the arrangement $\mathcal{A}(\mathcal{H})$ that
 2134 contains q in its interior (or that *is* q if q is a 0-cell of $\mathcal{A}(\mathcal{H})$), that is, the
 2135 positions $\sigma(H, q) \in \{-, 0, +\}$ of q with respect to all hyperplanes $H \in \mathcal{H}$.
 2136 In the k -SUM problem, the set \mathcal{H} is the set of $\Theta(n^k)$ hyperplanes with
 2137 equations of the form $x_{i_1} + x_{i_2} + \cdots + x_{i_k} = 0$. These equations can be
 2138 modified accordingly for k -LDT.

2139 We use Theorem 8.3 to design a prune and search algorithm for k -SUM as
 2140 follows: (1) construct an ε -net \mathcal{N} , (2) compute the cell C of $\mathcal{A}(\mathcal{N})$ containing
 2141 the input point q in its interior, (3) construct a simplex S inscribed in C and
 2142 containing q in its interior, (4) recurse on the hyperplanes of \mathcal{H} intersecting
 2143 the interior of S .

2144 Proceeding this way with a constant ε guarantees that at most a constant
 2145 fraction ε of the hyperplanes remains after the pruning step, and thus the
 2146 cumulative number of queries made to determine the enclosing cell at each
 2147 step is $O(n^2 \log n \log |\mathcal{H}|)$ when done in a brute-force way. However, we still
 2148 need to explain how to find a simplex S inscribed in C and containing q in
 2149 its interior. This procedure corresponds to the well-known *bottom vertex*
 2150 *decomposition* (or *triangulation*) of a hyperplane arrangement (see §6.4.1).

2151 **Finding a simplex**

2152 In order to simplify the exposition of the algorithm, we assume, without
 2153 loss of generality, that the input numbers q_i all lie in the interval $[-1, 1]$.
 2154 This assumption is justified by observing that we can normalize all the input
 2155 numbers by the largest absolute value of a component of q . One can then
 2156 see that every linear query on the normalized input can be implemented as
 2157 a linear query on the original input. A similar transformation can be carried
 2158 out for the k -LDT problem. This allows us to use bounding hyperplanes of
 2159 equations $x_i = \pm 1, i \in [n]$. We denote by \mathcal{B} this set of hyperplanes. Hence,
 2160 if we choose a subset \mathcal{N} of the hyperplanes, the input point is located in a
 2161 bounded cell of the arrangement $\mathcal{A}(\mathcal{N} \cup \mathcal{B})$. Note that $|\mathcal{N} \cup \mathcal{B}| = O(|\mathcal{N}|)$
 2162 for all interesting values of ε .

2163 We now explain how to construct S under this assumption. The algorithm
 2164 can be sketched as follows. (Recall that $\sigma(H, p)$ denotes the relative position
 2165 of p with respect to the hyperplane H .)

2166 **Algorithm 1** (Constructing S).

input A point q in $[-1, 1]^n$, a set \mathcal{I} of hyperplanes not containing q , and
 a set \mathcal{E} of hyperplanes in general position containing q , such that the
 cell

$$C = \{p: \sigma(H, p) = \sigma(H, q) \text{ or } \sigma(H, p) = 0 \text{ for all } H \in (\mathcal{I} \cup \mathcal{E})\}$$

2167 is a bounded polytope. The value $\sigma(H, q)$ is known for all $H \in (\mathcal{I} \cup \mathcal{E})$.

2168 **output** A simplex $S \in C$ that contains q in its interior (if it is not a
 2169 point), and all vertices of which are vertices of C .

- 2170 **0.** If $|\mathcal{E}| = n$, return q .
 2171 **1.** Determine a vertex ν of C .

- 2172 2. Let q' be the projection of q along $\vec{\nu}q$ on the boundary of C . Compute
 2173 $\mathcal{I}_\theta \subseteq \mathcal{I}$, the subset of hyperplanes in \mathcal{I} containing q' . Compute
 2174 $\mathcal{I}_\tau \subseteq \mathcal{I}_\theta$, a maximal subset of those hyperplanes such that $\mathcal{E}' = \mathcal{E} \cup \mathcal{I}_\tau$
 2175 is a set of hyperplanes in general position.
- 2176 3. Recurse on q' , $\mathcal{I}' = \mathcal{I} \setminus \mathcal{I}_\theta$, and \mathcal{E}' , and store the result in S' .
- 2177 4. Return S , the convex hull of $S' \cup \{\nu\}$.

2178 Step **0** is the base case of the recursion: when there is only one point
 2179 left, just return that point. This step uses no query.

2180 We can solve step **1** by using linear programming with the known values
 2181 of $\sigma(H, q)$ as linear constraints. We arbitrarily choose an objective function
 2182 with a gradient non-orthogonal to all hyperplanes in \mathcal{I} and look for the
 2183 optimal solution. The optimal solution being a vertex of the arrangement,
 2184 its coordinates are independent of q , and thus this step involves no query at
 2185 all.

2186 Step **2** prepares the recursive step by finding the hyperplanes containing
 2187 q' . This can be implemented as a ray-shooting algorithm that performs a
 2188 number of comparisons between projections of q on different hyperplanes
 2189 of \mathcal{I} without explicitly computing them. In §A.2.1, we prove that all such
 2190 comparisons can be implemented using $O(|\mathcal{I}|)$ linear queries. Constructing
 2191 \mathcal{E}' can be done by solving systems of linear equations that do not involve q .

2192 In step **3**, the input conditions are satisfied, that is, $q' \in [-1, 1]^n$, \mathcal{I}' is
 2193 a set of hyperplanes not containing q' , \mathcal{E}' is a set of hyperplanes in general
 2194 position containing q' , C' is a d -cell of C and is thus a bounded polytope.
 2195 The value $\sigma(H, q')$ differs from $\sigma(H, q)$ only for hyperplanes that have been
 2196 removed from \mathcal{I} , and for those $\sigma(H, q') = 0$, hence we know all necessary
 2197 values $\sigma(H, q')$ in advance.

2198 Since $|\mathcal{I}'| < |\mathcal{I}|$, $|\mathcal{E}'| > |\mathcal{E}|$, and $|\mathcal{I} \setminus \mathcal{I}'| - |\mathcal{E}' \setminus \mathcal{E}| \geq 0$, the complexity of
 2199 the recursive call is no more than that of the parent call, and the maximal
 2200 depth of the recursion is n . Thus, the total number of linear queries made
 2201 to compute S is $O(n|\mathcal{I}|)$.

2202 Hence given an input point $q \in [-1, 1]$, an arrangement of hyperplanes
 2203 $\mathcal{A}(\mathcal{N})$, and the value of $\sigma(H, q)$ for all $H \in (\mathcal{N} \cup \mathcal{B})$, we can compute
 2204 the desired simplex S by running Algorithm 1 on q , $\mathcal{I} = \{H \in (\mathcal{N} \cup$
 2205 $\mathcal{B}): \sigma(H, q) \neq 0\}$, and $\mathcal{E} \subseteq (\mathcal{N} \cup \mathcal{B}) \setminus \mathcal{I}$. This uses $O(n^3 \log n)$ linear queries.
 2206 Figure A.1 illustrates a step of the algorithm.

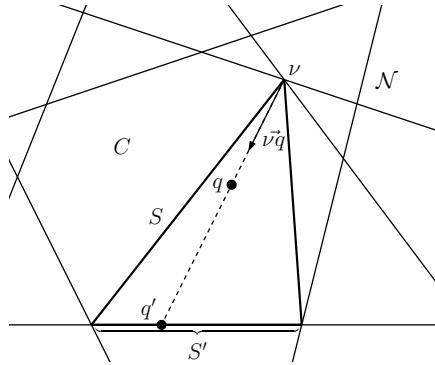


Figure A.1. Illustration of a step of Algorithm 1.

2207 **Assembling the pieces**

2208 Let us summarize the algorithm

2209 **Algorithm 2.**

2210 **input** $q \in [-1, 1]^n$

2211 1. Pick $O(n^2 \log n)$ hyperplanes of \mathcal{H} at random and locate q in this
2212 arrangement. Call C the cell containing q .

2213 2. Construct the simplex S containing q and inscribed in C , using Algo-
2214 rithm 1.

2215 3. For every hyperplane of \mathcal{H} containing S , output a solution.

2216 4. Recurse on hyperplanes of \mathcal{H} intersecting the interior of S .

2217 The query complexity of step 1 is $O(n^2 \log n)$, and that of step 2 is
2218 $O(n^3 \log n)$. Steps 3 and 4 do not involve any query at all. The recursion
2219 depth is $O(\log |\mathcal{H}|)$, with $|\mathcal{H}| = O(n^k)$, hence the total query complexity of
2220 this algorithm is $O(n^3 \log^2 n)$. This proves the first part of Theorem 1.

2221

2222 We can also consider the overall complexity of the algorithm in the RAM
2223 model, that is, taking into account the steps that do not require any query,
2224 but for which we still have to process the set \mathcal{H} . Note that the complexity
2225 bottleneck of the algorithm are steps 3-4, where we need to prune the list of
2226 hyperplanes according to their relative positions with respect to S . For this
2227 purpose, we simply maintain explicitly the list of all hyperplanes, starting

with the initial set corresponding to all k -tuples. Then the pruning step can be performed by looking at the position of each vertex of S relative to each hyperplane of \mathcal{H} . Because in our case hyperplanes have only k nonzero coefficients, this uses a number of integer arithmetic operations on $\tilde{O}(n)$ bits integers that is proportional to the number of vertices times the number of hyperplanes. (For the justification of the bound on the number of bits needed to represent vertices of the arrangement see §A.2.4.) Since we recurse on a fraction of the set, the overall complexity is $\tilde{O}(n^2|\mathcal{H}|) = \tilde{O}(n^{k+2})$. The next section is devoted to improving this running time.

A.1.2 Time Complexity

Proving the second part of Theorem 1 involves efficient implementations of the two most time-consuming steps of Algorithm 2. In order to efficiently implement the pruning step, we define an intermediate problem, that we call the *double k -SUM* problem.

Problem 15 (double k -SUM). Given two vectors $\nu_1, \nu_2 \in [-1, 1]^n$, where the coordinates of ν_i can be written down as fractions whose numerator and denominator lie in the interval $[-M, M]$, enumerate all $i \in [n]^k$ such that

$$\left(\sum_{j=1}^k \nu_{1,i_j} \right) \left(\sum_{j=1}^k \nu_{2,i_j} \right) < 0.$$

In other words, we wish to list all hyperplanes of \mathcal{H} intersecting the open line segment $\nu_1\nu_2$. We give an efficient output-sensitive algorithm for this problem.

Lemma A.1. *Double k -SUM can be solved in $O(n^{\lceil \frac{k}{2} \rceil} \log n \log M + Z)$ time in the word-RAM model, where Z is the size of the solution.*

Proof. If k is even, we consider all possible $\frac{k}{2}$ -tuples of numbers in ν_1 and ν_2 and sort their sums in increasing order. This takes time $O(n^{\frac{k}{2}} \log n)$ and yields two permutations π_1 and π_2 of $[n^{\frac{k}{2}}]$. If k is odd, then we sort both the $\lceil \frac{k}{2} \rceil$ -tuples and the $\lfloor \frac{k}{2} \rfloor$ -tuples. For simplicity, we will only consider the even case in what follows. The odd case carries through.

We let $N = n^{\frac{k}{2}}$. For $i \in [N]$ and $m \in \{1, 2\}$, let $\Sigma_{m,i}$ be the sum of the $\frac{k}{2}$ components of the i th $\frac{k}{2}$ -tuple in ν_m , in the order prescribed by π_m .

2254 We now consider the two $N \times N$ matrices M_1 and M_2 giving all possible
 2255 sums of two $\frac{k}{2}$ -tuples, for both ν_1 with the ordering π_1 and ν_2 with the
 2256 ordering π_2 .

2257 We first solve the k -SUM problem on ν_1 , by finding the sign of all pairs
 2258 $\Sigma_{1,i} + \Sigma_{1,j}$, $i, j \in [N]$. This can be done in time $O(N)$ by parsing the matrix
 2259 M_1 , just as in the standard k -SUM algorithm. We do the same with M_2 .

2260 The set of all indices $i, j \in [N]$ such that $\Sigma_{1,i} + \Sigma_{1,j}$ is positive forms
 2261 a staircase in M_1 . We sweep M_1 column by column in order of increasing
 2262 $j \in [N]$, in such a way that the number of indices i such that $\Sigma_{1,i} + \Sigma_{1,j} > 0$
 2263 is growing. For each new such value i that is encountered during the sweep,
 2264 we insert the corresponding $i' = \pi_2(\pi_1^{-1}(i))$ in a balanced binary search
 2265 tree.

2266 After each sweep step in M_1 — that is, after incrementing j and adding
 2267 the set of new indices i' in the tree — we search the tree to identify all the
 2268 indices i' such that $\Sigma_{2,i'} + \Sigma_{2,j'} < 0$, where $j' = \pi_2(\pi_1^{-1}(j))$. Since those
 2269 indices form an interval in the ordering π_2 when restricted to the indices
 2270 in the tree, we can search for the largest i'_0 such that $\Sigma_{2,i'_0} < -\Sigma_{2,j'}$ and
 2271 retain all indices $i' \leq i'_0$ that are in the tree. If we denote by z the number
 2272 of such indices, this can be done in $O(\log N + z) = O(\log n + z)$ time. Now
 2273 all the pairs i', j' found in this way are such that $\Sigma_{1,i} + \Sigma_{1,j}$ is positive and
 2274 $\Sigma_{2,i'} + \Sigma_{2,j'}$ is negative, hence we can output the corresponding k -tuples.
 2275 To get all the pairs i', j' such that $\Sigma_{1,i} + \Sigma_{1,j}$ is negative and $\Sigma_{2,i'} + \Sigma_{2,j'}$
 2276 positive, we repeat the sweeping algorithm after swapping the roles of ν_1
 2277 and ν_2 .

2278 Every matching k -tuple is output exactly once, and every $\frac{k}{2}$ -tuple is
 2279 inserted at most once in the binary search tree. Hence the algorithm runs
 2280 in the claimed time.

2281 Note that we only manipulate rational numbers that are the sum of at
 2282 most k rational numbers of size $O(\log M)$. □

2283 Now observe that a hyperplane intersects the interior of a simplex if and
 2284 only if it intersects the interior of one of its edges. Hence given a simplex S
 2285 we can find all hyperplanes of \mathcal{H} intersecting its interior by running the above
 2286 algorithm $\binom{n}{2}$ times, once for each pair of vertices (ν_1, ν_2) of S , and take the
 2287 union of the solutions. The overall running time for this implementation
 2288 will therefore be $\tilde{O}(n^2(n^{\lceil \frac{k}{2} \rceil} \log M + Z))$, where Z is at most the number of

intersecting hyperplanes and M is to be determined later. This provides an implementation of the pruning step in Meiser's algorithm, that is, step 4 of Algorithm 2.

Corollary A.2. *In the word-RAM model, given a simplex S , we can compute the Z k -SUM hyperplanes intersecting its interior in $\tilde{O}(n^2(n^{\lceil \frac{k}{2} \rceil} \log M + Z))$ time, where $\log M$ is proportional to the number of bits necessary to represent S .*

In order to detect solutions in step 3 of Algorithm 2, we also need to be able to quickly solve the following problem.

Problem 16 (multiple k -SUM). Given d points $\nu_1, \nu_2, \dots, \nu_d \in \mathbb{R}^n$, where the coordinates of ν_i can be written down as fractions whose numerator and denominator lie in the interval $[-M, M]$, decide whether there exists a hyperplane with equation of the form $x_{i_1} + x_{i_2} + \dots + x_{i_k} = 0$ containing all of them.

Here the standard k -SUM algorithm can be applied, taking advantage of the fact that the coordinates lie in a small discrete set.

Lemma A.3. *In the word-RAM model, k -SUM on n integers $\in [-V, V]$ can be solved in time $\tilde{O}(n^{\lceil \frac{k}{2} \rceil} \log V)$.*

Lemma A.4. *In the word-RAM model, multiple k -SUM can be solved in time $\tilde{O}(dn^{\lceil \frac{k}{2} \rceil + 2} \log M)$.*

Proof. Let $\mu_{i,j}$ and $\delta_{i,j}$ be the numerator and denominator of $\nu_{i,j}$ when written as an irreducible fraction. We define

$$\zeta_{i,j} = \nu_{i,j} \prod_{(i,j) \in [d] \times [n]} \delta_{i,j} = \frac{\mu_{i,j} \prod_{(i',j') \in [d] \times [n]} \delta_{i',j'}}{\delta_{i,j}}.$$

By definition $\zeta_{i,j}$ is an integer and its absolute value is bounded by $U = M^{n^2}$, that is, it can be represented using $O(n^2 \log M)$ bits. Moreover, if one of the hyperplanes contains the point $(\zeta_{i,1}, \zeta_{i,2}, \dots, \zeta_{i,n})$, then it contains ν_i . Construct n integers of $O(dn^2 \log M)$ bits that can be written $\zeta_{1,j} + U, \zeta_{2,j} + U, \dots, \zeta_{d,j} + U$ in base $2Uk+1$. The answer to our decision problem is “yes” if and only if there exists k of those numbers whose sum is kU, kU, \dots, kU .

2315 We simply subtract the number U, U, \dots, U to all n input numbers to obtain
2316 a standard k -SUM instance on n integers of $O(dn^2 \log M)$ bits. \square

2317 We now have efficient implementations of steps **3** and **4** of Algorithm 2
2318 and can proceed to the proof of the second part of Theorem 1.

2319 *Proof.* The main idea consists in modifying the first iteration of Algo-
2320 rithm 2, by letting $\varepsilon = \Theta(n^{-\frac{k}{2}})$. Hence we pick a random subset \mathcal{N} of
2321 $O(n^{k/2+2} \log n)$ hyperplanes in \mathcal{H} and use this as an ε -net. This can be
2322 done efficiently, as shown in §A.2.3.

Next, we need to locate the input q in the arrangement induced by \mathcal{N} . This can be done by running Algorithm 2 on the set \mathcal{N} . From the previous considerations on Algorithm 2, the running time of this step is

$$O(n|\mathcal{N}|) = \tilde{O}(n^{k/2+4}),$$

2323 and the number of queries is $O(n^3 \log^2 n)$.

2324 Then, in order to prune the hyperplanes in \mathcal{H} , we have to compute a
2325 simplex S that does not intersect any hyperplane of \mathcal{N} . For this, we ob-
2326 serve that the above call to Algorithm 2 involves computing a sequence of
2327 simplices for the successive pruning steps. We save the description of those
2328 simplices. Recall that there are $O(\log n)$ of them, all of them contain the
2329 input q and have vertices coinciding with vertices of the original arrange-
2330 ment $\mathcal{A}(\mathcal{H})$. In order to compute a simplex S avoiding all hyperplanes of
2331 \mathcal{N} , we can simply apply Algorithm 1 on the set of hyperplanes bounding
2332 the intersection of these simplices. The running time and number of queries
2333 for this step are bounded respectively by $n^{O(1)}$ and $O(n^2 \log n)$.

2334 Note that the vertices of S are not vertices of $\mathcal{A}(\mathcal{H})$ anymore. However,
2335 their coordinates lie in a finite set (see §A.2.4)

2336 **Lemma A.5.** *Vertices of S have rational coordinates whose fraction repre-
2337 sentations have their numerators and denominators absolute values bounded
2338 above by $C^{4n^5} n^{2n^5+n^3+\frac{n}{2}}$, where C is a constant.*

2339 We now are in position to perform the pruning of the hyperplanes in \mathcal{H}
2340 with respect to S . The number of remaining hyperplanes after the pruning
2341 is at most $\varepsilon n^k = O(n^{k/2})$. Hence from Corollary A.2, the pruning can be
2342 performed in time proportional to $\tilde{O}(n^{\lceil k/2 \rceil + 7})$.

2343 Similarly, we can detect any hyperplane of \mathcal{H} containing S using the
 2344 result of Lemma A.4 in time $\tilde{O}(n^{\lceil k/2 \rceil + 8})$. Note that those last two steps
 2345 do not require any query.

2346 Finally, it remains to detect any solution that may lie in the remaining
 2347 set of hyperplanes of size $O(n^{k/2})$. We can again fall back on Algorithm 2,
 2348 restricted to those hyperplanes. The running time is $\tilde{O}(n^{k/2+2})$, and the
 2349 number of queries is still $O(n^3 \log^2 n)$.

2350 Overall, the maximum running time of a step is $\tilde{O}(n^{\lceil \frac{k}{2} \rceil + 8})$, while the
 2351 number of queries is always bounded by $O(n^3 \log^2 n)$.

2352

□

2353 A.1.3 Query Size

2354 In this section, we consider a simple blocking scheme that allows us to
 2355 explore a tradeoff between the number of queries and the size of the queries.

2356 **Lemma A.6.** *For any integer $b > 0$, an instance of the k -SUM problem on
 2357 $n > b$ numbers can be split into $O(b^{k-1})$ instances on at most $k\lceil \frac{n}{b} \rceil$ numbers,
 2358 so that every k -tuple forming a solution is found in exactly one of the
 2359 subproblems. The transformation can be carried out in time $O(n \log n + b^{k-1})$.*

2360 *Proof.* Given an instance on n numbers, we can sort them in $O(n \log n)$
 2361 time. Partition the sorted sequence into b consecutive blocks B_1, B_2, \dots, B_b
 2362 of equal size. This partition can be associated with a partition of the real
 2363 line into b intervals, say I_1, I_2, \dots, I_b . Now consider the partition of \mathbb{R}^k
 2364 into grid cells defined by the k th power of the partition I_1, I_2, \dots, I_b . The
 2365 hyperplane of equation $x_1 + x_2 + \dots + x_k = 0$ hits $O(b^{k-1})$ such grid cells.
 2366 Each grid cell $I_{i_1} \times I_{i_2} \times \dots \times I_{i_k}$ corresponds to a k -SUM problem on the
 2367 numbers in the set $B_{i_1} \cup B_{i_2} \cup \dots \cup B_{i_k}$ (note that the indices i_j need not
 2368 be distinct). Hence each such instance has size at most $k\lceil \frac{n}{b} \rceil$. □

2369 Combining Lemma A.6 and Theorem 1 directly yields our second contribu-
 2370 tion.

2371 **Contribution 2.** *For any integer $b > 0$, there exists a $k\lceil \frac{n}{b} \rceil$ -linear decision
 2372 tree of depth $\tilde{O}(b^{k-4}n^3)$ solving the k -SUM problem. Moreover, this deci-
 2373 sion tree can be implemented as an $\tilde{O}(b^{\lfloor \frac{k}{2} \rfloor - 9}n^{\lceil \frac{k}{2} \rceil + 8})$ Las Vegas word-RAM
 2374 algorithm.*

2375 The following two corollaries are obtained by taking $b = \Theta(\text{polylog}(n))$,
 2376 and $b = \Theta(n^\alpha)$, respectively

2377 **Corollary A.7.** *There exists a $o(n)$ -linear decision tree of depth $\tilde{O}(n^3)$*
 2378 *solving the k -SUM problem. This decision tree can be implemented as a*
 2379 *$\tilde{O}(n^{\lceil \frac{k}{2} \rceil} + 8)$ Las Vegas word-RAM algorithm.*

2380 **Corollary A.8.** *For any α such that $0 < \alpha < 1$, there exists a $O(n^{1-\alpha})$ -*
 2381 *linear decision tree of depth $\tilde{O}(n^{3+(k-4)\alpha})$ solving the k -SUM problem. This*
 2382 *decision tree can be implemented as a $\tilde{O}(n^{(1+\alpha)\frac{k}{2}+8.5})$ Las Vegas word-RAM*
 2383 *algorithm.*

2384 Note that the latter query complexity improves on $\tilde{O}(n^{\frac{k}{2}})$ whenever
 2385 $\alpha < \frac{k-6}{2k-8}$ and $k \geq 7$. By choosing $\alpha = \frac{k-6}{2k-8} - \frac{\beta}{k-4}$ we obtain $O(n^{1-\frac{k-6}{2k-8} + \frac{\beta}{k-4}})$ -
 2386 linear decision trees of depth $\tilde{O}(n^{\frac{k}{2}-\beta})$ for any $k \geq 7$. Hence for instance,
 2387 we obtain $O(n^{\frac{3}{4}+\frac{\beta}{4}})$ -linear decision trees of depth $\tilde{O}(n^{4-\beta})$ for the 8SUM
 2388 problem.

2389 A.2 Missing Details

2390 This section regroups the missing details of the algorithms described in
 2391 the previous one (§A.1).

2392 In §A.2.1 we show how to keep the queries linear even though it is not
 2393 evident at first sight that they are because of all the algebra generated
 2394 by the recursive steps of the simplex construction. In §A.2.2 we show
 2395 that our linear decision tree can be implemented in the same time on the
 2396 algebraic computation tree model, even though some linear queries may
 2397 involve all members of the input, and hence have superconstant cost. In
 2398 §A.2.3 we explain how to efficiently implement uniform random sampling for
 2399 the construction of large ε -nets. In §A.2.4 we bound the size of the numbers
 2400 manipulated by the word-RAM implementation of our k -SUM algorithm.

2401 A.2.1 Keeping Queries Linear in Algorithm 1

2402 In Algorithm 1, we want to ensure that the queries we make in step 2
 2403 are linear and that the queries we will make in the recursion step remain
 2404 linear too.

2405 **Lemma A.9.** *Algorithm 1 can be implemented so that it uses $O(n|I|)$ linear*
 2406 *queries.*

2407 *Proof.* Let us first analyze what the queries of step **2** look like. In addition to
2408 the input point q we are given a vertex ν and we want to find the projection
2409 q' of q in direction $\nu\vec{q}$ on the hyperplanes of \mathcal{I}_θ . Let the equation of H_i be
2410 $\Pi_i(x) = c_i + d_i \cdot x = 0$ where c_i is a scalar and d_i is a vector. The projection
2411 of q along $\nu\vec{q}$ on a hyperplane H_i can thus be written¹ $\rho(q, \nu, H_i) = \nu + \lambda_i \nu\vec{q}$
2412 such that $\Pi_i(\nu + \lambda_i \nu\vec{q}) = c_i + d_i \cdot \nu + \lambda_i d_i \cdot \nu\vec{q} = 0$. Computing the closest
2413 hyperplane amounts to finding $\lambda_\theta = \min_{\lambda_i > 0} \lambda_i$. Since $\lambda_i = -\frac{c_i + d_i \cdot \nu}{d_i \cdot \nu\vec{q}}$ we
2414 can test whether $\lambda_i > 0$ using the linear query² $-\frac{d_i \cdot \nu\vec{q}}{c_i + d_i \cdot \nu} > ? 0$. Moreover,
2415 if $\lambda_i > 0$ and $\lambda_j > 0$ we can test whether $\lambda_i < \lambda_j$ using the linear query
2416 $\frac{d_i \cdot \nu\vec{q}}{c_i + d_i \cdot \nu} < ? \frac{d_j \cdot \nu\vec{q}}{c_j + d_j \cdot \nu}$. Step **2** can thus be achieved using $O(1)$ $(2k)$ -linear queries
2417 per hyperplane of \mathcal{N} .

2418 In step **4**, the recursive step is carried out on

$$\text{2419} \quad q' = \nu + \lambda_\theta \nu\vec{q} = \nu - \frac{c_\theta + d_\theta \cdot \nu}{d_\theta \cdot \nu\vec{q}} \nu\vec{q},$$

2420 hence, comparing λ'_i to 0 amounts to performing the query $-\frac{d_i \cdot \nu\vec{q}'}{c_i + d_i \cdot \nu'} > ? 0$,
2421 which is not linear in q . The same goes for comparing λ'_i to λ'_j with the
2422 query $\frac{d_i \cdot \nu\vec{q}'}{c_i + d_i \cdot \nu'} < ? \frac{d_j \cdot \nu\vec{q}'}{c_j + d_j \cdot \nu'}$.

2423 However, we can multiply both sides of the inequality test by $d_\theta \nu\vec{q}$ to
2424 keep the queries linear as shown below. We must be careful to take into
2425 account the sign of the expression $d_\theta \nu\vec{q}$, this costs us one additional linear
2426 query.

This trick can be used at each step of the recursion. Let $q^{(0)} = q$, then we have

$$q^{(s+1)} = \nu^{(s)} - \frac{c_{\theta_s} + d_{\theta_s} \cdot \nu^{(s)}}{d_{\theta_s} \cdot \nu\vec{q}^{(s)}} \nu\vec{q}^{(s)}$$

2427 and $(d_{\theta_s} \cdot \nu\vec{q}^{(s)}) q^{(s+1)}$ yields a vector whose components are linear in $q^{(s)}$.
2428 Hence, $(\prod_{k=0}^s d_{\theta_k} \cdot \nu\vec{q}^{(k)}) q^{(s+1)}$ yields a vector whose components are linear
2429 in q , and for all pairs of vectors d_i and $\nu^{(s+1)}$ we have that $(\prod_{k=0}^s d_{\theta_k} \cdot$
2430 $\nu\vec{q}^{(k)}) (d_i \cdot \nu\vec{q}^{(s+1)})$ is linear in q .

Hence at the s th recursive step of the algorithm, we will perform at

¹Note that we project from ν instead of q . We are allowed to do this since $\nu + \lambda_i \nu\vec{q} = q + (\lambda_i - 1) \nu\vec{q}$ and there is no hyperplane separating q from ν .

²Note that if $c_i + d_i \cdot \nu = 0$ then $\lambda_i = 0$, we can check this beforehand for free.

most $|\mathcal{N}|$ linear queries of the type

$$-\left(\prod_{k=0}^{s-1} d_{\theta_k} \cdot \nu \vec{q}^{(k)}\right) \frac{d_i \cdot \nu \vec{q}^{(s)}}{c_i + d_i \cdot \nu^{(s)}} ? 0$$

$|\mathcal{N}| - 1$ linear queries of the type

$$\left(\prod_{k=0}^{s-1} d_{\theta_k} \cdot \nu \vec{q}^{(k)}\right) \frac{d_i \cdot \nu \vec{q}^{(s)}}{c_i + d_i \cdot \nu^{(s)}} ? \left(\prod_{k=0}^{s-1} d_{\theta_k} \cdot \nu \vec{q}^{(k)}\right) \frac{d_j \cdot \nu \vec{q}^{(s)}}{c_j + d_j \cdot \nu^{(s)}}$$

and a single linear query of the type

$$d_{\theta_{s-1}} \cdot \nu \vec{q}^{(s-1)} ? 0.$$

2431 In order to detect all hyperplanes H_i such that $\lambda_i = \lambda_\theta$ we can afford
 2432 to compute the query $f(q) > g(q)$ for all query $f(q) < g(q)$ that we issue,
 2433 and vice versa.

2434 Note that, without further analysis, the queries can become n -linear as
 2435 soon as we enter the $\frac{n}{k}$ th recursive step. \square

2436 A.2.2 Algebraic Computation Trees

2437 The following theorem follows immediately from the analysis of the
 2438 linearity of queries

2439 **Theorem A.10.** *The algebraic computation tree complexity of k -LDT is*
 2440 $\tilde{O}(n^3)$.

2441 *Proof.* We go through each step of Algorithm 2. Indeed, each k -linear
 2442 query of step 1 can be implemented as $O(k)$ arithmetic operations, so step
 2443 1 has complexity $O(|\mathcal{N}|)$. The construction of the simplex in step 2 must
 2444 be handled carefully. What we need to show is that each n -linear query
 2445 we use can be implemented using $O(k)$ arithmetic operations. It is not
 2446 difficult to see from the expressions given in §A.2.1 that a constant number
 2447 of arithmetic operations and dot products suffice to compute the queries. A
 2448 dot product in this case involves a constant number of arithmetic operations
 2449 because the d_i are such that they each have exactly k non-zero components.
 2450 The only expression that involves a non-constant number of operations is
 2451 the product $\prod_{k=0}^s d_{\theta_k} \cdot \nu \vec{q}^{(k)}$, but this is equivalent to $(\prod_{k=0}^{s-1} d_{\theta_k} \cdot \nu \vec{q}^{(k)}) (d_{\theta_s} \cdot$

$\vec{\nu q}^{(s)}$) where the first factor has already been computed during a previous step and the second factor is of constant complexity. Since each query costs a constant number of arithmetic operations and branching operations, step 2 has complexity $O(n|\mathcal{N}|)$. Finally, steps 3 and 4 are free since they do not involve the input. The complexity of Algorithm 2 in this model is thus also $O(n^3 \log n \log |\mathcal{H}|)$. \square

2458 A.2.3 Uniform Random Sampling

2459 Theorem 8.3 requires us to pick a sample of the hyperplanes uniformly
 2460 at random. Actually the theorem is a little stronger; we can draw each
 2461 element of \mathcal{N} uniformly at random, only keeping distinct elements. This is
 2462 not too difficult to achieve for k -LDT when the $\alpha_i, i \in [k]$ are all distinct:
 2463 to pick a hyperplane of the form $\alpha_0 + \alpha_1 x_{i_1} + \alpha_2 x_{i_2} + \cdots + \alpha_k x_{i_k} = 0$
 2464 uniformly at random, we can draw each $i_j \in [n]$ independently and there
 2465 are n^k possible outcomes. However, in the case of k -SUM, we only have $\binom{n}{k}$
 2466 distinct hyperplanes. A simple dynamic programming approach solves the
 2467 problem for k -SUM. For k -LDT we can use the same approach, once for
 2468 each class of equal α_i .

2469 **Lemma A.11.** *Given $n \in \mathbb{N}$ and $(\alpha_0, \alpha_1, \dots, \alpha_k) \in \mathbb{R}^{k+1}$, m independent
 2470 uniform random draws of hyperplanes in \mathbb{R}^n with equations of the form
 2471 $\alpha_0 + \alpha_1 x_{i_1} + \alpha_2 x_{i_2} + \cdots + \alpha_k x_{i_k} = 0$ can be computed in time $O(mk^2 \log n)$
 2472 and preprocessing time $O(k^2 n)$.*

2473 *Proof.* We want to pick an assignment $a = \{(\alpha_1, x_{i_1}), \dots, (\alpha_k, x_{i_k})\}$ uniformly at random. Note that all x_i are distinct while the α_j can be equal.

2475 Without loss of generality, suppose $\alpha_1 \leq \alpha_2 \leq \cdots \leq \alpha_k$. There
 2476 is a bijection between assignments and lexicographically sorted k -tuples
 2477 $((\alpha_1, x_{i_1}), \dots, (\alpha_k, x_{i_k}))$.

2478 Observe that x_{i_j} can be drawn independently of $x_{i_{j'}}$ whenever $\alpha_j \neq \alpha_{j'}$.
 2479 Hence, it suffices to generate a lexicographically sorted $|\chi|$ -tuple of x_i for
 2480 each class χ of equal α_i .

Let $\omega(m, l)$ denote the number of lexicographically sorted l -tuples, where each element comes from a set of m distinct x_i . We have

$$\omega(m, l) = \begin{cases} 1 & \text{if } l = 0 \\ \sum_{i=1}^m \omega(i, l-1) & \text{otherwise.} \end{cases}$$

To pick such a tuple $(x_{i_1}, x_{i_2}, \dots, x_{i_l})$ uniformly at random we choose $x_{i_l} = x_o$ with probability

$$P(x_{i_l} = x_o) = \begin{cases} 0 & \text{if } o > m \\ \frac{\omega(o, l-1)}{\omega(m, l)} & \text{otherwise} \end{cases}$$

that we append to a prefix $(l - 1)$ -tuple (apply the procedure recursively), whose elements come from a set of o symbols. If $l = 0$ we just return the empty tuple.

Obviously, the probability for a given l -tuple to be picked is equal to $\frac{1}{\omega(m, l)}$.

Let X denote the partition of the α_i into equivalence classes, then the number of assignments is equal to $\prod_{\chi \in X} \omega(n, |\chi|)$. (Note that for k -SUM this is simply $\omega(n, k)$ since there is only a single class of equivalence.) For each equivalence class χ we draw independently a lexicographically sorted $|\chi|$ -tuple on n symbols using the procedure above. This yields a given assignment with probability $\frac{1}{\prod_{\chi \in X} \omega(n, |\chi|)}$. Hence, this corresponds to a uniform random draw over the assignments.

It is a well known fact that $\omega(n, k) = \binom{n+k-1}{k-1}$, hence each number we manipulate fits in $O(k \log n)$ bits, that is, $O(k)$ words. Moreover $\omega(n, k) = \omega(n-1, k) + \omega(n-1, k-1)$ so each $\omega(m, l)$ can be computed using a single addition on numbers of $O(k)$ words.

For given n and k , there are at most nk values $\omega(m, l)$ to compute, and for a given k -LDT instance, it must be computed only once. One way to perform the random draws is to compute the cumulative distribution functions of the discrete distributions defined above, then to draw x_{i_l} , we use binary search to find a generated random integer of $O(k)$ words in the cumulative distribution function. Computing the values $\omega(m, l)$ and all cumulative distributions functions can be done as a preprocessing step in $O(k^2 n)$ time. Assuming the generation of a random sequence of words takes linear time, performing a random draw takes time $O(k^2 \log n)$.

□

A.2.4 Proof of Lemma A.5

Theorem A.12 (Cramer's rule). *If a system of n linear equations for n unknowns, represented in matrix multiplication form $Ax = b$, has a unique*

solution $x = (x_1, x_2, \dots, x_n)^T$ then, for all $i \in [n]$,

$$x_i = \frac{\det(A_i)}{\det(A)}$$

where A_i is A with the i th column replaced by the column vector b .

Lemma A.13 (Meyer auf der Heide[114]). *The absolute value of the determinant of an $n \times n$ matrix $M = M_{i=1\dots n, j=1\dots n}$ with integer entries is an integer that is at most $C^n n^{\frac{n}{2}}$, where C is the maximum absolute value in M .*

Proof. The determinant of M must be an integer and is the volume of the hyperparallelepiped spanned by the row vectors of M , hence

$$|\det(M)| \leq \prod_{i=1}^n \sqrt{\sum_{j=1}^n M_{i,j}^2} \leq (\sqrt{nC^2})^n \leq C^n n^{\frac{n}{2}}.$$

□

Lemma A.14. *The determinant of an $n \times n$ matrix $M = M_{i=1\dots n, j=1\dots n}$ with rational entries can be represented as a fraction whose numerators and denominators absolute values are bounded above by $(ND^{n-1})^n n^{\frac{n}{2}}$ and D^{n^2} respectively, where N and D are respectively the maximum absolute value of a numerator and a denominator in M .*

Proof. Let $\delta_{i,j}$ denote the denominator of $M_{i,j}$. Multiply each row M_i of M by $\prod_j \delta_{i,j}$. Apply Lemma A.13. □

We can now proceed to the proof of Lemma A.5.

Proof. Coefficients of the hyperplanes of the arrangement are constant rational numbers, those can be changed to constant integers (because each hyperplane has at most k nonzero coefficients). Let C denote the maximum absolute value of those coefficients.

Because of Theorem A.12 and Lemma A.13, vertices of the arrangement have rational coordinates whose numerators and denominators absolute values are bounded above by $C^n n^{\frac{n}{2}}$.

Given simplices whose vertices are vertices of the arrangement, hyperplanes that define the faces of those simplices have rational coefficients whose numerators and denominators absolute values are bounded above by

2531 $C^{2n^3} n^{n^3 + \frac{n}{2}}$ by Theorem A.12 and Lemma A.14. (Note that some simplices
2532 might be not fully dimensional, but we can handle those by adding vertices
2533 with coordinates that are not much larger than that of already existing
2534 vertices).

2535 By applying Theorem A.12 and Lemma A.14 again, we obtain that
2536 vertices of the arrangement of those new hyperplanes (and thus vertices of
2537 S) have rational coefficients whose numerators and denominators absolute
2538 values are bounded above by $C^{4n^5} n^{2n^5 + n^3 + \frac{n}{2}}$. □

B

2539

Subquadratic Algorithms for Algebraic 3SUM

2540 with Luis Barba, Jean Cardinal, John Iacono, Stefan Langerman, and Noam Solomon

2543 The 3SUM problem asks if an input n -set of real numbers contains a
2544 triple whose sum is zero. We qualify such a triple as *degenerate* because
2545 the probability of finding one in a random input is zero. We consider
2546 the 3POL problem, an algebraic generalization of 3SUM where we replace
2547 the sum function by a constant-degree polynomial in three variables. The
2548 motivations are threefold. Raz, Sharir, and de Zeeuw gave an $O(n^{11/6})$ upper
2549 bound on the number of degenerate triples for the 3POL problem. We give
2550 algorithms for the corresponding problem of counting them. Grønlund and
2551 Pettie designed subquadratic algorithms for 3SUM. We prove that 3POL
2552 admits bounded-degree algebraic decision trees of depth $O(n^{12/7+\varepsilon})$, and we
2553 prove that 3POL can be solved in $O(n^2(\log \log n)^{3/2}/(\log n)^{1/2})$ time in the
2554 real-RAM model, generalizing their results. Finally, we shed light on the
2555 General Position Testing (GPT) problem: “Given n points in the plane, do
2556 three of them lie on a line?”, a key problem in computational geometry: we
2557 show how to solve GPT in subquadratic time when the input points lie on a
2558 small number of constant-degree polynomial curves. Many other geometric
2559 degeneracy testing problems reduce to 3POL.

2560 B.1 First Subquadratic Algorithms for 3POL

2561 The results in this section have been published in Paper B. Our main
2562 contribution in this paper is the first subquadratic algorithm for the 3POL

2563 problem (Problem 13).

2564 This section is divided into four subsections: In §B.1.1, we design a
 2565 bounded-degree algebraic decision tree of depth $O(n^{\frac{12}{7}+\epsilon})$ for explicit 3POL
 2566 (Contribution 3), and in §B.1.2, we adapt this decision tree to run in time
 2567 $O(n^2(\log \log n)^{\frac{3}{2}}/(\log n)^{\frac{1}{2}})$ in the real-RAM model (Contribution 4). In
 2568 §B.1.3, we generalize the decision tree from §B.1.1 to work for 3POL with
 2569 the same depth, up to constant factors (Contribution 5). Finally, in §B.1.4,
 2570 we give a real-RAM implementation of this second decision tree to solve
 2571 3POL as fast as explicit 3POL, up to constant factors (Contribution 6).

2572 Missing details are found in §B.2 and applications can be found in §B.3.

2573 B.1.1 Nonuniform Algorithm for Explicit 3POL

2574 We begin with the description of a nonuniform algorithm for explicit
 2575 3POL which we use later as a basis for other algorithms.

2576 We recall the definition of the explicit 3POL problem 

2577 **Problem 14** (explicit 3POL). Let $f \in \mathbb{R}[x, y]$ be a bivariate polynomial
 2578 of constant degree, given three sets A , B , and C , each containing n real
 2579 numbers, decide whether there exist $a \in A$, $b \in B$, and $c \in C$ such that
 2580 $c = f(a, b)$.

2581 We prove the following

2582 **Contribution 3.** *Explicit 3POL can be solved in $O(n^{12/7+\epsilon})$ time in the
 2583 bounded-degree algebraic decision tree model.*

2584 **Idea** We partition the sets A and B into small groups of consecutive
 2585 elements. That way, we can divide the $A \times B$ grid into cells with the
 2586 guarantee that each curve $c = f(x, y)$ intersects a small number of those
 2587 cells. For each such curve and each cell it intersects, we search c among the
 2588 values $f(a, b)$ for all (a, b) in a given intersected cell. We generalize Fredman's
 2589 trick [77] — and how it is used in Grønlund and Pettie's paper [91] — to
 2590 quickly obtain a sorted order on those values, which provides us a logarithmic
 2591 search time for each cell. Below is a sketch of the algorithm.

2592 **Algorithm 3** (Nonuniform algorithm for explicit 3POL).

2593 **input** $A = \{a_1 < \dots < a_n\}, B = \{b_1 < \dots < b_n\}, C = \{c_1 < \dots <$
 2594 $c_n\} \subset \mathbb{R}$.

2595 **output** *accept* if $\exists (a, b, c) \in A \times B \times C$ such that $c = f(a, b)$, *reject*
 2596 otherwise.¹

2597 1. Partition the intervals $[a_1, a_n]$ and $[b_1, b_n]$ into blocks A_i^* and B_j^* such
 2598 that $A_i = A \cap A_i^*$ and $B_j = B \cap B_j^*$ have size g .

2599 2. Sort the sets $f(A_i \times B_j) = \{f(a, b) : (a, b) \in A_i \times B_j\}$ for all A_i, B_j .
 2600 This is the only step that is nonuniform.

2601 3. For each $c \in C$,

2602 3.1. For each cell $A_i^* \times B_j^*$ intersected by the curve $c = f(x, y)$,

2603 3.1.1. Binary search for c in the sorted set $f(A_i \times B_j)$. If c is found, *accept*
 2604 and halt.

2605 4. *reject* and halt.

2606 Like in Grønlund and Pettie's $\tilde{O}(n^{3/2})$ decision tree for 3SUM [91], the key
 2607 is to give an efficient implementation of step 2.

2608 **$A \times B$ grid partitioning** Let $A = \{a_1 < a_2 < \dots < a_n\}$ and $B =$
 2609 $\{b_1 < b_2 < \dots < b_n\}$. For some positive integer g to be determined later,
 2610 partition the interval $[a_1, a_n]$ into n/g blocks $A_1^*, A_2^*, \dots, A_{n/g}^*$ such that each
 2611 block contains g numbers in A . Do the same for the interval $[b_1, b_n]$ with
 2612 the numbers in B and name the blocks of this partition $B_1^*, B_2^*, \dots, B_{n/g}^*$.
 2613 For the sake of simplicity, and without loss of generality, we assume here
 2614 that g divides n . We continue to make this assumption in the following
 2615 sections. To each of the $(n/g)^2$ pairs of blocks A_i^* and B_j^* corresponds
 2616 a cell $A_i^* \times B_j^*$. By definition, each cell contains g^2 pairs in $A \times B$. For
 2617 the sake of notation, we define $A_i = A \cap A_i^* = \{a_{i,1} < a_{i,2} < \dots < a_{i,g}\}$
 2618 and $B_j = B \cap B_j^* = \{b_{j,1} < b_{j,2} < \dots < b_{j,g}\}$. Figure B.1 depicts this
 2619 construction.

2620 The following two lemmas result from this construction:

2621 **Lemma B.1.** *For a fixed value $c \in C$, the curve $c = f(x, y)$ intersects $O(\frac{n}{g})$
 2622 cells. Moreover, those cells can be found in $O(\frac{n}{g})$ time.*

2623 *Proof.* Since f has constant degree, the curve $c = f(x, y)$ can be parti-
 2624 tioned into a constant number of xy -monotone arcs. Split the curve into

¹Note that it is easy to modify the algorithm to count or report the solutions. In the latter case, the algorithm becomes output sensitive.

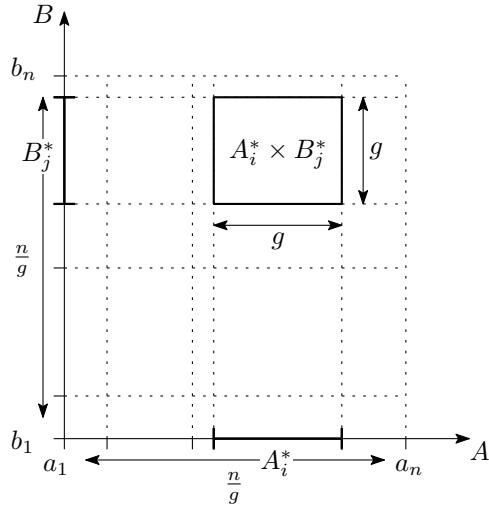


Figure B.1. The partitioning of $A \times B$. There are n/g columns A_i^* , n/g rows B_j^* , and $(n/g)^2$ cells $A_i^* \times B_j^*$. There are n^2 points in $A \times B$. Each column contains the ng points in $A_i \times B$, each row contains the ng points in $A \times B_j$, and each cell contains the g^2 points in $A_i \times B_j$.

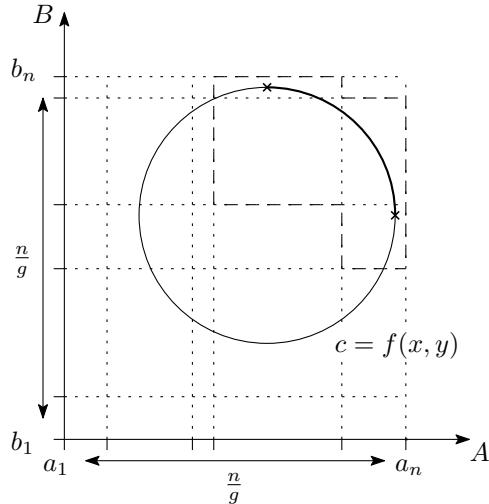


Figure B.2. An xy -monotone arc of the two-dimensional polynomial curve of equation $c = f(x, y)$. This arc intersects a staircase of at most $2\frac{n}{g} - 1$ cells in the grid. When f has constant degree, the defined curve can be partitioned into $O(1)$ such arcs.

2625 x -monotone pieces, then each x -monotone piece into y -monotone arcs. The
 2626 endpoints of the xy -monotone arcs are the intersections of $f(x, y) = c$ with
 2627 its derivatives $f'_x(x, y) = 0$ and $f'_y(x, y) = 0$. By Bézout's theorem, there
 2628 are $O(\deg(f)^2)$ such intersections and so $O(\deg(f)^2)$ xy -monotone arcs.
 2629 Figure B.2 shows that each such arc intersects $O(\frac{n}{g})$ cells since the cells
 2630 intersected by a xy -monotone arc form a staircase in the grid. This proves
 2631 the first part of the lemma.

2632 To prove the second part, notice that for each connected component of
 2633 $c = f(x, y)$ intersecting at least one cell of the grid either: (1) it intersects a
 2634 boundary cell of the grid, or (2) it is a (singular) point or contains vertical
 2635 and horizontal tangency points.² The cells intersected by $c = f(x, y)$ are
 2636 computed by exploring the grid from $O(\frac{n}{g})$ starting cells. Start with an
 2637 empty set. Find and add all boundary cells containing a point of the curve.
 2638 Finding those cells is achieved by solving the Tarski sentence $\exists x \exists y (c =$
 2639 $f(x, y) \wedge x \in A_i^* \wedge y \in B_j^*)$, for each cell $A_i^* \times B_j^*$ on the boundary. This
 2640 takes $O(\frac{n}{g})$ time. Find and add the cells containing singular points and
 2641 tangency points of $c = f(x, y)$. Finding those cells is achieved by first
 2642 finding the constant number of vertical and horizontal slabs $A_i^* \times \mathbb{R}$ and
 2643 $\mathbb{R} \times B_j^*$ containing such points:

$$\begin{aligned} 2644 \quad & \exists x \exists y (c = f(x, y) \wedge (f'_x(x, y) = 0 \vee f'_y(x, y) = 0) \wedge x \in A_i^*), \\ 2645 \quad & \exists x \exists y (c = f(x, y) \wedge (f'_x(x, y) = 0 \vee f'_y(x, y) = 0) \wedge y \in B_j^*). \end{aligned}$$

2647 This takes $O(\frac{n}{g})$ time. Then for each pair of vertical and horizontal slab
 2648 containing such a point, check that the cell at the intersection of the slabs
 2649 also contains such a point:

$$2650 \quad \exists x \exists y (c = f(x, y) \wedge (f'_x(x, y) = 0 \vee f'_y(x, y) = 0) \wedge x \in A_i^* \wedge y \in B_j^*).$$

2651 This takes $O(1)$ time. Note that we can always assume the constant-degree
 2652 polynomials we manipulate are square-free, as making them square-free
 2653 is trivial [159]: since $\mathbb{R}[x]$ and $\mathbb{R}[y]$ are unique factorization domains, let
 2654 $Q = P/\gcd(P, P'_x; x)$ and $\text{sf}(P) = Q/\gcd(P, P'_y; y)$, where $\gcd(P, Q; z)$ is the
 2655 greatest common divisor of P and Q when viewed as polynomials in $R[z]$
 2656 where R is a unique factorization domain and $\text{sf}(P)$ is the square-free part
 2657 of P . The set now contains, for each component of each type, at least one

²Note that vertical and horizontal lines fall in both categories.

2658 cell intersected by it. Initialize a list with the elements of the set. While
 2659 the list is not empty, remove any cell from the list, add each of the eight
 2660 neighbouring cells to the set and the list, if it contains a point of $c = f(x, y)$
 2661 — this can be checked with the same sentences as in the boundary case —
 2662 and if it is not already in the set. This costs $O(1)$ per cell intersected. The
 2663 set now contains all cells of the grid intersected by $c = f(x, y)$. \square

2664 **Lemma B.2.** *If the sets A, B, C can be preprocessed in $S_g(n)$ time so
 2665 that, for any given cell $A_i^* \times B_j^*$ and any given $c \in C$, testing whether
 2666 $c \in f(A_i \times B_j) = \{f(a, b) : (a, b) \in A_i \times B_j\}$ can be done in $O(\log g)$ time,
 2667 then, explicit 3POL can be solved in $S_g(n) + O(\frac{n^2}{g} \log g)$ time.*

2668 *Proof.* We need $S_g(n)$ preprocessing time plus the time required to search
 2669 each of the n numbers $c \in C$ in each of the $O(\frac{n}{g})$ cells intersected by
 2670 $c = f(x, y)$. Each search costs $O(\log g)$ time. We can compute the cells
 2671 intersected by $c = f(x, y)$ in $O(\frac{n}{g})$ time by Lemma B.1. \square

2672 **Remark** We do not give a $S_g(n)$ -time real-RAM algorithm for preprocess-
 2673 ing the input, but only a $S_g(n)$ -depth bounded-degree ADT. In fact, this
 2674 preprocessing step is the only nonuniform part of Algorithm 3. A real-RAM
 2675 implementation of this step is given in §B.1.2.

2676 **Preprocessing** All that is left to prove is that $S_g(n)$ is subquadratic for
 2677 some choice of g . To achieve this we sort the points inside each cell using
 2678 Fredman's trick [77]. Grønlund and Pettie [91] use this trick to sort the sets
 2679 $A_i + B_j = \{a + b : (a, b) \in A_i \times B_j\}$ with few comparisons: sort the set
 2680 $D = (\cup_i [A_i - A_i]) \cup (\cup_j [B_j - B_j])$, where $A_i - A_i = \{a - a' : (a, a') \in A_i \times A_i\}$
 2681 and $B_j - B_j = \{b - b' : (b, b') \in B_j \times B_j\}$, using $O(n \log n + |D|)$ comparisons,
 2682 then testing whether $a + b \leq a' + b'$ can be done using the free (already
 2683 computed) comparison $a - a' \leq b' - b$. We use a generalization of this trick to
 2684 sort the sets $f(A_i \times B_j)$. For each B_j , for each pair $(b, b') \in B_j \times B_j$, define
 2685 the curve $\gamma_{b,b'} = \{(x, y) : f(x, b) = f(y, b')\}$. Define the sets $\gamma_{b,b'}^0 = \gamma_{b,b'}$,
 2686 $\gamma_{b,b'}^- = \{(x, y) : f(x, b) < f(y, b')\}$, and $\gamma_{b,b'}^+ = \{(x, y) : f(x, b) > f(y, b')\}$.
 2687 The following lemma — illustrated by Figures B.3 and B.4 — follows by
 2688 definition:

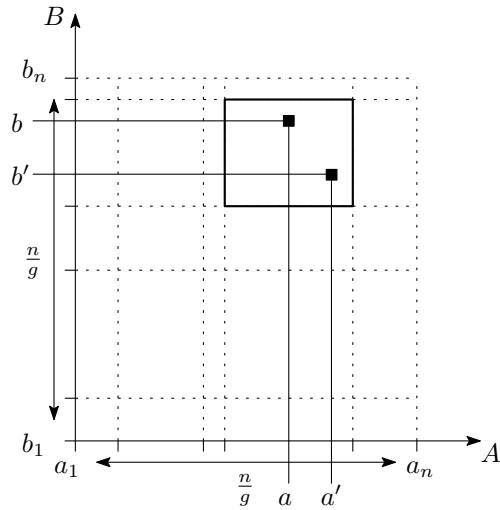


Figure B.3. For each cell, we sort the points it contains with comparisons. The points (a, b) and (a', b') are compared using the comparison $f(a, b) \leq f(a', b')$.

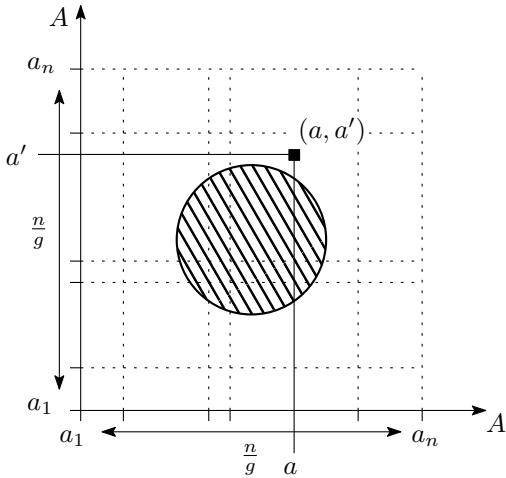


Figure B.4. The disk is the semi-algebraic set $\{(x, y): f(x, b) \leq f(y, b')\}$. Here (a, a') lies outside this semi-algebraic set which implies that $f(a, b) > f(a', b')$.

2689 **Lemma B.3.** *Given a cell $A_i^* \times B_j^*$ and two pairs $(a, b), (a', b') \in A_i \times$
 2690 B_j , deciding whether $f(a, b) < f(a', b')$ (respectively $f(a, b) = f(a', b')$ and
 2691 $f(a, b) > f(a', b')$) amounts to deciding whether the point (a, a') is contained
 2692 in $\gamma_{b,b'}^-$ (respectively $\gamma_{b,b'}^0$ and $\gamma_{b,b'}^+$).*

2693 There are $N := \frac{n}{g} \cdot g^2 = ng$ pairs $(a, a') \in \cup_i [A_i \times A_i]$ and there are N
 2694 pairs $(b, b') \in \cup_j [B_j \times B_j]$. Sorting the $f(A_i \times B_j)$ for all (A_i, B_j) amounts
 2695 to solving the following problem:

2696 **Problem 17** (Polynomial Batch Range Searching). Given N points and N
 2697 polynomial curves in \mathbb{R}^2 , locate each point with respect to each curve.

2698 We can now refine the description of step 2 in Algorithm 3

2699 **Algorithm 4** (Sorting the $f(A_i \times B_j)$ with a nonuniform algorithm).

2700 **input** $A = \{a_1 < a_2 < \dots < a_n\}, B = \{b_1 < b_2 < \dots < b_n\} \subset \mathbb{R}$

2701 **output** The sets $f(A_i \times B_j)$, sorted.

2702 **2.1.** Locate each point $(a, a') \in \cup_i [A_i \times A_i]$ w.r.t. each $\gamma_{b,b'}, (b, b') \in$
 2703 $\cup_j [B_j \times B_j]$.

2704 **2.2.** Sort the sets $f(A_i \times B_j)$ using the information retrieved in step 2.1.

2705 Note that this algorithm is nonuniform: step 2.2 costs at least quadratic
 2706 time in the real-RAM model, however, this step does not need to query the
 2707 input at all, as all the information needed to sort is retrieved during step 2.1.
 2708 Step 2.2 incurs no cost in our nonuniform model.

2709 To implement step 2.1, we use a modified³ version of the $N^{\frac{4}{3}} 2^{O(\log^* N)}$
 2710 algorithm of Matoušek [109] for Hopcroft's problem. In § B.2.1, we prove
 2711 the following upper bound:

2712 **Lemma B.4.** *Polynomial Batch Range Searching can be solved in $O(N^{\frac{4}{3}+\varepsilon})$
 2713 time in the real-RAM model when the input curves are the $\gamma_{b,b'}$.*

2714 **Analysis** Combining Lemma B.2 and Lemma B.4 yields a $O((ng)^{4/3+\varepsilon} +$
 2715 $n^2 \log g/g)$ -depth bounded-degree ADT for explicit 3POL. By optimizing
 2716 over g , we get $g = \Theta(n^{2/7-\varepsilon})$, and the previous expression simplifies to
 2717 $O(n^{12/7+\varepsilon})$, proving Contribution 3.

³The original algorithm relies on hierarchical cuttings which cannot be implemented in the bounded-degree ADT model.

2718 B.1.2 Uniform Algorithm for Explicit 3POL

2719 We again consider the explicit 3POL problem

2720 **Problem 14** (explicit 3POL). Let $f \in \mathbb{R}[x, y]$ be a bivariate polynomial
 2721 of constant degree, given three sets A , B , and C , each containing n real
 2722 numbers, decide whether there exist $a \in A$, $b \in B$, and $c \in C$ such that
 2723 $c = f(a, b)$.

2724 We build on the nonuniform algorithm described in §B.1.1 to prove the
 2725 following

2726 **Contribution 4.** *Explicit 3POL can be solved in $O\left(\frac{n^2(\log \log n)^{3/2}}{(\log n)^{1/2}}\right)$ time in
 2727 the real-RAM model.*

2728 We generalize again Grønlund and Pettie [91]. The algorithm we present
 2729 is derived from the first subquadratic algorithm in their paper.

2730 **Idea** We want the implementation of step 2 in Algorithm 3 to be uniform,
 2731 because then, the whole algorithm is. We use the same partitioning scheme
 2732 as before except we choose g to be much smaller. This allows to store
 2733 all permutations on g^2 items in a lookup table, where g is chosen small
 2734 enough to make the size of the lookup table $\Theta(n^\varepsilon)$. The preprocessing
 2735 part of the previous algorithm is replaced by $g^2!$ calls to an algorithm that
 2736 determines for which cells a given permutation gives the correct sorted order.
 2737 This preprocessing step stores a constant-size⁴ pointer from each cell to the
 2738 corresponding permutation in the lookup table. Search can now be done
 2739 efficiently: when searching a value c in $f(A_i \times B_j)$, retrieve the corresponding
 2740 permutation on g^2 items from the lookup table, then perform binary search
 2741 on the sorted order defined by that permutation. The sketch of the algorithm
 2742 is exactly Algorithm 3. The only differences with respect to §B.1.1 are the
 2743 choice of g and the implementation of step 2.

2744 **$A \times B$ grid partitioning** We use the same partitioning scheme as before,
 2745 hence Lemma B.1 and Lemma B.2 hold. We just need to find a replacement
 2746 for Lemma B.4.

⁴In the real-RAM and word-RAM models.

2747 **Preprocessing** For their simple subquadratic 3SUM algorithm, Grønlund
 2748 and Pettie [91] explain that for a permutation to give the correct sorted order
 2749 for a cell, that permutation defines a *certificate* — a set of inequalities — that
 2750 the cell must verify. They cleverly note — using Fredman’s Trick [77] as in
 2751 Chan [44] and Bremner et al. [33] — that the verification of a single certificate
 2752 by all cells amounts to solving a red/blue point dominance reporting problem.
 2753 We generalize their method. For each permutation $\pi: [g^2] \rightarrow [g]^2$, where
 2754 $\pi = (\pi_r, \pi_c)$ is decomposed into row and column functions $\pi_r, \pi_c: [g^2] \rightarrow [g]$,
 2755 we enumerate all cells $A_i^* \times B_j^*$ for which the following *certificate* holds:

$$2756 \quad f(a_{i,\pi_r(1)}, b_{j,\pi_c(1)}) \leq f(a_{i,\pi_r(2)}, b_{j,\pi_c(2)}) \leq \cdots \leq f(a_{i,\pi_r(g^2)}, b_{j,\pi_c(g^2)}).$$

2757 **Remark** Since some entries may be equal, to make sure each cell cor-
 2758 responds to exactly one certificate, we replace \leq symbols by choices of
 2759 $g^2 - 1$ symbols in $\{=, <\}$. Each permutation π gets a certificate for each
 2760 of those choices. This adds a 2^{g^2-1} factor to the number of certificates to
 2761 test, which will eventually be negligible. Note that some of those 2^{g^2-1}
 2762 certificates are equivalent. We need to skip some of them, as otherwise
 2763 we might output some cells more than once, and then there will be no
 2764 guarantee with respect to the output size. For example, the certificate
 2765 $f(a_{i,9}, b_{j,5}) = f(a_{i,6}, b_{j,7}) < \cdots < f(a_{i,4}, b_{j,4})$ is equivalent to the certificate
 2766 $f(a_{i,6}, b_{j,7}) = f(a_{i,9}, b_{j,5}) < \cdots < f(a_{i,4}, b_{j,4})$. Among equivalent certifi-
 2767 cates, we only consider the certificate whose permutation π precedes the
 2768 others lexicographically. In the previous example, $((6, 7), (9, 5), \dots, (4, 4)) \prec$
 2769 $((9, 5), (6, 7), \dots, (4, 4))$ hence we would only process the second certificate.
 2770 For the sake of simplicity, we will write inequality when we mean either
 2771 strict inequality or equation, and “ \leq ” when we mean either “ $<$ ” or “ $=$ ”.

2772 **Fredman’s Trick** This is where Fredman’s Trick comes into play. By
 2773 Lemma B.3, each inequality $f(a_{i,\pi_r(t)}, b_{j,\pi_c(t)}) \leq f(a_{i,\pi_r(t+1)}, b_{j,\pi_c(t+1)})$ of a
 2774 certificate can be checked by computing the relative position of the point
 2775 $(a_{i,\pi_r(t)}, a_{i,\pi_r(t+1)})$ with respect to the curve $\gamma_{b_{j,\pi_c(t)}, b_{j,\pi_c(t+1)}}$ in \mathbb{R}^2 . For a
 2776 given certificate, for each A_i and each B_j , define

$$2777 \quad p_i = \left((a_{i,\pi_r(1)}, a_{i,\pi_r(2)}), \dots, (a_{i,\pi_r(g^2-1)}, a_{i,\pi_r(g^2)}) \right),$$

$$2778 \quad q_j = \left(f(x, b_{j,\pi_c(1)}) \leq f(y, b_{j,\pi_c(2)}), \dots, f(x, b_{j,\pi_c(g^2-1)}) \leq f(y, b_{j,\pi_c(g^2)}) \right).$$

2780 A certificate is verified by a cell $A_i \times B_j$ if and only if, for all $t \in [g^2 - 1]$, the
 2781 point $p_{i,t}$ verifies the inequality $q_{j,t}$. Enumerating all cells $A_i \times B_j$ for which
 2782 the certificate holds therefore amounts to solving the following problem:

2783 **Problem 18** (Polynomial Dominance Reporting (PDR)). Given N k -tuples
 2784 p_i of points in \mathbb{R}^2 and N k -tuples q_j of bivariate polynomial inequalities of
 2785 degree at most Δ , output all pairs (p_i, q_j) where, for all $t \in [k]$, the point
 2786 $p_{i,t}$ verifies the inequality $q_{j,t}$.

2787 In § B.2.2, we explain how to solve PDR efficiently and prove the following:
 2788

2789 **Lemma B.5.** *We can output all ℓ such pairs in time $2^{O(k)} N^{2 - \frac{4}{\Delta^2 + 3\Delta + 2} + \varepsilon} + O(\ell)$.*

2791 We can now give a uniform implementation of step 2 in Algorithm 3:

2792 **Algorithm 5** (Sorting the $f(A_i \times B_j)$ with a uniform algorithm).

2793 **input** $A = \{a_1 < a_2 < \dots < a_n\}, B = \{b_1 < b_2 < \dots < b_n\} \subset \mathbb{R}$

2794 **output** The sets $f(A_i \times B_j)$, sorted.

2795 **2.1.** Initialize a table that will contain all $g^2!$ permutations on g^2 elements.

2796 **2.2.** For each permutation $\pi: [g^2] \rightarrow [g]^2$,

2797 **2.2.1.** Append permutation π to the table,

2798 **2.2.2.** For each choice of $g^2 - 1$ symbols in $\{=, <\}$,

2799 **2.2.2.1.** If there is any “=” symbol that corresponds to a lexicographically
 2800 decreasing pair of tuples of indices in π , skip this choice of symbols.

2801 **2.2.2.2.** Solve the PDR instance associated with A, B, π and the choice
 2802 of symbols.

2803 **2.2.2.3.** For each output pair (i, j) , store a pointer to the last entry in
 2804 the table.

2805 **Analysis** Plugging in $k = g^2 - 1$, $N = \frac{n}{g}$, iterating over all permutations
 2806 $(\sum_{\pi} \ell = (n/g)^2)$, and adding the binary search step we get that explicit
 2807 3POL can be solved in time

$$2808 (g^2!) 2^{g^2} 2^{O(g^2)} (n/g)^{2 - \frac{4}{\deg(f)^2 + 3\deg(f) + 2} + \varepsilon} + O((n/g)^2) + O(n^2 \log g/g).$$

2809 The first two terms correspond to the complexity of step 2 in Algorithm 3,
 2810 and the last term corresponds to the complexity of step 3 in Algorithm 3. To
 2811 get subquadratic time we can set $g = c_{\deg(f)} \sqrt{\log n / \log \log n}$, because then
 2812 for some appropriate choice of the constant factor $c_{\deg(f)}$, $(g^2)!2^{g^2}2^{O(g^2)} = n^\delta$
 2813 where $\delta < 4/(\deg(f)^2 + 3\deg(f) + 2) - \varepsilon$, making the first term negligible.
 2814 The complexity of the algorithm is dominated by $O(n^2 \log g/g) = O(n^2 (\log \log n)^{3/2}/(\log n)^{1/2})$. This proves Theorem 4.

2816 B.1.3 Nonuniform Algorithm for 3POL

2817 In this section, we extend the nonuniform algorithm given for explicit
 2818 3POL in §B.1.1 to work for the more general 3POL problem.

2819 We recall the definition of the 3POL problem.

2820 **Problem 13** (3POL). Let $F \in \mathbb{R}[x, y, z]$ be a trivariate polynomial of con-
 2821 stant degree, given three sets A , B , and C , each containing n real numbers,
 2822 decide whether there exist $a \in A$, $b \in B$, and $c \in C$ such that $F(a, b, c) = 0$.

2823 We prove the following:

2824 **Contribution 5.** *3POL can be solved in $O(n^{12/7+\varepsilon})$ time in the bounded-
 2825 degree algebraic decision tree model.*

2826 **Idea** The idea is the same as for explicit 3POL. Partition the plane into
 2827 $A_i^* \times B_j^*$ cells. Note that for a fixed $c \in C$, the curve $F(x, y, c)$ intersects $O(\frac{n}{g})$
 2828 cells $A_i^* \times B_j^*$. The algorithm is the following: (1) for each cell $A_i^* \times B_j^*$, sort
 2829 the real roots of the $F(a, b, z) \in \mathbb{R}[z]$ taking the union over all $(a, b) \in A_i \times B_j$,
 2830 (2) for each $c \in C$, for each cell $A_i^* \times B_j^*$ intersected by $F(x, y, c)$, binary
 2831 search on the sorted order computed in step (1) to find c . Step (2) costs
 2832 $O(n^2 \log g/g)$. It only remains to implement step (1) efficiently.

2833 **$A \times B$ partition** We use the same partitioning scheme as before. Hence,
 2834 counterparts of Lemma B.1 and Lemma B.2 hold

2835 **Lemma B.6.** *For a fixed $c \in C$, the curve $F(x, y, c) = 0$ intersects $O(\frac{n}{g})$
 2836 cells. Moreover, those cells can be computed in $O(\frac{n}{g})$ time.*

2837 **Lemma B.7.** *If the sets A, B, C can be preprocessed in $S_g(n)$ time so
 2838 that, for any given cell $A_i^* \times B_j^*$ and any given $c \in C$, testing whether*

2839 $c \in \{ z : \exists (a, b) \in A_i \times B_j \text{ such that } F(a, b, z) = 0 \}$ can be done in $O(\log g)$
2840 time, then, 3POL can be solved in $S_g(n) + O\left(\frac{n^2}{g} \log g\right)$ time.

2841 **Interleavings** Let $\mathcal{P} = (P_1, P_2, \dots, P_m)$ be a tuple of m (nonzero) uni-
2842 variate polynomials. Let $\{ p_{i,1} < \dots < p_{i,\Delta_i} \}$ be the set of real roots of P_i
2843 (without multiplicities). Let $I = ((i_1, j_1), \dots, (i_\Delta, j_\Delta))$ be a tuple of pairs of
2844 positive integers. We say that \mathcal{P} realizes I if and only if I is a permutation
2845 of $\{ (i, j) : i \in [m], j \in [\Delta_i] \}$, and for all $t \in [\Delta - 1]$, $p_{i_t, j_t} \leq p_{i_{t+1}, j_{t+1}}$.
2846 When used in this context, we call I an *interleaving*. Note that (1) the
2847 first condition implies $\Delta = \sum_{i=1}^m \Delta_i$, (2) a tuple of polynomials realizes at
2848 least one interleaving, (3) a tuple of polynomials realizes more than one
2849 interleaving if some of the polynomials have common real roots. We denote
2850 by $\mathcal{I}(\mathcal{P})$ the set of interleavings realized by \mathcal{P} .

2851 **$A \times A$ (b, b')-partitions** For a fixed pair $(b, b') \in B \times B$, we partition \mathbb{R}^2
2852 into (b, b') -cells that encode equivalence classes. Each cell \mathcal{C} is mapped to
2853 a unique interleaving I , and if we take any two points (a_1, a'_1) and (a_2, a'_2)
2854 inside \mathcal{C} , I is realized by both polynomial tuples $(F(a_1, b, z), F(a'_1, b', z))$ and
2855 $(F(a_2, b, z), F(a'_2, b', z))$. It is possible that a degenerate case arises where
2856 we cannot associate an interleaving to \mathcal{C} because one of the polynomials
2857 is the zero polynomial. We can easily tackle these degeneracies, because,
2858 if any point (a, a') is contained in such a cell, we can immediately answer
2859 the instance with the affirmative. Identifying the interleaving associated
2860 with each cell of each (b, b') -partition, then locating each (a, a') inside each
2861 (b, b') -partition provides the answers to all questions of the form “Is the k th
2862 real root of $F(a, b, z)$ greater than the ℓ th real root of $F(a', b', z)$?", for some
2863 $(a, b), (a', b') \in A_i \times B_j$. Those answers are all we need to binary search for c
2864 in the union of the real roots of the $F(a, b, z) \in \mathbb{R}[z]$ in time $O(\log g)$. Note
2865 again that in the nonuniform setting, we do not sort the roots explicitly, but
2866 we must be able to recover the order from the previous computation steps.

2867 **$\gamma_{b,b'}$ and δ_b curves** We consider the set of interleavings \mathcal{I} that the poly-
2868 nomial tuple $(F(x, b, z), F(y, b', z))$ realizes, where z is a variable, and x and y
2869 are parameters. We identify four types of event that can happen when the
2870 parameters x and y vary continuously (ignoring zero polynomials): (1) a real
2871 root of P_i and a real root of P_j that were previously distinct become equal,

for some P_i and P_j in \mathcal{P} , (2) a real root of P_i and a real root of P_j that were previously equal become distinct, for some P_i and P_j in \mathcal{P} , (3) a real root appears in one of the polynomials, and (4) a real root disappears in one of the polynomials. Note that many of those events can happen concurrently. By definition of an interleaving, those events are the only ones that can cause \mathcal{I} to change.

To handle events of the types (1) and (2), we redefine the curves $\gamma_{b,b'}$ from §B.1.1:

$$\gamma_{b,b'} = \{ (x, y) : \exists z \text{ such that } F(x, b, z) = F(y, b', z) = 0 \},$$

that is, $(a, a') \in \gamma_{b,b'}$ if and only if $F(a, b, z)$ and $F(a', b', z)$ have at least one common root.⁵ Note that this curve is defined by the equation

$$\text{res}(F(x, b, z), F(y, b', z); z) = 0,$$

that is, the set of pairs (x, y) for which the resultant (in z) of $F(x, b, z)$ and $F(y, b', z)$ vanishes. This resultant is a polynomial in $\mathbb{R}[x, y]$ of degree $O(\deg(F)^2)$ and can be computed in constant time [56]. The following lemma follows by continuity of the manipulated curve:

Lemma B.8. *Let (a_1, a'_1) and (a_2, a'_2) be two points in the plane such that there does not exist an interleaving that both $(F(a_1, b, z), F(a'_1, b', z))$ and $(F(a_2, b, z), F(a'_2, b', z))$ realize. Moreover, suppose that those two points belong to a connected region in the plane such that for any point (a, a') in that region, the number of real roots of $F(a, b, z)$ and $F(a', b', z)$ is fixed (and finite). Then the interior of any continuous path from (a_1, a'_1) to (a_2, a'_2) lying in this connected region must intersect $\gamma_{b,b'}$.*

Proof. Let I_1 be an interleaving realized by $(F(a_1, b, z), F(a'_1, b', z))$ and let I_2 be an interleaving realized by $(F(a_2, b, z), F(a'_2, b', z))$. Because the number of real roots of the polynomials $F(x, b, z)$ and $F(y, b', z)$ is fixed for any point (x, y) lying in the connected region, I_1 and I_2 differ by a nonzero number of swaps. Moreover, by contradiction, there is a swap that is common to every choice of I_1 and I_2 . Since there is a common swap, for some $i, j \in [\deg(F)]$ and without loss of generality, the i th root of $F(a_1, b, z)$ is smaller than the j th root of $F(a'_1, b', z)$ whereas the i th root of $F(a_2, b, z)$

⁵Note that Raz, Sharir, and de Zeeuw [133] use the same points and curves.

2903 is larger than the j th root of $F(a'_2, b', z)$. By continuity, on any continuous
 2904 path from (a_1, a'_1) and (a_2, a'_2) there is a point (a, a') such that the i th root
 2905 of $F(a, b, z)$ is equal to the j th root of $F(a', b', z)$. This point cannot be an
 2906 endpoint of the path, hence, the interior of the path intersects $\gamma_{b,b'}$. \square

2907 The contrapositive states that, if there exists a continuous path from
 2908 (a_1, a'_1) to (a_2, a'_2) whose interior does not intersect the curve $\gamma_{b,b'}$, then
 2909 there exists an interleaving realized by both $(F(a_1, b, z), F(a'_1, b', z))$ and
 2910 $(F(a_2, b, z), F(a'_2, b', z))$.

2911 To handle events of the types (3) and (4), we define the curve

$$2912 \quad \delta_b = \{ (x, z) : F(x, b, z) = 0 \},$$

2913 which lies in the xz -plane.

2914 **Lemma B.9.** *We can partition the x axis of the xz -plane into a constant
 2915 number of intervals so that for each interval the number of real roots of
 2916 $F(a, b, z)$ is fixed for all a in this interval.*

2917 *Proof.* We partition the xz -plane into a constant number of vertical slabs
 2918 and lines. The x coordinates of vertical tangency points and singular points
 2919 of δ_b are the values a for which a real root of $F(a, b, z) = 0$ appears or
 2920 disappears. The number of singular and vertical tangency points of δ_b is
 2921 quadratic in $\deg(F)$. For each of those points, draw a vertical line that
 2922 contains the point. Those vertical lines partition the xz -plane into slabs
 2923 and lines. The number of vertical lines we draw is constant because the
 2924 degree of F is constant. Figure B.5 depicts this drawing. The projection of
 2925 the vertical lines on the x axis produce the desired partition (with roughly
 2926 half of the intervals being singletons). Let us name those lines δ_b -lines for
 2927 further reference. \square

2928 We can do a symmetric construction for $F(y, b', z)$ in the zy -plane and
 2929 get horizontal $\delta_{b'}$ -lines.

2930 **Lemma B.10.** *We can partition the y axis of the zy -plane into a constant
 2931 number of intervals so that for each interval the number of real roots of
 2932 $F(a', b', z)$ is fixed for all a' in this interval.*

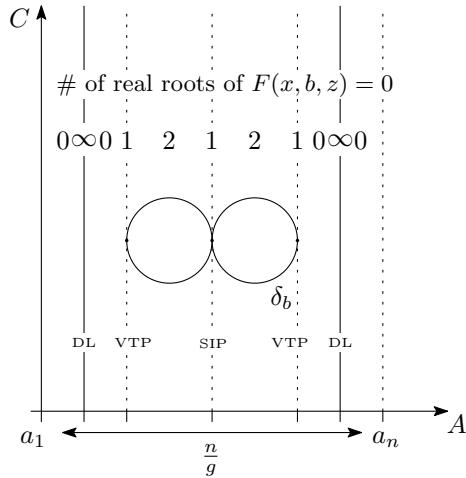


Figure B.5. The vertical tangency points (VTP), self-intersection points (SIP) and degenerate lines (DL) of δ_b partition the A axis into intervals. For all x of the same interval, the polynomial $F(x, b, z) \in \mathbb{R}[z]$ has a fixed number of real roots.

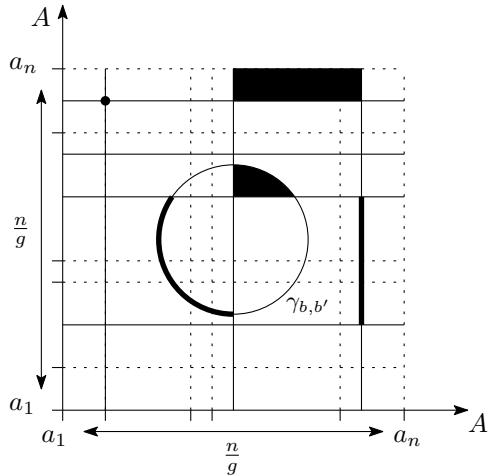


Figure B.6. Cells obtained after partitioning the plane using the curve $\gamma_{b,b'}$ and the δ_b and $\delta_{b'}$ -lines. The five darkened regions highlight examples of (b, b') -cells.

2933 **Cells of the (b, b') -partition** For a given $(b, b') \in B^2$, let $\Gamma_{b,b'}$ be the
 2934 set containing the curve $\gamma_{b,b'}$, the vertical δ_b -lines and the horizontal $\delta_{b'}$ -
 2935 lines. The arrangement $\mathcal{A}(\Gamma_{b,b'})$ of those constant-degree polynomial curves
 2936 partitions \mathbb{R}^2 into a constant-size set $\mathcal{C}(\Gamma_{b,b'})$ of (b, b') -cells. Let $P = \cup_{\gamma \in \Gamma_{b,b'}} \gamma$
 2937 and $\mathcal{C} = \emptyset$. Add all vertices of $\mathcal{A}(\Gamma_{b,b'})$ to \mathcal{C} . Add each connected component
 2938 of $P \setminus \mathcal{C}$ to \mathcal{C} . Add each connected component of $\mathbb{R}^2 \setminus P$ to \mathcal{C} . Finally
 2939 $\mathcal{C}(\Gamma_{b,b'}) = \mathcal{C}$. Those cells can be connected regions, pieces of the curve $\gamma_{b,b'}$,
 2940 pieces of the δ_b - and $\delta_{b'}$ -lines (vertical and horizontal line segments), and
 2941 intersections and self-intersection of those curves (vertices). This partitioning
 2942 scheme is depicted in figure B.6. By construction, all (b, b') -cells have the
 2943 following *invariant property*

2944 **Definition 13.** A (b, b') -cell has the invariant property if, for all points
 2945 (a, a') in that cell, (1) the number of real roots of $F(a, b, z)$ is fixed, (2) the
 2946 number of real roots of $F(a', b', z)$ is fixed, and (3) either, at least one of
 2947 $F(a, b, z)$ or $F(a', b', z)$ is the zero polynomial, or the sorted order of the real
 2948 roots of $F(a, b, z)$ and $F(a', b', z)$ is fixed, that is, $\mathcal{I}((F(a, b, z), F(a', b', z)))$
 2949 is fixed.

2950 **Lemma B.11.** All (b, b') -cells have the invariant property.

2951 *Proof.* First, (1) and (2) hold for all (b, b') -cells because of the partition
 2952 induced by the δ_b -lines and the $\delta_{b'}$ -lines. Second, (3) holds for all (b, b') -
 2953 cells that are not contained in $\gamma_{b,b'}$ since (1) and (2) hold for those cells and
 2954 because of the partition induced by $\gamma_{b,b'}$ (see Lemma B.8). Third, (3) holds
 2955 for all (b, b') -cells that are both contained in $\gamma_{b,b'}$ and some δ_b - or $\delta_{b'}$ -line
 2956 because one of the associated polynomials must be the zero polynomial.

2957 Finally, if a (b, b') -cell is contained in $\gamma_{b,b'}$ but not in any of the δ_b - or $\delta_{b'}$ -
 2958 lines, we make a simple observation. This cell has two distinct neighbouring
 2959 connected regions lying on each of its sides. We just showed that those two
 2960 neighbouring cells have the invariant property. The union of this piece of
 2961 $\gamma_{b,b'}$ with its two neighbouring cells is a connected region as in Lemma B.8.
 2962 Hence, the ordering of any two real roots cannot swap along the piece of
 2963 $\gamma_{b,b'}$, as this would otherwise contradict Lemma B.8. Hence, (3) holds for
 2964 those pieces of $\gamma_{b,b'}$. \square

2965 This lemma implies that, provided we compute in which (b, b') -cells each
 2966 (a, a') point lies, we only need to probe a single point per (b, b') -cell to reveal

2967 the sorted permutation associated with each cell of the $A \times B$ partition.

2968 **Preprocessing** Locate all points $(a, a') \in A_i \times A_i$ for all A_i with respect
 2969 to all $\gamma_{b,b'}$ curves, all vertical lines derived from δ_b and all horizontal lines
 2970 derived from $\delta_{b'}$ for all $(b, b') \in B_j \times B_j$ for all B_j . As in §B.1.1, this can
 2971 be done in a single batch using the algorithm described in §B.2.1, and the
 2972 following generalization of Lemma B.13:

2973 **Lemma B.12.** *Define*

$$2974 \hat{\gamma}_{a,a'} = \{ (x, y) : \text{res}(F(a, x, z), F(a', y, z); z) = 0 \},$$

$$2975 \hat{\delta}_a\text{-lines} = \{ (x, y) : \text{res}(F(a, x, z), F'_x(a, x, z); z) = 0 \},$$

$$2976 \hat{\delta}_{a'}\text{-lines} = \{ (x, y) : \text{res}(F(a', y, z), F'_y(a', y, z); z) = 0 \},$$

$$2977 2978 \hat{\Gamma}_{a,a'} = \hat{\gamma}_{a,a'} \cup \hat{\delta}_a\text{-lines} \cup \hat{\delta}_{a'}\text{-lines}.$$

2979 *Locating* (a, a') *with respect to* $\Gamma_{b,b'}$ *amounts to locating* (b, b') *with respect*
 2980 *to* $\hat{\Gamma}_{a,a'}$.

2981 This gives us the information needed for the binary search in step (2).

2982 **Analysis** The analysis mirrors the explicit case (described immediately
 2983 after Lemma B.4). Combining Lemma B.7 and Lemma B.4 yields a bounded-
 2984 degree ADT of depth $O((ng)^{4/3+\varepsilon} + n^2 \log g/g)$ for 3POL. By optimizing
 2985 over g , we get $g = \Theta(n^{2/7-\varepsilon})$, and the previous expression simplifies to
 2986 $O(n^{12/7+\varepsilon})$, proving Theorem 5.

2987 B.1.4 Uniform Algorithm for 3POL

2988 In this section, we combine the uniform algorithm for explicit 3POL
 2989 given in §B.1.2 with the nonuniform algorithm for 3POL given in §B.1.3 to
 2990 obtain a uniform subquadratic algorithm for 3POL.

2991 We recall the definition of the 3POL problem.

2992 **Problem 13** (3POL). Let $F \in \mathbb{R}[x, y, z]$ be a trivariate polynomial of con-
 2993 stant degree, given three sets A , B , and C , each containing n real numbers,
 2994 decide whether there exist $a \in A$, $b \in B$, and $c \in C$ such that $F(a, b, c) = 0$.

2995 We prove the following:

2996 **Contribution 6.** *3POL can be solved in $O(\frac{n^2(\log \log n)^{3/2}}{(\log n)^{1/2}})$ time in the real-*
 2997 *RAM model.*

2998 **Idea** In the uniform algorithm for explicit 3POL of §B.1.2, we partition
2999 the set $A \times B$ into very small sets $A_i \times B_j$, sort the sets $f(A_i \times B_j)$ using
3000 the dominance reporting algorithm of §B.2.2 then binary search on those
3001 sorted sets in order to find a matching c . Here we devise a similar scheme
3002 with the only difference that the sets to sort are the unions of the real roots
3003 of the univariate polynomials $F(a, b, z) \in \mathbb{R}[z]$ over all $(a, b) \in A_i \times B_j$. The
3004 main difficulty resides in implementing the equivalent of the certificates of
3005 §B.1.2 to reuse the dominance reporting algorithm of §B.2.2. We show how
3006 to implement those certificates using the $\gamma_{b,b'}$ and δ_b curves defined in §B.1.3.

3007 **$A \times B$ partition** We use the same partitioning scheme as all previous
3008 algorithms, hence Lemma B.6 and Lemma B.7 hold. We apply the same
3009 certificate verification scheme as in §B.1.2, hence, the dominance reporting
3010 algorithm of §B.2.2 and the analysis in §B.1.2 still apply.

3011 **Preprocessing** The preprocessing algorithm is essentially the same as
3012 Algorithm 5 with more complex certificates. We explain how to construct
3013 those new certificates. The first part of the explanation consists in general-
3014 lizing the definition of a certificate. The rest of the explanation focuses on
3015 the implementation of the verification of those certificates via Polynomial
3016 Dominance Reporting.

3017 **The certificates** For a fixed pair (a, b) , $F(a, b, z) \in \mathbb{R}[z]$ is a polynomial
3018 in z of degree at most $\deg(F)$. Hence, $F(a, b, z)$ has at most $\deg(F)$ real
3019 roots. For each cell $A_i^* \times B_j^*$, let

$$\text{3020} \quad A_i \times B_j = \{(a_{i,1}, b_{j,1}), (a_{i,1}, b_{j,2}), \dots, (a_{i,2}, b_{j,1}), (a_{i,2}, b_{j,2}), \dots, (a_{i,g}, b_{j,g})\}.$$

3021 Let $\rho: [g]^2 \rightarrow \{0, 1, \dots, \deg(F)\}$ be a function that maps a pair (k, l) to the
3022 number of real roots of $F(a_{i,k}, b_{j,l}, z)$. Let $\Sigma_\rho = \sum_{(i,j) \in [g]^2} \rho(i, j) \leq \deg(F)g^2$.
3023 Given a function ρ , let $\pi: [\Sigma_\rho] \rightarrow [g]^2 \times \{0, 1, \dots, \deg(F)\}$ be a permutation
3024 of the union of the real roots of all g^2 polynomials

$$\text{3025} \quad F(a_{i,1}, b_{j,1}, z), F(a_{i,1}, b_{j,2}, z), \dots, F(a_{i,2}, b_{j,1}, z), \dots, F(a_{i,g}, b_{j,g}, z),$$

3026 where the number of real roots of each polynomial is prescribed by ρ . De-
3027 compose $\pi = (\pi_r, \pi_c, \pi_s)$ into row, column and real root number functions

3028 $\pi_r, \pi_c: [\Sigma_\rho] \rightarrow [g]$, and $\pi_s: [\Sigma_\rho] \rightarrow \{0, 1, \dots, \deg(F)\}$. Let $\diamond(a, b, s)$ denote
3029 the s th real root of $F(a, b, z)$. To fix the permutation of the union of the real
3030 roots of all g^2 polynomials, we define the following interleaving certificate
3031 with $\Sigma_\rho - 1$ inequalities, for each possible function ρ and permutation π

$$\Phi_{\rho, \pi} = \diamond(a_{i, \pi_r(1)}, b_{j, \pi_c(1)}, \pi_s(1)) \leq \dots \leq \diamond(a_{i, \pi_r(\Sigma_\rho)}, b_{j, \pi_c(\Sigma_\rho)}, \pi_s(\Sigma_\rho)).$$

3032 To fix the number of real roots each of the g^2 polynomials can have, we
3033 define the following cardinality certificate for each function ρ

$$\Psi_\rho = \bigwedge_{(k,l) \in [g]^2} F(a_{i,k}, b_{j,l}, z) \text{ has } \rho(k, l) \text{ real roots.}$$

3034 For each possible function ρ and permutation π we define the certificate
3035 $\Upsilon_{\rho, \pi} = \Psi_\rho \wedge \Phi_{\rho, \pi}$ that fixes both the number of real roots each polynomial
3036 has and the permutation of those real roots. The total number of certificates
3037 $\Upsilon_{\rho, \pi}$ is $\sum_{\rho: [g]^2 \rightarrow \{0, 1, \dots, \deg(F)\}} \Sigma_\rho!$ which is of the order of $(g^2)^{O(g^2)}$.

3038 Finally, we need to handle the edge cases where a polynomial $F(a, b, z)$
3039 is the zero polynomial. In that case, $F(a, b, z)$ cancels for all $z \in \mathbb{R}$. Hence,
3040 all planar curves $F(x, y, c) = 0$ go through (a, b) and we can immediately
3041 accept the 3POL instance. To capture those edge cases, we will check the
3042 following certificate before running the main algorithm

$$\aleph = \bigvee_{(k,l) \in [g]^2} F(a_{i,k}, b_{j,l}, z) \text{ is the zero polynomial.}$$

3043 We can check if \aleph holds for any cell $A_i \times B_j$ in $O(n \log n)$ time. For each
3044 $b \in B$ binary search for a $a \in A$ that lies on a vertical line component of δ_b .

3045 If this certificate is verified we accept and halt. Otherwise we can safely
3046 run the main algorithm.

3047 **$A \times A$ (b, b')-partitions** For each B_j and for each $(b, b') \in B_j^2$ compute a
3048 partition of the $A \times A$ grid according to the (b, b') -cells defined by $\Gamma_{b,b'}$ —
3049 see §B.1.3. For each (b, b') -cell of that partition, pick a sample point (a, a') ,
3050 compute the interleaving $\mathcal{I}((F(a, b, z), F(a', b', z)))$. Store that information
3051 for future lookup. All this takes $O(ng)$ time.

3055 **PDR instance for Ψ_ρ** For a fixed pair (a, b) , suppose $F(a, b, z)$ has r real
 3056 roots. Then a must lie in one of the open intervals or be one of the breaking
 3057 points defined by the VTP, SIP and DL of δ_b that fixes the number of real
 3058 roots of $F(a, b, z)$ to r . Hence Ψ_ρ can be rewritten as follows

$$3059 \quad \Psi_\rho = \bigwedge_{(k,l) \in [g]^2} \left(\bigvee_{[u,v] \in \mathcal{I}_{\rho(k,l)}} u < a_{i,k} < v \right) \bigvee \left(\bigvee_{w \in \mathcal{B}_{\rho(k,l)}} a_{i,k} = w \right)$$

3060 where $\mathcal{I}_{\rho(k,l)}$ denotes the set of intervals fixing the number of real roots of
 3061 $F(a_{i,k}, b_{j,l}, z)$ to $\rho(k, l)$, and $\mathcal{B}_{\rho(k,l)}$ denotes the set of breaking points fixing
 3062 the number of real roots of $F(a_{i,k}, b_{j,l}, z)$ to $\rho(k, l)$.

3063 The PDR algorithm can only check conjunctions of polynomial inequali-
 3064 ties. However, we can transform Ψ_ρ into disjunctive normal form (DNF) by
 3065 splitting the certificate into distinct branches, each consisting of a conjunction
 3066 of polynomial inequalities. Since the number of intervals and breaking points
 3067 considered above is constant for each pair (k, l) , the number of branches to
 3068 test is $2^{O(g^2)}$.

3069 For each A_i we have thus a single vector of reals

$$3070 \quad p_i = (a_{i,1}, a_{i,1}, a_{i,2}, a_{i,2}, \dots, a_{i,g}, a_{i,g}),$$

3071 and for each B_j we have $2^{O(g^2)}$ vectors of linear inequalities

$$3072 \quad q_j = (x \bowtie_{u_{1,1}} u_{1,1}, x \bowtie_{v_{1,1}} v_{1,1}, x \bowtie_{u_{1,2}} u_{1,2}, x \bowtie_{v_{1,2}} v_{1,2}, \dots \\ 3073 \quad \dots, x \bowtie_{u_{g,g}} u_{g,g}, x \bowtie_{v_{g,g}} v_{g,g}),$$

3076 where each $(\bowtie_{u_{k,l}}, u_{k,l}, \bowtie_{v_{k,l}}, v_{k,l})$ is an element of

$$3077 \quad \{ (>, u, <, v) : (u, v) \in \mathcal{I}_{\rho(k,l)} \} \cup \{ (=, w, =, w) : w \in \mathcal{B}_{\rho(k,l)} \}.$$

3078 For a fixed function ρ , the sets of vectors p_i and q_j is a valid PDR
 3079 instance of size $N = (n/g) \cdot 2^{O(g)} = n2^{O(g)}$ and with parameter $k = 2g^2$
 3080 that will output all cells $A_i^* \times B_j^*$ such that $F(a_{i,k}, b_{j,l}, z) \in \mathbb{R}[z]$ has exactly
 3081 $\rho(k, l)$ real roots for all $(a_{i,k}, a_{j,l}) \in A_i \times B_j$.

3082 **PDR instance for $\Phi_{\rho,\pi}$** For fixed pairs (a, b) and (a', b') , suppose the s -th
 3083 real root of $F(a, b, z)$ is smaller or equal to the q -th real root of $F(a, b, z)$.

3084 Then, (a, a') must lie in a (b, b') -cell that orders the s -th root of $F(x, b, z)$
 3085 before the q -th root of $F(y, b', z)$ for all points (x, y) in that cell.

3086 Hence $\Phi_{\rho, \pi}$ can be rewritten as follows

$$3087 \quad \Phi_{\rho, \pi} = \bigwedge_{t \in [\Sigma_\rho - 1]} \bigvee_{C \in \mathcal{C}_{\rho, \pi, t}} (a_{i, \pi_r(t)}, a_{i, \pi_r(t+1)}) \in C$$

3088 where $\mathcal{C}_{\rho, \pi, t}$ denotes the set of (b, b') -cells fixing to $\rho(\pi_r(t), \pi_c(t))$ the number
 3089 of real roots of $F(a_{i, \pi_r(t)}, b_{j, \pi_c(t)}, z)$, fixing to $\rho(\pi_r(t+1), \pi_c(t+1))$ the number
 3090 of real roots of $F(a_{i, \pi_r(t+1)}, b_{j, \pi_c(t+1)}, z)$, and ordering the $\pi_s(t)$ -th root of
 3091 $F(a_{i, \pi_r(t)}, b_{j, \pi_c(t)}, z)$ before the $\pi_s(t+1)$ -th root of $F(a_{i, \pi_r(t+1)}, b_{j, \pi_c(t+1)}, z)$.

3092 The PDR algorithm can only check conjunctions of polynomial inequalities.
 3093 However, we can transform $\Phi_{\rho, \pi}$ in DNF as we did for Ψ_ρ . Again the
 3094 number of cells considered above is constant for each t , the description of
 3095 each cell is constant, hence, the number of branches to test is $2^{O(g^2)}$.

3096 For each A_i we have thus a single vector of 2-dimensional points

$$3097 \quad p_i = (\underbrace{\dots, (a_{i, \pi_r(1)}, a_{i, \pi_r(2)}), \dots, \dots, \dots,}_{\omega}, \underbrace{(a_{i, \pi_r(\Sigma_\rho - 1)}, a_{i, \pi_r(\Sigma_\rho)}), \dots,}_{\omega}),$$

3098 where ω is the size of the largest description of a (b, b') -cell C , and for each
 3099 B_j we have $2^{O(g^2)}$ vectors of polynomial inequalities,

$$3100 \quad q_j = (\underbrace{\dots, h_{1, \vartheta}(x, y) \bowtie_{1, \vartheta} 0, \dots, \dots, \dots,}_{\vartheta \in [\omega]}, \underbrace{h_{\Sigma_\rho - 1, \vartheta}(x, y) \bowtie_{\Sigma_\rho - 1, \vartheta} 0 \dots,}_{\vartheta \in [\omega]}),$$

3101 where each $(\dots, h_{t, \vartheta}(x, y) \bowtie_{t, \vartheta} 0, \dots)$ is an element of $\{ \text{desc}(C) : C \in \mathcal{C}_{\rho, \pi, t} \}$,
 3102 where $\text{desc}(C)$ is the description of the cell C given as a certificate of
 3103 belonging to C in the form of a Tarski sentence. The description of each
 3104 (b, b') -cell is padded with its last component so that it has length ω .

3105 For a fixed function ρ , for a fixed function π , the sets of vectors p_i
 3106 and q_j is a valid PDR instance of size $N = n2^{O(g)}$ and with parameter
 3107 $k = \Theta(g^2)$ that will output all cells $A_i^* \times B_j^*$ such that the number of real
 3108 roots of $F(a_{i, \pi_r(t)}, b_{j, \pi_c(t)}, z)$ is $\rho(\pi_r(t), \pi_c(t))$, the number of real roots of
 3109 $F(a_{i, \pi_r(t+1)}, b_{j, \pi_c(t+1)}, z)$ is
 3110 $\rho(\pi_r(t+1), \pi_c(t+1))$, and the $\pi_s(t)$ -th root of $F(a_{i, \pi_r(t)}, b_{j, \pi_c(t)}, z)$ comes
 3111 before the $\pi_s(t+1)$ -th root of $F(a_{i, \pi_r(t+1)}, b_{j, \pi_c(t+1)}, z)$, for all $t \in [\Sigma_\rho - 1]$.

3112 **PDR instance for $\Upsilon_{\rho,\pi}$** We can combine the certificates given above for
3113 Ψ_ρ and $\Phi_{\rho,\pi}$ to obtain the ones for $\Upsilon_{\rho,\pi}$: concatenate the p_i and q_j together
3114 (add a dummy y variable for the p_i and q_j of Ψ_ρ). For a fixed function ρ ,
3115 for a fixed function π , the sets of vectors p_i and q_j is a valid PDR instance
3116 of size $N = n2^{O(g)}$ and with parameter $k = \Theta(g^2)$ that will output all cells
3117 $A_i^* \times B_j^*$ such that $F(a_{i,k}, b_{j,l}, z) \in \mathbb{R}[z]$ has exactly $\rho(k, l)$ real roots for
3118 all $(a_{i,k}, a_{j,l}) \in A_i \times B_j$, and the $\pi_s(t)$ -th root of $F(a_{i,\pi_r(t)}, b_{j,\pi_c(t)}, z)$ comes
3119 before the $\pi_s(t+1)$ -th root of $F(a_{i,\pi_r(t+1)}, b_{j,\pi_c(t+1)}, z)$ for all $t \in [\Sigma_\rho - 1]$.
3120 The rest of the analysis in §B.1.2 applies. This proves Theorem 6.

3121 B.2 Subproblems

3122 In §B.2.1 we prove Lemma B.4 necessary for the implementation of our
3123 nonuniform algorithms. In §B.2.2 we prove Lemma B.5 necessary for the
3124 implementation of our uniform algorithms.

3125 B.2.1 Polynomial Batch Range Searching

3126 In this section we present a uniform algorithm that computes the relative
3127 position of M points with respect to N $\gamma_{b,b'}$ curves. We call such a problem
3128 an (M, N) -problem. When $M = N$ the complexity of the algorithm is
3129 $O(N^{4/3+\varepsilon})$. The algorithm gives the output in “concise form”: it outputs
3130 a set of $(\Pi_\alpha, \Gamma_\beta, \sigma)$ triples where Π_α is a subset of input points, Γ_β is a
3131 subset of input curves, and $\sigma \in \{-, 0, +\}$ indicates the relative position of
3132 all points in Π_α with respect to all curves in Γ_β . Note that if one is only
3133 interested in incident point-curve pairs, the algorithm can explicitly report
3134 all of them in $O(N^{4/3+\varepsilon})$ time, because there are at most $O(N^{4/3})$ such pairs
3135 and because they can easily be filtered from the output.

3136 **Tools** The proof of Lemma B.4 involves standard computational geometry
3137 tools: vertical decomposition of an arrangement of polynomial curves (see
3138 §??), ε -nets (see §8.1.3) and cuttings (see §8.1.5). and derandomization.

3139 *Proof of Lemma B.4.* Fix some constant $r \geq 2$. If $M < r^2$ or $N < r$, solve
3140 by brute-force in $O(M + N)$ time. Otherwise, consider the range space
3141 defined by $\gamma_{b,b'}$ curves and y -axis aligned trapezoidal patches whose top
3142 and bottom sides are pieces of $\gamma_{b,b'}$ curves. This range space has constant
3143 VC-dimension. Compute an $\frac{1}{r}$ -net of size $O(r \log r)$ for the input curves with

respect to this range space. Compute the vertical decomposition Ξ of the arrangement of this $\frac{1}{r}$ -net. This decomposition is a $\frac{1}{r}$ -cutting: it partitions \mathbb{R}^2 into $O(r^2 \log^2 r)$ cells of constant complexity each of which intersects at most $\frac{N}{r}$ input curves. Note that some of those cells are degenerate trapezoidal patches. The decomposition contains vertices, line segments, and curve segments as cells, each of which could contain input points and be intersected or contained by an input curve. Denote by Π_C the set of points contained in the cell $C \in \Xi$. Partition each Π_C into $\left\lceil \frac{|\Pi_C|}{Mr^{-2}} \right\rceil$ disjoint subsets of size at most $\frac{M}{r^2}$. All of this can be done in $O(M + N)$ time. The last step consists in solving $O(r^2 \log^2 r)$ $(\frac{M}{r^2}, \frac{N}{r})$ -problems, that is, solving the problem recursively for the points and curves intersecting each cell. Each recursive call is done by swapping the roles of the points and curves using a form of duality to be described below. The whole algorithm can be formally described as follows,

Algorithm 6 (Polynomial Batch Range Searching).

input A set Π of M points (a, a') , A set Γ of N curves $\gamma_{b,b'}$.

output A set of triples $(\Pi_\alpha, \Gamma_\beta, \sigma)$ covering $\Pi \times \Gamma$ such that, for any triple

$(\Pi_\alpha, \Gamma_\beta, \sigma)$, for all points (a, a') in Π_α and all curves γ in Γ_β , $(a, a') \in \gamma^\sigma$.

0. If $M < r^2$ or $N < r$, solve the problem by brute force in $O(M + N)$ time.

1. Compute an $\frac{1}{r}$ -net of size $O(r \log r)$ for the input curves.

2. Compute the vertical decomposition Ξ of the arrangement of this $\frac{1}{r}$ -net.

3. Denote by Π_C the set of points contained in the cell $C \in \Xi$. Partition each Π_C into $\left\lceil \frac{|\Pi_C|}{Mr^{-2}} \right\rceil$ disjoint subsets $\Pi_{C,i}$ of size at most $\frac{M}{r^2}$.

4. For each cell C of the vertical decomposition,

4.1. For each subset $\Pi_{C,i}$ of points contained in that cell,

4.1.1. Solve an $(\frac{N}{r}, \frac{M}{r^2})$ -problem on the curves intersecting that cell and the points in $\Pi_{C,i}$, swapping the roles of lines and curves via duality.

4.2. For each curve γ not intersecting C ,

4.2.1. Compute the location $\sigma_{C,\gamma}$ of any point in C with respect to γ .

3177 **4.3.** Output $(\{\gamma: \sigma_{C,\gamma} = -\}, \Pi_C, -)$.

3178 **4.4.** Output $(\{\gamma: \sigma_{C,\gamma} = +\}, \Pi_C, +)$.

3179 **4.5.** Output $(\{\gamma: \sigma_{C,\gamma} = 0\}, \Pi_C, 0)$.

3180 **Correctness** We want to locate each point with respect to each curve.
3181 When considering a curve-cell pair, there are two cases: either the curve
3182 intersects the cell, or it does not. For the first case we locate each point
3183 in the cell with respect to the curve in one of the recursive steps. For the
3184 second case, the relative position of all points in the cell with respect to
3185 the curve is the same, it suffices thus to locate one of those points with
3186 respect to the curve to get the location of all the points in $O(1)$ time. Each
3187 recursive call divides M and N by some constant strictly greater than one,
3188 hence, the base case is reached in each of the paths of the recursion tree
3189 and the algorithm always terminates.

3190 **Analysis** For c_1 some constant and bounding $c_1 r^2 \log^2 r$ above by $c_2 r^{2+\varepsilon}$
3191 for some large enough constant c_2 , the complexity $T(M, N)$ of an (M, N) -
3192 problem is thus

$$\text{3193} \quad T(M, N) \leq c_2 r^{2+\varepsilon} T\left(\frac{M}{r^2}, \frac{N}{r}\right) + O(M + N).$$

3195 The complexity $T(N, M)$ of a (N, M) -problem is the same as the complexity
3196 $T(M, N)$ of an (M, N) -problem by the following point-curve duality result
3197 whose proof is straightforward

3198 **Lemma B.13.** *Define*

$$\text{3199} \quad \hat{\gamma}_{a,a'} = \{(x, y) : f(a, x) = f(a', y)\},$$

3200 then, locating (a, a') with respect to $\gamma_{b,b'}$ amounts to locating (b, b') with
3201 respect to $\hat{\gamma}_{a,a'}$.

3202 By doing alternately one step in the primal with the points (a, a') and
3203 the curves $\gamma_{b,b'}$, then a second step with the dual points (b, b') and the dual

3204 curves $\hat{\gamma}_{a,a'}$, we get the following recurrence

$$\begin{aligned} \text{3205} \quad T(M, N) &\leq c_2^2 r^{4+\varepsilon} T\left(\frac{M}{r^3}, \frac{N}{r^3}\right) + c_2 r^{2+\varepsilon} O\left(\frac{M}{r^2} + \frac{N}{r}\right) + O(M + N) \\ \text{3206} \quad &\leq c_2^2 r^{4+\varepsilon} T\left(\frac{M}{r^3}, \frac{N}{r^3}\right) + O(M + N) \\ \text{3207} \end{aligned}$$

3208 Hence, for some large enough constant c_3 ,

$$\begin{aligned} \text{3209} \quad T(N, N) = T(N) &\leq c_3 r^{4+\varepsilon} T\left(\frac{N}{r^3}\right) + O(N) \\ \text{3210} \quad &\leq O\left(N^{\log_{r^3}(c_3 r^{4+\varepsilon})}\right) \\ \text{3211} \quad &\leq O(N^{\frac{4}{3}+\varepsilon}). \\ \text{3212} \end{aligned}$$

3213

□

3214 B.2.2 Polynomial Dominance Reporting

3215 We combine a standard dominance reporting algorithm [127] with 3216 Matoušek's algorithm [109] to prove Lemma B.5. For a pair of blue and red 3217 points in \mathbb{R}^k , we say that the blue point *dominates* the red point if for all 3218 indices $i \in [k]$ the i th coordinate of the blue point is greater or equal to the 3219 i th coordinate of the red point. The standard algorithm in [127] solves the 3220 following problem:

3221 **Problem 19.** Given N blue and M red points in \mathbb{R}^k , report all bichromatic 3222 dominating pairs.

3223 Our problem is significantly more complicated and general. Instead of 3224 blue points we have blue k -tuples p_i of 2-dimensional points, instead of red 3225 points we have red k -tuples q_j of bivariate polynomial inequalities, and we 3226 want to report all bichromatic pairs (p_i, q_j) such that, for all $t \in [k]$, the point 3227 $p_{i,t}$ verifies the inequality $q_{j,t}$. The standard algorithm essentially works by 3228 a combination of divide and conquer and prune and search, using a one- 3229 dimensional cutting (median selection) to split a problem into subproblems. 3230 We generalize the standard algorithm by using higher dimensional cuttings, in 3231 a way similar to Matoušek's algorithm [109]. For the analysis, we generalize 3232 Chan's analysis of the standard algorithm when k is not constant [44].

3233 *Proof of Lemma B.5.* We use the Veronese embedding [95, 96]. Since the
 3234 polynomials have constant degree, we can trade polynomial inequalities
 3235 for linear inequalities by lifting to a space of higher — but constant —
 3236 dimension. The degree of each polynomial is at most Δ . There are exactly
 3237 $d = \binom{\Delta+2}{2} - 1$ different bivariate monomials of degree at most Δ .⁶ To each
 3238 monomial we associate a variable in \mathbb{R}^d . By this association, points in the
 3239 plane are mapped to points in \mathbb{R}^d and bivariate polynomial inequalities are
 3240 mapped to d -variate linear inequalities.

3241 By abuse of notation, let p_i denote the tuple p_i where each 2-dimensional
 3242 point has been replaced by its d -dimensional counterpart, and let q_i denote
 3243 the tuple q_i where each bivariate polynomial inequality has been replaced
 3244 by its d -variate linear counterpart. We have N k -tuples p_i and M k -tuples
 3245 q_j . The algorithm checks each of the k components of the tuples in turn
 3246 and can be described recursively as follows for some positive integer $r > 1$:

3247

3248 **Algorithm 7** (Polynomial Dominance Reporting).

3249 **input** N k -tuples p_i of d -dimensional points, M k -tuples q_j of d -variate
 3250 linear inequalities.

3251 **output** All (p_i, q_j) pairs such that, for all $t \in [k]$, the point $p_{i,t}$ verifies
 3252 the inequality $q_{j,t}$.

- 3253 1. If $k = 0$, then output all pairs (p_i, q_j) and halt.
- 3254 2. If $N < r^d$ or $M < r$, solve the problem by brute force in $O((N + M)k)$
 3255 time.
- 3256 3. We now only consider the k th component of each input k -tuple and call
 3257 these *active* components. To each active d -variate linear inequality
 3258 corresponds a defining hyperplane in \mathbb{R}^d . Construct, as in [109], a
 3259 hierarchical cutting of \mathbb{R}^d using $O(r^d)$ simplicial cells such that each
 3260 simplicial cell is intersected by at most $\frac{M}{r}$ of the defining hyperplanes.
 3261 This construction also gives us for each simplicial cell of the cutting
 3262 the list of defining hyperplanes intersecting it. This takes $O(Mr^{d-1})$
 3263 time. Locate each active point inside the hierarchical cutting in time
 3264 $O(N \log r)$. Let S be a simplicial cell of the hierarchical cutting.
 3265 Denote by Π_S the set of active points in S . Partition each Π_S into

⁶Not including the independent monomial, namely, 1.

3266 $\left\lceil \frac{|\Pi_S|}{Nr^{d-2}} \right\rceil$ disjoint subsets of size at most $\frac{N}{r^d}$. For each simplicial cell,
 3267 find the active inequalities whose corresponding geometric object
 3268 (hyperplane, closed or open half-space) contains the cell. This takes
 3269 $O(Mr^d)$ time. The whole step takes $O(N \log r + Mr^d)$ time.

- 3270 4. For each of the $O(r^d)$ simplicial cells, recurse on the at most $\frac{N}{r^d} k$ -tuples
 3271 p_i whose active point is inside the simplicial cell and the at most $\frac{M}{r}$
 3272 k -tuples q_j whose active inequality's defining hyperplane intersects
 3273 the simplicial cell.
- 3274 5. For each of the $O(r^d)$ simplicial cells, recurse on the at most $\frac{N}{r^d} (k - 1)$ -prefixes of k -tuples p_i whose active point is inside the simplicial
 3275 cell and the $(k - 1)$ -prefixes of k -tuples q_j whose active inequality's
 3276 corresponding geometric object contains the simplicial cell.

3278 **Correctness** In each recursive call, either k is decremented or M and N
 3279 are divided by some constant strictly greater than one, hence, one of the
 3280 conditions in steps **1** and **2** is met in each of the paths of the recursion
 3281 tree and the algorithm always terminates. Step **5** is correct because it only
 3282 recurses on (p_i, q_j) pairs whose suffix pairs are dominating. The base case
 3283 in step **1** is correct because the only way for a pair (p_i, q_j) to reach this
 3284 point is to have had all k components checked in step **5**. The base case in
 3285 step **2** is correct by definition. Each dominating pair is output exactly once
 3286 because the recursive calls of step **4** and **5** partition the set of pairs (p_i, q_j)
 3287 that can still claim to be candidate dominating pairs.

3288 **Analysis** For $k, N, M \geq 0$, the total complexity $T_k(N, M)$ of comput-
 3289 ing the inclusions for the first k components, excluding the output cost
 3290 (steps **1** and **2**), is bounded, in general, by

$$3291 \quad T_k(N, M) \leq \underbrace{O(r^d) T_{k-1}(N, M)}_{\text{Step 5}} + \underbrace{O(r^d) T_k\left(\frac{N}{r^d}, \frac{M}{r}\right)}_{\text{Step 4}} + \underbrace{O(N + M)}_{\text{Step 3}},$$

3292 and, in particular, by $T_k(N, M) = 0$ when $k = 0$, $T_k(N, M) = O(Nk)$ when
 3293 $M < r$, and $T_k(N, M) = O(Mk)$ when $N < r^d$.

3294 By point-hyperplane duality, $T_k(N, M) = T_k(M, N)$, hence, we can ex-
 3295 ecute step **4** on dual linear inequalities and dual points to balance the

recurrence. For some constant $c_1 \geq 1$,

$$T_k(N, M) \leq c_1 r^{2d} T_{k-1}(N, M) + c_1 r^{2d} T_k\left(\frac{N}{r^{d+1}}, \frac{M}{r^{d+1}}\right) + c_1(N + M).$$

For simplicity, we ignore some problem-size reductions occurring in this balancing step.

Let $T_k(N) = T_k(N, N)$ denote the complexity of solving the problem when $M = N$, excluding the output cost. Hence, we have

$$T_k(N) \leq c_1 r^{2d} T_{k-1}(N) + c_1 r^{2d} T_k\left(\frac{N}{r^{d+1}}\right) + c_1 N, \quad (\text{B.1})$$

and, in particular, $T_k(N) = 0$ when $k = 0$, and $T_k(N) = O(k)$ when $N < r^{d+1}$.

To get rid of the parameter k and progress into the analysis of the recurrence, Chan makes an ingenious change of variable [44]. With hindsight, choose $b = r^{d+1}$ and let

$$T(N') = \max \left\{ T_k(N) : k \geq 0, N \geq 1, \text{ and } b^k N \leq N' \right\}. \quad (\text{B.2})$$

For the rest of the discussion, we shorten the notation to

$$T(N') = \max_{b^k N \leq N'} T_k(N).$$

By combining (B.1) and (B.2) we obtain

$$T(N') = \max_{b^k N \leq N'} T_k(N) \leq c_1 \max_{b^k N \leq N'} \left[r^{2d} T_{k-1}(N) + r^{2d} T_k\left(\frac{N}{r^{d+1}}\right) + N \right].$$

The maximum of a sum is always bounded by the sum of the maxima of its terms, hence,

$$T(N') \leq \max_{b^k N \leq N'} \left[c_1 r^{2d} T_{k-1}(N) \right] + \max_{b^k N \leq N'} \left[c_1 r^{2d} T_k\left(\frac{N}{r^{d+1}}\right) \right] + \max_{b^k N \leq N'} [c_1 N].$$

Looking at each term separately, by definition of $T(N')$, we have

$$\max_{b^k N \leq N'} T_{k-1}(N) = \max_{b^{k-1} N \leq \frac{N'}{b}} T_{k-1}(N) = T\left(\frac{N'}{b}\right) = T\left(\frac{N'}{r^{d+1}}\right),$$

$$\max_{b^k N \leq N'} T_k\left(\frac{N}{r^{d+1}}\right) = \max_{b^k \frac{N}{r^{d+1}} \leq \frac{N'}{r^{d+1}}} T_k\left(\frac{N}{r^{d+1}}\right) = T\left(\frac{N'}{r^{d+1}}\right),$$

$$\max_{b^k N \leq N'} N = \max_{N \leq \frac{N'}{b^k}} N = \frac{N'}{b^k} \leq N',$$

which, when combined with the previous inequality, produce the following recurrence

$$T(N') \leq 2c_1 r^{2d} T\left(\frac{N'}{r^{d+1}}\right) + c_1 N'.$$

Powers of r^{d+1} We claim that if N' is a power of r^{d+1} , then $T(N') \leq c_2[N'^{\alpha} - N']$ for some constants $\alpha > 1$ and $c_2 \geq 1$. We prove by induction that this (educated) guess is indeed correct. For $N' = 1$, we have

$$T(1) = \max_{b^k N \leq 1} T_k(N) = T_0(1) = 0 \leq c_2[1^\alpha - 1].$$

For $N' \geq r^{d+1}$ a power of r^{d+1} , and assuming the claim holds for all smaller powers:

$$\begin{aligned} T(N') &\leq 2c_1 r^{2d} c_2 \left[\left(\frac{N'}{r^{d+1}} \right)^\alpha - \frac{N'}{r^{d+1}} \right] + c_1 N' \\ &\leq c_2 N'^{\alpha} \left[\frac{2c_1 r^{2d}}{(r^{d+1})^\alpha} \right] - c_2 N' \left[2c_1 r^{d-1} - \frac{c_1}{c_2} \right]. \end{aligned}$$

We want

$$\frac{c_1 r^{2d}}{(r^{d+1})^\alpha} \leq \frac{1}{2} \quad \text{and} \quad 2c_1 r^{d-1} - \frac{c_1}{c_2} \geq 1.$$

For the first inequality, we can set the left hand side to be equal to $c_1 r^{-\epsilon'} = \frac{1}{2}$ with some small $\epsilon' = \frac{1+\log c_1}{\log r}$. Hence, $2d - \alpha(d+1) = -\epsilon'$, and for $\epsilon = \frac{\epsilon'}{d+1}$, we get $\alpha = \frac{2d}{d+1} + \epsilon$. The second inequality is equivalent to $2r^{d-1} \geq \frac{1}{c_1} + \frac{1}{c_2}$, which always holds since $r > 1, d \geq 1, c_1 \geq 1, c_2 \geq 1$.

We now have

$$T(N') = O\left(N'^{\frac{2d}{d+1} + \epsilon}\right),$$

where $\epsilon = \frac{1+\log c_1}{(d+1)\log r}$ can be chosen arbitrarily small by picking an arbitrarily large $r = (2c_1)^{1/\epsilon(d+1)}$.

Remark The choice $b = r^{d+1}$ gives a simpler analysis. Although giving more freedom to the value of b — as in Chan's paper — yields a slightly better relation between ϵ and r , namely $r > c_1^{1/\epsilon(d+1)}$, it does not get rid of the dependency of ϵ in r , unless $c_1 = 1$.

3348 **General case** When $N' \geq 2$ is not a power of r^{d+1} , we use the fact that
3349 $T(N') \leq T(N' + 1)$ by definition,

$$\begin{aligned} \text{3350} \quad T(N') &= T\left((r^{d+1})^{\log_{r^{d+1}} N'}\right) \\ \text{3351} \quad &\leq T\left((r^{d+1})^{\lfloor \log_{r^{d+1}} N' \rfloor + 1}\right) \\ \text{3352} \quad &= O\left((r^{d+1})^{(\lfloor \log_{r^{d+1}} N' \rfloor + 1)(\frac{2d}{d+1} + \epsilon)}\right) \\ \text{3353} \quad &= O\left((r^{d+1})^{\frac{2d}{d+1} + \epsilon} (r^{d+1})^{\lfloor \log_{r^{d+1}} N' \rfloor \frac{2d}{d+1} + \epsilon}\right) \\ \text{3354} \quad &= O\left((r^{d+1})^{\lfloor \log_{r^{d+1}} N' \rfloor \frac{2d}{d+1} + \epsilon}\right) \\ \text{3355} \quad &= O\left((r^{d+1})^{\log_{r^{d+1}} N' \frac{2d}{d+1} + \epsilon}\right) \\ \text{3356} \quad &= O\left(N'^{\frac{2d}{d+1} + \epsilon}\right). \\ \text{3357} \end{aligned}$$

3358 **Finally** We can now bound $T_k(N)$ using the upper bound for $T(N')$,

$$\text{3359} \quad T_k(N) \leq \max_{b^{k_i} N_i \leq b^k N} T_{k_i}(N_i) = T(b^k N) = O\left((b^k N)^{\frac{2d}{d+1} + \epsilon}\right) = 2^{O(k)} N^{\frac{2d}{d+1} + \epsilon}.$$

3360 Hence, $T_k(N) = 2^{O(k)} N^{\frac{2d}{d+1} + \epsilon_r}$, and since $d = \binom{\Delta+2}{2} - 1$, we have

$$\text{3361} \quad T_k(N) = 2^{O(k)} N^{2 - \frac{4}{\Delta^2 + 3\Delta + 2} + \epsilon_r}.$$

3362 To that complexity we add a constant time unit for each output pair in
3363 steps **1** and **2**. \square

3364 B.3 Applications

3365 To illustrate the expressive power of 3POL, we give some geometric
3366 applications in the sections that follow. We show the following:

- 3367 1. GPT can be solved in subquadratic time provided the input points lie
3368 on few parameterized constant-degree polynomial curves.

- 3369 2. In the plane, given three sets C_i of n unit circles and three points p_i
 3370 such that a circle $c \in C_i$ contains p_i , deciding whether there exists
 3371 $(a, b, c) \in C_1 \times C_2 \times C_3$ such that $a \cap b \cap c \neq \emptyset$ can be done in
 3372 subquadratic time.
- 3373 3. Given n input points in the plane, deciding whether any triple spans
 3374 a unit triangle can be done in subquadratic time, provided the input
 3375 points lie on few parameterized constant-degree polynomial curves.

3376 **B.3.1 GPT for Points on Curves**

3377 The following is a corollary of Theorem 2.16 in Raz, Sharir and de
 3378 Zeeuw [133]

3379 **Corollary B.14** (Raz, Sharir and de Zeeuw [133]). *Any n points on an irre-*
 3380 *ducible algebraic curve of degree d in \mathbb{C}^2 determine $\tilde{O}_d(n^{11/6})$ proper collinear*
 3381 *triples, unless the curve is a line or a cubic.*

3382 An interesting application of our results is the existence of subquadratic
 3383 nonuniform and uniform algorithms for the computational version of this
 3384 corollary.

3385 **Problem 20** (GPT on curves). Let C_1, C_2, C_3 be three (not necessarily
 3386 distinct) parameterized constant-degree polynomial curves in \mathbb{R}^2 , so that
 3387 each C_i can be written $(g_i(t), h_i(t))$ for some polynomials of constant degree
 3388 g_i, h_i . Given three n -sets $S_1 \subset C_1, S_2 \subset C_2, S_3 \subset C_3$, decide whether there
 3389 exist any collinear triple of points in $S_1 \times S_2 \times S_3$.

3390 **Theorem B.15.** *GPT on curves reduces linearly to 3POL.*

3391 *Proof.* For each set S_i , construct the set $T_i = \{t : p \in S_i, p = (g_i(t), h_i(t))\}$.
 3392 Testing whether there exists a collinear triple in $S_1 \times S_2 \times S_3$ amounts to
 3393 testing whether any determinant

$$\begin{vmatrix} g_1(t_1) & h_1(t_1) & 1 \\ g_2(t_2) & h_2(t_2) & 1 \\ g_3(t_3) & h_3(t_3) & 1 \end{vmatrix}$$

3395 equals zero. This determinant is a trivariate constant-degree polynomial in
 3396 $\mathbb{R}[t_1, t_2, t_3]$. Solving the original problem amounts thus to deciding whether
 3397 this polynomial cancels for any triple $(t_1, t_2, t_3) \in T_1 \times T_2 \times T_3$. \square

3398 Note that a similar polynomial predicate exists for testing collinearity in
 3399 higher dimension.

3400 **Lemma B.16.** *Let $p = (p_1, p_2, \dots, p_d)$, $q = (q_1, q_2, \dots, q_d)$, and $r =$*
 3401 *(r_1, r_2, \dots, r_d) be three points in \mathbb{R}^d , then p , q , and r are collinear if and*
 3402 *only if*

$$3403 \quad \left[\sum_{i=1}^d (p_i - r_i)(q_i - p_i) \right]^2 - \left[\sum_{i=1}^d (p_i - r_i)^2 \right] \left[\sum_{i=1}^d (q_i - p_i)^2 \right] = 0.$$

3404 *Proof.* Let $a = (p_1, p_2, \dots, p_d)$, $b = (q_1, q_2, \dots, q_d)$, and $c = (r_1, r_2, \dots, r_d)$
 3405 be three points in \mathbb{R}^d . The points p , q , and r are collinear if and only if
 3406 $r = p + \lambda(q - p)$ for some unique $\lambda \in \mathbb{R}$, that is

$$\begin{aligned} 3407 \quad & (p - r) + \lambda(q - p) = 0, \\ 3408 \quad \Rightarrow \quad & \forall i \in [d]: (p_i - r_i) + \lambda(q_i - p_i) = 0, \\ 3409 \quad \Rightarrow \quad & \sum_{i=1}^d [(p_i - r_i) + \lambda(q_i - p_i)]^2 = 0, \\ 3410 \quad \Rightarrow \quad & \sum_{i=1}^d \left[(q_i - p_i)^2 \lambda^2 + 2(p_i - r_i)(q_i - p_i)\lambda + (p_i - r_i)^2 \right] = 0, \\ 3411 \quad \Rightarrow \quad & \underbrace{\sum_{i=1}^d (q_i - p_i)^2}_{A} \lambda^2 + \underbrace{\left[2 \sum_{i=1}^d (p_i - r_i)(q_i - p_i) \right]}_{B} \lambda + \underbrace{\sum_{i=1}^d (p_i - r_i)^2}_{C} = 0, \\ 3412 \quad \Rightarrow \quad & \lambda = \frac{-B \pm \sqrt{B^2 - 4AC}}{2A}. \\ 3413 \end{aligned}$$

3414 For λ to exist and be unique $B^2 - 4AC$ must be zero. Hence, p , q , and
 3415 r are collinear if and only if

$$3416 \quad \left[2 \sum_{i=1}^d (p_i - r_i)(q_i - p_i) \right]^2 - 4 \left[\sum_{i=1}^d (p_i - r_i)^2 \right] \left[\sum_{i=1}^d (q_i - p_i)^2 \right] = 0.$$

3417 \square

3418 Moreover, the improvement that we obtain in the time complexity of
 3419 3POL can be exploited to boost the number of curves we pick the points
 3420 from.

3421 **Theorem B.17.** Let C_1, C_2, \dots, C_k be $k = o\left((\log n)^{\frac{1}{6}}/(\log \log n)^{\frac{1}{2}}\right)$ (not
3422 necessarily distinct) constant-degree polynomial curves in \mathbb{R}^d . Given k n -sets
3423 $S_{i_j} \subset C_{i_j}$, deciding whether there exists any collinear triple of points in any
3424 triple of sets $S_{i_1} \times S_{i_2} \times S_{i_3}$ can be solved in subquadratic time.

3425 *Proof.* Solve a 3POL instance for each choice of $S_{i_1} \times S_{i_2} \times S_{i_3}$. Since there
3426 are $o\left((\log n)^{\frac{1}{2}}/(\log \log n)^{\frac{3}{2}}\right)$ such choices, the theorem follows. \square

3427 B.3.2 Incidences on Unit Circles

3428 Raz, Sharir and Solymosi [137] mention the following problem as a special
3429 case of the framework they introduce. Let p_1, p_2, p_3 be three distinct points
3430 in the plane, and, for $i = 1, 2, 3$, let \mathcal{C}_i be a family of n unit circles (a circle
3431 of radius 1) that pass through p_i . Their goal is to obtain an upper bound
3432 on the number of *triple points*, which are points that are incident to a circle
3433 of each family. They prove:

3434 **Theorem B.18.** Let p_1, p_2, p_3 be three distinct points in the plane, and, for
3435 $i = 1, 2, 3$, let \mathcal{C}_i be a family of n unit circles that pass through p_i . Then the
3436 number of points incident to a circle of each family is $O(n^{11/6})$.

3437 They observe that the following dual formulation is equivalent to their
3438 original problem:

3439 **Theorem B.19.** Let C_1, C_2, C_3 be three unit circles in \mathbb{R}^2 , and, for each
3440 $i = 1, 2, 3$, let S_i be a set of n points lying on C_i . Then the number of unit
3441 circles, spanned by triples of points in $S_1 \times S_2 \times S_3$, is $O(n^{11/6})$.

3442 Our new algorithms indeed allow us to solve the decision version of their
3443 problems in subquadratic time.

3444 **Problem 21.** Let C_1, C_2, C_3 be three unit circles in \mathbb{R}^2 with centers c_i , and,
3445 for each $i = 1, 2, 3$, let $S_i = \{(x_{i,1}, y_{i,1}), (x_{i,2}, y_{i,2}), \dots, (x_{i,n}, y_{i,n})\}$ be a set
3446 of n points lying on C_i . Decide whether any triple of points $(p_1, p_2, p_3) \in$
3447 $S_1 \times S_2 \times S_3$ spans a unit circle.

3448 **Theorem B.20.** Problem 21 can be solved in $O(n^2(\log \log n)^{\frac{3}{2}}/(\log n)^{\frac{1}{2}})$
3449 time.

3450 *Proof.* Without loss of generality, assume all input points lie on the right
 3451 y -monotone arc of their respective circle. All other seven cases can be
 3452 handled similarly. We can also assume that no input point is the top or
 3453 bottom vertex of its circle, rotating the plane if necessary.

3454 Given three points p_1, p_2 , and p_3 , let

3455 $x = \|p_1 - p_2\|, X = x^2, y = \|p_1 - p_3\|, Y = y^2, z = \|p_2 - p_3\|, Z = z^2.$

3456 Testing if the three points p_1, p_2 , and p_3 span a unit circle amounts to
 3457 testing whether

3458
$$X^2 + Y^2 + Z^2 - 2XY - 2XZ - 2YZ + XYZ = 0.$$

3459 The fact that the input points lie on the right y -monotone arc of unit
 3460 circles of centers c_1, c_2, c_3 allows us to get down to a single variable per
 3461 point. Let $c_i = (c_i^x, c_i^y)$ and $t_{i,j} = \sqrt{\frac{1-x_{i,j}+c_i^x}{1+x_{i,j}-c_i^x}}$. Then the j th input point of
 3462 the i th circle can be expressed as

3463
$$p_{i,j} = (x_{i,j}, y_{i,j}) = c_i + \left(\frac{1-t_{i,j}^2}{1+t_{i,j}^2}, \frac{2t_{i,j}}{1+t_{i,j}^2} \right).$$

3464 Combining those two observations with some algebraic manipulations,
 3465 one can show that there exists some trivariate polynomial F of degree at
 3466 most 24 that cancels on t_1, t_2, t_3 when the points $c_1 + \left(\frac{1-t_1^2}{1+t_1^2}, \frac{2t_1}{1+t_1^2} \right)$, $c_2 +$
 3467 $\left(\frac{1-t_2^2}{1+t_2^2}, \frac{2t_2}{1+t_2^2} \right)$, and $c_3 + \left(\frac{1-t_3^2}{1+t_3^2}, \frac{2t_3}{1+t_3^2} \right)$ span a unit circle.

3468 Hence, the polynomial F together with the sets $\{t_{1,1}, t_{1,2}, \dots, t_{1,n}\}$,
 3469 $\{t_{2,1}, t_{2,2}, \dots, t_{2,n}\}$, and $\{t_{3,1}, t_{3,2}, \dots, t_{3,n}\}$ give an instance of 3POL we
 3470 can solve in subquadratic time with our new algorithms.

3471 Unfortunately, the computation $\sqrt{\cdot}$ is not allowed in our model, and so,
 3472 we cannot compute $t_{i,j}$. However, we can generalize the 3POL problem to
 3473 make it fit:

3474 **Problem 22** (Modified 3POL). Let $F \in \mathbb{R}[x, y, z]$ be a trivariate polyno-
 3475 mial of constant degree, given three sets A, B , and C , each containing n
 3476 real numbers, decide whether there exist $a \in A, b \in B$, and $c \in C$ such that

3477
$$\exists t_1, t_2, t_3 (t_1^2 = a \wedge t_2^2 = b \wedge t_3^2 = c \wedge F(t_1, t_2, t_3) = 0).$$

3478 Hence, the polynomial F together with the sets $\{t_{1,1}^2, t_{1,2}^2, \dots, t_{1,n}^2\}$,
 3479 $\{t_{2,1}^2, t_{2,2}^2, \dots, t_{2,n}^2\}$, and $\{t_{3,1}^2, t_{3,2}^2, \dots, t_{3,n}^2\}$ give an instance of this variant
 3480 of 3POL. Note that those sets are computable in our models.

3481 We can tweak our algorithms so that they work for this new version of
 3482 3POL. We prefix each decision we make on the first-order theory of the
 3483 reals with an existential quantifier and a condition of the type $t_i^2 = x$, with
 3484 x the square of t_i , when we reference t_i in the formula we test. This new
 3485 algorithm answers positively if and only if the original problem contains a
 3486 triple of points spanning a unit circle.

3487 In general, any constant-degree polynomial curve can be decomposed
 3488 in a constant number of pieces as above. Each point on this curve can be
 3489 given a parameterization that might involve roots of its coordinates. Those
 3490 can be taken care of by appropriately augmenting the Tarski sentences in
 3491 our algorithm with equations that encode those roots for free. \square

3492 B.3.3 Points Spanning Unit Triangles

3493 A similar problem, namely counting the number of input point triples
 3494 spanning an area S triangle (provided they lie on a few curves), can also
 3495 easily be reduced to 3POL. The polynomial to look at in this case is

$$3496 \quad F(x, y, z) = X^2 + Y^2 + Z^2 - 2XY - 2XZ - 2YZ + 16S^2.$$

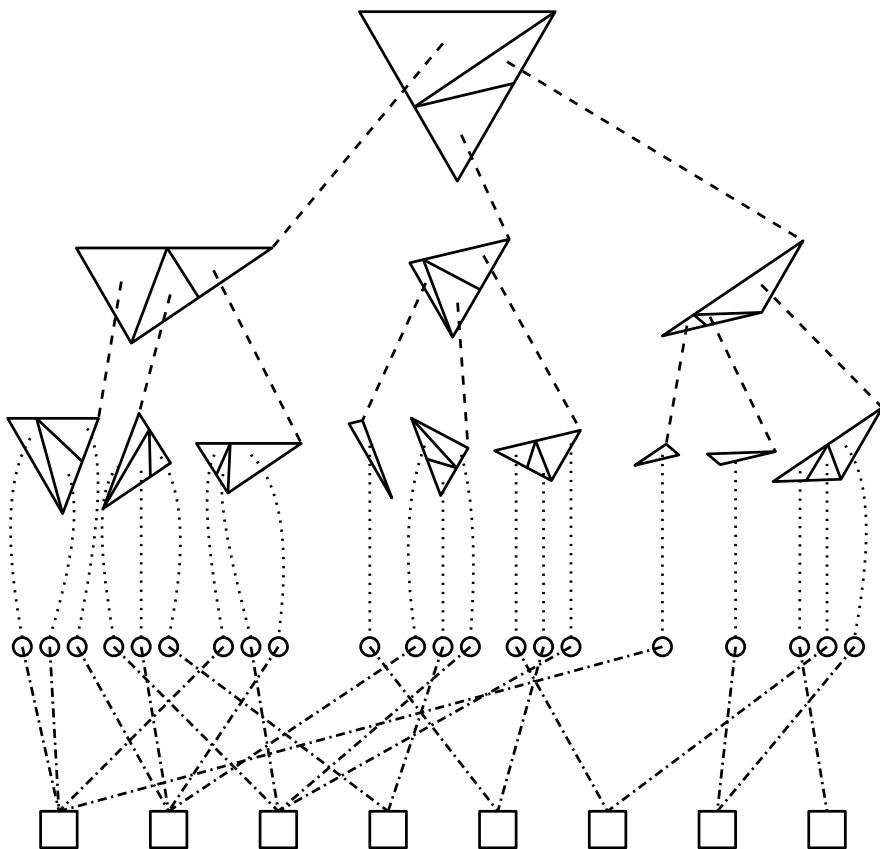
3497 Note that when the input points lie in the plane, the number of solutions
 3498 is more than quadratic [132, 135].

IV

3499

3500

Data Structures



C

3501

Subquadratic Encodings for Point Configurations

3502 with Jean Cardinal, Timothy Chan, John Iacono, and Stefan Langerman

3503 For many algorithms dealing with sets of points in the plane, the only
3504 relevant information carried by the input is the combinatorial configuration
3505 of the points: the orientation of each triple of points in the set (clockwise,
3506 counterclockwise, or collinear). This information is called the *order type* of
3507 the point set. In the dual, realizable order types and abstract order types are
3508 combinatorial analogues of line arrangements and pseudoline arrangements.
3509 Too often in the literature we analyze algorithms in the real-RAM model for
3510 simplicity, putting aside the fact that computers as we know them cannot
3511 handle arbitrary real numbers without some sort of encoding. Encoding an
3512 order type by the integer coordinates of a realizing point set is known to
3513 yield doubly exponential coordinates in some cases. Other known encodings
3514 can achieve quadratic space or fast orientation queries, but not both. In
3515 §C.1, we give a compact encoding for abstract order types that allows an
3516 efficient query of the orientation of any triple: the encoding uses $O(n^2)$
3517 bits and an orientation query takes $O(\log n)$ time in the word-RAM model
3518 with word size $w \geq \log n$. This encoding is space-optimal for abstract order
3519 types. We show how to shorten the encoding to $O(n^2(\log \log n)^2 / \log n)$
3520 bits for realizable order types, giving the first subquadratic encoding for
3521 those order types with fast orientation queries. In §C.2, we further refine
3522 our encoding to attain $O(\log n / \log \log n)$ query time at the expense of a
3523 negligibly larger space requirement. In the realizable case, we show that
3524 negligibly larger space requirement. In the realizable case, we show that

3526 all those encodings can be computed efficiently. In §C.3, we generalize our
 3527 results to the encoding of point configurations in higher dimension.

3528 **C.1 Encoding Order Types via Hierarchical
 3529 Cuttings**

3530 To make our statements clear, we use the concept of an encoding. We
 3531 recall its definition

3532 **Definition 2.** For fixed k and given a function $f : [n]^k \rightarrow [O(1)]$, we define
 3533 a $(S(n), Q(n))$ -encoding of f to be a string of $S(n)$ bits such that, given this
 3534 string and any $i \in [n]^k$, we can compute $f(i)$ in $Q(n)$ time in the word-RAM
 3535 model with word size $w \geq \log n$.

3536 In this section, we use this definition with f being some order type,¹
 3537 $k = 3$ and the codomain of f being $\{ -, 0, + \}$. For the rest of the discussion,
 3538 we assume the word-RAM model with word size $w \geq \log n$ and the standard
 3539 arithmetic and bitwise operators. We prove our main theorems for the
 3540 two-dimensional case:

3541 **Contribution 7.** *All abstract order types have an encoding using $O(n^2)$
 3542 bits of space and allowing for queries in $O(\log n)$ time.*

3543 **Contribution 8.** *All realizable order types have a $O(\frac{n^2(\log \log n)^2}{\log n})$ -bits en-
 3544 coding allowing for queries in $O(\log n)$ time.*

3545 **Contribution 17.** *In the real-RAM model and the constant-degree algebraic
 3546 decision tree model, given n real-coordinate input points in \mathbb{R}^2 we can compute
 3547 the encoding of their order type as in Theorems 7 and 8 in $O(n^2)$ time.*

3548 For instance, Theorem 8 implies that for any set of points $\{ p_1, \dots, p_n \}$,
 3549 there exists a string of $O(n^2(\log \log n)^2 / \log n)$ bits such that given this
 3550 string and any triple of indices $(a, b, c) \in [n]^3$ we can compute the value of
 3551 $\chi(a, b, c) = \nabla(p_a, p_b, p_c)$ in $O(\log n)$ time.

3552 Throughout the rest of this paper, we assume that we can access some
 3553 arrangement of pairwise intersecting lines or pseudolines that realizes the

¹Technically, we encode the orientation predicate of some realizing arrangement of the order type and skip the isomorphism. If desired, a canonical labeling of the arrangement can be produced in $O(n^2)$ time for abstract and realizable order types (see Lemma 7.1).

order type we want to encode. We thus exclusively focus on the problem of encoding the order type of a given arrangement. This does not pose a threat against the existence of an encoding. However, we have to be more careful when we bound the preprocessing time required to compute such an encoding. This is why, in Theorem 17, we specify the model of computation and how the input is given.

Idea We want to preprocess n pseudolines $\{p'_1, p'_2, \dots, p'_n\}$ in the plane so that, given three indices a , b , and c , we can compute their orientation, that is, whether the intersection $p'_a \cap p'_b$ lies above, below or on p'_c . Our data structure builds on cuttings as follows: Given a cutting Ξ and the three indices, we can locate the intersection of p'_a and p'_b with respect to Ξ . The location of this intersection is a cell of Ξ . The next step is to decide whether p'_c lies above, lies below, contains or intersects that cell. In the first three cases, we are done. Otherwise, we can answer the query by recursing on the subset of pseudolines intersecting the cell containing the intersection. We build on hierarchical cuttings to control the size of each such subproblem.

Intersection Location When the arrangement consists of straight lines, locating the intersection $p'_a \cap p'_b$ in Ξ is trivial if we know the real parameters of p'_a and p'_b and of the descriptions of the subcells of Ξ . However, in our model we are not allowed to store real numbers. To circumvent this annoyance, and to handle arrangements of pseudolines, we make an observation illustrated by Figure C.1.

Definition 14. Given a set of pseudolines, a pseudocircle is a simple closed curve such that each pseudoline properly intersects it in exactly two points. A pseudodisk is a region bounded by a pseudocircle.

Observation C.1. *Two pseudolines p'_a and p'_b intersect in the interior of a pseudodisk \mathcal{C} if and only if the intersections of p'_a and p'_b with \mathcal{C} alternate on the boundary of \mathcal{C} .*

We construct Ξ so that its cells are pseudodisks for the pseudolines that intersect them. Hence, this observation gives us a way to encode the location of the intersection of p'_a and p'_b in Ξ using only bits. We formalize this observation with the following definition:

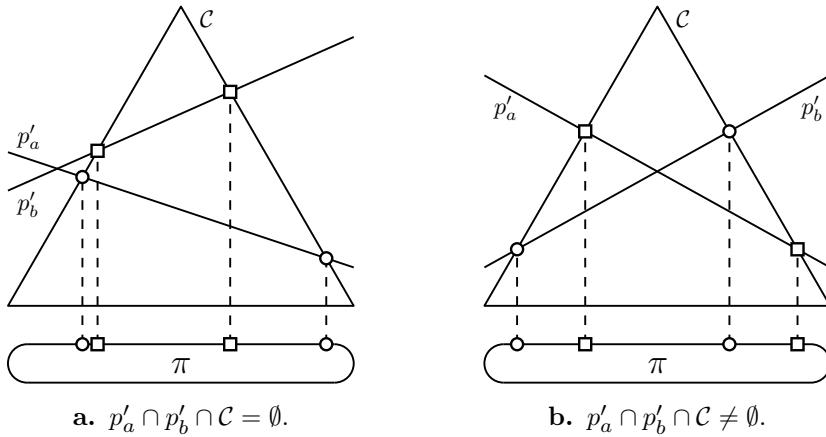


Figure C.1. Cyclic permutations (π).

3586 **Definition 15** (Cyclic Permutation). The *cyclic permutation* of a full-
 3587 dimensional cell of a cutting is the cyclic ordering of the pseudolines crossing
 3588 its boundary.

3589 For an arrangement of pseudolines, we use the standard vertical decompo-
 3590 sition (§6.4.2) to construct a hierarchical cutting (§8.2). This decompo-
 3591 sition partitions the space into trapezoidal cells. For an arrangement of lines, we
 3592 can use the standard bottom-vertex triangulation (§6.4.1) instead, which
 3593 allows us to generalize our results to higher dimensions in §C.3. In the plane,
 3594 the bottom-vertex triangulation partitions the space into triangular cells.
 3595 Note that non-full-dimensional cells are easier to encode. For a 0-dimensional
 3596 cell and a pseudoline, we store whether the pseudoline lies above, lies below,
 3597 or contains the 0-dimensional cell. For a 1-dimensional cell, a pseudoline
 3598 could also intersect the interior of the cell, but in only one point. The
 3599 pseudolines intersecting that cell define an (acyclic) permutation with poten-
 3600 tially several intersections at the same position. This information suffices to
 3601 answer location queries for those cells, and the space taken is not more than
 3602 that necessary for full-dimensional cells. When two pseudolines intersect
 3603 in a 1-dimensional cell or contain the same 0-dimensional cell, they appear
 3604 simultaneously in the cyclic permutation of an adjacent 2-dimensional cell if
 3605 they intersect its interior. If that is the case, the location of the intersection
 3606 of those two pseudolines in the cutting is the non-full-dimensional cell. A
 3607 constant number of bits can be added to the encoding each time we need to

3608 know the dimension of the cell we encode.

3609 **Encoding** We encode the order type of an arrangement via hierarchical
 3610 cuttings as defined in [48] (see §8.2). Given n pseudolines in the plane and
 3611 some fixed parameters r and ℓ , compute a ℓ -levels hierarchical cutting of
 3612 parameter r for those pseudolines as in Lemma 8.6. This hierarchical cutting
 3613 consists of ℓ levels labeled $0, 1, \dots, \ell - 1$. Level i has $O(r^{2i})$ cells. Each of
 3614 those cells is further partitioned into $O(r^2)$ subcells. The $O(r^{2(i+1)})$ subcells
 3615 of level i are the cells of level $i + 1$. Each cell of level i is intersected by at
 3616 most $\frac{n}{r^i}$ pseudolines, and hence each subcell is intersected by at most $\frac{n}{r^{i+1}}$
 3617 pseudolines.

3618 We compute and store a combinatorial representation of the hierarchical
 3619 cutting as follows: For each level of the hierarchy, for each cell in that level,
 3620 for each pseudoline intersecting that cell, for each subcell of that cell, we
 3621 store two bits to indicate the location of the pseudoline with respect to
 3622 that subcell, that is, whether the pseudoline lies above (00), lies below (01),
 3623 intersects the interior of that subcell (10), or contains the subcell (11). When
 3624 a pseudoline intersects the interior of a 2-dimensional subcell, we also store
 3625 the two indices of that pseudoline in the cyclic permutation of that subcell,
 3626 beginning at an arbitrary location in, say, clockwise order. If the intersected
 3627 subcell is 1-dimensional instead, we store the index of the pseudoline in the
 3628 acyclic permutation of that subcell, beginning at an arbitrary endpoint. If
 3629 two pseudolines intersect in the interior of a 1-dimensional subcell or on
 3630 the boundary of a 2-dimensional subcell, they share the same index in the
 3631 permutation of that subcell.

3632 This representation takes $O(\frac{n}{r^i} + \frac{n}{r^{i+1}} \log \frac{n}{r^{i+1}})$ bits per subcell of level i by
 3633 storing for each pseudoline its location and, when needed, the permutation
 3634 indices of its intersections with the subcell. At the last level of the hierarchy,
 3635 let $t = \frac{n}{r^\ell}$ denote an upper bound on the number of pseudolines intersecting
 3636 each subcell. For each of those $O(r^{2\ell}) = O(\frac{n^2}{t^2})$ subcells we store a pointer to
 3637 a lookup table of size $O(t^3)$ that allows to answer the query of the orientation
 3638 of any triple of pseudolines intersecting that subcell.

3639 Storing the permutation at each subcell would suffice to answer all queries
 3640 that do not reach the last level of the hierarchy. However, for those queries to
 3641 be answered efficiently, we need to have access to all bits belonging to a given
 3642 pseudoline without having to read the bits of the others. One solution is to

00	10	010	110	10	101	111	01	11	10	000	011	00	01
----	----	-----	-----	----	-----	-----	----	----	----	-----	-----	----	----

Figure C.2. A trace $\text{TR}(\mathcal{C}, p')$. The cell \mathcal{C} has eight subcells. Each subcell is intersected by at most four pseudolines. The pseudoline p' lies above two of them, lies below two of them, contains one of them, and intersects three of them at indices $(2, 6)$, $(5, 7)$, and $(0, 3)$.

3643 augment each subcell with a hash table that translates pseudoline indices of
 3644 the parent cell into pseudoline indices of the subcell. Another cleaner solution
 3645 is to use the Zone Theorem (Theorem 6.2): by constructing the hierarchical
 3646 cutting via decompositions of subsets of the input pseudolines, we can bound
 3647 the number of subcells of a given cell a given pseudoline intersects by $O(r)$.²
 3648 Hence, the number of bits stored for a single intersecting cell-pseudoline pair
 3649 at level i is bounded by $|\text{TR}_i| = O(r^2 + r \log \frac{n}{r^{i+1}})$. This bound allows us to
 3650 store all bits belonging to a given cell-pseudoline pair (\mathcal{C}, p') in a contiguous
 3651 block of memory, denoted by $\text{TR}(\mathcal{C}, p')$, whose address in memory is easy
 3652 to compute (as detailed later on). The overall number of bits stored stays
 3653 the same up to a constant factor. We call $\text{TR}(\mathcal{C}, p')$ the *trace* of p' in \mathcal{C} .
 3654 Figure C.2 depicts an example trace.

3655 For queries that reach the last level of the hierarchy, storing an individual
 3656 lookup table for each leaf would cost too much as soon as $t = \omega(1)$. However,
 3657 as long as t is small enough, each order type is shared by many leaves, and we
 3658 can thus save space. Formally, let $\nu(t)$ denote the number of order types of
 3659 size t , which is $\nu(t) = 2^{\Theta(t^2)}$ for abstract order types [73] and $\nu(t) = 2^{\Theta(t \log t)}$
 3660 for realizable order types [16, 87]. At most $\nu(t)$ distinct lookup tables are
 3661 needed to answer the queries on the subcells of the last level of the hierarchy.
 3662 Hence, the pointers have size $|\text{POINTER}| = O(\log \nu(t))$ and the total space
 3663 needed for the lookup tables is $O(t^3 \nu(t))$. For each leaf, we store a canonical
 3664 labeling of size $|\text{LABELING}| = O(t \log t)$ on the pseudolines that intersect it,
 3665 as in Lemma 7.1. We use that labeling to order the queries in the associated
 3666 lookup table.³

²A reason to prefer this solution is that it enables the query time reduction in the next section.

³Note that the use of this canonical labeling is not necessary if t is constant, because then we can afford to use one lookup table per leaf. For superconstant t , the canonical labeling is necessary to get the construction time down to $O(n^2)$ as explained later. Space

3667 **Layout** For completeness, we detail precisely how bits of the encoding are
 3668 laid out in memory to allow an efficient decoding. Indeed, the data structure
 3669 we encode is a tree and many space-efficient layouts exists for those when
 3670 their nodes each have the same size. Here however, node size shrinks as
 3671 we go down the hierarchy. We spend a few paragraphs showing it is still
 3672 possible to address all components in a time- and space-efficient way. As
 3673 this is fairly straightforward to adapt for the encodings in §C.2 and §C.3,
 3674 we do not give the details for those sections.

3675 The encoding is the concatenation of the parameters n , r , and t , the
 3676 cells of the hierarchy, and the lookup tables. We order the cells of the
 3677 hierarchy in a depth-first manner: a cell of level i is denoted by $\mathcal{C}_{0,j_1,j_2,\dots,j_i}$
 3678 with $j_1, j_2, \dots, j_i \in \{0, 1, \dots, O(r^2)\}$. The root cell is \mathcal{C}_0 and the cell
 3679 $\mathcal{C}_{0,j_1,j_2,\dots,j_i,j_{i+1}}$ is the $(j_{i+1} + 1)$ -th subcell of the cell $\mathcal{C}_{0,j_1,j_2,\dots,j_i}$. Cells are
 3680 then ordered lexicographically. For each leaf cell we store its pointer and
 3681 canonical labeling. For each internal cell \mathcal{C} we store the traces $\text{TR}(\mathcal{C}, p')$
 3682 in a certain permutation of the pseudolines p' . To order the pseudolines,
 3683 we use the first index of each pseudoline in the cyclic permutation of the
 3684 cell (for the root cell \mathcal{C}_0 this is the index given by the input permutation).
 3685 Note that those indices are between 0 and $2\lfloor \frac{N}{r^i} \rfloor - 1$ for a cell of level i . We
 3686 allocate twice the required space for traces and canonical labelings to avoid
 3687 defining a mapping between the indices of a cell and the indices of each of
 3688 its subcells.

3689 For the root cell of the hierarchy \mathcal{C}_0 representing the entire space and
 3690 containing all the intersections of the arrangement, $\text{ADDR}(\mathcal{C}_0)$ is the first
 3691 free address after the encoding of the parameters n , r , and t .

3692 The address $\text{ADDR}(\mathcal{C}_{0,j_1,j_2,\dots,j_{i-1},0})$ of the first subcell of $\mathcal{C}_{0,j_1,j_2,\dots,j_{i-1}}$ is
 3693 offset by twice the space taken by the traces for that cell

$$3694 \quad \text{ADDR}(\mathcal{C}_{0,j_1,j_2,\dots,j_{i-1},0}) = \underbrace{\text{ADDR}(\mathcal{C}_{0,j_1,j_2,\dots,j_{i-1}})}_{\text{Address of the parent cell}} + \underbrace{2 \left\lfloor \frac{n}{r^{i-1}} \right\rfloor |\text{TR}_{i-1}|}_{\text{Traces of the parent cell}} .$$

3695 Let c_h be the constant hidden by the $O(r^2)$ of the hierarchical cutting.

and query complexity are not affected by this design choice (other than constant factors).

3696 The space taken by a subtree rooted at a node of level i is bounded by
 3697

$$\begin{aligned}
 3698 |\text{SUBTREE}_i| = & \underbrace{2 \left\lfloor \frac{n}{r^i} \right\rfloor |\text{TR}_i|}_{\text{Traces at the root}} + \underbrace{2c_h \sum_{k=1}^{\ell-i-1} r^{2k} \left\lfloor \frac{n}{r^{i+k}} \right\rfloor |\text{TR}_{i+k}|}_{\text{Traces of the subtrees}} \\
 3699 & + \underbrace{c_h r^{2(\ell-i)} |\text{POINTER}|}_{\text{Pointers}} + \underbrace{2c_h r^{2(\ell-i)} |\text{LABELING}|}_{\text{Canonical labelings}}. \\
 3700
 \end{aligned}$$

3701 The address of any other subcell $\text{ADDR}(\mathcal{C}_{0,j_1,j_2,\dots,j_{i-1},j_i})$ is offset by the
 3702 space taken by the subtrees of its siblings $0, 1, \dots, j_i - 1$

$$3703 \text{ADDR}(\mathcal{C}_{0,j_1,j_2,\dots,j_{i-1},j_i}) = \underbrace{\text{ADDR}(\mathcal{C}_{0,j_1,j_2,\dots,j_{i-1},0})}_{\text{Address of the first sibling}} + \underbrace{j_i \cdot |\text{SUBTREE}_i|}_{\text{Left siblings subtrees}}.$$

3704 The address of the trace $\text{TR}(\mathcal{C}, p'_a)$, where $0 \leq a \leq 2 \lfloor \frac{N}{r^i} \rfloor - 1$ is the first
 3705 intersection index of p'_a in the cyclic permutation of the level- i cell \mathcal{C} , is

$$3706 \text{ADDR}(\mathcal{C}, p'_a) = \text{ADDR}(\mathcal{C}) + a \cdot |\text{TR}_i|.$$

3707 Since we do not store any trace for the leaves, define $|\text{TR}_\ell| = 0$. The
 3708 pointer and canonical labeling of the leaf $\mathcal{C}_{0,j_1,j_2,\dots,j_\ell}$ are concatenated at pos-
 3709 ition $\text{ADDR}(\mathcal{C}_{0,j_1,j_2,\dots,j_\ell})$ and the lookup tables are concatenated at position
 3710 $\text{ADDR}(\mathcal{C}_0) + |\text{SUBTREE}_0|$.

3711 This layout makes traversing the hierarchy from root to leaf efficient:
 3712 The address of the root cell is discovered after parsing the parameters n ,
 3713 r , and t . The address of the first subcell of a parent cell is computed in
 3714 constant time from the address of the parent cell. The size of the hierarchy is
 3715 computed in time proportional its height. The size of a subtree of level $i + 1$
 3716 is derived in constant time from the size of a subtree of level i . The address
 3717 of any subcell of level i is computed in constant time from its index, the
 3718 address of the parent cell, and the size of a subtree of level i . The address
 3719 of a trace is computed in constant time from the address of its cell and the
 3720 index of its pseudoline. The address of the pointer and canonical labeling of
 3721 a leaf is the address of the corresponding cell. The address of a lookup table
 3722 is computed in constant time given the address of the root cell, the size of
 3723 the hierarchy, the size of a lookup table, and the leaf pointer.

3724 During a traversal, pseudoline indices of the parent cell are mapped to
 3725 indices in the subcell in constant time by using the first intersection index

3726 of each pseudoline in that subcell. A final index mapping happens when
 3727 translating leaf indices to lookup table indices using the canonical labeling
 3728 of the leaf.

3729 **Space Complexity** We first prove a general bound on the space taken
 3730 by the encoding of a ℓ -level hierarchical cutting of parameter $r \geq 2$. For the
 3731 space taken by the lookup tables, their associated pointers and canonical
 3732 labelings at the leaves, and the parameters of the hierarchy n, r and t , the
 3733 analysis is immediate.

3734 Let $H_r^\ell(n) \in \mathbb{N}$ be the maximum amount of space (bits), over all ar-
 3735 rangements of n pseudolines, taken by the $\ell \in \mathbb{N}$ levels of a hierarchy with
 3736 parameter $r \in (1, +\infty)$.

3737 **Lemma C.2.** *For $r \geq 2$ and $t = \frac{n}{r^\ell}$ we have*

$$3738 \quad H_r^\ell(n) = O\left(\frac{n^2}{t}(\log t + r)\right).$$

3739 *Proof.* By definition, summing over all subcells, we have

$$3740 \quad H_r^\ell(n) = O\left(\sum_{i=0}^{\ell-1} \left(r^{2i} \cdot r^2 \cdot \left(\frac{n}{r^i} + \frac{n}{r^{i+1}} \log \frac{n}{r^{i+1}}\right)\right)\right).$$

3741 Note that we obtain the same bound by summing over all traces (of inter-
 3742 secting cell-pseudoline pairs)

$$3743 \quad H_r^\ell(n) = O\left(n \sum_{i=0}^{\ell-1} \left(r^i \cdot \left(r^2 + r \log \frac{n}{r^{i+1}}\right)\right)\right).$$

3744 Multiply any of the previous equations by $\frac{n}{tr^\ell} = 1$ to obtain

$$3745 \quad H_r^\ell(n) = O\left(\frac{n^2}{t} \sum_{i=0}^{\ell-1} \left(\frac{1}{r^{\ell-i-1}} \cdot \left(r + \log \frac{n}{r^{i+1}}\right)\right)\right).$$

3746 We use the equivalence $\frac{n}{r^{i+1}} = tr^{\ell-i-1}$ to replace the last term in the pre-
 3747 vious equation

$$3748 \quad H_r^\ell(n) = O\left(\frac{n^2}{t} \sum_{i=0}^{\ell-1} \left(\frac{1}{r^{\ell-i-1}} \cdot (r + \log t + (\ell - i - 1) \log r)\right)\right).$$

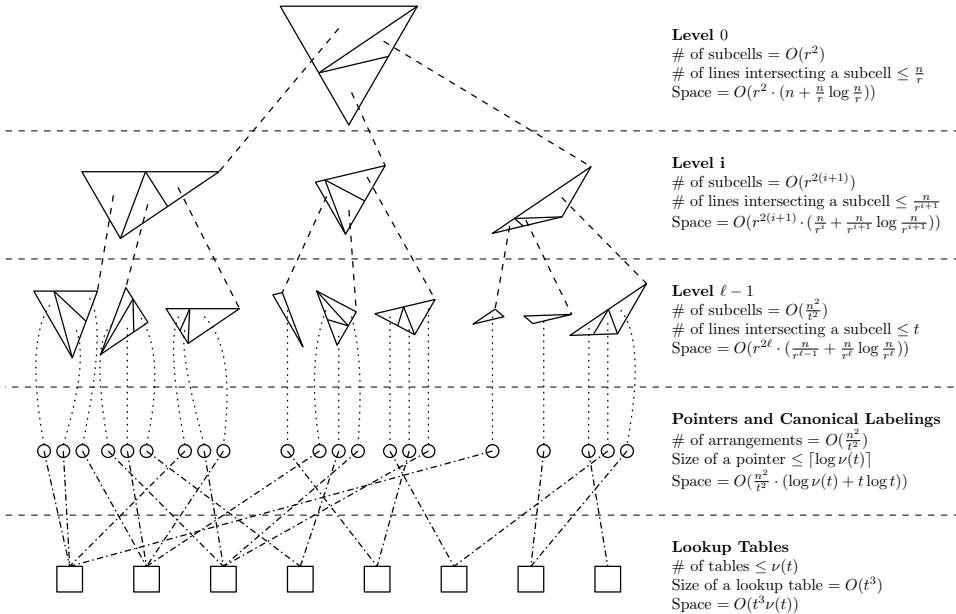


Figure C.3. Space analysis.

3749 We reverse the summation by redefining $i \leftarrow \ell - i - 1$ and group the terms

$$3750 \quad H_r^\ell(n) = O\left(\frac{n^2}{t} \left((\log t + r) \sum_{i=0}^{\ell-1} \frac{1}{r^i} + \log r \sum_{i=0}^{\ell-1} \frac{i}{r^i} \right)\right).$$

3751 Using the following inequalities (the geometric series and a multiple of its
3752 derivative):

$$3753 \quad \sum_{i=0}^k x^i \leq \frac{1}{1-x} \quad \text{and} \quad \sum_{i=0}^k ix^i \leq \frac{x}{(1-x)^2}, \quad \forall k \in \mathbb{N}, \forall x \in (0, 1),$$

3754 we conclude that

$$3755 \quad H_r^\ell(n) = O\left(\frac{n^2}{t} \left(\left(1 + \frac{1}{r-1}\right) (\log t + r) + \left(1 + \frac{2r-1}{r^2-2r+1}\right) \frac{\log r}{r} \right)\right),$$

3756 and the statement follows from $r \geq 2$. \square

3757 Figure C.3 sketches the different components of the encoding and shows
3758 the space taken by each of them. To that we must add the space taken by

3759 the parameters of the hierarchy n , r and t if those are not implicitly known
 3760 (here we assume the dimension $d = 2$ is implicitly known). We have thus
 3761 the following bound:

3762 **Lemma C.3.** *The space taken by the encoding described in §C.1 is proportional to*

$$3764 \quad \underbrace{\log ntr}_{\text{Parameters}} + \underbrace{\frac{n^2}{t}(\log t + r)}_{\text{Traces}} + \underbrace{\frac{n^2}{t^2}(\log \nu(t) + t \log t)}_{\text{Pointers and Canonical Labelings}} + \underbrace{t^3 \nu(t)}_{\text{Lookup Tables}} .$$

3765 We pick r constant for both abstract and realizable order types. We
 3766 have $\nu(t) = 2^{\Theta(n^2)}$ for abstract order types, hence we choose $t = \sqrt{\delta \log n}$
 3767 for small enough δ and the third term in Lemma C.3 dominates with n^2 .
 3768 Note how the quadratic bottleneck of this encoding is the storage of the
 3769 order type pointers at the leaves of the hierarchy. We have $\nu(t) = 2^{\Theta(n \log n)}$
 3770 for realizable order types, hence we choose $t = \delta \log n / \log \log n$ for small
 3771 enough δ and the second and third term in Lemma C.3 dominate with
 3772 $n^2(\log \log n)^2 / \log n$. This proves the space constraints in Theorems 7 and 8.

3773 **Correctness and Query Complexity** Given our encoding and three
 3774 pseudoline indices a, b, c we answer a query as follows: We start by decoding
 3775 the parameters n , r , and t . In our model, this can be done in $O(\log^* n +$
 3776 $\log^* r + \log^* t)$ time, where \log^* is the iterated logarithm (as in [109]).⁴ Let
 3777 $\mathcal{C} = \mathcal{C}_0$. First, find the subcell \mathcal{C}' of \mathcal{C} containing $p'_a \cap p'_b$ by testing for
 3778 each subcell whether p'_a and p'_b alternate in its cyclic permutation. This
 3779 can be done in $O(r^2)$ time by scanning $\text{TR}(\mathcal{C}, p'_a)$ and $\text{TR}(\mathcal{C}, p'_b)$ in parallel.
 3780 Note that non-full dimensional cells and subcells are easier to test. Next, if
 3781 p'_c does not properly intersect \mathcal{C}' , answer the query accordingly. If on the
 3782 other hand p'_c does properly intersect the subcell we recurse on \mathcal{C}' . This
 3783 can be tested by scanning $\text{TR}(\mathcal{C}, p'_c)$ in $O(r^2)$ time. Note that in case that

⁴Logarithmic space and constant decoding time is trivial when $w = \Theta(\log n)$. If w is too large, encode n in binary using $\lceil \log n + 1 \rceil$ bits, $\lceil \log n + 1 \rceil$ using $\lceil \log \lceil \log n + 1 \rceil + 1 \rceil$ bits, $\lceil \log \lceil \log n + 1 \rceil + 1 \rceil$ using $\lceil \log \lceil \log \lceil \log n + 1 \rceil + 1 \rceil + 1 \rceil$ bits, etc. until the number to encode is smaller than a constant which we encode in unary with 1's. Prepend a 1 to the largest number and 0 to all the others. Concatenate those numbers from smallest to largest. Total space is $O(\log n)$ bits and decoding n can be done in $O(\log^* n)$ time in the word-RAM model with $w \geq \log n$. As an alternative, logarithmic space and logarithmic decoding time is also trivially achievable with no constraint on w .

3784 the subcell is non-full-dimensional we can already answer the query. When
 3785 we reach the relative interior of a subcell of the last level of the hierarchy
 3786 without having found a satisfactory answer, we can answer the query by table
 3787 lookup in constant time. This works as long as each order type identifier
 3788 for at most t pseudolines fits in a constant number of words, which is the
 3789 case for the values of t we defined. The layout described earlier makes all
 3790 memory address computations of this query algorithm take constant time.
 3791 The total query time is thus proportional to $r^2 \log_r n$ in the worst case, which
 3792 is logarithmic since r is constant. This proves the query time constraints in
 3793 Theorems 7 and 8.

3794 With the hope of getting faster queries we could pick $r = \Theta(\log t)$ to
 3795 reduce the depth of the hierarchy, without changing the space requirements
 3796 by more than a constant factor. However, if no additional care is taken, this
 3797 would slow the queries down by a $\Theta(\log^2 t / \log \log t)$ factor because of the
 3798 scanning approach taken when locating the intersection $p'_a \cap p'_b$. We show
 3799 how to handle small but superconstant r properly in the next section.

3800 **Preprocessing Time** For a set of n points in the plane, or an arrangement
 3801 of n lines in the dual, we can construct the encoding of their order type in
 3802 quadratic time in the real-RAM and constant-degree algebraic computation
 3803 tree models. We prove Theorem 17.

3804 *Proof.* Using Lemma 8.6, a hierarchical cutting can be computed in $O(nr^\ell)$
 3805 time in the dual plane. All traces $\text{TR}(\mathcal{C}, p')$ can be computed from the
 3806 cutting in the same time. The lookup tables and leaf-table pointers can be
 3807 computed in $O(n^2 + t^3\nu(t))$ time as follows: For each subcell \mathcal{C} among the
 3808 $\frac{n^2}{t^2}$ subcells of the last level of the hierarchy, compute a canonical labeling
 3809 and representation of the lines intersecting \mathcal{C} in $O(t^2)$ time as in Lemma 7.1.
 3810 Insert the canonical representation in some trie in $O(t^2)$ time. If the canon-
 3811 ical representation was not already in the trie, create a lookup table with
 3812 the answers to all $O(t^3)$ queries on those lines and attach a pointer to that
 3813 table in the trie. This happens at most $\nu(t)$ times. In the encoding, store
 3814 the canonical labeling and this new pointer or the pointer that was already
 3815 in the trie for the subcell \mathcal{C} . All parts of the encoding can be concatenated
 3816 together in time proportional to the size of the encoding. \square

3817 C.2 Sublogarithmic Query Complexity

3818 We further refine the encoding introduced in the previous section so as
 3819 to reduce the query time by a $\log \log n$ factor. We do so using specificities of
 3820 the word-RAM model that allow us to preprocess computations on inputs of
 3821 small but superconstant size. This refinement is applicable to both abstract
 3822 and realizable order types, and leads to an improvement of our main theorems
 3823 for the two-dimensional case:

3824 **Contribution 9.** *All abstract order types have an encoding using $O(n^2)$
 3825 bits of space and allowing for queries in $O(\frac{\log n}{\log \log n})$ time.*

3826 **Contribution 10.** *All realizable order types have a $O(\frac{n^2 \log^\epsilon n}{\log n})$ -bits encoding
 3827 allowing for queries in $O(\frac{\log n}{\log \log n})$ time.*

3828 **Contribution 18.** *In the real-RAM model and the constant-degree algebraic
 3829 decision tree model, given n real-coordinate input points in \mathbb{R}^2 we can compute
 3830 the encoding of their order type as in Theorems 9 and 10 in $O(n^2)$ time.*

3831 **Idea** A natural idea is to pick r to be small but superconstant to reduce
 3832 the number of levels of the hierarchy and thus the query time. As already
 3833 pointed out, this has the drawback of increasing the time complexity of the
 3834 intersection location primitive from constant to $\Theta(r^2)$. Since this primitive
 3835 is used at each level of the hierarchy, the $\Theta(\log r)$ factor saved by having
 3836 less levels is lost.

3837 To get past this difficulty, the trick is to encode approximations of the
 3838 traces $\text{TR}(\mathcal{C}, p')$ to still allow constant intersection location. We call those
 3839 approximations *signatures* and denote them by $\text{SIG}(\mathcal{C}, p')$. We define those
 3840 signatures so that they approximately encode the cyclic permutation of the
 3841 intersections around each subcell. By carefully choosing some parameters,
 3842 we are able to fit two of those signatures in a single word of memory. We
 3843 can then precompute the output to all possible inputs for the intersection
 3844 location primitive and put them in a small table.

3845 Because of this size reduction, distinct pseudolines could be mapped to
 3846 identical signatures. Those ambiguous situations can be deterministically
 3847 handled using an additional lookup table. Because those situations rarely
 3848 arise, this table also is small.

Once we have located the intersection $p'_a \cap p'_b$, we still need to deal with p'_c . We change the layout of a trace to locate the subcell containing $p'_a \cap p'_b$ with respect to p'_c in constant time. In case p'_c properly intersects the subcell, we need to recurse on the subcell. To do that, we need to map the pseudoline indices a , b , and c to the indices those pseudolines have in that subcell. This is done with one indirection: For each of the three pseudolines we identify the index of the subcell in the list of subcells intersected by that pseudoline. This is implemented as a rank operation on a list of $O(r^2)$ bits. Given that index, we can find the intersection indices of that pseudoline in the cyclic permutation of the subcell in constant time.

In what follows, we describe five new structures: the signatures, the intersection oracle, the disambiguation table, the augmented traces, and the subcell mapper. The first reduces the bitsize of the original traces so that the second can be implemented in constant time. The third handles bad cases that arise because of this size reduction. The fourth defines a new layout for the traces that includes the signature. The fifth allows to implement the parent-to-subcell pseudoline index mapping in constant time using this new layout.

Signatures Fix a small constant α and define $r = \Theta(\log^\alpha n)$.⁵ We encode a ℓ -levels hierarchical cutting of parameter r . Note that we can construct a hierarchical cutting with superconstant r by constructing a hierarchical cutting with some appropriate constant parameter r' , and then skip levels that we do not need. As in §C.1, we store a combinatorial representation of this hierarchical cutting. We make some tweaks to this representation.

We augment the traces of §C.1 with a *signature*. The trace $\text{TR}(\mathcal{C}, p')$ of a cell-pseudoline pair is composed of two parts: The incidence bits that tell us for each subcell of the cell whether the pseudoline lies above, lies below, intersects or contains it, and the cyclic permutation bits used to locate the intersection of two pseudolines inside the cell. The first part uses $\Theta(r^2)$ bits. The second part uses $\Theta(r \log \frac{N}{r^{i+1}})$ bits for a cell of level i .

To construct the signature $\text{SIG}(\mathcal{C}, p')$, we keep the $\Theta(r^2) = \Theta(\log^{2\alpha} n)$ incidence bits because they fit in sublogarithmic space for sufficiently small

⁵The exact bound for how small α must be depends on hidden constant factors in the Zone Theorem and in the definition of the word size. In particular, we must have $\alpha \leq \frac{1}{2}$ when $w = \Theta(\log n)$.

3881 α . The second part would use superlogarithmic space if handled as before.
3882 We thus replace the $\Theta(r \log \frac{n}{r^{i+1}})$ bits of the cyclic permutation by a well
3883 chosen approximation.

3884 Let $\beta = 2^{\Theta(\log^\alpha n)}$ and denote by $n_i = n/r^i$ an upper bound on the
3885 number of pseudolines intersecting a cell of level i . For each subcell of level
3886 i , partition its cyclic permutation into $\beta_i \leq \min\{\beta, n_{i+1}\}$ blocks of at most
3887 $\lceil n_{i+1}/\beta_i \rceil$ intersections. For each pseudoline intersecting a cell we store
3888 the block indices that that pseudoline touches instead of storing the cyclic
3889 permutation indices.⁶ Hence, the second part of each signature only uses
3890 $O(r \log \beta) = O(\log^{2\alpha})$ bits.

3891 **Intersection Oracle** We construct a lookup table to compute in constant
3892 time, for any given cell of any given level, the subcell in which $q_i \cap q_j$ lies.
3893 For that we need a general observation on the precomputation of functions
3894 on small universes.

3895 **Observation C.4.** *In the word-RAM model with word size $w \geq \log n$, for
3896 any word-to-word function $f : [2^w] \rightarrow [2^w]$, we can build a lookup table of
3897 total bitsize $2^g h$ for all 2^g inputs $x \in [2^g]$ of bitsize $g \leq w$, mapping to images
3898 $y \in [2^h]$ of bitsize $h \leq w$, in time $2^g T(g)$ where $T(g)$ is the complexity of
3899 computing $f(x)$, $x \in [2^g]$. The image of bitsize h of any input of bitsize g can
3900 then be retrieved in $O(1)$ time by a single lookup (since inputs and outputs
3901 fit in a single word). In particular, the preprocessing time $2^g T(g)$ and the
3902 space $2^g h$ are sublinear as long as $T(g) = g^{O(1)}$ and $g + \log h = o(\log n)$.*

3903 In other words, any polynomial time computable word-to-word function
3904 can be precomputed in sublinear time and space for all inputs and outputs
3905 of sublogarithmic size.

3906 Since each pseudoline signature fits in $O(r^2 + r \log \beta) = O(\log^{2\alpha} n)$
3907 bits, and since the number of subcells of each cell is $O(\log^{2\alpha} n)$, we can
3908 choose an appropriate α so as to satisfy the requirements given above:
3909 take $\alpha < \frac{1}{2}$ so that two pseudoline signatures have a combined bitsize of
3910 $g = o(\log n)$. The output size is the bitsize of a subcell identifier, which is
3911 $h \leq 2\alpha \log \log n = o(\log n)$. In some cases, the block indices stored in the

⁶For β a power of two, this can be implemented by truncating the original cyclic permutation indices.

3912 signatures will not contain enough information to point to a unique output.
 3913 In those cases we store a special value that indicates ambiguity of the input.

3914 We can thus precompute the function that sends two pseudoline signatures
 3915 to either the subcell containing their intersection or to some special value in
 3916 case of an ambiguous input. Since we compute the function for all members
 3917 of its universe, we can implement the lookup table using direct addressing
 3918 into an array.

3919 Note that the output of this oracle is the same no matter what level or
 3920 cell we consider: for non-ambiguous inputs, all the information required to
 3921 locate $q_i \cap q_j$ is included in the input. We thus only need a single lookup
 3922 table, and the space needed is proportional to

$$3923 2^g h = 2^{\Theta(\log^{2\alpha} n)}.$$

3924 **Disambiguation** An input for the intersection oracle is ambiguous if and
 3925 only if at least one boundary intersection of each input pseudoline appears in
 3926 the same cyclic permutation block of the cell that contains their intersection.
 3927 Thus, ambiguous inputs rarely occur: less than the number of blocks times
 3928 the number of pairs of boundary intersections in a block, that is, less than
 3929 $\beta \cdot (n_{i+1}/\beta)^2 = \frac{n^2}{r^{2(i+1)}}/\beta$ times per subcell of level i of the hierarchy. When
 3930 $\beta \geq n_{i+1}$ all ambiguity is lifted so we can ignore those cases.

3931 The disambiguation table is a hash table storing the answer to all
 3932 ambiguous inputs for all cells of each level $i \in \{0, 1, \dots, \ell - 1\}$. This table
 3933 maps a triple of a cell $\mathcal{C}_{0,j_1,j_2,\dots,j_i}$, a pseudoline p'_a , and a pseudoline p'_b to
 3934 the index $j_{i+1} \in \{0, 1, \dots, O(r^2)\}$ of the subcell $\mathcal{C}_{0,j_1,j_2,\dots,j_i,j_{i+1}}$ containing
 3935 the intersection $p'_a \cap p'_b$. Summing over all subcells of each level, we obtain
 3936 that the number of entries in this table is bounded above by

$$3937 \sum_{i=0}^{\ell-1} r^{2i} \cdot r^2 \cdot \frac{n^2}{r^{2(i+1)}}/\beta = \frac{n^2 \ell}{2^{\Theta(\log^\alpha n)}}.$$

3938 The number of bits of each entry is at most $\ell \lceil \log cr^2 \rceil + 2 \log n = O(\log n)$
 3939 for the key and $\lceil \log cr^2 \rceil$ for the value. Both get absorbed by the $2^{-\Theta(\log^\alpha n)}$
 3940 factor in the number of entries, so we can keep the same expression for the
 3941 number of bits used by the disambiguation table.

3942 Since the keys and values fit in a constant number of words, we can
 3943 guarantee worst case constant query time using cuckoo hashing [124] or

Incidence bits	Block indices	Intersection indices
00 10 10 01 11 10 00 01	01 11 10 11 00 01	010 101 000
Signature		

Figure C.4. An augmented trace $\text{TR}'(\mathcal{C}, p')$ for the same cell-pseudoline pair as in Figure C.2.

3944 perfect hashing [78]. In both cases the construction of the table is randomized
 3945 and takes expected linear time in the number of entries. Perfect hashing
 3946 has the advantage that we can drop the entry keys since we only query this
 3947 table for existing entries. Note that this is the only part of the construction
 3948 that is randomized.

3949 **Augmented Traces** The augmented traces are simply the concatenation
 3950 of the signature and the first intersection indices as depicted in Figure C.4.
 3951 This layout allows for constant time access to the signature. Given a subcell
 3952 index we can test its incidence bits in constant time. Given an intersected
 3953 subcell index we can access the first intersection index of the pseudoline in
 3954 that subcell in constant time.

3955 As discussed earlier, the first part of the signature uses $O(r^2)$ bits.
 3956 We already noted that the second part of the signature uses $O(r \log \beta) =$
 3957 $O(\log^{2\alpha} n)$ bits. However, a better bound to use for the analysis of the total
 3958 space is $O(r \log \beta_i) = O(r \log n_{i+1})$ bits, which is proportional to the number
 3959 of bits needed for the first intersection indices.

3960 Summing over all pseudolines and all intersected cells of each level, the
 3961 space used for the augmented traces is proportional to

$$3962 n \sum_{i=0}^{\ell-1} r^i \cdot (r^2 + r \log n_{i+1}) = O\left(\frac{n^2}{t} (\log t + r)\right),$$

3963 as in Lemma C.2. Since $r = \Theta(\log^\alpha n)$ the bound becomes

$$3964 O\left(\frac{n^2}{t} (\log t + \log^\alpha n)\right).$$

3965 **Subcell Mapper** As before, let c_h be the constant hidden by the $O(r^2)$
 3966 of the hierarchical cutting, and let $c_z = 9.5$ be the constant in Theorem 6.2.

3967 We need a way to map a subcell index $0 \leq j_i \leq c_h r^2 - 1$ to the index
 3968 $0 \leq j'_i \leq \lfloor c_z r \rfloor - 2$ this subcell has in the list of subcells intersected by a
 3969 given pseudoline. If we can achieve this in constant time, then we can also
 3970 access the first intersection index this pseudoline has in this subcell.

3971 It is not hard to see that this mapping operation boils down to computing
 3972 $\text{rank}_{10}(j_i)$ where the data array is composed of the incidence bits of the given
 3973 pseudoline. Hence this can be solved by adding an auxiliary rank-select data
 3974 structure to each trace [24, 131]. Another solution is to reuse Observation C.4
 3975 to construct a lookup table for all $2^{\Theta(r^2)}$ possible incidence vectors.

3976 With the first solution, the space used by the traces stays the same up
 3977 to a constant factor. With the second solution, the space used is dominated
 3978 by the space taken by the intersection oracle. Hence, we do not bother
 3979 including it in the space analysis.

3980 All that is left to do is to properly solve the subproblems spawned by the
 3981 last level of the hierarchy. This is done exactly as in the previous section.

3982 **Leaves of the Hierarchy** As before, we have $O(\frac{n^2}{t^2})$ subproblems of size
 3983 t to encode. We follow the solution previously described to obtain: $O(\frac{n^2}{t^2})$
 3984 pointers of size $\lceil \log \nu(t) \rceil$, $O(\frac{n^2}{t^2})$ canonical labelings of size $\Theta(t \log t)$, and
 3985 $\nu(t)$ lookup tables of size $O(t^3)$. This is sufficient to take care of each of
 3986 those subproblems in constant time. The total space usage for those leaves
 3987 is unchanged and stays proportional to

$$3988 \quad \frac{n^2}{t^2} \cdot (t \log t + \log \nu(t)) + t^3 \nu(t).$$

3989 **Space Complexity** Summing all terms together and adding the space
 3990 taken by the parameters of the hierarchy n , r and t , we obtain:

3991 **Lemma C.5.** *The space taken by the encoding described in §C.2 is propor-*
 3992 *tional to*

$$3993 \quad \begin{aligned} & \underbrace{\log ntr}_{\text{Parameters}} + \underbrace{\frac{n^2}{t} (\log t + \log^\alpha n)}_{\text{Traces}} + \underbrace{2^{\Theta(\log^{2\alpha} n)}}_{\text{Intersection Oracle}} \\ & + \underbrace{\frac{n^2 \ell}{2^{\Theta(\log^\alpha n)}}}_{\text{Disambiguation Table}} + \underbrace{\frac{n^2}{t^2} (\log \nu(t) + t \log t) + t^3 \nu(t)}_{\text{Leaves}}. \end{aligned}$$

3997 As before, we take $t = \sqrt{\delta \log n}$ for abstract order types and $t =$
 3998 $\delta \log n / \log \log n$ for realizable ones. Taking δ to be sufficiently small, the
 3999 space taken by the leaves of the hierarchy is thus $\Theta(n^2)$ for abstract order
 4000 types and dominated by the term $\frac{n^2}{t} \log t$ in the case of realizable order types.
 4001 Setting $\alpha < \frac{1}{2}$ guarantees that the space taken by the intersection oracle is
 4002 subpolynomial, and that the space taken by the traces is subquadratic. The
 4003 space taken by the disambiguation table is in $O(\frac{n^2}{\log^c n})$ for all c and is thus
 4004 dominated by the other terms.

4005 For abstract order types, all those terms are subquadratic, except for
 4006 the pointers at the leaves. The total space usage for abstract order types is
 4007 thus dominated by this term and is quadratic. For realizable order types,
 4008 the total space is dominated by the term $\frac{n^2}{t} \log^\alpha n$. Indeed, we can take α
 4009 as small as desired to make the factor $\log^\alpha n = O(\log^\varepsilon n)$. This proves the
 4010 space constraints in Theorems 9 and 10. Unfortunately, the present solution
 4011 incurs a nonabsorbable extra $\log^\varepsilon n$ factor in the realizable case. Note that a
 4012 $\log \log \log n$ factor can be squeezed from the query time without increasing
 4013 the space usage by choosing $r = \Theta(\log \log n)$ instead.

4014 **Correctness and Query Complexity** Given a query p'_a, p'_b, p'_c and a
 4015 cell \mathcal{C} , the subcell \mathcal{C}' containing $p'_a \cap p'_b$ is found in constant time via the
 4016 intersection oracle and, if necessary, the disambiguation table. The location
 4017 of that subcell with respect to p'_c can then be retrieved by a single lookup
 4018 in the incidence bits of $\text{TR}'(\mathcal{C}, p'_c)$. In case of recursion, we can compute
 4019 the address of the traces $\text{TR}'(\mathcal{C}', p'_a)$, $\text{TR}'(\mathcal{C}', p'_b)$, and $\text{TR}'(\mathcal{C}', p'_c)$ in constant
 4020 time using the subcell mapper. The base case is handled in constant time as
 4021 before: using the pointers, canonical labelings and order type lookup tables.

4022 We now have a shallower decision tree of depth $\log_r \frac{n}{t} = O_\alpha(\frac{\log n}{\log \log n})$
 4023 and the work at each level takes constant time. This proves the query time
 4024 constraints in Theorems 9 and 10.

4025 **Preprocessing Time** We prove Theorem 18.

4026 *Proof.* As before, the hierarchical cutting and all traces $\text{TR}'(C, p')$ can be
 4027 computed in $O(nr^\ell)$ time (with or without rank-select data structures).
 4028 The lookup tables and leaf-table pointers can be computed in $O(n^2)$ time.

4029 The intersection oracle, the disambiguation table, and the optional subcell
 4030 mapper can be computed in subquadratic time. \square

4031 C.3 Higher-Dimensional Encodings

4032 We generalize our point configuration encoding to any dimension d . The
 4033 chirotope of a point set in \mathbb{R}^d consists of all orientations of simplices defined
 4034 by $d + 1$ points of the set [139]. The orientation of the simplex with $d + 1$
 4035 ordered vertices p_i with coordinates $(p_{i,1}, p_{i,2}, \dots, p_{i,d})$ is given by the sign
 4036 of the determinant

$$4037 \quad \begin{vmatrix} 1 & p_{1,1} & p_{1,2} & \dots & p_{1,d} \\ 1 & p_{2,1} & p_{2,2} & \dots & p_{2,d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & p_{d+1,1} & p_{d+1,2} & \dots & p_{d+1,d} \end{vmatrix}.$$

4038 We obtain the following generalized result:

4039 **Contribution 12.** *All realizable chirotopes of rank $k \geq 4$ have an encoding*
 4040 *using $O(\frac{n^{k-1}(\log \log n)^2}{\log n})$ bits of space and allowing for queries in $O(\frac{\log n}{\log \log n})$*
 4041 *time.*

4042 **Contribution 13.** *In the real-RAM model and the constant-degree algebraic*
 4043 *decision tree model, given n real-coordinate input points in \mathbb{R}^d we can compute*
 4044 *the encoding of their chirotope as in Theorem 12 in $O(n^d)$ time.*

4045 **Idea** In the primal, the orientation of a simplex whose vertices are ordered
 4046 can be interpreted as the location of its last vertex with respect to the
 4047 (oriented) hyperplane spanned by its first d vertices. In the dual, this
 4048 orientation corresponds to the location of the intersection of the first d
 4049 dual hyperplanes with respect to the last (oriented) dual hyperplane. In the
 4050 primal, degenerate simplices have orientation 0. In the dual, this corresponds
 4051 to linearly dependent subsets of $d + 1$ hyperplanes.

4052 We gave an encoding for the two-dimensional case in §C.1. With this
 4053 encoding, a query is answered by traversing the levels of some hierarchical
 4054 cutting, branching on the location of the intersection of two of the three
 4055 query lines. We generalize this idea to d dimensions. Now the cell considered
 4056 at the next level of the hierarchy depends on the location of the intersection

4057 of d of the $d + 1$ query hyperplanes. We will also have to take care of
 4058 degenerate cases.

4059 **Intersection Location** In §C.1, we solved the following two-dimensional
 4060 subproblem:

4061 **Problem 23.** Given a triangle and n lines in the plane, build a data struc-
 4062 ture that, given two of those lines, allows to decide whether their intersec-
 4063 tion lies in the interior of the triangle.

4064 In retrospective, we showed that there exists such a data structure using
 4065 $O(n \log n)$ bits that allows for queries in $O(1)$ time. We generalize this result.
 4066 Consider the following generalization of the problem in d dimensions:

4067 **Problem 24.** Given a convex body and n hyperplanes in \mathbb{R}^d , build a data
 4068 structure that, given d of those hyperplanes, allows to decide whether their
 4069 intersection is a vertex that lies in the interior of the convex body.

4070 Of course this problem can be solved using $O(n^d)$ space by explicitly
 4071 storing the answers to all possible queries. If the input hyperplanes are given
 4072 in an arbitrary order, this is best possible for $d = 1$. For $d \geq 2$, we show how
 4073 to reduce the space to $O(n^{d-1} \log n)$ by recursing on the dimension, taking
 4074 $d = 2$ as the base case.

4075 We encode the function $\mathcal{I}_{C,H}$ that maps a d -tuple of indices of input dual
 4076 hyperplanes H_i to 1 if their intersection is a vertex that lies in the interior
 4077 of a fixed convex body C , and to 0 otherwise.

$$4078 \quad \mathcal{I}_{C,H}: [n]^d \rightarrow \{0, 1\}: (i_1, i_2, \dots, i_d) \mapsto (H_{i_1} \cap H_{i_2} \cap \dots \cap H_{i_d}) \in C.$$

4079 We call this function the *intersection function* of (C, H) . We prove the
 4080 following:

4081 **Lemma C.6.** All intersection functions have a $O(n^{d-1} \log n)$ -bits $O(d)$ -
 4082 querytime encoding.

4083 *Proof.* Consider a convex body C and n hyperplanes H_i . At first, for sim-
 4084 plicity, assume C is d -dimensional and assume that any d hyperplanes H_{i_j}
 4085 meet in a single point. With those assumptions, we want a data structure
 4086 that can answer any query of the type

$$4087 \quad (H_{i_1} \cap H_{i_2} \cap H_{i_3} \cap H_{i_4} \cap \dots \cap H_{i_{d-1}} \cap H_{i_d}) \cap C \neq \emptyset.$$

4088 Note that this is equivalent to deciding whether

$$4089 (H_{i_1} \cap H_{i_2} \cap H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_{d-1}}) \cap (H_{i_d} \cap C) \neq \emptyset,$$

4090 where $H_{i_d} \cap C$ is a convex body of dimension $d - 1$ (or empty), and the
4091 number of hyperplanes we want to intersect it with is $d - 1$.

4092 We unroll the recursion until the convex body is of dimension two (or
4093 empty), and only two hyperplanes are left to intersect. We then notice that
4094 the decision we are left with is equivalent to

4095

$$4096 (H_{i_1} \cap (H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d})) \cap (H_{i_2} \cap (H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d})) \\ 4097 \cap (C \cap (H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d})) \neq \emptyset,$$

4099 which, if the three objects are non-empty, reads: “Given two lines and a
4100 convex body in some plane, do they intersect?”. We can answer this query
4101 if we have the encoding for $d = 2$ which is obtained by replacing *triangle* by
4102 *convex body* in the two-dimensional original problem. The total space taken
4103 is multiplied by n for each time we unroll the recursion times the space taken
4104 in two dimensions, which is proportional to $n^{d-2} \cdot n \log n = O(n^{d-1} \log n)$.
4105 Queries can then be answered in $O(d)$ time.

4106 Note that degenerate cases are likely to arise: empty convex bodies be-
4107 cause of nonintersecting hyperplanes, convex bodies that are higher dimen-
4108 sional because of linearly dependent hyperplanes, and convex bodies that
4109 are lower dimensional because C was not full-dimensional to start with.
4110 However, all those cases can be dealt with appropriately: If the query suffix
4111 $H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d}$ leads to an empty convex body $C \cap (H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d})$
4112 then the query point is not in C and we encode 0 for all the queries in C
4113 ending in this suffix. This information can be encoded in a table of size
4114 $O(n^{d-2})$. If the query suffix $H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d}$ leads to a convex body
4115 $C \cap (H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d})$ of dimension ≥ 3 then the intersection of all objects
4116 is not a 0-flat and we encode a 0 to follow the definition of $\mathcal{I}_{C,H}$. Again,
4117 this information can be encoded in a table of size $O(n^{d-2})$. If the query
4118 suffix $H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d}$ leads to a convex body $C \cap (H_{i_3} \cap H_{i_4} \cap \cdots \cap H_{i_d})$
4119 of dimension zero, one, or two, then we use the encoding of size $O(n \log n)$
4120 described in §C.1 as a base case. \square

4121 In the paragraphs that follow, we show how to plug this result in those

of the previous sections to obtain analogous results for the d -dimensional version of the problem (Theorem 12 and Theorem 13).

Encoding We build a hierarchical cutting as in §C.1 (this time in dimension d). Given n hyperplanes in \mathbb{R}^d and some fixed parameters r and ℓ , compute a ℓ -levels hierarchical cutting of parameter r for those hyperplanes as in Lemma 8.5. This hierarchical cutting consists of ℓ levels labeled $0, 1, \dots, \ell - 1$. Level i has $O(r^{di})$ cells. Each of those cells is further partitioned into $O(r^d)$ subcells. The $O(r^{d(i+1)})$ subcells of level i are the cells of level $i + 1$. Each cell of level i is intersected by at most $\frac{n}{r^i}$ hyperplanes, and hence each subcell is intersected by at most $\frac{n}{r^{i+1}}$ hyperplanes.

We compute and store a combinatorial representation of the hierarchical cutting as follows: For each level of the hierarchy, for each cell in that level, for each hyperplane intersecting that cell, for each subcell of that cell, we store two bits to indicate the location of the hyperplane with respect to that subcell, that is, whether the hyperplane lies above (00), lies below (01), intersects the interior of that subcell (10), or contains the subcell (11). When a hyperplane H_i intersects the interior of a subcell \mathcal{C}' , we also store the two $O(\log \frac{n}{r^{i+1}})$ -bits two-dimensional intersections indices this hyperplane has in each of the $O\left(\left(\frac{n}{r^{i+1}}\right)^{d-2}\right)$ query suffixes $(H_{i_3} \cap H_{i_4} \cap \dots \cap H_{i_d})$ with each of the H_{i_j} properly intersecting \mathcal{C}' .

This representation takes $O\left(\frac{n}{r^i} + \left(\frac{n}{r^{i+1}}\right)^{d-1} \log \frac{n}{r^{i+1}}\right)$ bits per subcell of level i by storing for each hyperplane its location and, when needed, the bits they hold in the encoding of the intersection function of the subcell. At the last level of the hierarchy, let $t = \frac{n}{r^\ell}$ denote an upper bound on the number of hyperplanes intersecting each subcell. For each of those $O(r^{d\ell}) = O(\frac{n^d}{t^\ell})$ subcells we store a pointer to a lookup table of size $O(t^{d+1})$ that allows to answer the query of the orientation of any triple of hyperplanes intersecting that subcell.

Using the Zone Theorem in higher dimensions (Theorem 6.3), we can have all bits belonging to a single cell-hyperplane pair in a contiguous block of memory with the same space bound.

Space Complexity As before, for the space taken by the lookup tables, their associated pointers and canonical labelings at the leaves, and the

parameters of the hierarchy n , r and t , the analysis is immediate. If not implicitly known, the dimension d can also trivially be added to the encoding.

For the space taken by the hierarchy, we generalize Lemma C.2 of §C.1. Let $H_r^\ell(n, d) \in \mathbb{N}$ be the maximum amount of space (bits), over all arrangements of n hyperplanes in \mathbb{R}^d , taken by the $\ell \in \mathbb{N}$ levels of a hierarchy with parameter $r \in (1, +\infty)$.

Lemma C.7. *For $r \geq 2$ we have*

$$H_r^\ell(n, d) = O\left(\frac{n^d}{t} \left(\log t + \frac{r}{t^{d-2}}\right)\right).$$

Proof. By definition we have

$$H_r^\ell(n, d) = O\left(\sum_{i=0}^{\ell-1} \left(r^{di} \cdot r^d \cdot \left(\frac{n}{r^i} + \left(\frac{n}{r^{i+1}}\right)^{d-1} \cdot \log \frac{n}{r^{i+1}}\right)\right)\right).$$

Using $t = \frac{n}{r^\ell}$, reversing the summation with $i \leftarrow \ell - i - 1$, and grouping the terms, we have

$$H_r^\ell(n, d) = O\left(\frac{n^d}{t} \left(\frac{r}{t^{d-2}} \sum_{i=0}^{\ell-1} \frac{1}{(r^{d-1})^i} + \log t \sum_{i=0}^{\ell-1} \frac{1}{r^i} + \log r \sum_{i=0}^{\ell-1} \frac{i}{r^i}\right)\right).$$

Using the geometric inequalities (see the proof of Lemma C.2) the statement follows from $r \geq 2$. \square

The final picture is almost the same as in Figure C.3. Summing all terms, we obtain

Lemma C.8. *The space taken by the encoding described in §C.3 is proportional to*

$$\underbrace{\log dntr}_{\text{Parameters}} + \underbrace{\frac{n^d}{t} \left(\log t + \frac{r}{t^{d-2}}\right)}_{\text{Traces}} + \underbrace{\frac{n^d}{t^d} (\log \nu_d(t) + t \log t)}_{\text{Leaves}} + \underbrace{\frac{t^{d+1}}{n} \nu_d(t)}_{\text{Lookup Tables}},$$

where $\nu_d(t) = 2^{\Theta(d^2 t \log t)}$ denotes the number of realizable rank- $(d+1)$ chirotopes of size t .

We pick r constant and choose $t = \delta \log n / \log \log n$ for small enough δ . The second term in Lemma C.8 dominates with $n^d (\log \log n)^2 / \log n$. This proves the space constraint in Theorem 12.

4180 **Correctness and Query Complexity** As before, a query is answered by
4181 traversing the hierarchy, which takes $O(\log n)$ time. The query time can be
4182 further improved using the method from §C.2 with $r = \Theta(t^{d-2} \log t)$. This
4183 proves the query time constraint in Theorem 12.

4184 **Preprocessing Time** We prove Theorem 13.

4185 *Proof.* The hierarchical cuttings can be computed in $O(n(r^\ell)^{d-1})$ time. The
4186 lookup table and leaf-table pointers can be computed in $O(n^d)$ time using
4187 the canonical labeling and representation for rank- $(d+1)$ chirotopes given
4188 in [17]. The intersection oracle, the disambiguation table, and the subcell
4189 mapper can be computed in $o(n^d)$ time. \square

D

4190

Encoding 3SUM

4191

4192 with Sergio Cabello, Jean Cardinal, John Iacono, Stefan Langerman, and Pat Morin

4193 Given three sets of n real numbers $A = \{a_1 < a_2 < \dots < a_n\}$, $B =$
4194 $\{b_1 < b_2 < \dots < b_n\}$, and $C = \{c_1 < c_2 < \dots < c_n\}$, we wish to build a
4195 discrete data structure (using bits, words, and pointers) such that, given any
4196 triple $(i, j, k) \in [n]^3$ it is possible to compute the sign of $a_i + b_j + c_k$ by only
4197 inspecting the data structure (we cannot consult A , B , or C). We refer to
4198 the map $\chi : [n]^3 \rightarrow \{-, 0, +\}, (i, j, k) \mapsto \text{sgn}(a_i + b_i + c_k)$ as the *3SUM type*
4199 of the instance $\langle A, B, C \rangle$.

4200 Obviously, one can simply construct a lookup table of size $O(n^3)$, such
4201 that triple queries can be answered in $O(1)$ time. In §D.1 we show that a
4202 minimal integer representation of a 3SUM instance may require $\Theta(n)$ bits
4203 per value, yielding $O(n)$ query time and $O(n^2)$ space. In §D.2 we show how
4204 to use an optimal $O(n \log n)$ bits of space with a polynomial query time.
4205 Finally, in §D.3 we show how to use $\tilde{O}(n^{3/2})$ space to achieve $O(1)$ -time
4206 queries.

4207 D.1 Representation by Numbers

4208 A first natural idea is to encode the real 3SUM instance by *rounding*
4209 its numbers to integers. We show a tight bound of $\Theta(n^2)$ bits for this
4210 representation.

4211 **Contribution 14.** *Every 3SUM instance has an equivalent integer instance*
4212 *where all values have absolute value at most $2^{O(n)}$. Furthermore, there exists*
4213 *an instance of 3SUM where all equivalent integer instances require numbers*

4214 at least as large as the n th Fibonacci number and where the standard binary
4215 representation of the instance requires $\Omega(n^2)$ bits.

4216 *Proof.* Every 3SUM instance $A = \{a_1 < a_2 < \dots < a_n\}$, $B = \{b_1 <$
4217 $b_2 < \dots < b_n\}$, and $C = \{c_1 < c_2 < \dots < c_n\}$ can be interpreted as
4218 the point $(a_1, \dots, a_n, b_1, \dots, b_n, c_1, \dots, c_n)$ in \mathbb{R}^{3n} . Let us use the variables
4219 x_1, \dots, x_n to encode the first n dimensions of \mathbb{R}^{3n} , y_1, \dots, y_n to encode the
4220 next n dimensions, and z_1, \dots, z_n for the remaining dimensions. Consider
4221 the subset of \mathbb{R}^{3n}

4222

$$\Delta = \{(x_1, \dots, x_n, y_1, \dots, y_n, z_1, \dots, z_n) : \quad$$

$$\quad x_i < x_{i+1}, y_j < y_{j+1}, z_k < z_{k+1} \forall i, j, k \in [n - 1]\}$$

4226 and the set Π of n^3 hyperplanes $x_i + y_j + z_k = 0$, where $i, j, k \in [n]$. Let \mathcal{A}
4227 be the arrangement defined by Π *inside* Δ . Instances of 3SUM correspond
4228 to points in Δ . Moreover, two 3SUM instances have the same 3SUM type
4229 if and only if they are in the same cell of \mathcal{A} .

4230 Consider an instance $\langle A, B, C \rangle$ and let $\sigma = \sigma(A, B, C)$ be the cell of \mathcal{A}
4231 that contains it. Then σ is the cell defined by the inequalities

$$\begin{aligned} \forall i, j, k \in [n] : & \begin{cases} x_i + y_j + z_k > 0 & \text{if } \chi(i, j, k) = +1, \\ x_i + y_j + z_k = 0 & \text{if } \chi(i, j, k) = 0, \\ x_i + y_j + z_k < 0 & \text{if } \chi(i, j, k) = -1. \end{cases} \\ \forall i, j, k \in [n - 1] : & \begin{cases} x_i - x_{i+1} < 0, \\ y_j - y_{j+1} < 0, \\ z_k - z_{k+1} < 0. \end{cases} \end{aligned}$$

4235 Let σ' be the subset of \mathbb{R}^{3n} defined by the following inequalities:

$$\begin{aligned} \forall i, j, k \in [n] : & \begin{cases} x_i + y_j + z_k \geq 1 & \text{if } \chi(i, j, k) = +1, \\ x_i + y_j + z_k = 0 & \text{if } \chi(i, j, k) = 0, \\ x_i + y_j + z_k \leq -1 & \text{if } \chi(i, j, k) = -1. \end{cases} \\ \forall i, j, k \in [n - 1] : & \begin{cases} x_i - x_{i+1} \leq 1, \\ y_j - y_{j+1} \leq 1, \\ z_k - z_{k+1} \leq 1. \end{cases} \end{aligned}$$

4238

4239 Clearly σ' is contained in σ . Moreover, for a sufficiently large $\lambda > 0$ the
 4240 scaled instance $\langle \lambda A, \lambda B, \lambda C \rangle$ belongs to σ' . Therefore, σ' is nonempty.

4241 Since σ' is defined by a collection of linear inequalities defining closed
 4242 halfspaces, there exists a point p in σ' defined by a subset of at most $3n$
 4243 inequalities, where the inequalities are actually equalities. Let us assume for
 4244 simplicity that exactly $3n$ equalities define the point p . Then, $p = (x, y, z)$
 4245 is the solution to a linear system of equations $M[x \ y \ z]^T = \delta$ where M
 4246 and δ have their entries in $\{-1, 0, 1\}$ and each row of M has at most three
 4247 non-zero entries. The solution p to this system of equations is an instance
 4248 equivalent to $\langle \lambda A, \lambda B, \lambda C \rangle$.

4249 Because of Cramer's rule, the system of linear equations has solution
 4250 with entries $\det(M_i) / \det(M)$, where M_i is the matrix obtained by replacing
 4251 the i th column of M by δ . We use the following simple bound on the de-
 4252 terminant. Since $\det(M) = \sum_{\pi} \text{sgn}(\pi) \prod_i m_{i,\pi(i)}$, where π iterates over the
 4253 permutations of $[3n]$, there are at most 3^{3n} summands where π gives non-
 4254 zero product $\prod_i m_{i,\pi(i)}$ (we have to select one non-zero entry per row), and
 4255 the product is always in $\{-1, 0, 1\}$. Therefore $|\det(M)| \leq 3^{3n}$. Similarly,
 4256 $|\det(M_i)| \leq 4^{3n}$ because each row of M_i has at most 4 non-zero entries. We
 4257 conclude that the solution to the system $M[x \ y \ z]^T = \delta$ are rationals that
 4258 can be expressed with $O(n)$ bits. This solution gives a 3SUM instance with
 4259 rationals that is equivalent to $\langle A, B, C \rangle$. Since all the rationals have the
 4260 common denominator ($\det(M)$), we can scale the result by $\det(M)$ and we
 4261 get an equivalent instance with integers, where each integer has $O(n)$ bits.

4262 The proof of the second statement is by implementing the Fibonacci
 4263 recurrence in each of the arrays A, B, C . This can be achieved by letting:

$$\begin{aligned} 4264 \quad a_i + b_1 + c_{n-i+1} &= 0, \text{ for } i \in [n] \\ 4265 \quad a_1 + b_i + c_{n-i+1} &= 0, \text{ for } i \in [n] \\ 4266 \quad a_{i-1} + b_{i-2} + c_{n-i+1} &< 0, \text{ for } i \in \{3, 4, \dots, n\}, \end{aligned}$$

4267 The first two sets of equations ensure that the two arrays A and B are
 4268 identical, while the array C contains the corresponding negated numbers,
 4269 in reverse order. From the inequalities in the third group, and depending
 4270 on the choice of the initial values a_1, a_2 , each array contains a sequence
 4271 growing at least as fast as the Fibonacci sequence. \square

4272 Note that this is a much smaller lower bound than for order types of

4273 points sets in the plane, the explicit representation of which can be shown
4274 to require exponentially many bits per coordinate [90].

4275 D.2 Space-Optimal Representation

4276 By considering the arrangement of hyperplanes defining the 3SUM prob-
4277 lem, we get an information-theoretic lower bound on the number of bits in a
4278 3SUM-type.

4279 **Lemma D.1.** *There are $2^{\Theta(n \log n)}$ distinct 3SUM-types of size n .*

4280 *Proof.* 3SUM-types of size n are in one-to-one correspondence with cells
4281 of the arrangement of n^3 hyperplanes in \mathbb{R}^{3n} . The number of such cells is
4282 $O(n^{9n})$ and at least $(n!)^2$. \square

4283 In order to reach this lower bound, we can simply encode the label of the
4284 cell of the arrangement in $\Theta(n \log n)$ bits. However, decoding the information
4285 requires to construct the whole arrangement which takes $n^{O(n)}$ time. An
4286 alternative solution is to store a vertex of the arrangement of hyperplanes
4287 $a_i + b_j + c_k \in \{-1, 0, 1\}$. There exists such a vertex that has the same
4288 3SUM-type as the input point, as shown in the proof of Contribution 14. To
4289 answer any query, either recompute the vertex from the basis then answer
4290 the query using arithmetic, or use linear programming. Hence we can build
4291 a data structure of $O(n \log n)$ bits such that triple queries can be answered
4292 in polynomial time.

4293 Note that we do not exploit much of the 3SUM structure here. In
4294 particular, the same essentially holds for k -SUM, and can also be generalized
4295 to a SUBSET SUM data structure of $O(n^2)$ bits, from which we can extract
4296 the sign of the sum of any subset of numbers.

4297 D.3 Subquadratic Space and Constant Query 4298 Time

4299 Our encoding is inspired by Grønlund and Pettie's $\tilde{O}(n^{3/2})$ non-uniform
4300 algorithm for 3SUM [91]. Our data structure stores three components, which
4301 we call the *differences*, the *staircase* and the *square neighbors*.

4302 **Differences.** Partition A and B into *blocks* of \sqrt{n} consecutive elements.
4303 Let D be the set of all differences of the form $a_{i_1} - a_{i_2}$ and $b_{j_1} - b_{j_2}$ where the

4304 items come from the same block. There are $O(n^{3/2})$ such differences. Sort
 4305 D and store a table indicating for each difference in D its rank among all
 4306 differences in D . This takes $O(\log n)$ bits for each of the $O(n^{3/2})$ differences,
 4307 for a total of $O(n^{3/2} \log n)$ bits.

4308 **Staircase.** Look at the table G formed by all sums of the form $a_i + b_j$. It is
 4309 monotonic in its rows and columns due to A and B being sorted. We view it
 4310 as being partitioned into a grid G' of size $\sqrt{n} \times \sqrt{n}$ where each *square* of the
 4311 grid is also of size $\sqrt{n} \times \sqrt{n}$. For each element $c_k \in C$, for each $i' \in [1, \sqrt{n}]$
 4312 we store the largest j' such that some elements of the square $G'[i', j']$ are $< c$,
 4313 denote this as $L[k, i']$. We also store, for each $c_k \in C$, for each $i' \in [1, \sqrt{n}]$
 4314 the smallest j' such that some elements of the square $G'[i', j']$ are $\geq c$, denote
 4315 this as $U[k, i']$. We thus store, in L and U , $2\sqrt{n}$ values of size $O(\log n)$ for
 4316 each of the n elements of C , for a total space usage of $O(n^{3/2} \log n)$ bits. We
 4317 call this the *staircase* as this implicitly classifies, for each $c \in C$, whether
 4318 each square has elements larger than c , smaller than c , or some larger and
 4319 some smaller; only $O(\sqrt{n})$ can be in the last case, which we refer to as the
 4320 *staircase* of c .

4321 **Square neighbors.** For each element $c_k \in C$, for each of the $O(\sqrt{n})$
 4322 squares on the staircase, we store the location of the predecessor and successor
 4323 of c_k . Those squares are the $G'[i', j']$ such that $L[k, i'] \leq j' \leq U[k, i']$, for
 4324 $i', j' \in [1, \sqrt{n}]$. This takes space $O(n^{3/2} \log n)$.

4325 To execute a query (a_i, b_j, c_k) , only a constant number of lookups in the
 4326 tables stored are needed. Let us define $i' = \lceil i/\sqrt{n} \rceil$ and $j' = \lceil j/\sqrt{n} \rceil$ to
 4327 be the indices of the cell of G' containing the point (a_i, b_j) . If $j' < L[k, i']$,
 4328 then we know $a_i + b_j < c_k$. If $j' > U[k, i']$, then we know $a_i + b_j > c_k$. If
 4329 neither of these is true, then the square $G'[i', j']$ is on the staircase of c_k and
 4330 thus using the square neighbors table we can determine the location of the
 4331 predecessor and successor of c_k in this square; suppose they are at $G[s_i, s_j]$
 4332 and $G[p_i, p_j]$ and thus $G[s_i, s_j] \leq c_k \leq G[p_i, p_j]$. One need only determine
 4333 how these two compare to $G[i, j] = a_i + b_j$ to answer the query. But this can
 4334 be done using the differences as follows: to compare $G[s_i, s_j]$ to $G[i, j]$ this
 4335 would be determining the sign of $(a_i + b_j) - (a_{s_i} + b_{s_j})$ which is equivalent to
 4336 determining the result of comparing $a_i - a_{s_i}$ and $b_j - b_{s_j}$, which since both
 4337 are in the same square, these differences are in D and the comparison can be

	1	2	10	14	17	22	32	33	40	91	92	97	98	110	120	127
1	2	3	11	15	18	23	33	34	41	92	93	98	99	111	121	128
11	12	13	21	25	28	33	43	44	51	102	103	108	109	121	131	138
13	14	15	23	27	30	35	45	46	53	104	105	110	111	123	133	140
19	20	21	29	33	36	41	51	52	59	110	111	116	117	129	139	146
24	25	26	34	38	41	46	56	57	64	115	116	121	122	134	144	151
34	35	36	44	48	51	56	66	67	74	125	126	131	132	144	154	161
51	52	53	61	65	68	73	83	84	91	142	143	148	149	161	171	178
57	58	59	67	71	74	79	89	90	97	148	149	154	155	167	177	184
59	60	61	69	73	76	81	91	92	99	150	151	156	157	169	179	186
114	115	116	124	128	131	136	146	147	154	205	206	211	212	224	234	241
119	120	121	129	133	136	141	151	152	159	210	211	216	217	229	239	246
127	128	129	137	141	144	149	159	160	167	218	219	224	225	237	247	254
128	129	130	138	142	145	150	160	161	168	219	220	225	226	238	248	255
133	134	135	143	147	150	155	165	166	173	224	225	230	231	243	253	260
138	139	140	148	152	155	160	170	171	178	229	230	235	236	248	258	265
142	143	144	152	156	159	164	174	175	182	233	234	239	240	252	262	269

Figure D.1. Illustration of the staircase and square neighbors of the constant query time encoding. Here the 16×16 table is partitioned into a 4×4 grid of squares of size 4×4 . If $c_k = 100$, the grey illustrates the squares that form the staircase, containing values both larger and smaller than 100. Predecessors and successors within each staircase square are shown in red and blue.

4338 obtained by examining their stored ranks. By doing this for the predecessor
 4339 and successor we will determine the relationship between $a_i + b_j$ and c_k .

4340

Bibliography

- [1] Amir Abboud, Kevin Lewi, and Ryan Williams. Losing weight by gaining edges. In *European Symposium on Algorithms (ESA 2014)*, pages 1–12. Springer, 2014.
- [2] Amir Abboud and Virginia Vassilevska Williams. Popular conjectures imply strong lower bounds for dynamic problems. In *FOCS*, pages 434–443. IEEE Computer Society, 2014.
- [3] Amir Abboud, Virginia Vassilevska Williams, and Oren Weimann. Consequences of faster alignment of sequences. In *ICALP (1)*, volume 8572 of *LNCS*, pages 39–51, 2014.
- [4] Amir Abboud, Virginia Vassilevska Williams, and Huacheng Yu. Matching triangles and basing hardness on an extremely popular conjecture. In *STOC*, pages 41–50. ACM, 2015.
- [5] Pankaj K. Agarwal and Jivr’i Matouvsek. On range searching with semialgebraic sets. *Discrete & Computational Geometry*, 11:393–418, 1994.
- [6] Alfred V. Aho, Kenneth Steiglitz, and Jeffrey D. Ullman. Evaluating polynomials at fixed sets of points. *SIAM Journal on Computing*, 4(4):533–539, 1975.
- [7] Oswin Aichholzer, Franz Aurenhammer, and Hannes Krasser. Enumerating order types for small point sets with applications. *Order*, 19(3):265–281, 2002.
- [8] Oswin Aichholzer, Franz Aurenhammer, and Hannes Krasser. On the crossing number of complete graphs. In *SoCG*, pages 19–24. ACM, 2002.
- [9] Oswin Aichholzer, Jean Cardinal, Vincent Kusters, Stefan Langerman, and Pavel Valtr. Reconstructing point set order types from

- 4367 radial orderings. *International Journal of Computational Geometry &*
4368 *Applications*, 26(3-4):167–184, 2016.
- 4369 [10] Oswin Aichholzer, Matias Korman, Alexander Pilz, and Birgit Vogten-
4370 huber. Geodesic order types. *Algorithmica*, 70(1):112–128, 2014.
- 4371 [11] Oswin Aichholzer and Hannes Krasser. The point set order type data
4372 base: A collection of applications and results. In *CCCG*, pages 17–20,
4373 2001.
- 4374 [12] Oswin Aichholzer and Hannes Krasser. Abstract order type extension
4375 and new results on the rectilinear crossing number. In *SoCG*, pages
4376 91–98. ACM, 2005.
- 4377 [13] Oswin Aichholzer, Vincent Kusters, Wolfgang Mulzer, Alexander Pilz,
4378 and Manuel Wettstein. An optimal algorithm for reconstructing point
4379 set order types from radial orderings. In *ISAAC*, pages 505–516.
4380 Springer, 2015.
- 4381 [14] Oswin Aichholzer, Tillmann Miltzow, and Alexander Pilz. Extreme
4382 point and halving edge search in abstract order types. *Computational
4383 Geometry*, 46(8):970–978, 2013.
- 4384 [15] Nir Ailon and Bernard Chazelle. Lower bounds for linear degeneracy
4385 testing. *J. ACM*, 52(2):157–171, 2005.
- 4386 [16] Noga Alon. The number of polytopes configurations and real matroids.
4387 *Mathematika*, 33(1):62–71, 1986.
- 4388 [17] Greg Aloupis, John Iacono, Stefan Langerman, Özgür Özkan, and
4389 Stefanie Wuhrer. The complexity of order type isomorphism. In
4390 *SODA*, pages 405–415. SIAM, 2014.
- 4391 [18] Amihood Amir, Timothy M. Chan, Moshe Lewenstein, and Noa Lewen-
4392 stein. On hardness of jumbled indexing. In *ICALP (1)*, volume 8572
4393 of *LNCS*, pages 114–125, 2014.
- 4394 [19] Ilya Baran, Erik D. Demaine, and Mihai Pătrașcu. Subquadratic
4395 algorithms for 3SUM. *Algorithmica*, 50(4):584–596, 2008.

- 4396 [20] Luis Barba, Jean Cardinal, John Iacono, Stefan Langerman, Aurélien
4397 Ooms, and Noam Solomon. Subquadratic algorithms for algebraic
4398 3SUM. *Discrete & Computational Geometry*, 61(4):698–734, Jun
4399 2019. URL: <https://doi.org/10.1007/s00454-018-0040-y>, doi:
4400 [10.1007/s00454-018-0040-y](https://doi.org/10.1007/s00454-018-0040-y).
- 4401 [21] Gill Barequet and Sariel Har-Peled. Polygon-containment and transla-
4402 tional min-Hausdorff-distance between segments sets are 3SUM-hard.
4403 *International Journal of Computational Geometry & Applications*,
4404 11(04):465–474, 2001.
- 4405 [22] Saugata Basu, Richard Pollack, and Marie-Françoise Roy. Computing
4406 roadmaps of semi-algebraic sets (extended abstract). In *STOC*, pages
4407 168–173. ACM, 1996.
- 4408 [23] Saugata Basu, Richard Pollack, and Marie-Françoise Roy. *Algorithms*
4409 *in real algebraic geometry*, volume 10 of *Algorithms and Computation*
4410 *in Mathematics*. Springer, 2006.
- 4411 [24] Tim Baumann and Torben Hagerup. Rank-select indices without tears.
4412 In *WADS*, 2019.
- 4413 [25] Michael Ben-Or. Lower bounds for algebraic computation trees. In
4414 *STOC*, pages 80–86. ACM, 1983.
- 4415 [26] Jon Louis Bentley, Dorothea Haken, and James B Saxe. A general
4416 method for solving divide-and-conquer recurrences. *ACM SIGACT News*,
4417 12(3):36–44, 1980.
- 4418 [27] Marshall W. Bern, David Eppstein, Paul E. Plassmann, and F. Frances
4419 Yao. Horizon theorems for lines and polygons. In *Discrete and Compu-
4420 tational Geometry*, volume 6 of *DIMACS Series in Discrete Mathemat-
4421 ics and Theoretical Computer Science*, pages 45–66. DIMACS/AMS,
4422 1990.
- 4423 [28] Anders Björner, Michel Las Vergnas, Bernd Sturmfels, Neil White, and
4424 Günter M Ziegler. Oriented matroids. In *Encyclopedia of Mathematics*,
4425 volume 46. Cambridge University Press, 1993.

- 4426 [29] Anselm Blumer, Andrzej Ehrenfeucht, David Haussler, and Manfred K.
4427 Warmuth. Learnability and the vapnik-chervonenkis dimension. *J.*
4428 *ACM*, 36(4):929–965, 1989.
- 4429 [30] Jürgen Bokowski, Susanne Mock, and Ileana Streinu. On the Folkman-
4430 Lawrence topological representation theorem for oriented matroids of
4431 rank 3. *European Journal of Combinatorics*, 22(5):601–615, 2001.
- 4432 [31] Jürgen Bokowski, Jürgen Richter-Gebert, and Werner Schindler. On
4433 the distribution of order types. *Computational Geometry*, 1(3):127–142,
4434 1992.
- 4435 [32] Peter Brass, William O. J. Moser, and János Pach. *Research problems*
4436 *in discrete geometry*. Springer, 2005.
- 4437 [33] David Bremner, Timothy M. Chan, Erik D. Demaine, Jeff Erickson,
4438 Ferran Hurtado, John Iacono, Stefan Langerman, Mihai Pătrașcu, and
4439 Perouz Taslakian. Necklaces, convolutions, and X+Y. *Algorithmica*,
4440 69(2):294–314, 2014.
- 4441 [34] Hervé Brönnimann, Bernard Chazelle, and Jirí Matoušek. Product
4442 range spaces, sensitive sampling, and derandomization. *SIAM J.*
4443 *Comput.*, 28(5):1552–1575, 1999.
- 4444 [35] Robert Creighton Buck. Partition of space. *The American Mathematical*
4445 *Monthly*, 50(9):541–544, 1943.
- 4446 [36] Peter Bürgisser, Michael Clausen, and Mohammad Amin Shokrollahi.
4447 *Algebraic complexity theory*, volume 315 of *Grundlehren der*
4448 *mathematischen Wissenschaften*. Springer, 1997.
- 4449 [37] Sergio Cabello, Jean Cardinal, John Iacono, Stefan Langerman, Pat
4450 Morin, and Aurélien Ooms. Encoding 3SUM. *ArXiv e-prints*, 2019.
4451 [arXiv:1903.02645 \[cs.DS\]](https://arxiv.org/abs/1903.02645).
- 4452 [38] Jean Cardinal, Timothy M. Chan, John Iacono, Stefan Langerman,
4453 and Aurélien Ooms. Subquadratic encodings for point configurations.
4454 In *Symposium on Computational Geometry*, volume 99 of *LIPICS*, pages
4455 20:1–20:14. Schloss Dagstuhl - Leibniz-Zentrum fuer Informatik, 2018.

- [4456] [39] Jean Cardinal, Samuel Fiorini, Gwenaël Joret, Raphaël M Jungers,
[4457] and J Ian Munro. An efficient algorithm for partial order production.
[4458] *SIAM journal on computing*, 39(7):2927–2940, 2010.
- [4459] [40] Jean Cardinal, Samuel Fiorini, Gwenaël Joret, Raphaël M Jungers,
[4460] and J Ian Munro. Sorting under partial information (without the
[4461] ellipsoid algorithm). *Combinatorica*, 33(6):655–697, 2013.
- [4462] [41] Jean Cardinal, John Iacono, and Aurélien Ooms. Solving k -SUM using
[4463] few linear queries. In *ESA*, volume 57 of *LIPICS*, pages 25:1–25:17,
[4464] 2016.
- [4465] [42] Marco L. Carmosino, Jiawei Gao, Russell Impagliazzo, Ivan Mihajlin,
[4466] Ramamohan Paturi, and Stefan Schneider. Nondeterministic extensions
[4467] of the strong exponential time hypothesis and consequences for non-
[4468] reducibility. In *ITCS*, pages 261–270. ACM, 2016.
- [4469] [43] Bob F Caviness and Jeremy R Johnson. *Quantifier elimination and*
[4470] *cylindrical algebraic decomposition*. Springer, 2012.
- [4471] [44] Timothy M. Chan. All-pairs shortest paths with real weights in
[4472] $O(n^3 / \log n)$ time. *Algorithmica*, 50(2):236–243, 2008.
- [4473] [45] Timothy M. Chan. More algorithms for all-pairs shortest paths in
[4474] weighted graphs. *SIAM J. Comput.*, 39(5):2075–2089, 2010.
- [4475] [46] Timothy M. Chan. More logarithmic-factor speedups for 3SUM, (me-
[4476] dian, +)-convolution, and some geometric 3SUM-hard problems. In
[4477] *SODA*, pages 881–897. SIAM, 2018.
- [4478] [47] Timothy M. Chan and Moshe Lewenstein. Clustered integer 3SUM
[4479] via additive combinatorics. In *STOC*, pages 31–40. ACM, 2015.
- [4480] [48] Bernard Chazelle. Cutting hyperplanes for divide-and-conquer. *Dis-
[4481] crete & Computational Geometry*, 9:145–158, 1993.
- [4482] [49] Bernard Chazelle, Herbert Edelsbrunner, Leonidas J. Guibas, and
[4483] Micha Sharir. A singly exponential stratification scheme for real
[4484] semi-algebraic varieties and its applications. *Theor. Comput. Sci.*,
[4485] 84(1):77–105, 1991.

- 4486 [50] Bernard Chazelle, Leonidas J. Guibas, and D. T. Lee. The power of
4487 geometric duality. *BIT*, 25(1):76–90, 1985.
- 4488 [51] Bernard Chazelle and Jirí Matoušek. On linear-time deterministic
4489 algorithms for optimization problems in fixed dimension. *J. Algorithms*,
4490 21(3):579–597, 1996.
- 4491 [52] Otfried Cheong, Ketan Mulmuley, and Edgar Ramos. Randomization
4492 and derandomization. In *Handbook of Discrete and Computational*
4493 *Geometry, 2nd Ed.*, pages 895–926. Chapman and Hall/CRC, 2004.
- 4494 [53] Kenneth L. Clarkson. A randomized algorithm for closest-point queries.
4495 *SIAM J. Comput.*, 17(4):830–847, 1988.
- 4496 [54] George E. Collins. Hauptvortrag: Quantifier elimination for real closed
4497 fields by cylindrical algebraic decomposition. In *Automata Theory*
4498 and *Formal Languages*, volume 33 of *LNCS*, pages 134–183. Springer,
4499 1975.
- 4500 [55] Thomas H Cormen, Charles E Leiserson, Ronald L Rivest, and Clifford
4501 Stein. *Introduction to algorithms (3rd Ed.)*. MIT press, 2009.
- 4502 [56] David Cox, John Little, and Donal O’shea. *Ideals, Varieties, and*
4503 *Algorithms: An Introduction to Computational Algebraic Geometry*
4504 and *Commutative Algebra*. Undergraduate Texts in Mathematics.
4505 Springer, 2007.
- 4506 [57] James H. Davenport and Joos Heintz. Real quantifier elimination is
4507 doubly exponential. *J. Symb. Comput.*, 5(1/2):29–35, 1988.
- 4508 [58] David P. Dobkin and Richard J. Lipton. On some generalizations of
4509 binary search. In *Symposium on Theory of Computing (STOC 1974)*,
4510 pages 310–316, 1974.
- 4511 [59] David P. Dobkin and Richard J. Lipton. A lower bound of the $\frac{1}{2}n^2$
4512 on linear search programs for the knapsack problem. *J. Comput. Syst.*
4513 *Sci.*, 16(3):413–417, 1978.
- 4514 [60] Yevgeniy Dodis, Mihai Pătraşcu, and Mikkel Thorup. Changing base
4515 without losing space. In *Proceedings of the 42nd ACM Symposium on*

- 4516 *Theory of Computing, STOC 2010, Cambridge, Massachusetts, USA,*
4517 5-8 June 2010, pages 593–602, 2010.
- 4518 [61] Herbert Edelsbrunner. *Algorithms in Combinatorial Geometry*, vol-
4519 umne 10. Springer Science & Business Media, 2012.
- 4520 [62] Herbert Edelsbrunner, Leonidas J. Guibas, János Pach, Richard Pol-
4521 lack, Raimund Seidel, and Micha Sharir. Arrangements of curves in
4522 the plane - topology, combinatorics and algorithms. *Theor. Comput.
4523 Sci.*, 92(2):319–336, 1992.
- 4524 [63] Herbert Edelsbrunner, Joseph O’Rourke, and Raimund Seidel. Con-
4525 structing arrangements of lines and hyperplanes with applications.
4526 *SIAM J. Comput.*, 15(2):341–363, 1986.
- 4527 [64] Herbert Edelsbrunner, Raimund Seidel, and Micha Sharir. On the zone
4528 theorem for hyperplane arrangements. *SIAM J. Comput.*, 22(2):418–
4529 429, 1993.
- 4530 [65] György Elekes and Lajos Rónyai. A combinatorial problem on poly-
4531 nomials and rational functions. *J. Comb. Theory, Ser. A*, 89(1):1–20,
4532 2000.
- 4533 [66] György Elekes and Endre Szabó. How to find groups? (and how to
4534 use them in Erdős geometry?). *Combinatorica*, 32(5):537–571, 2012.
- 4535 [67] David Eppstein. *Forbidden Configurations in Discrete Geometry*. Cam-
4536 bridge University Press, 2018.
- 4537 [68] Jeff Erickson. New lower bounds for Hopcroft’s problem. *Discrete &
4538 Computational Geometry*, 16(4):389–418, 1996.
- 4539 [69] Jeff Erickson. Lower bounds for linear satisfiability problems. *Chicago
4540 J. Theor. Comput. Sci.*, 1999.
- 4541 [70] Jeff Erickson. New lower bounds for convex hull problems in odd
4542 dimensions. *SIAM J. Comput.*, 28(4):1198–1214, 1999.
- 4543 [71] Hazel Everett, Ferran Hurtado, and Marc Noy. Stabbing information
4544 of a simple polygon. *Discrete Applied Mathematics*, 91(1-3):67–82,
4545 1999.

- 4546 [72] Esther Ezra and Micha Sharir. A nearly quadratic bound for the
4547 decision tree complexity of k-sum. In *Symposium on Computational*
4548 *Geometry*, volume 77 of *LIPics*, pages 41:1–41:15. Schloss Dagstuhl -
4549 Leibniz-Zentrum fuer Informatik, 2017.
- 4550 [73] Stefan Felsner. On the number of arrangements of pseudolines. In
4551 *SoCG*, pages 30–37. ACM, 1996.
- 4552 [74] Stefan Felsner and Pavel Valtr. Coding and counting arrangements
4553 of pseudolines. *Discrete & Computational Geometry*, 46(3):405–416,
4554 2011.
- 4555 [75] Jon Folkman and Jim Lawrence. Oriented matroids. *Journal of*
4556 *Combinatorial Theory, Series B*, 25(2):199–236, 1978.
- 4557 [76] Hervé Fournier. *Complexité et expressibilité sur les réels*. PhD thesis,
4558 École normale supérieure de Lyon, 2001.
- 4559 [77] Michael L. Fredman. How good is the information theory bound in
4560 sorting? *Theor. Comput. Sci.*, 1(4):355–361, 1976.
- 4561 [78] Michael L. Fredman, János Komlós, and Endre Szemerédi. Storing a
4562 sparse table with $O(1)$ worst case access time. *J. ACM*, 31(3):538–544,
4563 1984.
- 4564 [79] Ari Freund. Improved subquadratic 3SUM. *Algorithmica*, pages 1–19,
4565 2015.
- 4566 [80] Anka Gajentaan and Mark H. Overmars. On a class of $O(n^2)$ problems
4567 in computational geometry. *Comput. Geom.*, 5:165–185, 1995.
- 4568 [81] Omer Gold and Micha Sharir. Improved bounds for 3SUM, k-SUM,
4569 and linear degeneracy. In *ESA*, volume 87 of *LIPics*, pages 42:1–42:13.
4570 Schloss Dagstuhl - Leibniz-Zentrum fuer Informatik, 2017.
- 4571 [82] Jacob E. Goodman. Proof of a conjecture of Burr, Grünbaum, and
4572 Sloane. *Discrete Mathematics*, 32(1):27–35, 1980.
- 4573 [83] Jacob E. Goodman. Pseudoline arrangements. In *Handbook of Discrete*
4574 *and Computational Geometry, 2nd Ed.*, pages 97–128. Chapman and
4575 Hall/CRC, 2004.

- 4576 [84] Jacob E. Goodman and Richard Pollack. Proof of Grünbaum’s con-
4577 jecture on the stretchability of certain arrangements of pseudolines.
4578 *Journal of Combinatorial Theory, Series A*, 29(3):385–390, 1980.
- 4579 [85] Jacob E. Goodman and Richard Pollack. Multidimensional sorting.
4580 *SIAM Journal on Computing*, 12(3):484–507, 1983.
- 4581 [86] Jacob E. Goodman and Richard Pollack. Semispaces of configurations,
4582 cell complexes of arrangements. *Journal of Combinatorial Theory,*
4583 *Series A*, 37(3):257–293, 1984.
- 4584 [87] Jacob E. Goodman and Richard Pollack. Upper bounds for config-
4585 urations and polytopes in \mathbb{R}^d . *Discrete & Computational Geometry*,
4586 1:219–227, 1986.
- 4587 [88] Jacob E. Goodman and Richard Pollack. The complexity of point
4588 configurations. *Discrete Applied Mathematics*, 31(2):167–180, 1991.
- 4589 [89] Jacob E. Goodman and Richard Pollack. Allowable sequences and
4590 order types in discrete and computational geometry. In *New Trends*
4591 in *Discrete and Computational Geometry*, pages 103–134. Springer,
4592 1993.
- 4593 [90] Jacob E. Goodman, Richard Pollack, and Bernd Sturmfels. Coordinate
4594 representation of order types requires exponential storage. In *STOC*,
4595 pages 405–410. ACM, 1989.
- 4596 [91] Allan Grønlund and Seth Pettie. Threesomes, degenerates, and love
4597 triangles. *J. ACM*, 65(4):22:1–22:25, 2018.
- 4598 [92] Branko Grünbaum. *Convex Polytopes*. Springer, 2005.
- 4599 [93] Dan Halperin. Arrangements. In *Handbook of Discrete and Compu-*
4600 *tational Geometry, 2nd Ed.*, pages 529–562. Chapman and Hall/CRC,
4601 2004.
- 4602 [94] Heiko Harborth and Meinhard Möller. The Esther Klein problem in
4603 the projective plane. *J. Combin. Math. Combin. Comput.*, 15:171–179,
4604 1994.

- [95] Joe Harris. *Algebraic geometry: a first course*, volume 133. Springer, 2013.
- [96] Robin Hartshorne. *Algebraic geometry*, volume 52. Springer, 1977.
- [97] David Haussler and Emo Welzl. ε -nets and simplex range queries. *Discrete & Computational Geometry*, 2(1):127–151, 1987.
- [98] Monika Henzinger, Sebastian Krinnerger, Danupon Nanongkai, and Thatchaphol Saranurak. Unifying and strengthening hardness for dynamic problems via the online matrix-vector multiplication conjecture. In *STOC*, pages 21–30. ACM, 2015.
- [99] Max Hopkins, Daniel M. Kane, and Shachar Lovett. The power of comparisons for actively learning linear classifiers. *ArXiv e-prints*, 2019. [arXiv:1907.03816 \[cs.LG\]](https://arxiv.org/abs/1907.03816).
- [100] Alfredo Hubard, Luis Montejano, Emilio Mora, and Andrew Suk. Order types of convex bodies. *Order*, 28(1):121–130, 2011.
- [101] Guy Joseph Jacobson. *Succinct static data structures*. PhD thesis, Carnegie Mellon University, 1988.
- [102] Daniel M. Kane, Shachar Lovett, and Shay Moran. Generalized comparison trees for point-location problems. In *ICALP*, volume 107 of *LIPICS*, pages 82:1–82:13. Schloss Dagstuhl - Leibniz-Zentrum fuer Informatik, 2018.
- [103] Daniel M. Kane, Shachar Lovett, and Shay Moran. Near-optimal linear decision trees for k-sum and related problems. In *STOC*, pages 554–563. ACM, 2018.
- [104] Donald E. Knuth. *Axioms and Hulls*, volume 606 of *Lecture Notes in Computer Science*. Springer, 1992.
- [105] Tsvi Kopelowitz, Seth Pettie, and Ely Porat. Higher lower bounds from the 3SUM conjecture. In *SODA*, pages 1272–1287. SIAM, 2016.
- [106] Hsiang-Tsung Kung. Fast evaluation and interpolation. Technical report, Carnegie Mellon University, 1973.

- 4634 [107] Hsiang-Tsung Kung. A new upper bound on the complexity of derivative evaluation. Technical report, Carnegie Mellon University, 1973.
4635
- 4636 [108] Friedrich Levi. Die teilung der projektiven ebene durch gerade oder
4637 pseudogerade. *Ber. Math.-Phys. Kl. Sächs. Akad. Wiss.*, 78:256–267,
4638 1926.
- 4639 [109] Jirí Matoušek. Range searching with efficient hierarchical cutting.
4640 *Discrete & Computational Geometry*, 10:157–182, 1993.
- 4641 [110] Jirí Matoušek. Approximations and optimal geometric divide-and-
4642 conquer. *J. Comput. Syst. Sci.*, 50(2):203–208, 1995.
- 4643 [111] Jirí Matoušek. Derandomization in computational geometry. *J. Algorithms*,
4644 20(3):545–580, 1996.
- 4645 [112] Yoshitake Matsumoto, Sonoko Moriyama, Hiroshi Imai, and David
4646 Bremner. Matroid enumeration for incidence geometry. *Discrete &*
4647 *Computational Geometry*, 47(1):17–43, 2012.
- 4648 [113] Stefan Meiser. Point location in arrangements of hyperplanes. *Information*
4649 and Computation
- 4650 [114] Friedhelm Meyer auf der Heide. A polynomial linear search algorithm
4651 for the n -dimensional knapsack problem. *J. ACM*, 31(3):668–676,
4652 1984.
- 4653
- 4654 [115] John Milnor. On the Betti numbers of real varieties. *Proceedings of*
4655 *the American Mathematical Society*, 15(2):275–280, 1964.
- 4656 [116] Bhubaneswar Mishra. Computational real algebraic geometry. In
4657 *Handbook of Discrete and Computational Geometry*, 2nd Ed., pages
4658 743–764. Chapman and Hall/CRC, 2004.
- 4659 [117] Joseph S. B. Mitchell and Joseph O'Rourke. Computational geometry
4660 column 42. *Int. J. Comput. Geometry Appl.*, 11(5):573–582, 2001.
- 4661 [118] Nikolai E Mnëv. On manifolds of combinatorial types of projective
4662 configurations and convex polyhedra. In *Soviet Math. Doklady*, volume 32,
4663 pages 335–337, 1985.

- [119] Nikolai E Mnëv. The universality theorems on the classification problem of configuration varieties and convex polytopes varieties. In *Topology and geometry—Rohlin seminar*, pages 527–543. Springer, 1988.
- [120] H. Nassajian Mojarrad, T. Pham, C. Valculescu, and F. de Zeeuw. Schwartz-Zippel bounds for two-dimensional products. *ArXiv e-prints*, 2016. arXiv:1507.08181 [math.CO].
- [121] Jaroslav Nešetřil and Pavel Valtr. A Ramsey property of order types. *Journal of Combinatorial Theory, Series A*, 81(1):88–107, 1998.
- [122] János Pach and Micha Sharir. On the number of incidences between points and curves. *Combinatorics, Probability & Computing*, 7(1):121–127, 1998.
- [123] János Pach and Micha Sharir. Combinatorial geometry with algorithmic applications – the Alcalá lectures. *AMS, Providence*, 2009.
- [124] Rasmus Pagh and Flemming Friche Rodler. Cuckoo hashing. *J. Algorithms*, 51(2):122–144, 2004.
- [125] Mihai Pătrașcu. Towards polynomial lower bounds for dynamic problems. In *Symposium on Theory of Computing (STOC 2010)*, pages 603–610. ACM, 2010.
- [126] Mihai Pătrașcu and Ryan Williams. On the possibility of faster SAT algorithms. In *Proceedings of the Twenty-First Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2010, Austin, Texas, USA, January 17-19, 2010*, pages 1065–1075, 2010.
- [127] Franco P. Preparata and Michael Ian Shamos. *Computational Geometry - An Introduction*. Texts and Monographs in Computer Science. Springer, 1985.
- [128] Michael O. Rabin. Proving simultaneous positivity of linear forms. *J. Comput. Syst. Sci.*, 6(6):639–650, 1972.
- [129] Prabhakar Raghavan. Probabilistic construction of deterministic algorithms: Approximating packing integer programs. *J. Comput. Syst. Sci.*, 37(2):130–143, 1988.

- [130] Rajeev Raman, Venkatesh Raman, and Srinivasa Rao Satti. Succinct dynamic data structures. In *WADS*, volume 2125 of *Lecture Notes in Computer Science*, pages 426–437. Springer, 2001.
- [131] Rajeev Raman, Venkatesh Raman, and Srinivasa Rao Satti. Succinct indexable dictionaries with applications to encoding k -ary trees, prefix sums and multisets. *ACM Trans. Algorithms*, 3(4):43, 2007.
- [132] Orit E. Raz and Micha Sharir. The number of unit-area triangles in the plane: Theme and variations. In *SoCG*, volume 34 of *LIPICS*, pages 569–583, 2015.
- [133] Orit E. Raz, Micha Sharir, and Frank de Zeeuw. Polynomials vanishing on cartesian products: The Elekes-Szabó theorem revisited. In *SoCG*, volume 34 of *LIPICS*, pages 522–536, 2015.
- [134] Orit E. Raz, Micha Sharir, and Frank de Zeeuw. The Elekes-Szabó theorem in four dimensions. *ArXiv e-prints*, 2016. arXiv:1607.03600 [math.CO].
- [135] Orit E. Raz, Micha Sharir, and Ilya D. Shkredov. On the number of unit-area triangles spanned by convex grids in the plane. *Comput. Geom.*, 62:25–33, 2017.
- [136] Orit E. Raz, Micha Sharir, and János Solymosi. Polynomials vanishing on grids: The Elekes-Rónyai problem revisited. In *SoCG*, page 251. ACM, 2014.
- [137] Orit E. Raz, Micha Sharir, and János Solymosi. On triple intersections of three families of unit circles. *Discrete & Computational Geometry*, 54(4):930–953, 2015.
- [138] Jürgen Richter. Kombinatorische realisierbarkeitskriterien für orientierte matroide. *Mitt. Math. Sem. Univ. Giessen*, 194:1–112, 1989.
- [139] Jürgen Richter-Gebert and Günter M. Ziegler. Oriented matroids. In *Handbook of Discrete and Computational Geometry, 2nd Ed.*, pages 129–151. Chapman and Hall/CRC, 2004.
- [140] Gerhard Ringel. Teilungen der ebene durch geraden oder topologische geraden. *Mathematische Zeitschrift*, 64(1):79–102, 1956.

- [141] Norbert Sauer. On the density of families of sets. *Journal of Combinatorial Theory, Series A*, 13(1):145–147, 1972.
- [142] Abraham Seidenberg. Constructions in algebra. *Transactions of the AMS*, 197:273–313, 1974.
- [143] Saharon Shelah. A combinatorial problem; stability and order for models and theories in infinitary languages. *Pacific Journal of Mathematics*, 41(1):247–261, 1972.
- [144] Jack Snoeyink. Point location. In *Handbook of Discrete and Computational Geometry, 2nd Ed.*, pages 767–785. Chapman and Hall/CRC, 2004.
- [145] Joel Spencer. *Ten lectures on the probabilistic method, 2nd Ed.*, volume 64. SIAM, 1994.
- [146] J. Michael Steele and Andrew Yao. Lower bounds for algebraic decision trees. *J. Algorithms*, 3(1):1–8, 1982.
- [147] William L. Steiger and Ileana Streinu. A pseudo-algorithmic separation of lines from pseudo-lines. *Information Processing Letters*, 53(5):295–299, 1995.
- [148] Volker Strassen. Die berechnungskomplexität von elementarsymmetrischen funktionen und von interpolationskoeffizienten. *Numerische Mathematik*, 20(3):238–251, 1973.
- [149] Ileana Streinu. Clusters of stars. In *SoCG*, pages 439–441. ACM, 1997.
- [150] Alfred Tarski. *A Decision Method for Elementary Algebra and Geometry*, pages 24–84. Springer Vienna, Vienna, 1998.
- [151] René Thom. Sur l’homologie des variétés algébriques. In *Differential and Combinatorial Topology (A Symposium in Honor of Marston Morse)*, pages 255–265, 1965.
- [152] VN Vapnik and A Ya Chervonenkis. On the uniform convergence of relative frequencies of events to their probabilities. *Theory of Probability & Its Applications*, 16(2):264–280, 1971.

- 4754 [153] Sebastiano Vigna. Broadword implementation of rank/select queries.
4755 In *WEA*, volume 5038 of *Lecture Notes in Computer Science*, pages
4756 154–168. Springer, 2008.
- 4757 [154] Ryan Williams. A new algorithm for optimal 2-constraint satisfaction
4758 and its implications. *Theor. Comput. Sci.*, 348(2-3):357–365, 2005.
- 4759 [155] Andrew Yao. A lower bound to finding convex hulls. *J. ACM*, 28(4):780–
4760 787, 1981.
- 4761 [156] Andrew Chi-Chih Yao. On parallel computation for the knapsack
4762 problem. *J. ACM*, 29(3):898–903, 1982.
- 4763 [157] Andrew Chi-Chih Yao and F. Frances Yao. A general approach to
4764 d-dimensional geometric queries. In *STOC*, pages 163–168. ACM,
4765 1985.
- 4766 [158] Chee K. Yap. Robust geometric computation. In *Handbook of Discrete*
4767 and *Computational Geometry*, 2nd Ed., pages 927–952. Chapman and
4768 Hall/CRC, 2004.
- 4769 [159] David Yun. On square-free decomposition algorithms. In *SYMSACC*,
4770 pages 26–35, 1976.

4771

List of Contributions and Open 4772 Questions

4773	1	Contribution	40
4774	2	Contribution	41
4775	3	Contribution	43
4776	4	Contribution	43
4777	5	Contribution	43
4778	6	Contribution	43
4779	7	Contribution	45
4780	8	Contribution	45
4781	9	Contribution	45
4782	10	Contribution	45
4783	11	Contribution (Contributions 17 and 18)	45
4784	12	Contribution	46
4785	13	Contribution	46
4786	14	Contribution	47
4787	15	Contribution	48
4788	16	Contribution	48
4789	1	Open Question	52
4790	2	Open Question	52
4791	3	Open Question	52
4792	4	Open Question	52
4793	5	Open Question	53
4794	6	Open Question	53
4795	7	Open Question	53
4796	8	Open Question	53
4797	9	Open Question	53
4798	10	Open Question	54
4799	11	Open Question	54

4800	12	Open Question	54
4801	13	Open Question	54
4802	14	Open Question	54
4803	1	Contribution	77
4804	2	Contribution	85
4805	3	Contribution	94
4806	4	Contribution	101
4807	5	Contribution	104
4808	6	Contribution	111
4809	7	Contribution	131
4810	8	Contribution	131
4811	17	Contribution	131
4812	9	Contribution	142
4813	10	Contribution	142
4814	18	Contribution	142
4815	12	Contribution	149
4816	13	Contribution	149
4817	14	Contribution	155

Solving k -SUM using few linear queries

Jean Cardinal^{*1}, John Iacono^{†2}, and Aurélien Ooms^{‡1}

1 Université libre de Bruxelles (ULB)

Brussels, Belgium

{jcardin,aureooms}@ulb.ac.be

2 New York University

New York, United States of America

icalp2016submission@johniacono.com

Abstract

The k -SUM problem is given n input real numbers to determine whether any k of them sum to zero. The problem is of tremendous importance in the emerging field of complexity theory within P , and it is in particular open whether it admits an algorithm of complexity $O(n^c)$ with $c < \lceil \frac{k}{2} \rceil$. Inspired by an algorithm due to Meiser (1993), we show that there exist linear decision trees and algebraic computation trees of depth $O(n^3 \log^3 n)$ solving k -SUM. Furthermore, we show that there exists a randomized algorithm that runs in $\tilde{O}(n^{\lceil \frac{k}{2} \rceil + 8})$ time, and performs $O(n^3 \log^3 n)$ linear queries on the input. Thus, we show that it is possible to have an algorithm with a runtime almost identical (up to the +8) to the best known algorithm but for the first time also with the number of queries on the input a polynomial that is independent of k . The $O(n^3 \log^3 n)$ bound on the number of linear queries is also a tighter bound than any known algorithm solving k -SUM, even allowing unlimited total time outside of the queries. By simultaneously achieving few queries to the input without significantly sacrificing runtime vis-à-vis known algorithms, we deepen the understanding of this canonical problem which is a cornerstone of complexity-within- P .

We also consider a range of tradeoffs between the number of terms involved in the queries and the depth of the decision tree. In particular, we prove that there exist $o(n)$ -linear decision trees of depth $o(n^4)$.

1998 ACM Subject Classification F.2.2 Nonnumerical Algorithms and Problems

Keywords and phrases k -SUM problem, linear decision trees, point location, ε -nets

Digital Object Identifier 10.4230/LIPIcs.xxx.yyy.p

1 Introduction

The k -SUM problem is defined as follows: given a collection of real numbers, and a constant k , decide whether any k of them sum to zero. It is a fixed-parameter version of the subset-sum problem, a standard NP -complete problem. The k -SUM problem, and in particular the special case of 3SUM, has proved to be a cornerstone of the fine-grained complexity program aiming at the construction of a complexity theory for problems in P . In particular, there are deep connections between the complexity of k -SUM, the Strong

* Supported by the “Action de Recherche Concertée” (ARC) COPHYMA, convention number 4.110.H.000023.

† Research partially completed while on sabbatical at the Algorithms Research Group of the Département d’Informatique at the Université Libre de Bruxelles with support from a Fulbright Research Fellowship, the Fonds de la Recherche Scientifique — FNRS, and NSF grants CNS-1229185, CCF-1319648, and CCF-1533564.

‡ Supported by the Fund for Research Training in Industry and Agriculture (FRIA).



© Jean Cardinal, John Iacono, and Aurélien Ooms;
licensed under Creative Commons License CC-BY

Conference title on which this volume is based on.

Editors: Billy Editor and Bill Editors; pp. 1–16

 Leibniz International Proceedings in Informatics
Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

Exponential Time Hypothesis [28, 12], and the complexity of many other major problems in P [19, 8, 26, 27, 5, 2, 22, 24, 1, 3, 13].

It has been long known that the k -SUM problem can be solved in time $O(n^{\frac{k}{2}} \log n)$ for even k , and $O(n^{\frac{k+1}{2}})$ for odd k . Erickson [17] proved a near-matching lower bound in the k -linear decision tree model. In this model, the complexity is measured by the depth of a decision tree, every node of which corresponds to a query of the form $q_{i_1} + q_{i_2} + \dots + q_{i_k} \leq^? 0$, where q_1, q_2, \dots, q_n are the input numbers. In a recent breakthrough paper, Grønlund and Pettie [22] showed that in the $(2k - 2)$ -linear decision tree model, where queries test the sign of weighted sums of up to $2k - 2$ input numbers, only $O(n^{\frac{k}{2}} \sqrt{\log n})$ queries are required for odd values of k . In particular, there exists a 4-linear decision tree for 3SUM of depth $\tilde{O}(n^{\frac{3}{2}})$ (here the notation \tilde{O} ignores polylogarithmic factors), while every 3-linear decision tree has depth $\Omega(n^2)$ [17]. This indicates that increasing the size of the queries, defined as the maximum number of input numbers involved in a query, can yield significant improvements on the depth of the minimal-height decision tree. Ailon and Chazelle [4] slightly extended the range of query sizes for which a nontrivial lower bound could be established, elaborating on Erickson's technique.

It has been well established that there exist nonuniform polynomial-time algorithms for the subset-sum problem. One of them was described by Meiser [25], and is derived from a data structure for point location in arrangements of hyperplanes using the bottom vertex decomposition. This algorithm can be cast as the construction of a linear decision tree in which the queries have non-constant size.

1.1 Our results.

Our first contribution, in Section 3, is through a careful implementation of Meiser's basic algorithm idea [25] to show the existence of an n -linear decision tree of depth $\tilde{O}(n^3)$ for k -SUM. Although the high-level algorithm itself is not new, we refine the implementation and analysis for the k -SUM problem. Meiser presented his algorithm as a general method of point location in \mathcal{H} given n -dimensional hyperplanes that yielded a depth $\tilde{O}(n^4 \log |\mathcal{H}|)$ linear decision tree; when viewing the k -SUM problem as a point location problem, \mathcal{H} is $O(n^k)$ and thus Meiser's algorithm can be viewed as giving a $\tilde{O}(n^4)$ -height linear decision tree. We reduce this height to $\tilde{O}(n^3)$ for the particular hyperplanes that an instance k -SUM problem presents when viewed as a point location problem. While the original algorithm was cast as a nonuniform polynomial-time algorithm, we show that it can be implemented in the linear decision tree model. This first result also implies the existence of nonuniform RAM algorithms with the same complexity, as shown in Appendix B.

There are two subtleties to this result. The first is inherent to the chosen complexity model: even if the number of queries to the input is small (in particular, the degree of the polynomial complexity is invariant on k), the time required to *determine which queries should be performed* may be arbitrary. In a naïve analysis, we show it can be trivially bounded by $\tilde{O}(n^{k+2})$. In Section 4 we present an algorithm to choose which decisions to perform whereby the running time can be reduced to $\tilde{O}(n^{\frac{k}{2}+8})$. Hence, we obtain an $\tilde{O}(n^{\frac{k}{2}+8})$ time randomized algorithm in the RAM model expected to perform $\tilde{O}(n^3)$ linear queries on the input¹.

¹ Grønlund and Pettie [22] mention the algorithms of auf Der Heyde [6] and Meiser [25], and state “(...) it was known that all k -LDT problems can be solved by n -linear decision trees with depth $O(n^5 \log n)$ [25], or with depth $O(n^4 \log(nK))$ if the coefficients of the linear function are integers with absolute value at most K [6]. Unfortunately these decision trees are not efficiently constructible. The time required to

	# blocks	query size	# queries	time
Theorem 3	1	n	$\tilde{O}(n^3)$	$\tilde{O}(n^{\lceil \frac{k}{2} \rceil + 8})$
Theorem 11	b	$k \lceil \frac{n}{b} \rceil$	$\tilde{O}(b^{k-4} n^3)$	$\tilde{O}(b^{\lfloor \frac{k}{2} \rfloor - 9} n^{\lceil \frac{k}{2} \rceil + 8})$
Corollary 12	$b = \frac{k}{\varepsilon}$	εn	$\tilde{O}(n^3)$	$\tilde{O}(n^{\lceil \frac{k}{2} \rceil + 8})$
Corollary 13	$b = O(n^\varepsilon)$	$O(n^{1-\varepsilon})$	$\tilde{O}(n^{3+(k-4)\varepsilon})$	$\tilde{O}(n^{(1+\varepsilon)\frac{k}{2} + 8.5})$

■ **Table 1** Complexities of our new algorithms for the k -SUM problem. The query size is the maximum number of elements of the input that can be involved in a single linear query. The number of blocks is a parameter that allows us to change the query size (see Section 5). The origin of the constant in the exponent of the time complexity is due to Lemma 9. We conjecture it can be reduced, though substantial changes in the analysis will likely be needed to do so.

The second issue we address is that the linear queries in the above algorithm may have size n , that is, they may use all the components of the input. The lower bound of Erickson shows that if the queries are of minimal size, the number of queries cannot be a polynomial independent of k such as what we obtain, so non-minimal query size is clearly essential to a drastic reduction in the number of queries needed. This gives rise to the natural question as to what is the relation between query size and number of queries. In particular, one natural question is whether queries of size less than n would still allow the problem to be solved using a polynomial independent of k number of queries. We show that this is possible; in Section 5, we introduce a range of algorithms exhibiting an explicit tradeoff between the number of queries and their size. Using a blocking scheme, we show that we can restrict to εn -linear decision trees, for arbitrary ε . We also give a range of tradeoffs for $n^{1-\varepsilon}$ -linear decision trees. Although the proposed algorithms still involve nonconstant-size queries, this is the first time such tradeoffs are explicitly tackled. Table 1 summarizes our results.

2 Definitions and previous work

2.1 Definitions

We consider the k -SUM problem for $k = O(1)$. In what follows, we use the notation $[n] = \{1, 2, \dots, n\}$.

► **Problem (k -SUM).** Given an input vector $q \in \mathbb{R}^n$, decide whether there exists a k -tuple $(i_1, i_2, \dots, i_k) \in [n]^k$ such that $\sum_{j=1}^k q_{i_j} = 0$.

The problem amounts to deciding in n -dimensional space, for each hyperplane H of equation $x_{i_1} + x_{i_2} + \dots + x_{i_k} = 0$, whether q lies on, above, or below H . Hence this indeed amounts to locating the point q in the arrangement formed by those hyperplanes. We emphasize that the set of hyperplanes depends only on k and n and not on the actual input vector q .

Linear degeneracy testing (k -LDT) is a generalization of k -SUM where we have arbitrary rational coefficients² and an independent term in the equations of the hyperplanes.

► **Problem (k -LDT).** Given an input vectors $q \in \mathbb{R}^n$ and $\alpha \in \mathbb{Q}^n$ and constant $c \in \mathbb{Q}$ decide whether there exists a k -tuple $(i_1, i_2, \dots, i_k) \in [n]^k$ such that $c + \sum_{j=1}^k \alpha_j q_{i_j} = 0$.

determine which comparisons to make is exponential.” We prove that the trees can have depth $\tilde{O}(n^3)$ and that the whole algorithm can run in randomized polynomial-time.

² The usual definition of k -LDT allows arbitrary real coefficients. However, the algorithm we provide for Lemma 8 needs the vertices of the arrangement of hyperplanes to have rational coordinates.

Our algorithms apply to this more general problem with only minor changes.

The s -linear decision tree model is a standard model of computation in which several lower bounds for k -SUM have been proven. In the decision tree model, one may ask well-defined questions to an oracle that are answered “yes” or “no.” For s -linear decision trees, a well-defined question consists of testing the sign of a linear function on at most s numbers q_{i_1}, \dots, q_{i_s} of the input q_1, \dots, q_n and can be written as

$$c + \alpha_1 q_{i_1} + \dots + \alpha_s q_{i_s} \stackrel{?}{\leq} 0$$

Each question is defined to cost a single unit. All other operations can be carried out for free but may not examine the input vector q . We refer to n -linear decision trees simply as linear decision trees.

In this paper, we consider algorithms in the standard integer RAM model, but in which the input $q \in \mathbb{R}^n$ is accessible *only* via a linear query oracle. Hence we are not allowed to manipulate the input numbers directly. The complexity is measured in two ways: by counting the total number of queries, just as in the linear decision tree model, and by measuring the overall running time, taking into account the time required to determine the sequence of linear queries. This two-track computation model, in which the running time is distinguished from the query complexity, is commonly used in results on comparison-based sorting problems where analyses of both runtime and comparisons are of interest (see for instance [29, 10, 11]).

2.2 Previous Results.

The seminal paper by Gajentaan and Overmars [19] showed the crucial role of 3SUM in understanding the complexity of several problems in computational geometry. It is an open question whether an $O(n^{2-\varepsilon})$ algorithm exists for 3SUM. Such a result would have a tremendous impact on many other fundamental computational problems [19, 8, 26, 27, 5, 2, 22, 24, 1, 3, 13]. It has been known for long that k -SUM is $W[1]$ -hard. Recently, it was shown to be $W[1]$ -complete by Abboud et al. [1].

In Erickson [17], it is shown that we cannot solve 3SUM in subquadratic time in the 3-linear decision tree model:

► **Theorem 1** (Erickson [17]). *The optimal depth of a k -linear decision tree that solves the k -LDT problem is $\Theta(n^{\lceil \frac{k}{2} \rceil})$.*

The proof uses an adversary argument which can be explained geometrically. As we already observed, we can solve k -LDT problems by modeling them as point location problems in an arrangement of hyperplanes. Solving one such problem amounts to determining which cell of the arrangement contains the input point. The adversary argument of Erickson [17] is that there exists a cell having $\Omega(n^{\lceil \frac{k}{2} \rceil})$ boundary facets and in this model point location in such a cell requires testing each facet.

Ailon and Chazelle [4] study s -linear decision trees to solve the k -SUM problem when $s > k$. In particular, they give an additional proof for the $\Omega(n^{\lceil \frac{k}{2} \rceil})$ lower bound of Erickson [17] and generalize the lower bound for the s -linear decision tree model when $s > k$. Note that the exact lower bound given by Erickson [17] for $s = k$ is $\Omega((nk^{-k})^{\lceil \frac{k}{2} \rceil})$ while the one given by Ailon and Chazelle [4] is $\Omega((nk^{-3})^{\lceil \frac{k}{2} \rceil})$. Their result improves therefore the lower bound for $s = k$ when k is large. The lower bound they prove for $s > k$ is the following

► **Theorem 2** (Ailon and Chazelle [4]). *The depth of an s -linear decision tree solving the k -LDT problem is*

$$\Omega\left((nk^{-3})^{\frac{2k-s}{2\lceil \frac{s-k+1}{2} \rceil}(1-\varepsilon_k)}\right),$$

where $\varepsilon_k > 0$ tends to 0 as $k \rightarrow \infty$.

This lower bound breaks down when $k = \Omega(n^{1/3})$ or $s \geq 2k$ and the cases where $k < 6$ give trivial lower bounds. For example, in the case of 3SUM with $s = 4$ we only get an $\Omega(n)$ lower bound.

As for upper bounds, Baran et al. [7] gave subquadratic Las Vegas algorithms for 3SUM on integer and rational numbers in the circuit RAM, word RAM, external memory, and cache-oblivious models of computation. The idea of their approach is to exploit the parallelism of the models, using linear and universal hashing.

More recently, Grønlund and Pettie [22] proved the existence of a linear decision tree solving the 3SUM problem using a strongly subquadratic number of linear queries. The classical quadratic algorithm for 3SUM uses 3-linear queries while the decision tree of Grønlund and Pettie uses 4-linear queries and requires $O(n^{3/2}\sqrt{\log n})$ of them. Moreover, they show that their decision tree can be used to get better upper bounds for k -SUM when k is odd.

They also provide two subquadratic 3SUM algorithms. A deterministic one running in $O(n^2/(\log n/\log\log n)^{2/3})$ time and a randomized one running in $O(n^2(\log\log n)^2/\log n)$ time with high probability. These results refuted the long-lived conjecture that 3SUM cannot be solved in subquadratic time in the RAM model.

Freund [18] and Gold and Sharir [20] later gave improvements on the results of Grønlund and Pettie [22]. Freund [18] gave a deterministic algorithm for 3SUM running in $O(n^2\log\log n/\log n)$ time. Gold and Sharir [20] gave another deterministic algorithm for 3SUM with the same running time and shaved off the $\sqrt{\log n}$ factor in the decision tree complexities of 3SUM and k -SUM given by Grønlund and Pettie.

Auf der Heide [6] gave the first point location algorithm to solve the knapsack problem in the linear decision tree model in polynomial time. He thereby answers a question raised by Dobkin and Lipton [15, 16], Yao [30] and others. However, if one uses this algorithm to locate a point in an arbitrary arrangement of hyperplanes the running time is increased by a factor linear in the greatest coefficient in the equations of all hyperplanes. On the other hand, the complexity of Meiser's point location algorithm is polynomial in the dimension, logarithmic in the number of hyperplanes and does not depend on the value of the coefficients in the equations of the hyperplanes. A useful complete description of the latter is also given by Bürgisser et al. [9] (Section 3.4).

3 Query complexity

In this section and the next, we prove the following first result.

► **Theorem 3.** *There exist linear decision trees of depth at most $O(n^3\log^3 n)$ solving the k -SUM and the k -LDT problems. Furthermore, for the two problems there exists an $\tilde{O}(n^{\lceil \frac{k}{2} \rceil} + 8)$ Las Vegas algorithm in the RAM model expected to perform $O(n^3\log^3 n)$ linear queries on the input.*

3.1 Algorithm outline

For a fixed set of hyperplanes \mathcal{H} and given input vertex q in \mathbb{R}^n , Meiser's algorithm allows us to determine the cell of the arrangement $\mathcal{A}(\mathcal{H})$ that contains q in its interior (or that is q if q is a 0-cell of $\mathcal{A}(\mathcal{H})$), that is, the positions $\sigma(H, q) \in \{-, 0, +\}$ of q with respect to all hyperplanes $H \in \mathcal{H}$. In the k -SUM problem, the set \mathcal{H} is the set of $\Theta(n^k)$ hyperplanes with

equations of the form $x_{i_1} + x_{i_2} + \dots + x_{i_k} = 0$. These equations can be modified accordingly for k -LDT.

We use standard results on ε -nets. Using a theorem due to Haussler and Welzl [23], it is possible to construct an ε -net \mathcal{N} for the range space defined by hyperplanes and simplices using a random uniform sampling on \mathcal{H} .

► **Theorem 4** (Haussler and Welzl [23]). *If we choose $O(\frac{n^2}{\varepsilon} \log^2 \frac{n}{\varepsilon})$ hyperplanes of \mathcal{H} uniformly at random and denote this selection \mathcal{N} then for any simplex intersected by more than $\varepsilon|\mathcal{H}|$ hyperplanes of \mathcal{H} , with high probability, at least one of the intersecting hyperplanes is contained in \mathcal{N} .*

The contrapositive states that if no hyperplane in \mathcal{N} intersects a given simplex, then with high probability the number of hyperplanes of \mathcal{H} intersecting the simplex is at most $\varepsilon|\mathcal{H}|$.

We can use this to design a prune and search algorithm as follows: (1) construct an ε -net \mathcal{N} , (2) compute the cell C of $\mathcal{A}(\mathcal{N})$ containing the input point q in its interior, (3) construct a simplex S inscribed in C and containing q in its interior, (4) recurse on the hyperplanes of \mathcal{H} intersecting the interior of S .

Proceeding this way with a constant ε guarantees that at most a constant fraction ε of the hyperplanes remains after the pruning step, and thus the cumulative number of queries made to determine the enclosing cell at each step is $O(n^2 \log^2 n \log |\mathcal{H}|)$ when done in a brute-force way. However, we still need to explain how to find a simplex S inscribed in C and containing q in its interior. This procedure corresponds to the well-known *bottom vertex decomposition* (or *triangulation*) of a hyperplane arrangement [21, 14].

3.2 Finding a simplex

In order to simplify the exposition of the algorithm, we assume, without loss of generality, that the input numbers q_i all lie in the interval $[-1, 1]$. This assumption is justified by observing that we can normalize all the input numbers by the largest absolute value of a component of q . One can then see that every linear query on the normalized input can be implemented as a linear query on the original input. A similar transformation can be carried out for the k -LDT problem. This allows us to use bounding hyperplanes of equations $x_i = \pm 1, i \in [n]$. We denote by \mathcal{B} this set of hyperplanes. Hence, if we choose a subset \mathcal{N} of the hyperplanes, the input point is located in a bounded cell of the arrangement $\mathcal{A}(\mathcal{N} \cup \mathcal{B})$. Note that $|\mathcal{N} \cup \mathcal{B}| = O(|\mathcal{N}|)$ for all interesting values of ε .

We now explain how to construct S under this assumption. The algorithm can be sketched as follows. (Recall that $\sigma(H, p)$ denotes the relative position of p with respect to the hyperplane H .)

Algorithm 1 (Constructing S).

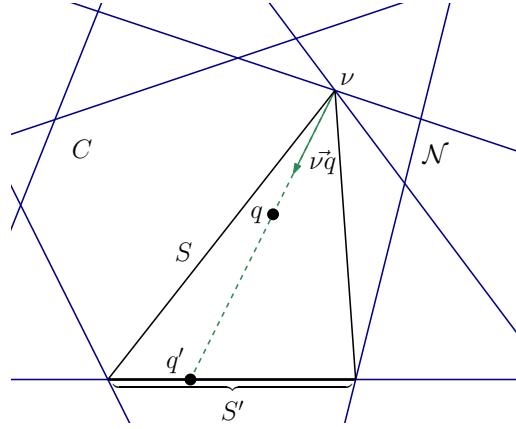
input A point q in $[-1, 1]^n$, a set \mathcal{I} of hyperplanes not containing q , and a set \mathcal{E} of hyperplanes in general position containing q , such that the cell

$$C = \{p: \sigma(H, p) = \sigma(H, q) \text{ or } \sigma(H, p) = 0 \text{ for all } H \in (\mathcal{I} \cup \mathcal{E})\}$$

is a bounded polytope. The value $\sigma(H, q)$ is known for all $H \in (\mathcal{I} \cup \mathcal{E})$.

output A simplex $S \in C$ that contains q in its interior (if it is not a point), and all vertices of which are vertices of C .

0. If $|\mathcal{E}| = n$, return q .
1. Determine a vertex ν of C .



■ **Figure 1** Illustration of a step of Algorithm 1.

2. Let q' be the projection of q along $\nu\vec{q}$ on the boundary of C . Compute $\mathcal{I}_\theta \subseteq \mathcal{I}$, the subset of hyperplanes in \mathcal{I} containing q' . Compute $\mathcal{I}_\tau \subseteq \mathcal{I}_\theta$, a maximal subset of those hyperplanes such that $\mathcal{E}' = \mathcal{E} \cup \mathcal{I}_\tau$ is a set of hyperplanes in general position.
3. Recurse on q' , $\mathcal{I}' = \mathcal{I} \setminus \mathcal{I}_\theta$, and \mathcal{E}' , and store the result in S' .
4. Return S , the convex hull of $S' \cup \{\nu\}$.

Step **0** is the base case of the recursion: when there is only one point left, just return that point. This step uses no query.

We can solve step **1** by using linear programming with the known values of $\sigma(H, q)$ as linear constraints. We arbitrarily choose an objective function with a gradient non-orthogonal to all hyperplanes in \mathcal{I} and look for the optimal solution. The optimal solution being a vertex of the arrangement, its coordinates are independent of q , and thus this step involves no query at all.

Step **2** prepares the recursive step by finding the hyperplanes containing q' . This can be implemented as a ray-shooting algorithm that performs a number of comparisons between projections of q on different hyperplanes of \mathcal{I} without explicitly computing them. In Appendix A, we prove that all such comparisons can be implemented using $O(|\mathcal{I}|)$ linear queries. Constructing \mathcal{E}' can be done by solving systems of linear equations that do not involve q .

In step **3**, the input conditions are satisfied, that is, $q' \in [-1, 1]^n$, \mathcal{I}' is a set of hyperplanes not containing q' , \mathcal{E}' is a set of hyperplanes in general position containing q' , C' is a d -cell of C and is thus a bounded polytope. The value $\sigma(H, q')$ differs from $\sigma(H, q)$ only for hyperplanes that have been removed from \mathcal{I} , and for those $\sigma(H, q') = 0$, hence we know all necessary values $\sigma(H, q')$ in advance.

Since $|\mathcal{I}'| < |\mathcal{I}|$, $|\mathcal{E}'| > |\mathcal{E}|$, and $|\mathcal{I} \setminus \mathcal{I}'| - |\mathcal{E}' \setminus \mathcal{E}| \geq 0$, the complexity of the recursive call is no more than that of the parent call, and the maximal depth of the recursion is n . Thus, the total number of linear queries made to compute S is $O(n|\mathcal{I}|)$.

Hence given an input point $q \in [-1, 1]$, an arrangement of hyperplanes $\mathcal{A}(\mathcal{N})$, and the value of $\sigma(H, q)$ for all $H \in (\mathcal{N} \cup \mathcal{B})$, we can compute the desired simplex S by running Algorithm 1 on q , $\mathcal{I} = \{H \in (\mathcal{N} \cup \mathcal{B}): \sigma(H, q) \neq 0\}$, and $\mathcal{E} \subseteq (\mathcal{N} \cup \mathcal{B}) \setminus \mathcal{I}$. This uses $O(n^3 \log^2 n)$ linear queries. Figure 1 illustrates a step of the algorithm.

3.3 Assembling the pieces

Let us summarize the algorithm

Algorithm 2.

input $q \in [-1, 1]^n$

1. Pick $O(n^2 \log^2 n)$ hyperplanes of \mathcal{H} at random and locate q in this arrangement. Call C the cell containing q .
2. Construct the simplex S containing q and inscribed in C , using Algorithm 1.
3. For every hyperplane of \mathcal{H} containing S , output a solution.
4. Recurse on hyperplanes of \mathcal{H} intersecting the interior of S .

The query complexity of step 1 is $O(n^2 \log^2 n)$, and that of step 2 is $O(n^3 \log^2 n)$. Steps 3 and 4 do not involve any query at all. The recursion depth is $O(\log |\mathcal{H}|)$, with $|\mathcal{H}| = O(n^k)$, hence the total query complexity of this algorithm is $O(n^3 \log^3 n)$. This proves the first part of Theorem 3.

We can also consider the overall complexity of the algorithm in the RAM model, that is, taking into account the steps that do not require any query, but for which we still have to process the set \mathcal{H} . Note that the complexity bottleneck of the algorithm are steps 3-4, where we need to prune the list of hyperplanes according to their relative positions with respect to S . For this purpose, we simply maintain explicitly the list of all hyperplanes, starting with the initial set corresponding to all k -tuples. Then the pruning step can be performed by looking at the position of each vertex of S relative to each hyperplane of \mathcal{H} . Because in our case hyperplanes have only k nonzero coefficients, this uses a number of integer arithmetic operations on $\tilde{O}(n)$ bits integers that is proportional to the number of vertices times the number of hyperplanes. (For the justification of the bound on the number of bits needed to represent vertices of the arrangement see Appendix D.) Since we recurse on a fraction of the set, the overall complexity is $\tilde{O}(n^2 |\mathcal{H}|) = \tilde{O}(n^{k+2})$. The next section is devoted to improving this running time.

4 Time complexity

Proving the second part of Theorem 3 involves efficient implementations of the two most time-consuming steps of Algorithm 2. In order to efficiently implement the pruning step, we define an intermediate problem, that we call the *double k -SUM* problem.

► **Problem (double k -SUM).** *Given two vectors $\nu_1, \nu_2 \in [-1, 1]^n$, where the coordinates of ν_i can be written down as fractions whose numerator and denominator lie in the interval $[-M, M]$, enumerate all $i \in [n]^k$ such that*

$$\left(\sum_{j=1}^k \nu_{1,i_j} \right) \left(\sum_{j=1}^k \nu_{2,i_j} \right) < 0.$$

In other words, we wish to list all hyperplanes of \mathcal{H} intersecting the open line segment $\nu_1 \nu_2$. We give an efficient output-sensitive algorithm for this problem.

► **Lemma 5.** *The double k -SUM problem can be solved in time $O(n^{\lceil \frac{k}{2} \rceil} \log n \log M + Z)$, where Z is the size of the solution.*

Proof. If k is even, we consider all possible $\frac{k}{2}$ -tuples of numbers in ν_1 and ν_2 and sort their sums in increasing order. This takes time $O(n^{\frac{k}{2}} \log n)$ and yields two permutations π_1 and π_2 of $[n^{\frac{k}{2}}]$. If k is odd, then we sort both the $\lceil \frac{k}{2} \rceil$ -tuples and the $\lfloor \frac{k}{2} \rfloor$ -tuples. For simplicity, we will only consider the even case in what follows. The odd case carries through.

We let $N = n^{\frac{k}{2}}$. For $i \in [N]$ and $m \in \{1, 2\}$, let $\Sigma_{m,i}$ be the sum of the $\frac{k}{2}$ components of the i th $\frac{k}{2}$ -tuple in ν_m , in the order prescribed by π_m .

We now consider the two $N \times N$ matrices M_1 and M_2 giving all possible sums of two $\frac{k}{2}$ -tuples, for both ν_1 with the ordering π_1 and ν_2 with the ordering π_2 .

We first solve the k -SUM problem on ν_1 , by finding the sign of all pairs $\Sigma_{1,i} + \Sigma_{1,j}$, $i, j \in [N]$. This can be done in time $O(N)$ by parsing the matrix M_1 , just as in the standard k -SUM algorithm. We do the same with M_2 .

The set of all indices $i, j \in [N]$ such that $\Sigma_{1,i} + \Sigma_{1,j}$ is positive forms a staircase in M_1 . We sweep M_1 column by column in order of increasing $j \in [N]$, in such a way that the number of indices i such that $\Sigma_{1,i} + \Sigma_{1,j} > 0$ is growing. For each new such value i that is encountered during the sweep, we insert the corresponding $i' = \pi_2(\pi_1^{-1}(i))$ in a balanced binary search tree.

After each sweep step in M_1 — that is, after incrementing j and adding the set of new indices i' in the tree — we search the tree to identify all the indices i' such that $\Sigma_{2,i'} + \Sigma_{2,j'} < 0$, where $j' = \pi_2(\pi_1^{-1}(j))$. Since those indices form an interval in the ordering π_2 when restricted to the indices in the tree, we can search for the largest i'_0 such that $\Sigma_{2,i'_0} < -\Sigma_{2,j'}$ and retain all indices $i' \leq i'_0$ that are in the tree. If we denote by z the number of such indices, this can be done in $O(\log N + z) = O(\log n + z)$ time. Now all the pairs i', j' found in this way are such that $\Sigma_{1,i} + \Sigma_{1,j}$ is positive and $\Sigma_{2,i'} + \Sigma_{2,j'}$ is negative, hence we can output the corresponding k -tuples. To get all the pairs i', j' such that $\Sigma_{1,i} + \Sigma_{1,j}$ is negative and $\Sigma_{2,i'} + \Sigma_{2,j'}$ positive, we repeat the sweeping algorithm after swapping the roles of ν_1 and ν_2 .

Every matching k -tuple is output exactly once, and every $\frac{k}{2}$ -tuple is inserted at most once in the binary search tree. Hence the algorithm runs in the claimed time.

Note that we only manipulate rational numbers that are the sum of at most k rational numbers of size $O(\log M)$. ◀

Now observe that a hyperplane intersects the interior of a simplex if and only if it intersects the interior of one of its edges. Hence given a simplex S we can find all hyperplanes of \mathcal{H} intersecting its interior by running the above algorithm $\binom{n}{2}$ times, once for each pair of vertices (ν_1, ν_2) of S , and take the union of the solutions. The overall running time for this implementation will therefore be $\tilde{O}(n^2(n^{\lceil \frac{k}{2} \rceil} \log M + Z))$, where Z is at most the number of intersecting hyperplanes and M is to be determined later. This provides an implementation of the pruning step in Meiser's algorithm, that is, step 4 of Algorithm 2.

► **Corollary 6.** *Given a simplex S , we can compute all k -SUM hyperplanes intersecting its interior in $\tilde{O}(n^2(n^{\lceil \frac{k}{2} \rceil} \log M + Z))$ time, where M is proportional to the number of bits necessary to represent S .*

In order to detect solutions in step 3 of Algorithm 2, we also need to be able to quickly solve the following problem.

► **Problem (multiple k -SUM).** *Given d points $\nu_1, \nu_2, \dots, \nu_d \in \mathbb{R}^n$, where the coordinates of ν_i can be written down as fractions whose numerator and denominator lie in the interval $[-M, M]$, decide whether there exists a hyperplane with equation of the form $x_{i_1} + x_{i_2} + \dots + x_{i_k} = 0$ containing all of them.*

Here the standard k -SUM algorithm can be applied, taking advantage of the fact that the coordinates lie in a small discrete set.

► **Lemma 7.** *k -SUM on n integers $\in [-V, V]$ can be solved in time $\tilde{O}(n^{\lceil \frac{k}{2} \rceil} \log V)$.*

► **Lemma 8.** *Multiple k -SUM can be solved in time $\tilde{O}(dn^{\lceil \frac{k}{2} \rceil + 2} \log M)$.*

Proof. Let $\mu_{i,j}$ and $\delta_{i,j}$ be the numerator and denominator of $\nu_{i,j}$ when written as an irreducible fraction. We define

$$\zeta_{i,j} = \nu_{i,j} \prod_{(i,j) \in [d] \times [n]} \delta_{i,j} = \frac{\mu_{i,j} \prod_{(i',j') \in [d] \times [n]} \delta_{i',j'}}{\delta_{i,j}}.$$

By definition $\zeta_{i,j}$ is an integer and its absolute value is bounded by $U = M^{n^2}$, that is, it can be represented using $O(n^2 \log M)$ bits. Moreover, if one of the hyperplanes contains the point $(\zeta_{i,1}, \zeta_{i,2}, \dots, \zeta_{i,n})$, then it contains ν_i . Construct n integers of $O(dn^2 \log M)$ bits that can be written $\zeta_{1,j} + U, \zeta_{2,j} + U, \dots, \zeta_{d,j} + U$ in base $2Uk + 1$. The answer to our decision problem is “yes” if and only if there exists k of those numbers whose sum is kU, kU, \dots, kU . We simply subtract the number U, U, \dots, U to all n input numbers to obtain a standard k -SUM instance on n integers of $O(dn^2 \log M)$ bits. ◀

We now have efficient implementations of steps 3 and 4 of Algorithm 2 and can proceed to the proof of the second part of Theorem 3.

Proof. The main idea consists of modifying the first iteration of Algorithm 2, by letting $\varepsilon = n^{-\frac{k}{2}}$. Hence we pick a random subset \mathcal{N} of $O(n^{k/2+2} \log^2 n)$ hyperplanes in \mathcal{H} and use this as an ε -net. This can be done efficiently, as shown in Appendix C.

Next, we need to locate the input q in the arrangement induced by \mathcal{N} . This can be done by running Algorithm 2 on the set \mathcal{N} . From the previous considerations on Algorithm 2, the running time of this step is

$$O(n|\mathcal{N}|) = \tilde{O}(n^{k/2+4}),$$

and the number of queries is $O(n^3 \log^3 n)$.

Then, in order to be able to prune the hyperplanes in \mathcal{H} , we have to compute a simplex S that does not intersect any hyperplane of \mathcal{N} . For this, we observe that the above call to Algorithm 2 involves computing a sequence of simplices for the successive pruning steps. We save the description of those simplices. Recall that there are $O(\log n)$ of them, all of them contain the input q and have vertices coinciding with vertices of the original arrangement $\mathcal{A}(\mathcal{H})$. In order to compute a simplex S avoiding all hyperplanes of \mathcal{N} , we can simply apply Algorithm 1 on the set of hyperplanes bounding the intersection of these simplices. The running time and number of queries for this step are bounded respectively by $n^{O(1)}$ and $O(n^2 \log n)$.

Note that the vertices of S are not vertices of $\mathcal{A}(\mathcal{H})$ anymore. However, their coordinates lie in a finite set (see Appendix D)

► **Lemma 9.** *Vertices of S have rational coordinates whose fraction representations have their numerators and denominators absolute value bounded by $O(C^{4n^5} n^{4n^5+2n^3+n})$, where C is a constant.*

We now are in position to perform the pruning of the hyperplanes in \mathcal{H} with respect to S . The number of remaining hyperplanes after the pruning is at most $\varepsilon n^k = O(n^{k/2})$. Hence from Corollary 6, the pruning can be performed in time proportional $\tilde{O}(n^{\lceil k/2 \rceil + 7})$.

Similarly, we can detect any hyperplane of \mathcal{H} containing S using the result of Lemma 8 in time $\tilde{O}(n^{\lceil k/2 \rceil} + 8)$. Note that those last two steps do not require any query.

Finally, it remains to detect any solution that may lie in the remaining set of hyperplanes of size $O(n^{k/2})$. We can again fall back on Algorithm 2, restricted to those hyperplanes. The running time is $\tilde{O}(n^{k/2+2})$, and the number of queries is still $O(n^3 \log^3 n)$.

Overall, the maximum running time of a step is $\tilde{O}(n^{\lceil \frac{k}{2} \rceil} + 8)$, while the number of queries is always bounded by $O(n^3 \log^3 n)$. ◀

5 Query size

In this section, we consider a simple blocking scheme that allows us to explore a tradeoff between the number of queries and the size of the queries.

► **Lemma 10.** *For any integer $b > 0$, an instance of the k -SUM problem on $n > b$ numbers can be split into $O(b^{k-1})$ instances on at most $k \lceil \frac{n}{b} \rceil$ numbers, so that every k -tuple forming a solution is found in exactly one of the subproblems. The transformation can be carried out in time $O(n \log n + b^{k-1})$.*

Proof. Given an instance on n numbers, we can sort them in time $O(n \log n)$, then partition the sorted sequence into b consecutive blocks B_1, B_2, \dots, B_b of equal size. This partition can be associated with a partition of the real line into b intervals, say I_1, I_2, \dots, I_b . Now consider the partition of \mathbb{R}^k into grid cells defined by the k th power of the partition I_1, I_2, \dots, I_b . The hyperplane of equation $x_1 + x_2 + \dots + x_k = 0$ hits $O(b^{k-1})$ such grid cells. Each grid cell $I_{i_1} \times I_{i_2} \times \dots \times I_{i_k}$ corresponds to a k -SUM problem on the numbers in the set $B_{i_1} \cup B_{i_2} \cup \dots \cup B_{i_k}$ (note that the indices i_j need not be distinct). Hence each such instance has size at most $k \lceil \frac{n}{b} \rceil$. ◀

Combining Lemma 10 and Theorem 3 directly yields the following.

► **Theorem 11.** *For any integer $b > 0$, there exists a $k \lceil \frac{n}{b} \rceil$ -linear decision tree of depth $\tilde{O}(b^{k-4} n^3)$ solving the k -SUM problem. Moreover, this decision tree can be implemented as an $\tilde{O}(b^{\lfloor \frac{k}{2} \rfloor - 9} n^{\lceil \frac{k}{2} \rceil} + 8)$ Las Vegas algorithm.*

The following two corollaries are obtained by taking $b = \frac{k}{\varepsilon}$, and $b = O(n^\varepsilon)$, respectively

► **Corollary 12.** *For any constant $\varepsilon > 0$ such that $b = \frac{k}{\varepsilon}$ and $\frac{n}{b}$ are positive integers, there exists an εn -linear decision tree of depth $\tilde{O}(n^3)$ solving the k -SUM problem. Moreover, this decision tree can be implemented as an $\tilde{O}(n^{\lceil \frac{k}{2} \rceil} + 8)$ Las Vegas algorithm.*

► **Corollary 13.** *For any ε such that $0 < \varepsilon < 1$, there exists an $O(n^{1-\varepsilon})$ -linear decision tree of depth $\tilde{O}(n^{3+(k-4)\varepsilon})$ solving the k -SUM problem. Moreover, this decision tree can be implemented as an $\tilde{O}(n^{(1+\varepsilon)\frac{k}{2} + 8.5})$ Las Vegas algorithm.*

Note that the latter query complexity improves on $\tilde{O}(n^{\frac{k}{2}})$ whenever $\varepsilon < \frac{1}{2}$ and $k \geq \frac{3-4\varepsilon}{1/2-\varepsilon}$. Hence for instance, we obtain an $O(n^{\frac{3}{4}})$ -linear decision tree of depth $\tilde{O}(n^4)$ for the 8SUM problem, and $o(n)$ -linear decision trees of depth $o(n^4)$ for any k .

References

- 1 Amir Abboud, Kevin Lewi, and Ryan Williams. Losing weight by gaining edges. 2015. arXiv:1311.3054 [cs.CC].

- 2 Amir Abboud, Virginia Vassilevska Williams, and Oren Weimann. Consequences of faster alignment of sequences. In *International Colloquium on Automata, Languages, and Programming (ICALP 2014)*, volume 8572 of *Lecture Notes in Computer Science*, pages 39–51. Springer, 2014.
- 3 Amir Abboud, Virginia Vassilevska Williams, and Huacheng Yu. Matching triangles and basing hardness on an extremely popular conjecture. In *Symposium on Theory of Computing (STOC 2015)*, pages 41–50, 2015.
- 4 Nir Ailon and Bernard Chazelle. Lower bounds for linear degeneracy testing. *J. ACM*, 52(2):157–171, 2005.
- 5 Amihood Amir, Timothy M. Chan, Moshe Lewenstein, and Noa Lewenstein. On hardness of jumbled indexing. In *International Colloquium on Automata, Languages, and Programming (ICALP 2014)*, volume 8572 of *Lecture Notes in Computer Science*, pages 114–125. Springer, 2014.
- 6 Meyer Friedhelm auf der Heide. A polynomial linear search algorithm for the n -dimensional knapsack problem. *J. ACM*, 31(3):668–676, 1984.
- 7 Ilya Baran, Erik D. Demaine, and Mihai Patrascu. Subquadratic algorithms for 3SUM. *Algorithmica*, 50(4):584–596, 2008.
- 8 Gill Barequet and Sariel Har-Peled. Polygon-containment and translational min-hausdorff-distance between segment sets are 3SUM-hard. In *Symposium on Discrete Algorithms (SODA 1999)*, pages 862–863, 1999.
- 9 Peter Bürgisser, Michael Clausen, and Mohammad Amin Shokrollahi. *Algebraic complexity theory*, volume 315 of *Grundlehren der mathematischen Wissenschaften*. Springer, 1997.
- 10 Jean Cardinal, Samuel Fiorini, Gwenaël Joret, Raphaël M Jungers, and J Ian Munro. An efficient algorithm for partial order production. *SIAM journal on computing*, 39(7):2927–2940, 2010.
- 11 Jean Cardinal, Samuel Fiorini, Gwenaël Joret, Raphaël M Jungers, and J Ian Munro. Sorting under partial information (without the ellipsoid algorithm). *Combinatorica*, 33(6):655–697, 2013.
- 12 Marco Carmosino, Jiawei Gao, Russell Impagliazzo, Ivan Mikhailin, Ramamohan Paturi, and Stefan Schneider. Nondeterministic extensions of the strong exponential time hypothesis and consequences for non-reducibility. *Electronic Colloquium on Computational Complexity (ECCC 2015)*, 22:148, 2015.
- 13 Timothy M. Chan and Moshe Lewenstein. Clustered integer 3SUM via additive combinatorics. In Rocco A. Servedio and Ronitt Rubinfeld, editors, *Symposium on Theory of Computing (STOC 2015)*, pages 31–40. ACM, 2015.
- 14 Kenneth L. Clarkson. A randomized algorithm for closest-point queries. *SIAM J. Comput.*, 17(4):830–847, 1988.
- 15 David P. Dobkin and Richard J. Lipton. On some generalizations of binary search. In *Symposium on Theory of Computing (STOC 1974)*, pages 310–316, 1974.
- 16 David P. Dobkin and Richard J. Lipton. A lower bound of $\frac{1}{2}n^2$ on linear search programs for the knapsack problem. *J. Comput. Syst. Sci.*, 16(3):413–417, 1978.
- 17 Jeff Erickson. Lower bounds for linear satisfiability problems. *Chicago Journal of Theoretical Computer Science*, 1999.
- 18 Ari Freund. Improved subquadratic 3SUM. *Algorithmica*, 2015. To appear.
- 19 Anka Gajentaan and Mark H. Overmars. On a class of $O(n^2)$ problems in computational geometry. *Comput. Geom.*, 5:165–185, 1995.
- 20 O. Gold and M. Sharir. Improved bounds for 3SUM, k -SUM, and linear degeneracy. *ArXiv e-prints*, 2015. arXiv:1512.05279 [cs.DS].

- 21 Jacob E. Goodman and Joseph O'Rourke, editors. *Handbook of Discrete and Computational Geometry, Second Edition*. Chapman and Hall/CRC, 2004.
- 22 Allan Grønlund and Seth Pettie. Threesomes, degenerates, and love triangles. In *Foundations of Computer Science (FOCS 2014)*, pages 621–630. IEEE, 2014.
- 23 David Haussler and Emo Welzl. ϵ -nets and simplex range queries. *Discrete & Computational Geometry*, 2(1):127–151, 1987.
- 24 T. Kopelowitz, S. Pettie, and E. Porat. Higher lower bounds from the 3SUM conjecture. *ArXiv e-prints*, 2014. arXiv:1407.6756 [cs.DS].
- 25 S. Meiser. Point location in arrangements of hyperplanes. *Information and Computation*, 106(2):286–303, 1993.
- 26 Joseph S. B. Mitchell and Joseph O'Rourke. Computational geometry column 42. *Int. J. Comput. Geometry Appl.*, 11(5):573–582, 2001.
- 27 Mihai Patrascu. Towards polynomial lower bounds for dynamic problems. In *Symposium on Theory of Computing (STOC 2010)*, pages 603–610, 2010.
- 28 Mihai Patrascu and Ryan Williams. On the possibility of faster SAT algorithms. In *Symposium on Discrete Algorithms (SODA 2010)*, pages 1065–1075, 2010.
- 29 William L. Steiger and Ileana Streinu. A pseudo-algorithmic separation of lines from pseudo-lines. *Information Processing Letters*, 53(5):295–299, 1995.
- 30 Andrew Chi-Chih Yao. On parallel computation for the knapsack problem. *J. ACM*, 29(3):898–903, 1982.

A Keeping queries linear in Algorithm 1

In Algorithm 1, we want to ensure that the queries we make in step 2 are linear and that the queries we will make in the recursion step remain linear too.

► **Lemma 14.** *Algorithm 1 can be implemented so that it uses $O(n|I|)$ linear queries.*

Proof. Let us first analyze what the queries of step 2 look like. In addition to the input point q we are given a vertex ν and we want to find the projection q' of q in direction $\vec{\nu}q$ on the hyperplanes of \mathcal{I}_θ . Let the equation of H_i be $\Pi_i(x) = c_i + d_i \cdot x = 0$ where c_i is a scalar and d_i is a vector. The projection of q along $\vec{\nu}q$ on a hyperplane H_i can thus be written³ $\rho(q, \nu, H_i) = \nu + \lambda_i \vec{\nu}q$ such that $\Pi_i(\nu + \lambda_i \vec{\nu}q) = c_i + d_i \cdot \nu + \lambda_i d_i \cdot \vec{\nu}q = 0$. Computing the closest hyperplane amounts to finding $\lambda_\theta = \min_{\lambda_i > 0} \lambda_i$. Since $\lambda_i = -\frac{c_i + d_i \cdot \nu}{d_i \cdot \vec{\nu}q}$ we can test whether $\lambda_i > 0$ using the linear query⁴ $-\frac{d_i \cdot \vec{\nu}q}{c_i + d_i \cdot \nu} > ? 0$. Moreover, if $\lambda_i > 0$ and $\lambda_j > 0$ we can test whether $\lambda_i < \lambda_j$ using the linear query $\frac{d_i \cdot \vec{\nu}q}{c_i + d_i \cdot \nu} < ? \frac{d_j \cdot \vec{\nu}q}{c_j + d_j \cdot \nu}$. Step 2 can thus be achieved using $O(1)$ $(2k)$ -linear queries per hyperplane of \mathcal{N} .

In step 4, the recursive step is carried out on $q' = \nu + \lambda_\theta \vec{\nu}q = \nu - \frac{c_\theta + d_\theta \cdot \nu}{d_\theta \cdot \vec{\nu}q} \vec{\nu}q$ hence comparing λ'_i to 0 amounts to performing the query $-\frac{d_i \cdot \vec{\nu}q'}{c_i + d_i \cdot \nu'} > ? 0$, which is not linear in q . The same goes for comparing λ'_i to λ'_j with the query $\frac{d_i \cdot \vec{\nu}q'}{c_i + d_i \cdot \nu'} < ? \frac{d_j \cdot \vec{\nu}q'}{c_j + d_j \cdot \nu'}$.

However, we can multiply both sides of the inequality test by $d_\theta \vec{\nu}q$ to keep the queries linear as shown below. We must be careful to take into account the sign of the expression $d_\theta \vec{\nu}q$, this costs us one additional linear query.

³ Note that we project from ν instead of q . We are allowed to do this since $\nu + \lambda_i \vec{\nu}q = q + (\lambda_i - 1) \vec{\nu}q$ and there is no hyperplane separating q from ν .

⁴ Note that if $c_i + d_i \cdot \nu = 0$ then $\lambda_i = 0$, we can check this beforehand for free.

This trick can be used at each step of the recursion. Let $q^{(0)} = q$, then we have

$$q^{(s+1)} = \nu^{(s)} - \frac{c_{\theta_s} + d_{\theta_s} \cdot \nu^{(s)}}{d_{\theta_s} \cdot \vec{\nu}^{(s)}} \vec{\nu}^{(s)}$$

and $(d_{\theta_s} \cdot \vec{\nu}^{(s)})q^{(s+1)}$ yields a vector whose components are linear in $q^{(s)}$. Hence, $(\prod_{k=0}^s d_{\theta_k} \cdot \vec{\nu}^{(k)})q^{(s+1)}$ yields a vector whose components are linear in q , and for all pairs of vectors d_i and $\nu^{(s+1)}$ we have that $(\prod_{k=0}^s d_{\theta_k} \cdot \vec{\nu}^{(k)})(d_i \cdot \vec{\nu}^{(s+1)})$ is linear in q .

Hence at the s th recursive step of the algorithm, we will perform at most $|\mathcal{N}|$ linear queries of the type

$$-\left(\prod_{k=0}^{s-1} d_{\theta_k} \cdot \vec{\nu}^{(k)}\right) \frac{d_i \cdot \vec{\nu}^{(s)}}{c_i + d_i \cdot \nu^{(s)}} \stackrel{?}{>} 0$$

$|\mathcal{N}| - 1$ linear queries of the type

$$\left(\prod_{k=0}^{s-1} d_{\theta_k} \cdot \vec{\nu}^{(k)}\right) \frac{d_i \cdot \vec{\nu}^{(s)}}{c_i + d_i \cdot \nu^{(s)}} \stackrel{?}{<} \left(\prod_{k=0}^{s-1} d_{\theta_k} \cdot \vec{\nu}^{(k)}\right) \frac{d_j \cdot \vec{\nu}^{(s)}}{c_j + d_j \cdot \nu^{(s)}}$$

and a single linear query of the type

$$d_{\theta_{s-1}} \cdot \vec{\nu}^{(s-1)} \stackrel{?}{<} 0.$$

In order to detect all hyperplanes H_i such that $\lambda_i = \lambda_\theta$ we can afford to compute the query $f(q) > g(q)$ for all query $f(q) < g(q)$ that we issue, and vice versa.

Note that, without further analysis, the queries can become n -linear as soon as we enter the $\frac{n}{k}$ th recursive step. ◀

B Algebraic computation trees

We consider *algebraic computation trees*, whose internal nodes are labeled with arithmetic ($r \leftarrow o_1 \text{ op } o_2$, $\text{op} \in \{+, -, \times, \div\}$) and branching ($z : 0$) operations. We say that an algebraic computation tree T *realizes* an algorithm A if the paths from the root to the leaves of T correspond to the execution paths of A on all possible inputs $q \in \mathbb{R}^n$, where n is fixed. A leaf is labeled with the output of the corresponding execution path of A . Such a tree is *well-defined* if any internal node labeled $r \leftarrow o_1 \text{ op } o_2$ has outdegree 1 and is such that either $o_k = q_i$ for some i or there exists an ancestor $o_k \leftarrow x \text{ op } y$ of this node, and any internal node labeled $z : 0$ has outdegree 3 and is such that either $z = q_i$ for some i or there exists an ancestor $z \leftarrow x \text{ op } y$ of this node. In the algebraic computation tree model, we define the complexity $f(n)$ of an algorithm A to be the minimum depth of a well-defined computation tree that realizes A for inputs of size n .

In the algebraic computation tree model, we only count the operations that involve the input, that is, members of the input or results of previous operations involving the input. The following theorem follows immediately from the analysis of the linearity of queries

► **Theorem 15.** *The algebraic computation tree complexity of k -LDT is $\tilde{O}(n^3)$.*

Proof. We go through each step of Algorithm 2. Indeed, each k -linear query of step 1 can be implemented as $O(k)$ arithmetic operations, so step 1 has complexity $O(|\mathcal{N}|)$. The construction of the simplex in step 2 must be handled carefully. What we need to show is that each n -linear query we use can be implemented using $O(k)$ arithmetic operations. It is not difficult to see from the expressions given in Appendix A that a constant number

of arithmetic operations and dot products suffice to compute the queries. A dot product in this case involves a constant number of arithmetic operations because the d_i are such that they each have exactly k non-zero components. The only expression that involves a non-constant number of operations is the product $\prod_{k=0}^s d_{\theta_k} \cdot \vec{v}q^{(k)}$, but this is equivalent to $(\prod_{k=0}^{s-1} d_{\theta_k} \cdot \vec{v}q^{(k)})(d_{\theta_s} \cdot \vec{v}q^{(s)})$ where the first factor has already been computed during a previous step and the second factor is of constant complexity. Since each query costs a constant number of arithmetic operations and branching operations, step 2 has complexity $O(n|\mathcal{N}|)$. Finally, steps 3 and 4 are free since they do not involve the input. The complexity of Algorithm 2 in this model is thus also $O(n^3 \log^2 n \log |\mathcal{H}|)$. \blacktriangleleft

C Uniform random sampling

Theorem 4 requires us to pick a sample of the hyperplanes uniformly at random. Actually the theorem is a little stronger; we can draw each element of \mathcal{N} uniformly at random, only keeping distinct elements. This is not too difficult to achieve for k -LDT when the $\alpha_i, i \in [k]$ are all distinct: to pick a hyperplane of the form $\alpha_0 + \alpha_1 x_{i_1} + \alpha_2 x_{i_2} + \dots + \alpha_k x_{i_k} = 0$ uniformly at random, we can draw each $i_j \in [n]$ independently and there are n^k possible outcomes. However, in the case of k -SUM, we only have $\binom{n}{k}$ distinct hyperplanes. A simple dynamic programming approach solves the problem for k -SUM. For k -LDT we can use the same approach, once for each class of equal α_i .

► **Lemma 16.** *Given $n \in \mathbb{N}$ and $(\alpha_0, \alpha_1, \dots, \alpha_k) \in \mathbb{R}^{k+1}$, a uniform random sample \mathcal{N} of hyperplanes in \mathbb{R}^n with equations of the form $\alpha_0 + \alpha_1 x_{i_1} + \alpha_2 x_{i_2} + \dots + \alpha_k x_{i_k} = 0$ can be computed in time $O(k|\mathcal{N}|)$ and preprocessing time $O(kn)$.*

Proof. We want to pick an assignment $a = \{(\alpha_1, x_{i_1}), (\alpha_2, x_{i_2}), \dots, (\alpha_k, x_{i_k})\}$ uniformly at random. Note that all x_i are distinct while the α_j can be equal.

Without loss of generality, suppose $\alpha_1 \leq \alpha_2 \leq \dots \leq \alpha_k$. There is a bijection between assignments and lexicographically sorted k -tuples $((\alpha_1, x_{i_1}), (\alpha_2, x_{i_2}), \dots, (\alpha_k, x_{i_k}))$.

Observe that x_{i_j} can be drawn independently of $x_{i_{j'}}$ whenever $\alpha_j \neq \alpha_{j'}$. Hence, it suffices to generate a lexicographically sorted $|\chi|$ -tuple of x_i for each class χ of equal α_i .

Let $\omega(m, l)$ denote the number of lexicographically sorted l -tuples, where each element comes from a set of m distinct x_i . We have

$$\omega(m, l) = \begin{cases} 1 & \text{if } l = 0 \\ \sum_{i=1}^m \omega(i, l-1) & \text{otherwise.} \end{cases}$$

To pick such a tuple $(x_{i_1}, x_{i_2}, \dots, x_{i_l})$ uniformly at random we choose $x_{i_l} = x_o$ with probability

$$P(x_{i_l} = x_o) = \begin{cases} 0 & \text{if } o > m \\ \frac{\omega(o, l-1)}{\omega(m, l)} & \text{otherwise} \end{cases}$$

that we append to a prefix $(l-1)$ -tuple (apply the procedure recursively), whose elements come from a set of o symbols. If $l = 0$ we just return the empty tuple.

Obviously, the probability for a given l -tuple to be picked is equal to $\frac{1}{\omega(m, l)}$.

Let X denote the partition of the α_i into equivalence classes, then the number of assignments is equal to $\prod_{\chi \in X} \omega(n, |\chi|)$. (Note that for k -SUM this is simply $\omega(n, k)$ since there is only a single class of equivalence.) For each equivalence class χ we draw independently a lexicographically sorted $|\chi|$ -tuple on n symbols using the procedure above. This yields

a given assignment with probability $\frac{1}{\prod_{x \in X} \omega(n, |X|)}$. Hence, this corresponds to a uniform random draw over the assignments.

For given n and k , there are at most nk values $\omega(m, l)$ to compute, and for a given k -LDT instance, it must be computed only once. Once those values have been computed, making a random draw takes time $O(k)$.

D Proof of Lemma 9

► **Theorem 17** (Cramer's rule). *If a system of n linear equations for n unknowns, represented in matrix multiplication form $Ax = b$, has a unique solution $x = (x_1, x_2, \dots, x_n)^T$ then, for all $i \in [n]$,*

$$x_i = \frac{\det(A_i)}{\det(A)}$$

where A_i is A with the i th column replaced by the column vector b .

► **Lemma 18.** *The absolute value of the determinant of an $n \times n$ matrix M with integer entries is $O((Cn)^n)$ where C is the maximum absolute value in M .*

Proof. The determinant of such a matrix is the sum of $2n!$ integers with maximum absolute value C^n .

► **Lemma 19.** *The determinant of an $n \times n$ matrix M with rational entries can be represented as a fraction whose numerators and denominators absolute values are bounded by $O((ND^{n-1}n)^n)$ and $O(D^{n^2})$ respectively, where N and D are respectively the maximum absolute value of a numerator and a denominator.*

Proof. Multiply each row M_i of M by $\prod_j d_{i,j}$. Apply Lemma 18.

We can now proceed to the proof of Lemma 9.

Proof. Coefficients of the hyperplanes of the arrangement are constant rational numbers, those can be changed to constant integers (because each hyperplane has at most k nonzero coefficients). Let C denote the maximum absolute value of those coefficients.

Because of Theorem 17 and Lemma 18, vertices of the arrangement have rational coordinates whose numerators and denominators absolute values are bounded by $O(C^n n^n)$.

Given simplices whose vertices are vertices of the arrangement, hyperplanes that define the faces of those simplices have rational coefficients whose numerator and denominator are bounded by $O(C^{2n^3} n^{2n^3+n})$ by Theorem 17 and Lemma 19. (Note that some simplices might be not fully dimensional, but we can handle those by adding vertices with coordinates that are not much larger than that of already existing vertices).

By applying Theorem 17 and Lemma 19 again, we obtain that vertices of the arrangement of those new hyperplanes (and thus vertices of S) have rational coefficients whose numerator and denominator are bounded by $O(C^{4n^5} n^{4n^5+2n^3+n})$.